The Evidence Base for Savant

A review of the research behind 283 learning mechanisms in an adaptive learning system

607 distinct sources · 551 verified against their DOI or official page with a verbatim quote

Curriculum and placement · Retention and review · Adaptation · Skill tracks and games · Engagement, accessibility and wellbeing

Savant · savantedu.com

Compiled 25 October 2026 · citation and consistency audit 27 October 2026

Abstract

Savant is an adaptive learning app covering school subjects from preschool to grade 12, test preparation, languages and three skill tracks (chess, piano and basketball). This report sets out the research behind every rule in the app that affects how a person learns: each threshold, timing, ordering, gate, adaptation and screen behaviour, 283 mechanisms in all. For each mechanism it gives where it lives in the code, what it does with its exact numbers, the published evidence for and against it with a verbatim quote from each source, a graded verdict (strong, moderate, mixed, design choice, standard, gap or product rule), and whether the code matches the evidence.

The evidence base is 607 distinct sources: peer-reviewed studies, meta-analyses, official standards and regulations, and academic textbooks. Every source was checked against its DOI record or official page; 551 have a verbatim quote confirmed in the abstract, full text or page, and 56 are title-level only and never count toward a mechanism's minimum of three sources. Of the 283 mechanisms, 50 rest on strong evidence, 104 on moderate evidence, 33 on mixed evidence, and 77 are reasoned design choices whose general principle is supported but whose exact value no study has set. 15 follow a law or standard. The review also produced 45 evidence-backed improvement proposals, 32 of which were built, almost all behind individual on/off switches.

On 27 October 2026 the whole report was audited. Each of its 1,611 citation occurrences was checked against the cited source's abstract or official text: first by two independent language models, then by a stronger model in a final review, with planted errors mixed in to test the checkers (every model caught all of them). 235 occurrences were corrected, 62 of them by removing a source that did not show what it was cited for. As a result 17 mechanisms fell below three quoted sources. A second search added 23 verified replacement sources, and 4 mechanisms still rest on fewer than three quoted sources, each marked. 31 citation occurrences remain undecided, mostly because no abstract or open full text could be obtained. The report was also searched for internal contradictions, and 13 were fixed. The section "Citation and consistency audit" lists every change.

The research shows which direction is better for each choice: retrieval over rereading, spacing over massing, explanations over bare marks, mastery before moving on. It rarely sets Savant's exact numbers, and no study has tested these mechanisms working together. The report therefore marks the limits of the evidence throughout, and names the in-app experiments (a holdout group, a switch per change, outcomes on later unaided checks) that can confirm the whole system with real learners.

How to read this report

- **Mechanism labels** take the form 01-M7: section 01, mechanism 7. Proposals take the form 01-P3.
- **Citations.** A number in square brackets, such as [12], points to the numbered reference list at the end. The short label before it (for example roediger2006 [1]) is the label used in the working files.
- **Before and after 25 October 2026.** The sections were written from the code as it stood on 25 October 2026. Changes built from this review that day are recorded on each mechanism's **Status** line. Where a mechanism changed, its **What** line describes the earlier behaviour unless it says otherwise, and the Status line gives the current behaviour.
- **Code locations** (`file:line`) are as read on 25 October 2026; the changes may have moved lines.

Contents

Abstract 2
How to read this report 3
Summary 7
How the sources were checked 11
Citation and consistency audit (27 October 2026) 12
How the citations were checked 12
What the citation check found 12
Second pass: fuller text and replacement sources 37
Contradictions 40
How it all fits together 42
Section 01. Curriculum, levels and placement 45
01 at a glance 45
01-A. Placement: the difficulty ladder (`services/placement.js`) 48
01-B. Adaptive placement: the subject walk and the unit walk 52
01-C. Gap-only learning paths (`services/learner-path.js`, `services/gap-scan.js`) 56
01-D. The learning path and level layout (`config/curriculum.js`, `config/level-layout.js`) 58
01-E. Lesson structure (`services/lessons.js`, `public/js/session.js`) 64
01-F. Checkpoints (`services/block-checkpoints.js`, `services/sessions.js`) 73
01-G. Session planning (`services/sessions.js` planLesson, `services/learner.js` sessionParams, `content/pool.js`) 76
01-H. Mastery and advancement (`services/mastery.js`) 82
01-I. Grade-fit (`config/grade-skills.js`) 85
01-J. Objective coverage (`services/coverage.js`) 86
01-K. Remediation after a failed checkpoint (`services/remediation.js`) 89
01 proposals (each backed by 3 or more sources; status on each) 91
Section 02. Retention and review 99
02 at a glance 99
02-A. Retention enforcers: what they are and what they cost 102
02-B. The calendar 106
02-C. Which questions an enforcer asks 110
02-D. What counts as kept, and stakes 118

02-E. The daily review and spaced review 119
02-F. Consolidation before the Final 125
02-G. Retrieval inside lessons and sittings 125
02-H. Feedback 127
02-I. Sleep and day boundaries 129
02-J. The forgetting model and its recall estimates 129
The eglington2020 [28] question: should the enforcer's selection weights shift? 132
Other proposals (3 or more sources each; status marked on each) 134
Corrections to RETENTION-ENFORCER-EVIDENCE.md found here 138
Section 03. Scorecard and adaptation 139
The 10 weak rules: what the new research says 139
03 at a glance 141
03-A. The scorecard gate and the insertion contract 143
03-B. Style quiz, survey and onboarding 150
03-C. How the scorecard drives teaching (ADAPTIVE_TEACHING, default on) 152
03-D. The effort monitor (ADAPT_EFFORT, default on) 164
03-E. Coach guards 166
03-F. Grading leniency and the injection check 169
03-G. Objective items 171
03-H. The young-learner oral mode (YOUNG_ORAL, default on) 173
03-I. The lesson sandbox (LESSON_SANDBOX, default on) 174
03-J. Owner base vs optimized (product behaviour, no evidence needed) 175
03 proposals (each backed by 3 or more sources; status on each) 176
Section 04. Skill tracks (chess, piano, basketball) and learning games 180
04 at a glance 180
04-A. How a stage is proven (skill-coverage, skill-record, skill-gate) 182
04-B. How much and when to practise (skill-adaptation, skill-dose) 189
04-C. The stage checks: items, takes and counts (content/skill-stage-checks.js, skill-checks.js) 192
04-D. Piano: score-following and the sittings (public/engine/piano, skill-piano-bb-sittings) 197
04-E. Basketball: counts, scenes, tracking and the court (skill-checks, skill-piano-bb-sittings, public/engine/basketball) 206
04-F. Learning games (game-lab, game-payloads, GAME_HOLDS) 213
04 proposals (each backed by 3 or more sources; status on each) 216
Section 05. Engagement, UX and wellbeing 224

05 at a glance 224
05-A. Streaks, XP and rewards 226
05-B. Study sessions, length and breaks 232
05-C. Feedback, rationales, hints and the coach 235
05-D. Accessibility, mobile and visual design 243
05-E. Calendar, reminders, errors and empty states 247
05-F. Owner skips 249
05-G. Privacy, consent and minors 249
Code-versus-evidence gaps found 254
05 proposals (each backed by 3 or more sources; status on each) 254
What was not done 257
Open items 258
References 260

Summary

This article collects the research behind every learning mechanism in Savant into one place. A mechanism is any rule, number, order, timing, gate or screen behaviour that affects how someone learns. It merges five evidence sections written on 2026-10-25: curriculum and placement, retention and review, the scorecard and adaptation, the skill tracks and games, and engagement and wellbeing. The authors' findings are kept as written. Only inconsistencies and duplicated material were fixed, and every source now has one number in a single reference list at the end.

	Count
Mechanisms	**283** (01: 64, 02: 50, 03: 61, 04: 58, 05: 50)
With 3 or more quoted, non-title-level sources	**277** of 283
Flagged as fewer than 3	6: 03-M60 (product rule, 0 needed), 03-M61 (product rule, 0 needed). After the 27 October 2026 audit removed sources that did not support them: 03-M38, 03-M45, 04-M22, 04-M23 (2 each)
Verdict: strong	50
Verdict: moderate	104
Verdict: mixed	33
Verdict: design choice	77
Verdict: standard	15
Verdict: gap	2
Verdict: product rule	2
Distinct sources, deduplicated by DOI or URL	**607**
Verified by DOI and a verbatim quote	509
Verified by official URL and a verbatim quote	42
Title-level only (real DOI or URL, no quote; never counted toward the 3)	56
Distinct verified sources (DOI or official URL checked: the three rows above)	**607**
Not verified (cited once as title-level, no DOI or URL)	0
Proposals	45: 32 built on 25 October 2026 (fully or in part), 13 not built
Changes live on 25 October 2026 from this article	36

How to read the verdicts.

- **Strong:** several studies agree, or a meta-analysis (a study that pools many studies) supports it.
- **Moderate:** good support, but from fewer studies, or from a setting unlike Savant's.
- **Mixed:** the studies disagree. Each mixed verdict says what conflicts.
- **Design choice:** the general idea is supported, but no study sets the exact number Savant uses.

- **Standard:** a law, regulation or accessibility standard sets the rule, and the code is judged against it.
- **Gap:** there was no mechanism in the code when the section was written. The evidence said one would help.
- **Product rule:** an owner decision about the product, such as the owner's base version, so no learning evidence is needed.

What was improved on 25 October 2026 (overnight, 24 to 25 October). 36 changes backed by the research went live: 35 behind on/off switches, plus one new check that has no switch. Each passed its tests. The visible ones also passed a browser check at 375 and 1280 px. The learning changes do not apply to the owner's base version. The owner does get the safety rule `LEVEL_SELF_CHECK` and four interface fixes: `UX_TOUCH44`, `UX_STYLE_COPY`, `UX_TOAST_STAY` and `UX_REPORT_DIALOG`. `UX_HIDE_STATS` also applies to the owner, if the owner turns it on. Every switch defaults to on, and setting it to `off` and restarting turns the change off. The rollback steps for each are in `curriculum-build/LIVE-CHANGES-2026-10-24.md`. **The app on port 6767 needs a restart to load them.**

Switch	What changed	From	Mechanisms
`PLACEMENT_SKILL_BELIEF`	Placement gaps and strengths from a belief, with an "unsure" list (computed; not yet shown to learners)	01-P6	01-M6
`INTERLEAVE_SESSIONS`	Interleaved practice body after the first practice level	01-P1	01-M43
`CHECKPOINT_SPRT`	Checkpoint pass decisions by likelihood ratio; the Final needs 34/40	01-P3	01-M21
`CHALLENGE_SPRT`	Challenge Exam bands by likelihood ratio	01-P3	01-M23
`RETAIN_REAL_DELAY`	"Retained" needs a check 7 or more days after proof	01-P5	01-M58, 01-M59
`LEVEL_SELF_CHECK`	No unverified AI level or lesson without the verifier key (safety rule)	01-P10	01-M31, 01-M34
`CURVE_REVIEWS`	Lesson reviews from the forgetting curve	02-P4	01-M41, 01-M52, 02-M37, 02-M38
`OBJ_ITEMS_3_OPTIONS`	Objective items may have 3 options	03-P11	03-M53
`ENFORCER_DIFFICULTY_WINDOW_95`	Enforcer difficulty window widened to 0.95	02-P1	02-M17
`ENFORCER_DEFER_SITTING_RETENTION`	No duplicate end-of-sitting retention set while a study session runs	02-P3	01-M47, 02-M43
`ENFORCER_LEFT_OUT_FLAG`	"Left out last time" rotation flag	02-P5	02-M18
`SCORECARD_DROP_RUSH_TYPED`	No typed-answer switch for rushers	03-P1	03-M28, 05-M21
`SCORECARD_SOFT_STRETCH`	Softer challenge high side (+0.05, cap 0.25)	03-P2	03-M21, 03-M22
`SCORECARD_LENGTHEN_NEW_ONLY`	Fuller teaching only for new objectives	03-P3	03-M19
`SCORECARD_OBSERVED_PREFS`	Observed behaviour replaces the stated preference; scorecard term gated	03-P4	03-M15

Switch	What changed	From	Mechanisms
SCORECARD_WORKED_HINT_TEXT	Hint-dependency directive: a worked example, then try first	03-P5	03-M30
SCORECARD_SUBGOAL_TEXT	Self-efficacy directive: sub-goals and specific feedback	03-P6	03-M31
(no switch) reliability.js	Split-half reliability check, Spearman-Brown corrected (a check; needs real data to run)	03-P7	03-M8
RAPID_ITEM_RELATIVE	Item-relative rapid-answer line for long items	03-P8	02-M28, 03-M39, 03-M41, 05-M27
COACH_LEAK_CHOICE	The coach leak guard catches the key's choice number or letter	03-P10	03-M43
CHESS_PRACTICE_TARGET	Practice puzzles at a 0.75 expected score	04-P1	04-M21
PIANO_FEEDBACK_FADE	Bandwidth and on-request feedback on warm piano takes	04-P2	04-M36, 04-M37
PIANO_TEMPO_LADDER	Tempo ladder on warm piano sets	04-P3	04-M36
PIANO_HOTSPOT_LOOP	Hotspot loop on the worst 1-2 bars	04-P3	04-M36
SCENE_RELIABILITY_GATE	Scene-set split-half reliability gates the basketball stages	04-P5	04-M9, 04-M10
SKILL_LOAD_CAPS	Youth weekly load caps and rest days in basketball	04-P7	04-M11, 04-M16, 04-M51
SKILL_TWO_DAY_TEXT	The two-day rule described as a delayed test	04-P9	04-M2, 04-M41
CHESS_NOTHING_TO_FIND	"Nothing to find" positions in CH2 and CH3 tactics	04-P12	04-M19
STREAK_LOCAL_GRACE	Streak on the local day, one grace day a week	05-P2	05-M1
BREAK_NUDGE	Optional break line in the study strip	05-P1	05-M13
UX_TOUCH44	44 px touch targets on coarse pointers and phones	05-P4	05-M29
UX_HIDE_STATS	Learners or parents can hide streak and XP	05-P5	05-M4
UX_STYLE_COPY	Honest learning-style copy	05-P6	05-M9
UX_TOAST_STAY	Error toasts stay until dismissed	05 (gap list)	05-M33
UX_REPORT_DIALOG	"Report a problem" dialog instead of window.prompt	05 (gap list)	05-M26
UX_NEXT_STUDY_PLANNER	"When will you study next?" planner (sends nothing)	05-P3	05-M37, 05-M38

Also on 25 October 2026, separate from this article's proposals:

- **Grade fit against the published standards** (switch GRADE_SKILLS). 90 grade ranges were checked against CCSS and the CEFR, and 10 were changed. Every math, English and Reading draw from preschool to grade 12 now carries its school grade. In a sample of 102,400 items, every kind of above-grade item that was found went to 0. For example, 1,038 of 14,930 math items for preschool to grade 8 had been above grade.

- **Fact re-verification fixes** (no switch). Piano scale parts now show each scale's own note spelling: before, 59 of 102 items showed a wrong name. Grading was unchanged. The basketball aim cue now says "toward the back of the ring".
- **UI fixes** (no switch). The study strip has 44 px touch targets on phones. The placement page shows its grade subtitle at 375 px.

What research can prove, and what it can't. Each mechanism here either has research behind it, or is clearly labelled: design choice, mixed, title-level only, or fewer than three sources. That is as far as published research can go. Most of the studies were run on other learners, in other settings and with other materials. They show which direction is better: retrieval beats rereading, spaced practice beats massed practice, explanations beat bare marks. They rarely give Savant's exact numbers, and they never test these 283 mechanisms working together. **That the whole combination is the best choice for each learner can only be shown with real learner data, through A/B tests inside the app:** a holdout group, a switch per change, and outcomes measured on later, unaided checks. Many mechanisms are already built for this. The 10% holdout (01-M35), the switches added on 25 October 2026 and the shadow modes record what would change. The open items list what needs that data.

How the sources were checked

How the sources were checked. This is the same method in all five sections.

- **Only credible sources count:** peer-reviewed papers, meta-analyses, official standards and regulations, and textbooks from academic publishers. Blogs, vendor marketing and AI-written text never count.
- **Every DOI was checked** on Crossref (or DataCite or OpenAlex), and the title, first author and year had to match. A standard or law is cited by its official URL.
- **Each claim has a verbatim quote** under 25 words, from the abstract, the full text or the official page. A script checked that each quote appears in the stored text. Sources reused from `scorecard-sources.json` and `retention-sources.json` were already DOI-verified. A reused source was quoted again when it was used for a new claim.
- **"Title-level only"** means the DOI or URL is real, but no abstract or text could be fetched. Such a source never counts toward a mechanism's three sources.
- **Deduplication.** Sources are counted once by DOI, or by URL when there is no DOI. Two DOIs for the same paper (a JSTOR copy and the publisher's) count once. So do two ids for the same page. Each WCAG success criterion, ICO standard and GDPR or COPPA article is its own page, so each counts separately.
- **Citations.** A number in square brackets, such as [12], points to the reference list. The short id before it (for example roediger2006 [1]) is the authors' own label. It is kept so this article can be checked against the section files.
- **Where:** `file:line` locations are as the authors read them on 2026-10-25. The 25 October 2026 changes may have moved lines.
- **Mechanism labels** take the form 01-M7: section 01, mechanism 7. Proposals take the form 01-P3.

Citation and consistency audit (27 October 2026)

After the five sections were merged, the whole article was audited in two ways. Every citation was checked against its source, and the text was searched for contradictions it did not itself explain.

How the citations were checked

- **Every citation occurrence.** The article cites a source 1,611 times, counting each source once per line it appears on. Each occurrence was checked against four things: the passage it supports, the source's reference, the verbatim quote the article gives, and the source's own abstract or official page. Abstracts were fetched from OpenAlex, Semantic Scholar, Crossref and Europe PMC, and official pages were fetched directly. 535 of the 584 sources had an abstract or page text. For the other 49, the check used the quote verified earlier and the reference.
- **What "supports" means.** The source must show what the article uses it for. That covers the topic, the direction of the effect, the population and setting, every number attributed to it, and the strength of the wording. "May" must not become "does", one study must not become "studies show", and a correlation must not become a cause. Where a mechanism's verdict is "design choice", a source only has to support the general principle, not the exact number.
- **Two independent checkers.** Each occurrence was first rated by two different language models working separately. The first was GLM 5.3 (Qwen 3.8 Max for 5 passages). The second was DeepSeek V4.1, Qwen 3.8 Max or Claude Sonnet 5.5. Both said "supported" for 1,211 occurrences, and 396 were questioned by at least one.
- **A final review of every occurrence by a stronger model.** Claude Sonnet 5.5 reviewed the questioned occurrences. A spot check of 100 occurrences that both checkers had passed found that the stronger model still questioned some. So every remaining occurrence was reviewed as well, by GPT-6.1 Sol. In all, 308 occurrences were finally reviewed by Sonnet and 1,303 by Sol. Each final reviewer decided from the source text and proposed the smallest correction.
- **Planted-error controls.** 24 planted errors were mixed into the work: real passages with their source swapped for an unrelated one. Every model used caught every planted error: GLM 24/24, Qwen 24/24, Sonnet 24/24 and Sol 24/24.
- **Guards on every correction.** A correction was applied only if any new quoted wording appears verbatim in the source's text, and any new number appears in the source. Every correction was then read by hand. Where a proposed correction was itself wrong or went too far, it was rewritten from the source or not applied.

What the citation check found

- **1,307 of 1,611** occurrences (81.1%) were confirmed as written.
- **Second check of every corrected line.** All 265 citation lines that changed were checked again against their sources by two models (Sol and GLM 5.3), and every line either

questioned was read by hand. 6 more corrections came from this second check; they are in the table below.

- **235** occurrences (14.6%) were corrected. In 62 of them the citation was removed, because the source does not show what it was cited for. The others were reworded to match the source: a narrower claim, the right setting, population or number, or a fuller quote.
- **31** occurrences (1.9%) are still undecided after a second pass with fuller text (see below). They are left as they were. Each of these sources was checked earlier by DOI or official URL with a quote, or is marked title-level only.
- **38** questioned occurrences were kept unchanged after review. 16 only had a source label whose year differed from the reference; the year corrections below settle those. 7 were already corrected through another citation on the same line or by a hand edit. For the other 15, the source text was read by hand and supports the claim as written; the final reviewer had asked for more than the claim needs, or proposed quoted wording that is not in the source.
- **Effect on the counts.** After the removals, 17 mechanisms had fewer than three quoted sources, and the second pass below left two more (03-M43 and 05-M5) in the same position. Replacement sources were added to 15 of those 19. 4 mechanisms (03-M38, 03-M45, 04-M22, 04-M23) still have two. They are marked on their verdict line, and in the summary and section counts. Two proposals (01-P8 and 02-P5) now rest on two sources, and each says so.
- **Publication years.** 22 reference years were corrected to the year of the printed issue. These are Crossref's "published-print" years; the earlier year was the online-first date. 7 source labels were renamed to match their reference's year.

Every correction:

Citation	Source	Problem found	Correction
01-P5	[159]	The meta-analysis supports spacing retrieval episodes over massed practice, but finds no significant overall advantage for expanding over uniform intervals (and says it does not support progressively increasing intervals as a general rule).	Citation removed.
03-M25	[137]	The meta-analysis supports meaningful, learner-paced segmentation of multimedia instruction, with small-to-medium benefits for retention and transfer and reduced cognitive load. It does not show that shorter app sections or 20% fewer items improve persistence.	Now reads: - rey2019 [137] — meaningful, learner-paced segmentation of multimedia instruction modestly improves retention and transfer and reduces cognitive load (above).
03-M31	[287]	Schunk (1983) supports feedback about progress, but the cited study does not establish the use of proximal subgoals; it studied attributional feedback during subtraction training.	Now reads: feedback on progress
03-M43	[33]	The review addresses balancing information giving and withholding, so it is relevant to a coach that withholds answers.	Citation removed.
03-M46	[33]	The review concerns balancing assistance and withholding to support learning, and says that optimal conditions remain an open problem. It does not	Now reads: - koedinger2007 [33] (above; reviews balancing assistance and withholding, not the

Citation	Source	Problem found	Correction
		establish that assisted answers are weaker evidence in tutoring assessment.	evidential weight of assisted answers).
04-M14	[259]	The meta-analysis found an overall spacing advantage across 63 studies, but it also found that task type and intertrial interval moderated the effect.	Now reads: Spacing showed an overall advantage across tasks, moderated by task type and interval (lee1988 [374], donovan1999 [259]).
Scorecard weak rules (section 03)	[271]	The abstract supports the characterization of assessment cut-offs as rules of thumb, but it does not support the additional claim that they are often misattributed.	Now reads: cut-offs are rules of thumb
Scorecard weak rules (section 03)	[110]	The source supports worked examples in tutored problem solving and adaptive fading, but the supplied abstract does not specifically establish that up-front worked examples are better supported.	Now reads: Worked examples are better supported
Section 03 proposals	[294]	The source is a relevant meta-analysis of achievement emotions and academic performance, but the supplied abstract/page text does not show the specific frustration correlation of rho = -0.02.	Now reads: add camachomorles2021 [294] on frustration and academic performance
04-M33	[447]	Lerch et al. review music performance analysis and discuss measuring performance parameters such as timing, but the source text does not specifically address piano pedalling or grading pedal changes against chord onsets.	Now reads: lerch2020 [447] (general background on performance-timing analysis; as M32).
04-M26	[425]	The source studies recall of presented chess positions and long-term memory templates, not calculating or generating a sequence of moves without moving the pieces.	Now reads: - gobet1996 [425] (as M25; supports internal chess-position memory structures, not move calculation).
04-P2	[496]	Salmoni et al. discuss the guidance hypothesis, but the abstract does not specifically support reduced feedback frequency; it also contains substantive discussion, so calling this source “title-level only” understates what it says.	Changed to: salmoni1984 [496] discusses the guidance hypothesis; it does not by itself establish benefits of reduced feedback frequency. (revised again on the second check; see below)
04-P7	[408]	The IOC abstract supports sustainable, healthy youth sports participation in general, but it does not establish the specific age-based weekly caps, rest-day numbers, or that those limits reduce overuse injury and burnout.	Now reads: - bergeron2015 [408] (general youth-athlete development guidance; not evidence for the specific caps).
05-M8	[516]	The source is not title-level only: its abstract describes a conceptual dashboard model for process-oriented feedback grounded in learning regulation and learner goals.	Now reads: sedrakyan2020 [516]: proposes dashboard feedback grounded in learning regulation and learner goals.
05-M5	[509]	The source is about behaviors, motivations, and emotions associated with gaming the system, but it does not show that making help less rewarding prevents help from becoming a free substitute for trying or supports the XP discount.	Citation removed.

Citation	Source	Problem found	Correction
05-M32	[127]	The source accurately supports the claim that interesting but irrelevant additions can reduce recall and transfer, but it does not study motion, animation, reduced-motion settings, or accessibility. Its relevance to this mechanism is therefore indirect.	Now reads: (Four experiments, students reading with interesting but irrelevant additions; indirect background evidence, not a motion study)
03-M7	[305]	The source supports a random-effects approach for pooling results from related studies and accounting for heterogeneity, but the abstract does not specifically establish normal pooling of log-odds estimates.	Now reads: random-effects pooling of estimates from related studies
03-M6	[305]	The source supports random-effects pooling and accounting for heterogeneity between clinical-trial effects, but it does not directly support empirical-Bayes shrinkage of individual learners or the stated method-of-moments implementation.	Now reads: - dersimonian1986 [305] — "We discuss a random effects approach to combining evidence from a series of experiments comparing two treatments." — clinical-trial random-effects modeling of between-study heterogeneity, not learner-level…
04-M41	[376]	Walker et al. supports sleep-dependent improvement in a practiced motor skill after an overnight interval, but it does not establish that a next-day cold attempt is the appropriate lesson score and is not evidence from children; the article's own verdict…	Now reads: - walker2002 [376] (as M02; supports sleep-dependent motor-skill improvement after practice, not specifically a child day-2 scoring rule).
01-M18	[95]	The source supports a region-of-proximal-learning model of study-time allocation, but it does not support the described linear grade-band difficulty progression or a steady increase in challenge.	Now reads: *Setting:* study-time allocation (region of proximal learning); not evidence for a linear difficulty progression.
03-M23	[293]	The source supports restoring equilibrium through thought, reflection, and problem solving when a learner encounters an impasse, but it does not specifically establish restoring the learner's sense of control.	Now reads: support restoring equilibrium when a learner is stuck
03-M31	[332]	The review supports the broader point that social-cognitive-theory-based educational programmes can enhance self-efficacy in higher education, but it does not specifically support the success-first or mastery-experience mechanism described here.	Now reads: Context: vandinther2011 [332] found that educational programmes based on social cognitive theory could enhance self-efficacy — review, higher education.
03-M34	[335]	The source compares rereading with a single reading and finds rereading usually adds little or no benefit; it does not compare rereading with retrieval or establish that the recap is retrieval rather than rereading.	Now reads: Rereading often added little or no benefit in experiments with educational texts (callender2009 [335]).
04-P1	[78]	The source is not merely title-level: its abstract reports that arithmetic practice items were sampled at a mean success probability of 0.75, with tasks described as challenging but not too difficult.	Changed to: klinkenberg2011 [78] reports arithmetic practice items sampled at mean success probability 0.75; this is not direct evidence about chess. (revised again on the second check; see below)

Citation	Source	Problem found	Correction
01-M8	[58]	Source [58] (wainer1987) is about testlet construction and item bundling in CAT, not about variable-length stopping rules or posterior-based termination criteria.	Citation removed.
01-M11	[75]	Source [75] is a systematic literature review about simulated learners in educational technology generally; it does not discuss this specific placement algorithm or the 926/928 result, and cannot support the characterization of this particular claim as…	Now reads: is a simulation, not a field result.
01-M6	[67]	The quote describes why cognitive tasks become popular (low between-subject variability), but does not state the paper's main finding relevant to this mechanism: that such tasks have poor reliability for measuring individual differences.	Now reads: "low between-subject variability causes low reliability for individual differences, destroying replicable correlations with other factors" (reliabilities ranged 0 to .82 across seven classic cognitive tasks)
01-M27	[119]	Porter 2002 describes measurement tools for researchers to assess curriculum alignment (whether instruction matches standards/materials), not evidence that telling students where a lesson sits in the sequence aids learning or functions as an advance organizer.	Citation removed.
01-M35	[149]	The source discusses confidence interval width in the context of large-scale RCTs (mean N ≈ 8,600 students per trial), not small holdouts. The wide intervals (mean 0.30 SDs) were found in studies with large samples.	Now reads: lortieforgues2019 [149] found that even large educational trials often gave wide, uninformative confidence intervals (mean width 0.30 SDs), so the holdout comparison needs enough learners before it is read.
01-M57	[71]	Baker et al. (2008) does not warn against setting BKT parameters by hand. The paper criticizes *fitting* parameters from performance data (which can lead to identifiability and degeneracy problems) and proposes contextual estimation as a better *fitting*…	Now reads: The parameters are design choices: they are set by hand, not fitted from performance data.
01-M39	[156]	Way 1998 addresses test security and item-pool integrity in computerized testing (preventing exposure to protect measurement validity), not learning effectiveness or the benefits of fresh questions for retry attempts.	Citation removed.
01-M63	[186]	The setting says 'children and adults' but the quoted finding (no difference when lag controlled) only applies to college students. For Grade 6 children, delayed feedback was superior even when lag was controlled.	Now reads: metcalfe2009 [186]. *Setting:* college students, vocabulary.
01-M60	[115]	Wrong context. Hausknecht2007 measures practice effects (score improvement from retaking) in selection/hiring contexts (job applicants, admissions testing), not remediation effectiveness in learning checkpoints.	Citation removed.
03-M19	[137]	Rey et al. 2019 studied the segmenting effect (breaking continuous material into learner-paced	Citation removed.

Citation	Source	Problem found	Correction
		segments), not brevity or content length. Segmenting is about presentation structure and pacing, not making teaching shorter or fuller.	
03-M25	[324]	Source describes classification of MOOC engagement trajectories but does not provide evidence that early payoffs or short sections help low-persistence learners; it is descriptive research about patterns, not evidence for the intervention strategy.	Citation removed.
03-M24	[59]	The article says this applies to 'low-stakes tests' but the abstract explicitly states rapid guessing indicates disengaged test taking 'regardless whether it occurs in high-stakes or low-stakes testing contexts.'	Now reads: — both high- and low-stakes contexts.
03-M22	[94]	wilson2019 argues for optimal learning at 85% accuracy (moderate difficulty, 'neither too easy nor too hard'), not for the low-challenge 'comfortable difficulty' approach the mechanism implements.	Citation removed.
03-M29	[266]	Source shows proctor notifications during test-taking increased engagement after rapid-guessing was detected, but the article feature is about proactively formatting learning content with warnings. Different context (proctored tests vs.	Now reads: wise2019 [266] — "after a proctor notification was triggered, test-taking engagement" (rose) during tests — K-12 test takers; whether formatting warnings in learning content work similarly is untested.
03-M33	[87]	The quote and source are about adapting to learners' *level*, not about interest-based personalization. major2021 explicitly found that 'more personalised approaches which adapt or adjust to learners' level led to significantly greater impact (an effect size…	Citation removed.
03-M39	[339]	The abstract confirms soland2021 compares 'threshold-setting methods' but does not specify they are 'item-relative' methods.	Now reads: kong2007 [338] and soland2021 [339] compare threshold-setting methods
03-M36	[14]	The article claims the code matches Brunmair's 'caution for new material and words,' but the source warns about 'expository texts and words'—not about whether material is new vs. proven.	Now reads: partly matches brunmair2019 [14], which found interleaving less helpful in some conditions, especially for expository texts and words; blocking while an objective is new is a reasoned reading of that caution, not something the…
03-M38	[32]	Jansen 2013 is about experiencing success in math (Math Garden) reducing anxiety and raising perceived competence; it fits the 'favour high success' claim, but no abstract is available to confirm a specific 85% threshold.	Now reads: - jansen2013 [32] (above; it supports the general principle that experiencing success helps, not the exact cut-off).
03-M38	[28]	Eglington & Pavlik 2020 is about optimizing spaced-practice scheduling for memory retention. It says optimal difficulty differed from prior recommendations but supported desirable difficulty, which cuts against a 'favour high success' framing for an 85%…	Citation removed.

Citation	Source	Problem found	Correction
03-M41	[341]	Year is wrong (source is 2016, not 2017); quote is truncated at a critical point, omitting the specific thresholds (6.25%, 12.5%, 12.5% for easy, moderately difficult, and difficult tests); and the source studies group-level aggregated scores in low-stakes…	Now reads: rios2017 [341] — "the true mean was underestimated by around 0.20 SDs if the total amount of careless responses exceeded 6.25%, 12.5%, and 12.5% for easy, moderately difficult, and difficult tests" — group-based low-stakes assessments.
03-M48	[350]	The quote and the '35 systems' figure are accurate. But the source is a review of computational short-answer grading. It says nothing about forgiving case, accents, '$', punctuation or a leading 'x =', and nothing about scoring the construct rather than the…	Now reads: review of 35 systems; background on automatic short-answer grading, not direct support for notation leniency.
03-M56	[363]	Quote uses 'Children' but the abstract says 'Preschoolers'	Now reads: "Preschoolers appear to progress along a continuum from scribbling to conventional spelling"
03-M50	[105]	Kane (2013) is about argument-based validation of score interpretations and uses. It does not address construct-irrelevant variance or grading leniency, or strictness on a required form.	Citation removed.
03-M57	[364]	Reed & Vaughn is a synthesis of retell protocols for reading comprehension in K-12. It says nothing about adult or parent confirmation, or about weighting an unverified report below direct evidence.	Citation removed.
03-M57	[105]	Kane's argument-based validation supports the general principle that a score use needs adequate evidence for the claim made, and that more ambitious claims need more support. It does not address adult attestation or weighting, so it is only indirect support.	Now reads: kane2013 [105] (above; general principle that a claim needs adequate supporting evidence).
03-M57	[151]	Messick's unified validity framework supports the general idea that score meaning needs multiple kinds of evidence. It says nothing about unverified adult reports or a 0.25 weight, so it is only indirect support.	Now reads: messick1995 [151] (above; general principle that score meaning needs sufficient evidence).
03-M58	[367]	The source is a real LSE evidence review on children's data and privacy online, so the topic matches. But the fetched page text is only an OpenAlex JavaScript notice, and the quoted sentence could not be confirmed in it.	Now reads: — "The relation between privacy and data online is becoming highly complex" — children's privacy (background only; it does not address voice recordings).
04-M1	[375]	Adams 1961 is a review of warm-up decrement: performance drops after a rest and recovers once practice resumes. That is related to the idea that warm and cold performance differ.	Now reads: background only: review of warm-up decrement, showing performance after a rest differs from warmed-up performance (title-level; does not directly test cold-sitting validity).
04-M19	[412]	Einstellung study shows familiar features trigger a first solution that blocks a better one, even in experts. It does not address deciding whether a tactic exists or no-mate positions. Related to the 'find vs.	Now reads: Pop.: expert chess players, eye tracking. Shows familiar cues can trigger a first

Citation	Source	Problem found	Correction
			solution that blocks a better one; it does not test no-tactic items.
04-M19	[413]	Source shows response latencies improve criterion validity of a chess test. It says nothing about deciding whether a tactic exists or unannounced items. The article itself notes it is about response times, which are not used here.	Citation removed.
04-M19	[414]	Burns shows blitz performance correlates with overall rating (up to 81% variance), which is correlational evidence that fast recognition matters.	Now reads: Blitz results track rating closely, which suggests fast pattern recognition accounts for much of skill (correlational).
04-M11	[401]	The meta-analysis (r = 0.61) relates accumulated task-relevant practice duration to musical achievement. It did not test targeting weaknesses or allocating practice by distance from target.	Now reads: Pop.: meta-analysis, 13 studies, N = 788. Shows that accumulated practice time predicts musical achievement (r = 0.61); it does not test weakness targeting.
01-M26	[115]	Direction reversed. The source found that 'effects were larger...when identical forms were used'—meaning repetition produced larger practice effects. This contradicts the article's principle that checks should never repeat; the source supports the opposite.	Citation removed.
01-M26	[116]	Wrong topic and year. The source (Smith & Karpicke 2013, not 2014) is about question format effects (short-answer vs. multiple-choice), not about repeating vs. varying practice items.	Citation removed.
01-M44	[115]	Source is about retesting/practice effects in personnel selection (job applicants retaking cognitive ability tests), not about instructional design or learning variety.	Citation removed.
01-M13	[83]	This source studies equity and fairness concerns in knowledge tracing algorithms across student populations, not the effectiveness or appropriateness of the Bayesian belief-updating method itself.	Citation removed.
01-P3	[65]	The source describes EAP (Expected A Posteriori) ability estimation in adaptive testing, not SPRT (Sequential Probability Ratio Test) classification decisions.	Now reads: babcock2012 [51].
01-P8	[133]	Little et al. (2012) [133] studied whether multiple-choice tests with plausible distractors can trigger productive retrieval processes and foster learning.	Now reads: - **Sources:** rodriguez2005 [136], haladyna2002 [132]. (Two after the 27 October 2026 audit: little2012 [133] tests whether plausible distractors trigger retrieval, not how many options an item needs.)
01-P4	[76]	meyers2008 [76] studies how item difficulty parameters change when item position changes (a threat to parameter invariance in IRT), not methods for online calibration of difficulty from live student responses.	Citation removed.

Citation	Source	Problem found	Correction
01-P4	[75]	kaser2024 [75] is a systematic review of simulated learners in educational technology and their validation, not about calibrating item difficulty from real student responses.	Citation removed.
01-P4	[79]	doroudi2017 [79] proposes the Robust Evaluation Matrix method for offline evaluation of instructional policies using student model simulators, not online calibration from live student data.	Citation removed.
01-P9	[161]	Source is a meta-analysis of computer-based scaffolding (hint/support structures for complex problem-solving), not adaptive difficulty adjustment or tuning item distribution by realized accuracy. Different mechanism entirely.	Now reads: wilson2019 [94], wauters2010 [55], klinkenberg2011 [78], tetzlaff2025 [61]
01-P6	[67]	Wrong year (cited as hedge2018 but source is 2017), and source addresses test-retest reliability of cognitive experimental tasks across sessions, not placement testing or skill classification from small item samples within a single assessment.	Now reads: parsons2019 [66], morris1983 [63], vandesande2013 [81], baker2008 [71]
02-M14	[208]	The quote is accurate, but the design is mislabeled. Wissman et al. (2012) is a survey of 374 undergraduates, with no lab component.	Now reads: (survey of 374 undergraduates)
02-M22	[205]	The quote says 'substantial benefits of learning to a higher initial criterion during the first session do not persist' but the article presents it as 'to a higher initial criterion during the first session do not persist' — omitting 'substantial benefits of…	Now reads: "substantial benefits of learning to a higher initial criterion during the first session do not persist across relearning sessions"
02-M37	[235]	The abstract does support the universal formula claim at 95% retention, but the article's parenthetical states '(the authors' own repetition data (abstract claim is the author's own))' which is confusing and redundant.	Now reads: (the authors' own repetition data)
02-M31	[235]	The article cites this as wozniak1995 but the source was published in 1994.	Now reads: (the authors' own repetition data)
02-M26	[229]	The quote discusses retrospective evaluations being based on 'snapshots' of experience generally, but does not specifically mention that the ending or final moments matter, which is the core principle behind this mechanism (ending on easier questions makes…	Now reads: **fredrickson1993** [229]: "Retrospective evaluations appear to be determined by a weighted average of "snapshots" of the actual affective experience" (adults, film clips; foundational for peak-end effects but does not specify ending in…
02-M47	[27]	Abstract confirms 'middle-school foreign-language course' but does not specify Spanish; all other details (16.5%, massed study comparison, time-matched) are accurate.	Now reads: (middle-school foreign-language course, time-matched)
02-P4	[27]	The abstract states 'middle-school foreign-language course' but the article specifies 'Spanish' which is not confirmed in the provided abstract.	Now reads: (middle-school foreign-language course, time-matched)
02-M49	[191]	The population 'college students' and material 'foreign-vocabulary pairs' are not confirmed in the	Now reads: (lab study)

Citation	Source	Problem found	Correction
		abstract; the abstract uses 'subjects' and 'items' without specifying student population or vocabulary materials.	
02-M49	[259]	The quote is incomplete—cuts off mid-sentence at 'performed significantly' without completing the comparison or outcome.	Now reads: performed significantly" (better than those in massed practice; meta-analysis, 63 studies of practice distribution)
Corrections to RETENTION-ENFORCER-EVIDEN	[28]	The article states that 'difficulty_window' now reads '1.0 up to 0.95' and attributes this range to eglington2020 [28].	Now reads: with `difficulty_window_basis` naming eglington2020.
02-P5	[193]	Cepeda 2009 is about optimal spacing intervals for material that IS practiced (distributed practice effect), not about retrieval-induced forgetting or prioritizing topics left out of practice sessions. It does not support a rotation flag mechanism.	Now reads: It was built on anderson1994 [262] and murayama2014 [263]. A third source, cepeda2009 [193], was removed in the 27 October 2026 audit because it concerns spacing of material that is practised, not rotating topics, so this change now rests…
02-P5	[193]	The quote refers to 'cumulative reviews' meaning repeated spaced practice of the same material over time, not rotating through topics to avoid leaving some unpracticed.	Citation removed.
02-P1	[26]	The abstract states 'the optimal reviewing schedule is given by the recall probability itself' but does not explicitly derive or state the formula need = 1 - p_H.	Now reads: **need = 1 - p_H**: this operationalizes the principle that review priority should track memory strength (tabibian2019 [26]).
03-M9	[306]	The article states '(no DOI)' but the source has an arXiv DOI: https://doi.org/10.48550/arXiv.2507.11538 (shown on the arXiv page as 'arXiv-issued DOI via DataCite').	Now reads: jaroslawicz2025 [306] —
03-M2	[301]	The quote is accurate and the source addresses reliability of individual-difference measurement in cognitive tasks, but the paper's main argument is about low cross-task correlations and the need for hierarchical trial-level models, not directly about minimum…	Now reads: rouder2019 [301] — "correlations among tasks that seemingly have processes in common are often low" — cognitive tasks; the paper asks whether trial-level noise and low individual variability cause this, which is the reason to ask for…
03-M10	[307]	The quote describes incrementally adding constraints to test LLM capability, not resolving conflicts between directives.	Now reads: — LLM benchmark (design only; it tests following added constraints, not resolving conflicting ones).
02-M36	[239]	The study's efficiency gain was specifically due to sleep consolidation between sessions (12-hour delay), not immediate reteaching; the mechanism implements	Citation removed.

Citation	Source	Problem found	Correction
		reteaching 'soon' but this source required sleep interleaving for the benefit	
03-M43	[38]	The article implies the safeguards in GPT Tutor were specifically about not giving the answer first, but the abstract only states there were 'prompts designed to safeguard learning' and 'safeguards' without specifying what they were.	Now reads: high school maths; a tutor with guardrails (prompts designed to safeguard learning) largely mitigated the harm.
03-M43	[34]	The source is about help-seeking behavior and whether students use on-demand help deliberately, not specifically about whether tutors should avoid stating the answer first.	Citation removed.
04-M14	[23]	Cepeda 2008 studied spacing effects for declarative learning (teaching facts), not motor skills; mismatch with this block's motor-learning context	Citation removed.
04-M26	[414]	The article says 'Search explains the variance that recognition does not,' but Burns (2004) found the opposite: fast processes like recognition account for up to 81% of variance in chess skill, with search being less important than traditionally thought.	Citation removed.
04-M23	[217]	The source studies unsuccessful retrieval attempts (memory testing with trivia and word associations), not skill performance against strong opponents. Playing chess endgames against perfect defense is not retrieval practice.	Citation removed.
04-M25	[411]	Source is about practice time, starting age, and handedness in chess expertise development, not about memory structures or piece/board knowledge that would support testing beginners on moves and setup.	Citation removed.
04-M25	[41]	Source is a motor learning framework about physical skill acquisition and practice conditions, not about chess cognition, memory, or piece knowledge.	Citation removed.
04-M36	[418]	Ash & Holding (1990) compared part-task sequencing methods (backward vs forward vs whole) for keyboard learning. It does not address errorless or guided learning vs independent timed performance, which is the mechanism's principle.	Citation removed.
04-M33	[387]	The source demonstrates successful audio-based pedal transcription (91.86% F1 score on MAESTRO), but the article uses it to claim 'audio remains imperfect, so MIDI only is right'—the opposite of what Kong et al. show.	Citation removed.
04-M43	[459]	Agresti & Coull 1998 discusses frequentist confidence intervals (Wilson score, adjusted Wald) for binomial proportions, not Bayesian Beta-posterior probability calculations.	Citation removed.
04-M43	[383]	Hopkins 2000 is about reliability and reproducibility of measurements (typical error, within-subject variation, retest correlation), not about Bayesian	Citation removed.

Citation	Source	Problem found	Correction
		inference, Beta posteriors, or binomial proportion estimation for performance gates.	
04-M44	[461]	The quote is accurate, but Kubatko 2007 is about possession estimation and general basketball analytics methods, not zone-based shooting analysis or uncertainty quantification.	Citation removed.
04-M48	[389]	The upper bound is misstated: the source reports ankle angle error of 7.4°, not 7°. The actual range from the abstract is hip 4.0°, knee 5.6°, ankle 7.4°.	Now reads: Joint-angle errors of 4.0-7.4°
04-M56	[14]	The source studies interleaving different categories of learning material (e.g., different types of math problems), not mixing easy and hard difficulty levels within one material type.	Citation removed.
04-M58	[43]	Source is about gamification (adding game elements to instruction) rather than learning games or serious games. While related, these are distinct interventions in the literature.	Now reads: sailer2020 [43]: "the effect of gamification on cognitive learning outcomes was stable in a subsplit analysis of studies employing high methodological rigor".
04-P3	[444]	The source shows that as pianists learned, they practiced difficult bars decreasingly and structural bars increasingly; experts practiced difficult bars less, not more. This contradicts the hotspot loop concept of repeatedly drilling the worst bars.	Citation removed.
04-P6	[473]	Source [473] is the technical OpenPose algorithm paper (Part Affinity Fields method), not a validation study. It does not provide error rates, agreement statistics, or validation of pose estimation against manual counts or ground truth—unlike sources [388]…	Citation removed.
04-P7	[464]	Lloyd 2012 is about a physical development model for training young athletes to improve performance, not about load caps or rest days to prevent overuse injury and burnout. Wrong focus and direction.	Citation removed.
05-M11	[509]	Gaming behavior (active system manipulation like rapid guessing or hint abuse) is conceptually different from idle inactivity.	Citation removed.
04-M10	[398]	The article claims the study reports split-half 0.92, but the abstract does not mention split-half reliability at all. The abstract only discusses diagnostic validity (discriminating performance levels) and external validity (correlation with game data).	Now reads: video 4v4 test
01-M17	[93]	Year is wrong: source is from 2009, not 2010. Grade level '4th-grade' not confirmed in abstract (only 'children').	Now reads: taylor2010 [93]. *Setting:* children, mathematics.
04-P2	[436]	Source discusses that complex skills (high processing load) benefit from conditions that reduce load, but does not explicitly advocate for fading feedback	Now reads: wulf2002 [436] (as M29). Complex skills (high processing load) benefit from

Citation	Source	Problem found	Correction
		rather than other reduction strategies. The article infers a fading approach from the general principle.	conditions that reduce the load to manageable levels.
05-M14	[22]	Cepeda 2006 is a meta-analysis of distributed (spaced) practice effects, not retrieval practice or testing effects. It does not support the use of a retention-enforcer prompt/test at the end of study sessions.	Citation removed.
04-P5	[398]	The Rösch 2021 abstract discusses diagnostic and external validity of a basketball decision-making assessment, not reliability thresholds or split-half reliability standards. No mention of 0.8 thresholds, Spearman-Brown, or reliability coefficient standards.	Citation removed.
05-M25	[219]	Metcalfe (2017) is about errorful learning—making errors during problem attempts and then receiving corrective feedback—particularly that high-confidence errors are corrected more readily.	Citation removed.
05-M21	[74]	wise2005 [74] measures rush behavior (answering before reading) via response time, but does not show that typed answers prevent rushing or force reading. The source describes the problem, not this intervention.	Citation removed.
05-M16	[342]	Source is about feedback interventions generally, not praise; does not support the praise-specific mechanism.	Citation removed.
05-M40	[554]	WCAG 3.3.1 Error Identification applies to input errors (form validation failures, user mistakes), not empty states. An empty dashboard or 'Nothing to lock in...' message is a no-content condition, not an automatically detected input error.	Citation removed.
05-M40	[555]	WCAG 3.3.3 Error Suggestion applies to input errors (form validation failures with known corrections), not empty states. Empty states are not input errors.	Citation removed.
05 proposals (each backed by 3 or more s	[116]	The source was published in 2013, not 2014, and its result is more qualified: it found little or no short-answer advantage over multiple choice in three experiments, but found an advantage in a fourth when retrieval success was improved.	Now reads: smith2014 [116] found little or no advantage over multiple choice in three experiments, but an advantage in a fourth when retrieval success was improved.
01-M3	[60]	The source supports mastery learning through variation in instructional methods and learning time, but it does not support sequential classification, stopping rules, or using consecutive failing bands to determine placement.	Citation removed.
01-M5	[65]	The source concerns EAP ability estimation in computerized adaptive testing and the relationship between posterior standard deviation and measurement error. It does not support the specific principle that thin evidence should be shrunk toward a prior.	Citation removed.

Citation	Source	Problem found	Correction
01-M13	[71]	The source says Bayesian Knowledge Tracing uses four parameters per skill fitted from student performance data; it does not support the specific claim that this implementation's SLIP and LUCK knobs are global and not fitted per unit.	Citation removed.
01-M2	[57]	The source supports the possibility that item order affects item and test characteristics, including item difficulty, but it does not establish that fatigue causes a learner to appear less able or that the hardest items are affected in that specific way.	Now reads: *Setting:* IRT methodology; supports position effects, not learner fatigue.
01-M15	[89]	The source supports substantial variation in initial performance and similar learning rates, but the article's wording that "missing prior practice" is what separates students is stronger than the abstract, which does not establish that missing prior practice…	Now reads: fits koedinger2023 [89]'s finding that students differ mainly in initial performance, not in learning rate.
01-M31	[134]	The source discusses the literature's lack of focus on generating questions with controlled difficulty, but the cited mechanism concerns validation of item structure and answer correctness. The validators described do not enforce controlled difficulty.	Citation removed.
01-M36	[152]	Walkington et al. shows that personalized interest contexts changed speed and accuracy on matched algebra problems, but that does not directly establish that the style changed problem difficulty.	Now reads: which suggests an "interest" style may affect performance even when the underlying problems are matched
01-M39	[60]	The source supports Bloom's general principle of varying instructional methods and learning time to match students' needs, but the supplied abstract and quote do not establish the more specific sequence of corrective instruction followed by a second…	Now reads: Bloom's model of varying instruction and learning time to match students' needs.
01-M39	[4]	The source supports positive examination effects of mastery-learning programs, but it does not address fresh versus repeated forms, item exposure, or score inflation from repeated identical questions.	Citation removed.
01-M21	[51]	The study concerns simulated computerized adaptive tests and trait-estimation accuracy; it does not establish the accuracy or classification error of fixed-length mastery checkpoints.	Now reads: "longer CATs yielded more accurate trait estimation" (a simulation; it supports the general principle that longer tests measure more accurately, not the checkpoint lengths themselves)
01-M42	[61]	The source supports giving more instructional assistance to low-prior-knowledge learners and less assistance to high-prior-knowledge learners, but it does not directly show that weaker learners should receive easier items at d−1.5 or that ready learners…	Now reads: - tetzlaff2025 [61] (supports adaptive assistance by prior knowledge). - "Low prior knowledge learners learn better from high-assistance instruction (d = 0.505); high prior knowledge learners learn better from low-assistance instruction (d…
01-M42	[162]	The review identifies contingency, fading, and transfer of responsibility as characteristics of scaffolding and reports that scaffolding can be	Now reads: - vandepol2010 [162]. *Setting:* review of scaffolding

Citation	Source	Problem found	Correction
		effective, but it does not support the specific use of easier or harder items at d−1.5 and d+1.5.	characteristics, not item-difficulty spacing.
01-M46	[169]	The source supports correcting errors after feedback, but specifically finds that high-confidence errors are easier to correct than low-confidence errors. It does not establish that an end-of-level recall question benefits all missed skills equally.	Now reads: - metcalfe2011 [169] (especially for high-confidence errors).
01-M23	[6]	The source supports adapting instructional techniques to learner expertise, but it does not directly study or establish that learners should skip ahead through levels.	Now reads: "can lose their effectiveness and even have negative consequences when used with more experienced learners" (supporting expertise-sensitive instruction rather than level skipping)
01-M54	[176]	The Grade 8 standard explicitly includes interpreting unit rate as slope, but it does not establish that slope instruction begins in Grade 8; that claim requires comparison with earlier-grade standards.	Now reads: Slope is addressed in grade 8
01-M58	[23]	The source supports that the optimal interstudy gap increases with the delay before the final test, but it does not establish that the optimal gap is generally "well beyond a day".	Now reads: shows that the optimal gap generally lengthens for longer retention delays
01-M59	[183]	The source describes 46 studies of group-based applications of mastery learning, but the abstract does not specifically characterize them as studies conducted in schools.	Now reads: *Setting:* meta-analysis of 46 studies of group-based applications of mastery learning.
01-P2	[6]	The source supports that instructional techniques can lose effectiveness or become counterproductive for more experienced learners, but it does not establish that advanced learners should specifically skip this blank step.	Now reads: Expertise reversal suggests the blank-step support may be less effective, or even counterproductive, for advanced learners
01-P3	[18]	Lewis et al. supports reducing average test length by half without sacrificing classification accuracy, but it does not establish that accuracy will be better, nor does the abstract specifically support fewer false fails for learners exactly at the pass mark.	Now reads: the same classification accuracy with fewer items
01-P7	[173]	Aleven et al. supports the claim that learners often fail to use on-demand help effectively, but it does not show that learners skip the corrective step that defines Bloom's mastery learning specifically, nor that they skip it simply because it is left to…	Now reads: while learners often fail to use optional help effectively (aleven2003 [173])
01-P1	[12]	Rohrer et al. report a delayed-test effect of d=0.79, but the d≈0.4 lower bound is the immediate-test effect, not a delayed-test result.	Now reads: test gains in mathematics of d≈0.42 a day later and d≈0.79-0.83 about a month later (rohrer2015 [12], rohrer2020 [13])
01-P1	[14]	The source supports caution for words and expository texts, but the supplied abstract does not specifically establish that the effect is smaller for those materials.	Now reads: The effect may be less reliable for vocabulary and expository material (brunmair2019 [14])

Citation	Source	Problem found	Correction
01-M61	[17]	The quoted result is supported, but the source abstract does not establish a college setting; it describes two experiments with subjects studying prose passages and taking tests.	Now reads: *Setting:* experimental study.
01-P5	[180]	The abstract supports successive relearning as a combination of retrieval and spaced practice across multiple sessions and reports improved course-exam learning, but it does not establish larger long-term retention gains than same-session recall or directly…	Citation removed.
02-M16	[214]	The source reports 220 million students' memory-behavior logs and deployment in an online language-learning app, but it does not describe the app specifically as a vocabulary app.	Now reads: (220 million student memory-behavior logs from an online language-learning app)
How it all fits together	[5]	Bloom's source describes mastery-learning experiments and formative corrective testing, but it is not a meta-analysis and the supplied abstract does not report the claimed large gains.	Changed to: A meta-analysis (kulik1990 [4]) and Bloom's experiments (bloom1984 [5]) report large gains. (revised again on the second check; see below)
How it all fits together	[7]	Sweller and Cooper studied worked examples in algebra with Year 9, Year 11, and university mathematics students; the source supports worked examples for novice algebra problem solving, but not the broader claim that this is one of the best-supported rules for…	Now reads: Worked examples first is well supported for novice problem-solving material
How it all fits together	[13]	Rohrer et al. (2020) reports d = 0.83, which is slightly above the stated upper bound of 0.8.	Now reads: interleaved practice gives test gains over blocked practice of about d = 0.42 a day later and d = 0.79 to 0.83 about a month later
How it all fits together	[14]	The meta-analysis says interleaving should be used with caution for expository texts and words, but the supplied abstract does not specifically establish that the gain is smaller for vocabulary and reading.	Now reads: with caution for words and expository texts
How it all fits together	[37]	Kulik and Fletcher's meta-analysis supports the effectiveness of intelligent tutoring systems, but the supplied abstract does not establish that step-based tutoring comes close to a human tutor; that comparison is supported by VanLehn.	Now reads: Step-based tutoring comes close to a human tutor (vanlehn2011 [36]), and intelligent tutors raised test scores by a median 0.66 SD in 50 controlled evaluations (kulik2016 [37])
How it all fits together	[38]	Bastani et al. show that safeguards mitigate the learning harm of unfettered GPT access, but the supplied abstract does not specifically say that guardrails stop an AI from giving the answer.	Now reads: guardrails can mitigate the learning harm of unfettered AI access
How it all fits together	[39]	Soderstrom and Bjork support the claim that observed performance can be an unreliable guide to long-term learning, but the abstract does not directly establish that a delayed test is better.	Now reads: Performance during practice can be an unreliable guide to long-term learning
How it all fits together	[41]	Guadagnoli and Lee propose matching task difficulty to skill level in motor learning, but the supplied	Now reads: Practice difficulty is matched to the learner's skill level

Citation	Source	Problem found	Correction
		abstract does not state that practice is dosed by distance from the target.	
02-M34	[193]	The source supports the 150% recall improvement, verbal materials, and test delays up to 6 months, but the supplied abstract does not identify the participants as adults.	Now reads: (laboratory studies of foreign vocabulary, facts, and visual-object names, with test delays up to 6 months)
02-M34	[191]	The source supports the 200% improvement from long intervals between repeated tests, but the supplied abstract does not specify college students or foreign-vocabulary pairs.	Now reads: (laboratory experiment on repeated retrieval and long-term retention)
02-M35	[10]	The quoted finding and the meta-analysis description are supported, but the supplied abstract does not establish that the literature was mostly adult laboratory studies.	Now reads: (meta-analysis of testing vs restudy)
02-M35	[166]	The quoted result is directly supported, but the supplied abstract does not specifically identify the participants as college students; it describes students in experimental tests.	Now reads: (students in laboratory experiments)
02-M35	[133]	The quoted finding is directly supported, but the supplied abstract does not specifically identify the participants as college students or explicitly label the setting as a laboratory.	Now reads: (participants in two experiments)
02-M7	[202]	The source clearly supports the claim that data-preprocessing and outlier-detection decisions affect learning-analytics outcomes and prediction accuracy, but the provided abstract does not establish that the data were specifically from a university LMS.	Now reads: (real-world student learning-activity data (LAK 2022, peer-reviewed proceedings))
02-M8	[23]	The abstract reports 20–40% of a 1-week test delay declining to 5–10% of a 1-year test delay, not tests up to 1 year as the endpoint of the stated range without specifying the endpoint; the evidence item's numbers are otherwise accurate.	Now reads: to about 5 to 10% of a 1-year test delay
02-M8	[193]	Source [193] establishes that optimal gaps have a nonmonotonic effect and that an optimal gap can improve recall by up to 150%, but it does not establish that the particular 3–4-day gaps are on the side of the curve that is penalized least.	Now reads: which is the side the curve penalises least: past the optimum, accuracy declines only gradually (cepeda2009 [193]).
02-M5	[24]	The source describes students learning conceptual material, not key-term definitions. Its cited recommendation of 3 initial correct recalls followed by 3 relearning sessions at widely spaced intervals is accurate.	Now reads: conceptual material
02-M2	[193]	The source supports the importance of cumulative reviews and reports test delays up to 6 months, but the article's parenthetical is imprecise: the studies used foreign vocabulary, facts, and names of visual objects, not only verbal materials, and the source…	Now reads: (laboratory studies, foreign vocabulary, facts, and names of visual objects, test delays up to 6 months)
L2637	[12]	Rohrer et al. support interleaved mathematics practice using different problem types, but the study does not	Now reads: Interleaved mathematics practice is supported

Citation	Source	Problem found	Correction
		specifically test placing confusable concepts side by side.	(brunmair2019 [14]; rohrer2015 [12])
02-P5	[262]	The source supports retrieval-induced forgetting in category-exemplar lists, but the provided abstract does not establish that the participants were college students.	Now reads: (category-exemplar lists)
02-P5	[153]	The source supports enhanced recall 24 hours later for initially nontested prose material, but the provided abstract does not establish that the participants were college students.	Now reads: (prose)
02-M49	[209]	The source supports the quoted spacing result and the flashcard context, but the abstract describes participants learning GRE-type word pairs via a web-based study program; it does not identify them as college students or as learning foreign vocabulary.	Now reads: (participants, GRE-type word-pair flashcards)
02-M47	[214]	The source supports the 12.6% improvement figure and the 220 million behavior logs, but it describes MaiMemo as an online language-learning app rather than specifically a vocabulary app.	Now reads: (220 million review logs from an online language-learning app)
Scorecard weak rules (section 03)	[270]	Koo and Li's cutoffs apply specifically to intraclass correlation coefficients, whereas the article applies them to a split-half correlation. The numerical thresholds are therefore related but not directly validated for the reported statistic.	Now reads: for ICCs, below 0.5 is poor and 0.5-0.75 moderate, judged on the 95% CI.
Scorecard weak rules (section 03)	[66]	Parsons et al. support bootstrapped split-half estimation and routine reliability reporting, but the abstract does not support the specific claim that a Spearman-Brown correction is needed for a split-half.	Now reads: Spearman-Brown correction (brown1910 [273], spearman1910 [274] title-level) is needed for a split-half; parsons2019 [66] recommends reporting split-half reliability.
Scorecard weak rules (section 03)	[283]	Baker et al. study gaming the system, which involves exploiting feedback and help and is associated with reduced learning; they do not show that heavy hint use itself is partly gaming.	Now reads: gaming the system exploits feedback and help and is associated with reduced learning
Scorecard weak rules (section 03)	[286]	Bandura and Schunk's experiment supports proximal goals improving mastery, self-efficacy, and interest, but the supplied abstract does not directly establish persistence. The other cited sources are the more direct support for the persistence claim.	Now reads: Persistence: botelho2019 [295], ventura2013ce [296].
The eglington 2020 [28] question: should	[28]	The source is about model-based optimal scheduling and supports the general idea, but the supplied abstract does not establish the specific threshold comparisons or show that 0.94 was among the best while 0.90–0.95 receives a lower weight.	Now reads: Broadly consistent with eglington2020 [28]?
The eglington 2020 [28] question: should	[26]	The abstract assigns the recall-probability result to the theoretical optimal-control analysis. The natural experiment tested the resulting MEMORIZE algorithm against heuristic schedules; it did not itself establish that the optimal reviewing rate equals…	Now reads: The theory in tabibian2019 [26] puts the optimal reviewing rate on the recall probability itself, while its large natural experiment found that the

Citation	Source	Problem found	Correction
			resulting algorithm outperformed heuristic schedules.
02-M46	[252]	The study found that sleep benefited restudied items but did not improve recall of retrieved items; it does not establish the app's design decision or prove that an evening suggestion should not be proposed.	Now reads: Given bauml2014 [252], sleep is not shown to provide an extra recall boost after retrieval practice.
02-P1	[28]	Eglington et al. support the general principle that item-level performance can guide adaptive practice scheduling, but the source does not establish the specific 0.50–0.95 plateau or the proposed difficulty weights.	Now reads: Add eglington2020 [28] to its `basis` as general support for adaptive, item-level scheduling.
02-P1	[94]	Wilson et al. derive an approximately 85% optimum for specified stochastic-gradient-descent learning algorithms and demonstrate it in neural networks; they do not establish that 0.85 is only a model result in contrast to an empirical 0.70–0.94 efficient…	Now reads: 0.70-0.94 was one efficient region in eglington2020 [28], while 0.85 comes from a model of learning algorithms, not a band tested with learners (wilson2019 [94]).
03-M20	[4]	The meta-analysis supports mastery learning generally and reports stronger effects for weaker students, but it does not specifically support shorter units for learners who stop early or less practice once an objective is mastered.	Citation removed.
03-M30	[283]	The source supports the general point that learners may game an intelligent tutoring system by exploiting feedback and help to obtain correct answers, and that gaming is associated with poorer learning.	Now reads: - baker2004 [283] — gaming the system can involve exploiting feedback and help to obtain correct answers, and is associated with lower post-test performance.
03-M38	[94]	The source supports an approximately 85% optimal training accuracy, but its evidence comes from binary-classification learning algorithms, artificial neural networks, and biologically plausible animal-learning models; it does not establish that human studies…	Now reads: - wilson2019 [94] (supports an approximately 85% training-accuracy sweet spot in AI and animal-learning models).
03-M39	[328]	The source reports that items with more text were associated with increased rapid guessing; it does not establish that more text causes or directly draws rapid guessing.	Now reads: items with more text are associated with more rapid guessing
03-M44	[344]	The source supports the approximately .70 sigma learning gain and the use of dialogue, prompts, feedback, and adaptive tutoring, but its abstract specifies Newtonian qualitative physics and computer literacy; it does not specify a college population.	Now reads: Newtonian qualitative physics and computer literacy
03-M44	[36]	The source reports that intelligent tutoring systems overall had an effect size of 0.76 versus 0.79 for human tutoring.	Now reads: intelligent tutoring systems about as effective
03-M45	[36]	The review discusses tutoring interaction granularity, including step-based and substep-based systems, but it does not show that keeping help aligned to the	Citation removed.

Citation	Source	Problem found	Correction
		specific step being assessed is more effective; it explicitly says the expected effectiveness…	
03-M42	[342]	Kluger and DeNisi support keeping feedback attention on the task rather than the self, but this abstract does not specifically establish that the note should be non-evaluative. That aspect is supported by Shute.	Now reads: The note is task-focused, which is what kluger1996 [342] points to; its non-evaluative form is supported by shute2008 [16].
03-M46	[34]	The source is a review of help seeking with intelligent tutoring systems and reports mixed learning effects; it does not directly support the assessment claim that assisted answers are weaker evidence or should receive less proof weight.	Now reads: - aleven2016help [34] — review of help seeking with intelligent tutoring systems; it reports that help can be useful under some circumstances but did not improve learning outcomes in one classroom study.
03-M46	[347]	The quoted statement supports that help provides scaffolding on the current problem, but the available source text does not establish the specific Reading Tutor setting or the stronger wording that a helped answer 'partly reflects the help.'	Now reads: Reading Tutor data; help can affect performance on the current problem.
03-M48	[349]	The source studies flawed multiple-choice item-writing features in medical examinations, not lenient short-answer formatting or notation.	Now reads: disadvantage some medical students" — medical students; indirect support (flawed multiple-choice items, not answer-format leniency).
02-M45	[248]	The source shows that feedback reduced lure intrusions, not that it cancelled the multiple-choice lure risk. The verdict overstates reduction as elimination.	Now reads: "both immediate and delayed feedback increased the proportion of correct responses and reduced, rather than eliminated, the proportion of intrusions"
02-M45	[249]	The source supports a spacing account in which delayed feedback strengthens initially correct responses, but it reports minimal effects on error correction. It therefore does not directly support a delayed recap of missed answers as beneficial.	Now reads: "delaying feedback strengthens initially correct responses due to the spacing of encoding opportunities" (college students, lab; it found little effect of the delay on correcting errors)
03-M17	[61]	The source says the evidence is less clear for younger students and some fields of study, specifically humanities and language learning. “Languages” is broader and omits the humanities qualification.	Now reads: younger students and some fields of study, including humanities and language learning
03-M15	[316]	This source concerns the limits of introspective reports about higher-order cognitive processes, not specifically whether stated learning preferences are a poor basis for tailoring instruction. It also notes that verbal reports can sometimes be accurate.	Citation removed.
03-M8	[270]	Koo and Li provide ICC-specific reliability thresholds, not thresholds for a Pearson split-half correlation. The quoted moderate range and confidence-interval qualification are accurate, but applying them directly to the proposed r statistic is not…	Now reads: ICC reliability guideline (judged on the 95% CI)

Citation	Source	Problem found	Correction
03-M7	[281]	Gart et al. supports the small-sample bias caveat, but its abstract does not justify the specific 10+ chances threshold or establish that this source is the reason for that threshold.	Now reads: The small-sample caveat is relevant to keeping a minimum-evidence threshold; the specific 10+ threshold is a design choice.
04-M2	[378]	The source reports a retention test 2 days later, not a day-later test.	Now reads: This shows that a 2-day-later retention test reverses the in-session ranking.
04-M2	[39]	The review supports the general distinction between in-session performance and long-term learning, but it does not specifically establish that a second-day test is the relevant or better index.	Now reads: A delayed test is a better index of learning than performance during practice (soderstrom2015 [39], abushanab2013 [378]); holding it on a second day is a design choice.
04-M7	[383]	The source supports the existence of within-subject random variation and measurement noise, but it does not establish that one miss in this specific metric is normally noise.	Now reads: A single miss can reflect trial-to-trial measurement noise.
04-M7	[24]	The abstract reports 533 students, but does not identify them as college students. It supports spaced relearning as improving long-term retention at relatively low practice cost, but does not specifically describe relearning after decay as efficient.	Now reads: Pop.: 533 students. Relearning at widely spaced intervals improves long-term retention at relatively minimal practice cost.
04-M3	[382]	The source studies binomial interval coverage and recommends alternatives for small n; it does not specifically establish that small-n single tests are unreliable.	Now reads: Small-n Wald intervals can have erratic coverage; Wilson or Jeffreys intervals are recommended for small n.
04-M12	[217]	The abstract supports unsuccessful retrieval attempts improving later learning with verbal materials, but it does not establish the stated adult population. It also does not study a five-day interval; that timing is a design choice.	Now reads: Pop.: verbal materials.
04-M22	[414]	Burns (2004) studies blitz chess performance, search and recognition mechanisms, and their relationship to overall chess rating; it does not examine puzzle-type tasks or show that they measure chess skill reliably.	Citation removed.
04-M23	[420]	The title supports a connection to endgame analysis, but it does not establish that the source concerns endgame tablebases or perfect-defence play-outs.	Now reads: endgame analysis
03-M56	[364]	The source supports the claim that retell has not demonstrated adequacy as a progress-monitoring instrument, but it does not specifically support making a tap check decide the pass. That design choice is an inference.	Now reads: which supports treating retell as supplementary rather than as the sole progress-monitoring measure
03-M52	[358]	The source identifies biases in LLM judges, but does not address prompt injection or support the injection-check and input-fencing principle this mechanism concerns.	Citation removed.

Citation	Source	Problem found	Correction
03-M52	[356]	The source reports JudgeDeceiver outperforms manually crafted prompt-injection attacks and jailbreak attacks, and that the tested defenses are insufficient. It does not specifically say optimized injections beat “simple detection.”	Now reads: shi2024judge [356] shows JudgeDeceiver outperforms manually crafted injection attacks and that tested defenses are insufficient.
03-M59	[369]	The source is a W3C Working Draft, not a finalized web standard; the cited text supports CSP's general purpose but not the stronger status label.	Now reads: — W3C Working Draft.
03-M54	[358]	The source identifies self-enhancement bias as a limitation of LLM-as-a-judge, but it does not show that this bias is why different model families differ or that using different families resolves it.	Now reads: self-enhancement bias is a limitation of LLM-as-a-judge
04-M42	[6]	The source supports adapting instructional techniques to learner expertise, but it does not specifically argue for skipping the early rungs of a lesson ladder for learners who show mastery.	Now reads: supports adapting instruction for learners who show mastery
04-M39	[428]	The review supports general sensorimotor synchronization and tapping variability, but its abstract does not establish that young children synchronize less precisely or that a wide timing window is age-appropriate; it also does not study the 75% threshold.	Now reads: - repp2005 [428] (as M27), a general review of sensorimotor synchronization and tapping variability.
04-M41	[239]	The study supports spacing learning sessions across sleep and testing retention later, but it studied foreign-vocabulary relearning in adults rather than a motor or piano skill, so it is only indirect support for a next-day cold performance test.	Now reads: - mazza2016 [239] (as M02; indirect evidence from foreign-vocabulary learning for spacing sessions across sleep).
Section 03 proposals	[34]	The source concerns help-seeking behavior and on-demand help in intelligent tutoring systems, but it does not specifically support opening with a fully worked example or the sequencing of a learner attempt before a hint.	Citation removed.
Section 03 proposals	[299]	The source supports faster completion for three-option items, but its abstract reports that three-choice questions are slightly easier than four-choice questions and does not establish no loss in psychometric quality.	Citation removed.
Section 03 proposals	[360]	The review discusses distractor quality and optimal numbers of distractors, but the abstract does not specifically show that allowing three options yields more items per minute with no loss in quality.	Citation removed.
04-M56	[417]	The source concerns computerized adaptive testing and selecting easier or more difficult test items; it does not address learning games, mixed-difficulty sequencing, or motivation. Its relevance is limited to adaptive variation in item difficulty.	Now reads: - eggen2006 [417]: computerized adaptive testing can select easier or more difficult item sets.
04-M36	[400]	The source reports significant correlations between the percentage of correctly performed practice trials and next-day retention-test rankings; it does not	Now reads: - duke2009 [400] (as M11). The percentage of correctly performed trials was significantly

Citation	Source	Problem found	Correction
		establish that the share of correct trials causally predicted next-day quality.	related to next-day retention-test quality.
04-M43	[382]	Brown et al. discuss binomial proportion interval estimation and recommend Wilson or equal-tailed Jeffreys intervals for small samples, but the supplied text does not specifically support an exact posterior probability gate under a uniform Beta(1,1) prior.	Now reads: brown2001 [382] (for related small-sample binomial interval guidance).
04-M35	[385]	The abstract supports YIN as an F0 estimator for speech or musical sounds, but does not address polyphonic difficulty or chord detection.	Now reads: - decheveigne2002 [385] (as M04; supports YIN fundamental-frequency estimation for speech and music, not polyphonic chord detection).
04-M35	[387]	The abstract supports piano-note transcription with precise onset and offset estimates, but does not establish that polyphonic audio is hard or address this app's chord-checking design.	Now reads: - kong2021 [387] (as M04; supports piano-note transcription, not the broader claim that polyphonic audio is hard).
04-P9	[379]	The source supports the claim that sleep does not enhance motor-sequence performance, but the abstract does not establish that its participants were children. Using it as direct evidence about children overstates the population match.	Now reads: - rickard2008 [379] (as M02; indirect motor-sequence evidence in its studied population).
05-M1	[45]	The source says that missing one opportunity did not materially affect habit formation; it does not establish that one miss categorically does not break a habit.	Now reads: one missed opportunity did not materially affect habit formation
05-M11	[74]	The source supports using unusually fast response times as an indicator of low effort, but it does not directly support treating idle gaps or inactivity as non-study time. Its population is computer-based test examinees, not learners in study sessions.	Now reads: "unmotivated examinees will answer too quickly (i.e., before they have time to read and fully consider the item)" (Computer-based test examinees; about fast answers, so treating idle time as non-study is a related extension, not something…
04-P4	[393]	The abstract reports that students using a model during self-evaluation improved more than those self-evaluating without a model in several performance areas. It does not establish that self-evaluation helped only when a model was present.	Now reads: Junior high instrumentalists who used a model during self-evaluation improved more in several performance areas than those self-evaluating without a model.
04-P4	[392]	The study found poor correlations between undergraduate musicians' self-evaluations and faculty and peer evaluations. The abstract does not identify the self-ratings as unanchored or attribute their poor agreement to a lack of anchors.	Now reads: Undergraduate musicians' self-ratings correlated poorly with faculty and peer ratings.
05-M3	[44]	Deci et al. identify engagement-contingent rewards as undermining intrinsic motivation, but the article's classification of +2 for a wrong answer as an engagement-contingent reward is not established by the cited source; it is a reward tied to an incorrect…	Now reads: "+2 for a wrong answer" may raise a related concern about contingent rewards (deci1999 [44] flags engagement-contingent rewards)
05-M46	[570]	The source supports the claim that people are uncertain about the consequences of privacy-related	Now reads: - acquisti2015 [570]: "people's uncertainty about the

Citation	Source	Problem found	Correction
		behavior, but the abstract does not establish that minimising data is superior to relying on consent.	consequences of privacy-related behaviors and their own preferences over those consequences" (review of privacy behaviour; supports uncertainty about privacy consequences, not that…
05-M46	[561]	Reyes et al. found widespread COPPA problems involving third-party SDKs in children's Android apps, including SDK transmission of identifiers and PII; the source does not show that an SDK sent camera frames or audio from this app, so calling it the exact…	Now reads: third-party SDK risk is among the failures reyes2018 [561] found in children's apps
05-M25	[17]	The abstract supports the explanation-feedback effect, prose-passage setting, and two-day delay, but it does not identify the participants as college students.	Now reads: (Prose passages, test two days later)
05-M45	[566]	The source supports concern about camera proctoring and higher test anxiety, but its population was university students and its setting was online exam proctoring, not children or general camera and microphone access in a learning app.	Now reads: (1,760 Dutch university students across four survey waves studying online exam proctoring)
05-M45	[567]	The source specifically analyzes facial recognition and facial detection in compulsory schools; the article broadens this to camera sensing generally and to the privacy and consent design choice without stating that limitation.	Now reads: school-based facial recognition raises social and governance concerns
05-M26	[533]	The review supports the general link between trust and appropriate reliance on automation, but it does not show that bad learning items erode trust.	Now reads: lee2004 [533] (trust influences reliance on automation)
05-M26	[534]	The source concerns workers' trust in AI technology; the reported learning-item check is not described as involving AI, so this source does not support the mechanism.	Citation removed.
05 proposals (each backed by 3 or more s	[553]	The ICO code warns against using nudges to lead children to provide unnecessary personal data or turn off privacy protections; it does not support a general warning against nudges for children as such.	Now reads: nudges that undermine children's privacy choices (ico-aadc-13 [553], ico-aadc-5 [50])
05 proposals (each backed by 3 or more s	[48]	WCAG 2.5.8 sets a 24 CSS-pixel minimum target size subject to exceptions, not a 44 px target; the 44 px standard is covered by the separately cited WCAG 2.5.5.	Now reads: wcag-2.5.8 [48] (24 CSS-pixel minimum), wcag-2.5.5 [535] (44 CSS-pixel enhanced target)
05 proposals (each backed	[166]	The study does not show that short-answer retrieval is generally best with feedback. It found that result in its second experiment, with specific study materials and a delayed test; without feedback, multiple choice outperformed short answer.	Now reads: kang2007 [166] (short answer performed best with feedback in its second experiment)

Citation	Source	Problem found	Correction
by 3 or more s			
05 proposals (each backed by 3 or more s	[6]	The expertise-reversal review supports tailoring instructional support to learner expertise, but does not establish that younger or low-prior-knowledge learners will feel frustrated; its stated concern is that techniques effective for novices can lose…	Now reads: mismatched support for learners at different expertise levels (kalyuga2003 [6] cuts both ways)
How it all fits together	[4], [5]	kulik1990 reports positive effects, not large ones; the large gains are from Bloom's experiments. (found on the second check)	Now reads: A meta-analysis of 108 evaluations found positive effects, strongest for weaker students (kulik1990 [4]), and Bloom's experiments reported large gains (bloom1984 [5]).
How it all fits together	[29], [30]	efron1975 and gelman2012 support shrinkage and partial pooling, not false-discovery control. (found on the second check)	Now reads: It uses shrinkage and partial pooling (efron1975 [29]; gelman2012 [30]) and false-discovery control (benjamini1995 [300]) so that noisy readings don't steer teaching.
How it all fits together	[48]	One WCAG success criterion was cited as if it covered all of WCAG 2.2. (found on the second check)	Now reads: Accessibility follows WCAG 2.2, for example its target-size criterion (wcag-2.5.8 [48]).
How it all fits together	[49], [50]	One COPPA rule and one code standard were cited as if they covered the whole of each. (found on the second check)	Now reads: Privacy for children follows COPPA and the UK children's code, for example COPPA's consent rule and the code's standard on detrimental use of data (coppa-312.5a1 [49]; ico-aadc-5 [50]).
02-M45	[248]	The verdict said feedback cancels the lure risk; the source says it reduces it. (found on the second check)	Now reads: Feedback reduces, though it does not remove, multiple choice's lure risk.
03-M46	[34]	The classroom result is about feedback on help seeking, not about help itself. (found on the second check)	Now reads: aleven2016help [34] — review of help seeking with intelligent tutoring systems; in its classroom study, feedback on help seeking made students use on-demand help more deliberately but did not improve learning outcomes.
01-P9	[78]	The abstract shows what the source supports (adaptive difficulty at a mean success probability of 0.75) but not this proposal's 80–85% target. (found on the second pass)	Now reads: wauters2010 [55], klinkenberg2011 [78] (adaptive difficulty, with items sampled at a mean success probability of 0.75; it does not set the 80–85% target), tetzlaff2025 [61].
04-M47	[460]	Art. 2.5.7 of the cited rules gives the no-charge semicircle radius as 1.30 m, measured to the outer edge of the semicircle line, so the bare figure of 1.25	Now reads: the no-charge semicircle (radius 1.30 m to the outer edge of its line, Art. 2.5.7; 1.25 m to the inner edge)

Citation	Source	Problem found	Correction
		m was not what the source states. (found on the second pass)	
04-P1	[78]	The earlier correction dropped pelanek2016 [77] from the line, which the review had not meant to remove. (found on the second pass)	Now reads: klinkenberg2011 [78] reports arithmetic practice items sampled at a mean success probability of 0.75, so tasks were challenging but not too difficult; this is not direct evidence about chess or persistence.
04-P2	[496]	The earlier correction dropped winstein1990 [497] from the line, which the review had not meant to remove. (found on the second pass)	Now reads: salmoni1984 [496] discusses the guidance hypothesis (knowledge of results as guidance); it does not by itself establish benefits of reduced feedback frequency. Also: winstein1990 [497]: title-level only (reduced feedback frequency).

Second pass: fuller text and replacement sources

- **Fuller text for the undecided citations.** 128 occurrences had first been left undecided, because the source had no abstract on file. Abstracts or open full text were then fetched for 46 sources, from publishers' own landing pages, PubMed, ERIC records and open-access copies. Where a publisher's site would not give an automated reader its abstract, none was taken. The same strong reviewer then judged those occurrences again against the new text, with the same guards on every correction.
- **Result.** 71 of the 128 are now confirmed as written. 22 were corrected (they are in the table above) and 4 were kept after review. 31 are still undecided: for 26 no abstract or open full text could be obtained, and for 5 the text was read and still did not settle whether the quote supports the claim. Seven of them were settled by hand against full text: the AERA/APA/NCME Standards and the FIBA rules (both read in full), the full text of the achievement-emotions meta-analysis, and the publisher's description of a book. Six were confirmed, and one led to a hand correction (in the table above). Those 31 are left as they were, and each source was checked earlier by DOI or official URL, or is marked title-level only.
- **Replacement sources.** After the removals, 17 mechanisms had fewer than three quoted sources, and the second pass left two more (03-M43 and 05-M5) in the same position. A search for replacements found 23 papers that do support part of what the mechanisms cite them for, and they were added to 15 mechanisms (table below). For each one, the DOI was resolved on Crossref and the authors, title and year matched; a quote of under 25 words was copied from the abstract and found verbatim; a strong model judged whether the abstract supports the use made of it; and the result was read by hand. Most of these papers support the principle, not Savant's exact setting, and each entry has a Scope line that says so. Nothing was added where no honest source was found.
- **Independent re-check.** Claude Sonnet 5.5, a different model from the reviewer above, then re-checked all 23 added sources and the 29 corrections from the second pass against

their abstracts (52 items). It flagged 10. Nine led to a change: seven scope sentences were made narrower or more exact, and two corrections that had dropped a citation (04-P1, 04-P2) were repaired. For the tenth, the abstract on file turned out to belong to a different paper, which also meant a second correction made on a similar basis was wrong; both (01-M50, 03-M8) were undone and left as written, undecided.

Occurrences still undecided:

Where	Source	Why
01-M10	pelanek2016 [77]	no abstract or open full text available
01-M49	pelanek2016 [77]	no abstract or open full text available
01-M50	wood1976 [174]	no abstract or open full text available
01-M56	anderson1972 [139]	no abstract or open full text available
01-M59	keller1968 [100]	no abstract or open full text available
01-P4	pelanek2016 [77]	no abstract or open full text available
02-M17	pyc2009 [215]	no abstract or open full text available
02-M40	kapler2015 [244]	no abstract or open full text available
Scorecard weak rules (section 03)	brown1910 [273]	text read; the quote could not be tied to it
Scorecard weak rules (section 03)	spearman1910 [274]	no abstract or open full text available
Scorecard weak rules (section 03)	haldane1956 [279]	no abstract or open full text available
Scorecard weak rules (section 03)	anscombe1956 [280]	text read; the quote could not be tied to it
03-M7	haldane1956 [279]	no abstract or open full text available
03-M7	anscombe1956 [280]	text read; the quote could not be tied to it
03-M8	brown1910 [273]	text read; the quote could not be tied to it
03-M8	spearman1910 [274]	no abstract or open full text available
Section 03 proposals	haldane1956 [279]	no abstract or open full text available
Section 03 proposals	anscombe1956 [280]	text read; the quote could not be tied to it
04-M3	wald1945 [384]	no abstract or open full text available
04-M21	pelanek2016 [77]	no abstract or open full text available
04-M25	chase1973 [426]	no abstract or open full text available
04-M43	gelman2013 [381]	no abstract or open full text available
04-M54	millman1965 [487]	no abstract or open full text available
04-P1	pelanek2016 [77]	no abstract or open full text available
05-M19	stankov2012 [526]	no abstract or open full text available
05-M29	vatavu2015 [537]	no abstract or open full text available
05-M38	gollwitzer2006 [551]	no abstract or open full text available
05-M46	hoepman2014 [569]	no abstract or open full text available
05 proposals (each backed by 3 or more s	castleman2015 [550]	no abstract or open full text available

Where	Source	Why
05 proposals (each backed by 3 or more s	gollwitzer2006 [551]	no abstract or open full text available
05 proposals (each backed by 3 or more s	vatavu2015 [537]	no abstract or open full text available

Sources added (the reference numbers continue the list):

Mechanism	Source	What it supports, and what it does not show
01-M26	lievens2007 [585]	Shows that retest scores were less saturated with g and more associated with memory than initial scores, in a high-stakes cognitive ability test. It did not separate identical from varied items or test a no-repeat window.
01-M39	calamia2012 [586]	Defines practice effects as including memory for specific items, and reports that the use of alternate forms was associated with the size of the change in many cases. It did not test freshly written questions on a retry.
01-M39	lievens2007 [585]	Shows that retesting the same measure biases scores. It did not test freshly written questions.
03-M20	rohrer2006 [198]	Backs only the 'less practice once mastered' half: extra practice added nothing in this task. It did not study learners who stop early.
03-M20	rohrer2005over [261]	Backs only the 'less practice once mastered' half. It did not test adaptive block sizes.
03-M57	zell2014 [587]	These are self-reports. Read across to a grown-up's unverified report, it shows only that unverified reports track performance loosely. It did not study an adult confirming a child's skill.
03-M57	bernes2021 [588]	Shows that informant ratings can line up only weakly with directly measured performance in these groups. It did not study an adult confirming a skill, or any weight such as 0.25.
04-M25	gobet1996b [589]	Shows that board recall differs by skill even for random positions. It did not test beginners' knowledge of piece moves or board set-up.
04-M25	astudillo2026 [590]	Shows that novices reconstruct real positions better than random ones, and links real-position reconstruction to rating. It did not test knowing how each piece moves or setting up the full board.
04-M26	chabris2003 [591]	Shows that grandmasters' blunder frequency and size did not differ between rapid and blindfold play, despite the extra visualisation load. It did not test typing whole lines as a training method.
04-M33	liang2018 [592]	Supports that sustain-pedal onset and offset can be measured and classified by technique. It did not use MIDI CC64 or grade pedalling against chord onsets.
04-M36	poolton2005 [593]	Supports minimising errors early in learning. It did not test note-by-note waiting, slow practice or independent play in time.
04-M36	furuya2013 [594]	Supports practising below top speed before playing at tempo. It did not test waiting at each note or errorless learning.
04-M43	thall1994 [595]	Supports deciding on the posterior probability that a binary success rate is promising. It did not use a Beta(1,1) prior or a learner gate.
04-M43	tuyl2008 [596]	Notes that Beta(1,1) is among the usual non-informative priors for a binomial rate, and favours the Bayes–Laplace prior over Jeffreys in its sensitivity analysis. It did not study pass gates.

Mechanism	Source	What it supports, and what it does not show
04-M44	shortridge2014 [597]	Supports smoothing noisy local shooting rates in basketball. It did not set a minimum number of shots per zone or show per-zone intervals.
04-M44	wongtoi2023 [598]	Supports describing shooting by court region. It did not study small-sample noise, minimum shots or intervals.
04-M56	hambleton1974 [599]	Shows that item order changes performance and stress. It did not test mixed difficulty in a game.
04-M56	wang2021 [600]	Shows that item-group difficulty and the order of difficulty groups affected test performance in this sample. It did not test mixed difficulty in a game or motivation.
05-M16	skipper2012 [601]	Shows that person praise worsens the response to a later failure. It did not test warmth or sincerity.
05-M16	morris2014 [602]	Shows that ambiguous praise, like effort praise, led to better persistence than trait praise after failure. It did not test warmth or sincerity.
05-M40	lazar2006 [603]	Shows that computer frustration is common and costly. It did not test empty states.
05-M40	taipalus2025 [604]	Shows that novices value a plain statement of what is wrong plus a hint on how to fix it, the pattern an empty state follows. It did not test empty states.
03-M43	sinha2021 [605]	Shows that attempting problems before being taught helped on average. It found the opposite trend for second to fifth graders, and it did not test a tutor that withholds the answer to the question on screen.
03-M43	kapur2014 [606]	Supports letting learners attempt a problem before being told. It did not study an AI coach or withholding the answer to a question on screen.
05-M5	roll2011 [607]	Shows that feedback on help-seeking errors improved help-seeking and the gain carried over. It did not test an XP discount or whether help replaces trying.

Mechanisms that stay under three. 4 mechanisms still have fewer than three quoted sources, and each is marked on its verdict line: 03-M38 (a search found only game studies in which easier conditions raised engagement but slowed learning, which does not support reading 85% as strong, so none was added); 03-M45 (the one candidate describes step-level hints but reports no learning outcome, which is the reason the earlier review was removed, so none was added); 04-M22 (no paper with an abstract on file shows that puzzle-type tasks measure chess strength; the one candidate showed that problem solving varies with opening specialisation); 04-M23 (no paper tests endgame play-outs against perfect defence; the one candidate found only that serious study predicts chess skill).

Contradictions

The text was searched for statements that contradict each other without saying why. The search ran once over the whole article, and once over every pair of sections side by side. Each candidate was then checked against both passages, and against the code where the passages disagreed on a number. 13 contradictions were confirmed and fixed. Two more candidates, in 01-M40 and 05-M12, were not contradictions: the code's profile formula gives 6 to 18 items, and its clamp at 5 matters only when the scorecard shortens a session, as the text says. The other candidates were already explained in the text.

Where	What conflicted	Fix
01-M6	The What line gave the raw-accuracy cut-offs (60%/80%) while the Status line said gaps and strengths now come from a belief.	What line marked as the behaviour before 25 October 2026.
01-M21	The Final was said to need 36 of 40 in one line and 34 of 40 in the Status line.	36/40 marked as the earlier rule; 34/40 is current.
02-P1 / 02-M17	A table headed "what the code does now" showed d = 0.65 for recall 0.90-0.95, a step removed on 25 October.	Table heading now says it shows the code before that change.
03-M15	The observed-preference weight was given as capped at 0.75, while the Status line says 1.0.	0.75 marked as the earlier cap.
03-M21	The stretch clamp was given as 0.05-0.35, while the Status line says the cap is now 0.25.	0.35 marked as the earlier cap.
03-M38	The summary text uses a 60% cut and the difficulty delta a 50% cut; both are what the code does, but the article did not say they differ.	A note that the two cuts differ in the code.
03-M53	Objective items were said to need exactly 4 options, while the Status line says 3 or 4 are now allowed.	Marked as the rule before 25 October 2026.
04-M19	CH2 and CH3 were said to still have no "nothing to find" positions, while the Status line says they were added.	Marked as the gap before 25 October 2026, closed that day.
04-M21	Practice puzzles were said to give about 50% expected success, while the Status line says they now aim at 0.75.	50% marked as the earlier behaviour.
04-M36	The tempo ladder and hotspot loop were said to be not built, while the Status line says both were built.	Marked as the state before 25 October 2026.
04-M50	The Status line said both the B03 and B07 aim cues now read "toward the back of the ring"; the fact re-verification shows B07 reads "the back of the ring, not your elbow".	Each drill's actual wording given.
Summary table / 01-P6	The summary listed the placement belief "with an 'unsure' list" as live, while the open items say unsure skills are not shown yet.	The summary now says the list is computed but not yet shown.
05-P1 / 05-M13	The built proposal said the break line comes after about 25 minutes "or 50 for older stages"; the code has one setting, 25 minutes by default.	The proposal now says what was built.

Most of these had one cause. The sections describe the code as it stood on 25 October 2026, and the changes built that day are recorded on each mechanism's Status line. Where a description and its Status line disagreed, the description now says it is the earlier behaviour.

How it all fits together

This is one learner's path through Savant. For each step it names the mechanisms involved, how each one sets up the next, and the strongest research behind each link.

1. Placement finds where to start. A new learner answers a short adaptive test, not a fixed quiz. The app keeps a belief about where the learner is on the grade's path. It asks the question nearest the middle of that belief, and stops when the belief is narrow enough, or the time budget runs out (01-M7 to 01-M12). This is computerized adaptive testing, which is long established (owen1975 [2]; weiss1984 [3]). Its result feeds everything after it. Placement marks units as known, and seeds starting mastery on a cautious, shrunk scale (01-M5). The per-skill gaps and strengths are now set by a belief rather than a raw percentage (01-M6, switch `PLACEMENT_SKILL_BELIEF`).

2. The path teaches only the gaps. The learning path steps over units the learner has shown they know, and records why (01-M13, 01-M14). This follows mastery learning: move on when a unit is mastered, not when time is up. A meta-analysis of 108 evaluations found positive effects, strongest for weaker students (kulik1990 [4]), and Bloom's experiments reported large gains (bloom1984 [5]). It also avoids the expertise-reversal effect, where teaching what a learner already knows can hurt (kalyuga2003 [6]). An earlier-grade gap scan is offered, never forced (01-M15).

3. Each lesson teaches, then checks. A lesson runs: topic, worked example, a harder example, a summary, then "Check what stuck" (01-M24, 02-M41). Worked examples first is well supported for novice problem-solving material (sweller1985 [7]; renkl2014 [8]; atkinson2000 [9]). The closing check is retrieval practice. Testing what was just learned beats reading it again (roediger2006 [1]; rowland2014 [10]; adesope2017 [11]). The check questions never repeat the examples (01-M26), so they test understanding, not memory of the page. Lessons written by AI pass a verification gate before any learner sees them (01-M31, 01-M34). On 25 October 2026 that gate was extended to run without the paid key (switch `LEVEL_SELF_CHECK`).

4. Practice sessions mix review, focus and stretch. A session opens with about 30% due review, one step easier, so it starts with success (01-M41). The focus work then mixes earlier topics with the new one (01-M43, switch `INTERLEAVE_SESSIONS`). In mathematics, interleaved practice gives test gains over blocked practice of about d = 0.42 a day later and d = 0.79 to 0.83 about a month later (rohrer2015 [12]; rohrer2020 [13]), with caution for words and expository texts (brunmair2019 [14]). Feedback after each answer gives the right answer and the reason why (02-M45, 05-M15, 05-M17). Explanation feedback beats a bare "right/wrong" (vanderkleij2015 [15]; shute2008 [16]; butler2013 [17]). On 25 October 2026, the reviews inside a lesson also started to come from the forgetting model (switch `CURVE_REVIEWS`). One scheduler now decides what is due everywhere.

5. Checkpoints confirm mastery. Every 25 and 100 levels, and at the end of a grade, a checkpoint samples every skill introduced so far, on a difficulty ramp (01-M21, 01-M37). Every retry uses new questions (01-M39). This is the "test, correct, retest" cycle of mastery learning (kulik1990 [4]; bloom1984 [5]). On 25 October 2026 the pass decisions became likelihood-ratio rules (switches `CHECKPOINT_SPRT` and `CHALLENGE_SPRT`), which control how often a learner is wrongly passed or failed (lewis1990 [18]; spray1996 [19]). A learner who is truly at 90% now passes the Final 90% of the time; before, it was 63%. A failed checkpoint leads to a diagnosis built from the actual wrong answers (01-M60 to 01-M64).

6. Retention enforcers keep what was learned. When a study session or a day ends, the app offers a short "retention enforcer": a sitting of unaided questions with feedback (02-M1, 02-M2). It then schedules more at growing gaps, from semi-weekly to monthly (02-M8 to 02-M11). Two of the best-replicated findings in learning research support this: retrieval practice (karpicke2008 [20]; yang2021 [21]) and spacing (cepeda2006 [22]; cepeda2008 [23]). Retrieving to a set standard and then relearning at widely spaced intervals ("successive relearning") is the schedule this research recommends (rawson2011 [24]). Which topics get asked is chosen by a forgetting model: topics predicted to be fading come first (02-M16, 02-M47). This follows tested adaptive schedulers (settles2016 [25]; tabibian2019 [26]; lindsey2014 [27]). On 25 October 2026 its difficulty window was widened to 0.95 recall (switch `ENFORCER_DIFFICULTY_WINDOW_95`), to match eglington2020 [28]. "Retained" now needs a real 7-day delay (switch `RETAIN_REAL_DELAY`), so graduation rests on durable learning, not same-day memory (01-M58, 01-M59).

7. The scorecard adapts the teaching. The scorecard watches how a learner behaves: pace, persistence, hint use, confidence. It turns a reading into a directive only when the evidence is strong enough (03-M1 to 03-M12). It uses shrinkage and partial pooling (efron1975 [29]; gelman2012 [30]) and false-discovery control (benjamini1995 [300]) so that noisy readings don't steer teaching. A stated learning style is never used to match the teaching, because that "meshing" idea has not survived testing (pashler2008 [31]; 03-M14, 05-M9). What does steer the teaching is the learner's own record on each objective (03-M17). Worked examples come first while an objective is new, and fade once it is proven (kalyuga2003 [6]). The 25 October 2026 changes made several directives better supported: a softer challenge setting (jansen2013 [32]), worked-example hints (koedinger2007 [33]; aleven2016help [34]), and an item-relative "too fast to count" line (wise2006 [35]; 03-M39 to 03-M42). The coach asks rather than tells (03-M43, 05-M23). Step-based tutoring comes close to a human tutor (vanlehn2011 [36]), and intelligent tutors raised test scores by a median 0.66 SD in 50 controlled evaluations (kulik2016 [37]), and guardrails can mitigate the learning harm of unfettered AI access (bastani2025 [38]).

8. Skill tracks apply the same rules to the body and the ear. In chess, piano and basketball, a stage is proven only by "cold" sittings on two different days (04-M1, 04-M2). Performance during practice can be an unreliable guide to long-term learning (soderstrom2015 [39]). Practice difficulty is matched to the learner's skill level (04-M11), following deliberate practice (ericsson1993 [40]) and the challenge-point idea (guadagnoli2004 [41]). Varied and random practice orders help retention (shea1979 [42]; 04-M29). Feedback is now faded as the learner improves (switch `PIANO_FEEDBACK_FADE`). Youth load caps protect young players (switch `SKILL_LOAD_CAPS`).

9. Engagement and wellbeing keep the learner coming back without pressure. Streaks, XP and progress bars help some learners. The evidence on them is mixed: gamification has small positive effects on learning (sailer2020 [43]), but rewards can undermine interest in some forms (deci1999 [44]). Savant keeps them informational. On 25 October 2026 a learner can hide them (switch `UX_HIDE_STATS`), and a streak now forgives one missed day a week (switch `STREAK_LOCAL_GRACE`), because one missed day does not break a habit (lally2009 [45]). Choice and autonomy support motivation (patall2008 [46]; ryan2000 [47]). Accessibility follows WCAG 2.2, for example its target-size criterion (wcag-2.5.8 [48]). Privacy for children follows COPPA and the UK children's code, for example COPPA's consent rule and the code's standard on detrimental use of data (coppa-312.5a1 [49]; ico-aadc-5 [50]).

How the parts reinforce each other.

- Placement and the gap-only path send the learner's time to what is not yet known. The forgetting model and the enforcers then protect what is known, so each new lesson builds on kept knowledge rather than forgotten knowledge.
- Retrieval runs through every step: lesson checks, review inside sessions, checkpoints, enforcers and skill sittings. Each retrieval both measures and strengthens memory (roediger2006 [1]; rowland2014 [10]). The same answers feed the forgetting model and the scorecard, so measuring and teaching use one record.
- "Kept" means the same thing everywhere: first try, unaided, not rushed, after a real delay (02-M28, 01-M58, 04-M1). So placement, coverage, checkpoints and the skill gate all count the same kind of evidence.
- The effort monitor removes rushed answers before they count (03-M39 to 03-M42). This keeps the mastery, coverage and scorecard numbers honest.
- Engagement features invite a return visit without pressure (05-M1 to 05-M9). Spacing only works if the learner comes back.

Section 01. Curriculum, levels and placement

This section covers how Savant places a learner, how a grade's path is laid out, how a lesson is built, how a session is planned, how mastery moves a learner on, how grade-fit and objective coverage are enforced, how a failed checkpoint is remediated, and the rules the AI lesson writers work under.

In numbers: 64 mechanisms; 64 with 3 or more quoted sources; 19 strong, 27 moderate, 6 mixed, 12 design choice; 171 distinct sources cited or listed; 13 mechanisms built or changed on 25 October 2026.

01 at a glance

Mechanism	Verdict	Quoted sources	Status
01-M1 Placement length scales with the claimed tier	design choice	4	existing
01-M2 Strictly ascending ladder, rotating through the grade's skills	mixed	5	existing
01-M3 Band scoring: hold at 70%, half credit at 50-70%, stop after two failing bands	design choice	3	existing
01-M4 Placement lands on a lesson, never on a checkpoint	moderate	4	existing
01-M5 Seeded mastery is shrunk and capped; the first review comes early	design choice	3	existing
01-M6 Per-skill "gaps" under 60% and "strengths" at 80% or better	mixed	4	built on 25 October 2026
01-M7 A Bayesian belief over position, with a guessing floor and a slip ceiling	strong	5	existing
01-M8 Ask nearest the belief's middle; stop when the belief is narrow, or at the cap or the clock	strong	3	existing
01-M9 Mode time budgets: Quick 20, Balanced 45, Accurate 90 minutes	design choice	4	existing
01-M10 Item positions calibrated by judges and simulated testers; "reliable" shift 0.3	mixed	3	existing; open proposal
01-M11 In-grade unit placement	moderate	4	existing
01-M12 Quick check of skipped units; a unit missed twice is reopened	design choice	4	existing
01-M13 One evidence model per unit; nothing is skipped without the unit's own evidence	strong	3	existing
01-M14 The path steps over known units and records why	strong	5	existing
01-M15 Earlier-grade gap scan, offered and never forced	design choice	4	existing
01-M16 Level rhythm: lesson every 10, Master every 25, Grand Master every 100, Graduation at the end	moderate	4	existing

Mechanism	Verdict	Quoted sources	Status
01-M17 Practice levels between lessons	mixed	4	existing
01-M18 Grade-band difficulty, linear inside the band	design choice	3	existing; open proposal
01-M19 Each level's focus skill, and the cumulative "introduced" set	strong	4	existing
01-M20 Retention levels (retrieval only)	strong	4	existing
01-M21 Checkpoint lengths 15, 25 and 40; pass marks 75%, 75% and 90%	moderate	3	built on 25 October 2026; open proposal
01-M22 Written checkpoints in English	moderate	3	existing
01-M23 The Challenge Exam: a 25-level clean run, then five bands of seven, each band perfect	moderate	4	built on 25 October 2026
01-M24 The order of a lesson: topic, worked example, a harder example, summary, then "Check what stuck"	strong	5	existing; open proposal
01-M25 Examples at difficulty d−1 then d; checks at d and d±0.5	moderate	4	existing; open proposal
01-M26 Checks never repeat the examples, or anything asked recently	strong	3	existing
01-M27 The lesson says where it sits: topic N of M, which levels practise it, and the next lesson	moderate	3	existing
01-M28 A lesson level is cleared by working through it, not by its check score	design choice	4	existing
01-M29 The lesson writer is told the topic, its neighbours and the last 4 lessons	design choice	4	existing
01-M30 The AI lesson writer's rules	moderate	5	existing
01-M31 Validating lessons and items, including the key gate	moderate	4	built on 25 October 2026; open proposal
01-M32 How a block lesson is laid out	strong	5	existing
01-M33 Checks may not be answerable from the text verbatim; they test understanding on new material	moderate	4	existing
01-M34 Fact research and the verification gate	strong	6	built on 25 October 2026
01-M35 A 10% holdout gets the shared plan	strong	4	existing
01-M36 A fixed blueprint; tailoring changes only how questions are presented	moderate	3	existing
01-M37 Checkpoint sampling covers every introduced skill, on a difficulty ramp	strong	4	existing
01-M38 Checkpoint formats and styles are varied	moderate	5	existing
01-M39 Every retry uses freshly written questions, with no cooldown	strong	4	existing
01-M40 Session length: 6 to 18 items from the profile	design choice	4	existing

Mechanism	Verdict	Quoted sources	Status
01-M41 About 30% due review, first, at d−1	strong	4	built on 25 October 2026
01-M42 Support and stretch shares	moderate	4	existing; open proposal
01-M43 Interleaving within the focus items, and the order of roles	strong	5	built on 25 October 2026
01-M44 Question families, shape rotation and no-repeat history	moderate	3	existing
01-M45 Typed answers preferred for recall-leaning learners	mixed	5	existing
01-M46 One recall question at the end of each level, from what was missed	moderate	4	existing
01-M47 The end-of-session recall set	moderate	4	built on 25 October 2026
01-M48 Next-session difficulty moves ±0.5	design choice	4	existing; open proposal
01-M49 Mastery is an exponential moving average, weighted by difficulty	moderate	3	existing
01-M50 Coach-assisted answers count for less	design choice	3	existing
01-M51 Review intervals and ease	moderate	4	existing
01-M52 Due reviews: weakest and most overdue first	moderate	4	built on 25 October 2026
01-M53 The level-up gate: average 0.8 or better, weakest 0.6 or better, 12 or more items	moderate	4	existing
01-M54 Which skills fit which grade, and re-tagging by rule	strong	4	built on 25 October 2026
01-M55 Difficulty caps per grade and skill, and the school grade passed to the bank	strong	4	built on 25 October 2026
01-M56 Objectives come from each lesson's own "covers" sentence	moderate	4	existing
01-M57 Knowledge tracing per objective, and what counts as proof	strong	4	existing
01-M58 Retained or lapsed: a delayed, unaided check	moderate	4	built on 25 October 2026
01-M59 The coverage gate holds graduation until every objective is retained	moderate	3	built on 25 October 2026
01-M60 Two doors: retry now, or "Help me learn"	moderate	3	existing; open proposal
01-M61 Diagnosis from the actual wrong answers	strong	4	existing
01-M62 Five drill questions per area; an area holds at the pass mark	moderate	4	existing
01-M63 Diagnosis starts as soon as the quiz is lost	mixed	4	existing
01-M64 After a failed Final, remediation is planned from the whole record	moderate	4	existing

01-A. Placement: the difficulty ladder (`services/placement.js`)

The difficulty ladder is used for grades without a unit bank. Grades with a bank use the adaptive walk in section B.

01-M1. Placement length scales with the claimed tier

- **Where:** `config/curriculum.js:745-747`, used by `services/placement.js:26-28`.
- **What:** 20 questions at tier 1, rising 7.5 per tier to 50 at tier 5. Fixed length, with no precision-based stop.
- **Evidence:**
  - babcock2012 [51]. *Setting:* simulation of CAT termination rules.
    - "longer CATs yielded more accurate trait estimation, but there were diminishing returns with a very large number of items"
  - ackerman2009 [52]. *Setting:* 239 first-year university students on SAT batteries of 3.5 to 5.5 hours.
    - "subjective fatigue increased with increasing time-on-task. However, mean performance increased in the longer test length conditions"
  - debeer2014 [53]. *Setting:* PISA 2009 reading, 15-year-olds.
    - "A decrease in examinee effort during the PISA reading assessment was found consistently across countries"
  - borgonovi2016 [54]. *Setting:* PISA 2012, 15-year-olds.
    - "academic endurance, defined as students' ability to maintain their baseline rates of successful test completion for the duration of a low-stakes cognitive test"
- **Verdict:** design choice. More items buy precision with diminishing returns, while effort drops over a low-stakes test. No study fixes 20 to 50.
- **Match:** partly. The trade-off is respected by the cap. The gap is that the ladder stops on length, not on precision. The adaptive walk (M8) already does the better thing; this ladder is the fallback.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M2. Strictly ascending ladder, rotating through the grade's skills

- **Where:** `services/placement.js:17-38`.
- **What:**
  - Difficulty climbs linearly from 1 to one point above the claimed tier's target (`{1:3, 2:4.5, 3:6, 4:7.5, 5:9}`), with a cap of 10.
  - Item *i* goes to skill `i % skills.length`, so every skill is sampled at several difficulties.
- **Evidence:**
  - weiss1984 [3]. *Setting:* review of CAT applications in education.
    - "Each of these applications makes use of item response theory to select test questions from an item pool"

    - wauters2010 [55]. *Setting:* review of adaptive item-based learning environments.
        - "adaptive item sequencing by matching the difficulty of the items to the learner's knowledge level"
    - hartig2012 [56]. *Setting:* PISA, multilevel IRT.
        - "A significant negative effect of item position on performance was found in all countries"
    - debeer2013 [57]. *Setting:* IRT methodology; supports position effects, not learner fatigue.
        - "changes in item order may affect item and test characteristics"
    - wainer1987 [58]. *Setting:* measurement essay.
        - "a bundle of items that can be arranged either hierarchically or linearly, thus maintaining the efficiency of an adaptive test"
- **Verdict:** mixed.
    - Content balancing across skills is standard.
    - A fixed ascending order confounds difficulty with position and fatigue: the hardest items always come last, when effort is lowest. Adaptive selection is more efficient.
- **Match:** partly. Skill rotation matches. Ascending order means a tired learner looks less able at the top of the ladder.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M3. Band scoring: hold at 70%, half credit at 50-70%, stop after two failing bands

- **Where:** `services/placement.js:44-54`, `services/placement.js:71-123`.
- **What:**
    - Five difficulty bands. A band at 70% or better is held, which credits its top. A band at 50-70% credits its middle.
    - Two consecutive bands under 50% end the walk, so one lucky hard answer or one slip cannot decide the level.
- **Evidence:**
    - lewis1990 [18]. *Setting:* simulated certification exam.
        - "average test lengths can be reduced by half without sacrificing classification accuracy"
    - spray1996 [19]. *Setting:* simulated pass/fail classification.
        - "the SPRT procedure required fewer test items than the sequential Bayes procedure to achieve the same level of classification accuracy"
    - wise2017 [59]. *Setting:* review of low-stakes testing.
        - "very short response time—termed rapid guessing—has been shown to indicate disengaged test taking"
- **Verdict:** design choice. Sequential classification with a stopping rule is well supported. The 70/50 cut-offs and the "two bands" rule are reasoned values.
- **Match:** yes in spirit.

    - The gap: with roughly 4 to 10 items a band, a 70% band is noisy, and nothing models guessing on 4-option items.
    - The adaptive walk (M7) models guessing explicitly. The ladder does not.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M4. Placement lands on a lesson, never on a checkpoint

- **Where:** `config/level-layout.js:490-508`, `services/placement.js:56-69`.
- **What:**
    - The placed level is rounded to the nearest lesson at or below it, or else ahead of it.
    - A learner is never dropped onto a quiz they have not been taught for, nor onto a bare practice level.
- **Evidence:**
    - kalyuga2003 [6]. *Setting:* review, cognitive load theory.
        - "Instructional techniques that are highly effective with inexperienced learners can lose their effectiveness and even have negative consequences when used with more experienced learners."
    - tetzlaff2025 [61]. *Setting:* meta-analysis of the expertise reversal effect.
        - "Low prior knowledge learners learn better from high-assistance instruction (d = 0.505)."
    - renkl2014 [8]. *Setting:* theory review, example-based learning.
        - "Learning from examples is a very effective means of initial cognitive skill acquisition."
    - pan2023pretest [62]. *Setting:* review of the pretesting literature.
        - "can benefit learning if there is an opportunity to study the correct answers afterwards"
- **Verdict:** moderate. Starting new material with instruction rather than a test fits the guidance and worked-example literature. A test without later study is the case pretesting does not support.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M5. Seeded mastery is shrunk and capped; the first review comes early

- **Where:** `services/placement.js:142-161`.
- **What:**
    - Seeded mastery is `min(0.7, accuracy × 0.85)`. The placement items are added to `attempts`, so later evidence keeps its weight.
    - The first review is due in 1 day, or in 3 days for a skill at 80% or better.
- **Evidence:**
    - morris1983 [63]. *Setting:* statistics review.
        - "reviews the state of multiparameter shrinkage estimators with emphasis on the empirical Bayes viewpoint"

    - casella1985 [64]. *Setting:* statistics tutorial.
        - "The empirical Bayes model is much richer than either the classical or the ordinary Bayes model and often provides superior estimates of parameters."
    - cepeda2008 [23]. *Setting:* 1,354 adults, retention up to a year.
        - "the optimal gap declined from about 20 to 40% of a 1-week test delay to about 5 to 10% of a 1-year test delay"
- **Verdict:** design choice. Shrinking thin evidence toward a prior is the textbook remedy. The 0.7 cap, the 0.85 factor and the 1-day or 3-day gaps are reasoned values.
- **Match:** yes.
    - The shrinkage is a fixed multiplier, not one scaled by how many items each skill got.
    - A skill seen twice and a skill seen ten times shrink alike. See proposal P6.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M6. Per-skill "gaps" under 60% and "strengths" at 80% or better

- **Where:** `services/placement.js:125-139`.
- **What:** before 25 October 2026, each skill's raw placement accuracy marked it as a gap (under 60%) or a strength (80% or better). Since then a belief sets gaps, strengths and an "unsure" list (see Status).
- **Evidence:**
    - parsons2019 [66]. *Setting:* methods paper, cognitive-behavioural measures.
        - "subpar reliability has severe implications for measures' validity and interpretation"
    - hedge2018 [67]. *Setting:* reliability of cognitive tasks.
        - "low between-subject variability causes low reliability for individual differences, destroying replicable correlations with other factors" (reliabilities ranged 0 to.82 across seven classic cognitive tasks)
    - brown1978 [68]. *Setting:* elementary arithmetic bugs.
        - "explaining why a student is making a mistake as opposed to simply identifying the mistake"
    - kalyuga2004 [69]. *Setting:* rapid knowledge tests.
        - "a highly significant correlation (up to.92) between performance on the rapid measure and traditional measures of knowledge"
- **Verdict:** mixed.
    - Rapid diagnostic tests can be valid when they are built for the purpose (kalyuga2004 [69]).
    - A raw percentage from 4 to 12 incidental items per skill is unreliable. It labels a gap without saying why (brown1978 [68]).
- **Match:** partly. A skill can be called a gap on 1 of 3 right. Proposal P6 makes this Bayesian, with an "unsure" outcome.

- **Status: Built on 25 October 2026** (switch `PLACEMENT_SKILL_BELIEF`, from 01-P6). Gaps, strengths and a new "unsure" list come from a Bayes belief: a strength needs P(known) >= 0.9, a gap P(known) <= 0.15. The owner keeps the percentage rule.

01-B. Adaptive placement: the subject walk and the unit walk

01-M7. A Bayesian belief over position, with a guessing floor and a slip ceiling

- **Where:** `services/adaptive-placement.js:24-41`, `services/adaptive-placement.js:69-96`.
- **What:**
  - Belief sits on a grid in steps of 0.05 grade.
  - P(correct) = luck + (1 − slip − luck)·logistic((x − pos + 0.3)/0.2), where luck = guess + 0.15 (capped at 0.6) and slip = 0.10.
  - The posterior comes from summed log likelihoods.
- **Evidence:**
  - owen1975 [2]. *Setting:* adaptive mental testing (statistical method).
    - "A new kind of Bayesian motivated procedure is introduced which leads to a strongly consistent estimator."
  - bock1982 [65]. *Setting:* CAT simulation.
    - "is shown to have unusually good properties for computerized adaptive testing"
  - barton1981 [70]. *Setting:* four achievement-test data sets.
    - "The fourth parameter increased the likelihood in only two of the four sets."
  - baker2008 [71]. *Setting:* Cognitive Tutor data.
    - "existing methods for determining these parameters are prone to the Identifiability Problem"
  - vanderlinden1998 [72]. *Setting:* CAT method.
    - "This paper suggests several item selection criteria for adaptive testing which are all based on the use of the true posterior."
- **Verdict:** strong. Grid (EAP-style) Bayesian estimation with a guessing floor is the established CAT method.
- **Match:** yes.
  - The upper asymptote (slip) is added. Barton and Lord found it helps only sometimes, so it is harmless rather than necessary.
  - The parameters were tuned on simulated testers, not fitted to real responses (see M10).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M8. Ask nearest the belief's middle; stop when the belief is narrow, or at the cap or the clock

- **Where:** `services/adaptive-placement.js:44-67` (MODES), `services/adaptive-placement.js:176`, `services/adaptive-placement.js:205-211`.
- **What:**

Mode	Most questions	Fewest questions	Mass within ±1 grade to stop
Quick	32	18	0.94
Balanced	70	20	0.97
Accurate	130	40	0.995

- Questions come in grade blocks, re-aimed after every answer.

- **Evidence:**
  - babcock2012 [51]. *Setting:* CAT simulation.
    - "Fixed-length CATs performed either slightly worse or comparable to their variable-length termination counterparts"
  - lewis1990 [18]. *Setting:* simulated certification test.
    - "shorter tests for individuals who have clearly mastered (or clearly not mastered) a given subject and longer tests for those individuals"
  - eggen1999 [73]. *Setting:* simulated classification CAT.
    - "Wald's (1947) sequential probability ratio test can be implemented as an adaptive test for classifying examinees into categories."
- **Verdict:** strong for the principle (variable-length, posterior-based stopping). The thresholds are design choices, tuned in simulation.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M9. Mode time budgets: Quick 20, Balanced 45, Accurate 90 minutes

- **Where:** `services/adaptive-placement.js:50-55`, `services/unit-placement.js:32-47`, `services/gap-scan.js:31-40`.
- **What:** one budget per mode, shared by the subject walk, the in-grade walk and the gap scan. The owner set the limits.
- **Evidence:**
  - ackerman2009 [52]. *Setting:* university students.
    - "subjective fatigue increased with increasing time-on-task"
  - debeer2014 [53]. *Setting:* PISA.
    - "A decrease in examinee effort during the PISA reading assessment was found consistently across countries"
  - wise2005 [74]. *Setting:* low-stakes computer-based testing.
    - "unmotivated examinees will answer too quickly (i.e., before they have time to read and fully consider the item)"
  - babcock2012 [51]. *Setting:* CAT simulation.
    - "there were diminishing returns with a very large number of items"

- **Verdict:** design choice.
- **Match:** yes. The learner picks the precision they will sit for.
    - Gap: 90 minutes is long for a child, and the walk does not watch for rapid guessing late in a long session. The scorecard measures rapid guessing, but placement does not use it.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M10. Item positions calibrated by judges and simulated testers; "reliable" shift 0.3

- **Where:** `services/adaptive-placement.js:24-41`, `services/unit-placement.js:80-117`.
- **What:**
    - Each question has a calibrated point, the place where a learner answers it reliably (about 9 in 10). The judges set it (`scripts/placement-second-judge.js`).
    - The slope is centred 0.3 grade below that point.
    - Uncalibrated unit questions sit at 0.75 of their unit.
- **Evidence:**
    - kaser2024 [75]. *Setting:* systematic review of simulated learners.
        - "almost half of the studies using simulated learners do not provide any evidence that their modeling addresses the most fundamental question in simulation design"
    - meyers2008 [76]. *Setting:* state test equating.
        - "item parameter invariance is threatened when an item appears in a different location on the live test"
    - pelanek2016 [77]. *Setting:* review, Elo in adaptive systems. Title-level only.
    - klinkenberg2011 [78]. *Setting:* Math Garden, on-the-fly item calibration. Title-level only.
    - doroudi2017 [79]. *Setting:* offline policy evaluation.
        - "existing methods for doing so can overestimate the performance of new teaching methods"
- **Verdict:** mixed.
    - Calibrating from expert judgement and simulation is a reasonable cold start.
    - The literature calibrates item difficulty from real responses, and warns that simulation-validated policies overestimate.
- **Match:** partly. No step updates item positions from live learner responses. See P4.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 01-P4, calibrating item positions from live responses (needs data).

01-M11. In-grade unit placement

- **Where:** `services/unit-placement.js:31-61`, `services/unit-placement.js:167-230`.
- **What:**

Mode	Most questions	Fewest questions	Confidence to stop	Clock
Quick	14	6	0.85	6 min

Mode	Most questions	Fewest questions	Confidence to stop	Clock
Balanced	30	10	0.95	13 min
Accurate	80	16	0.995	25 min

- Unit slope width 0.2 and reliable shift 0.3 (in units). Rounding goes 0.15 past the middle.
- The learner starts at the placed unit's first lesson.

- **Evidence:**
  - lewis1990 [18]. Mastery testing with Bayesian decision theory.
    - "A theoretical framework for mastery testing based on item response theory and Bayesian decision theory is described."
  - spray1996 [19]. Placement as classification.
    - "Examples of such applications include licensure and certification, college selection, and placement into entry-level or developmental college courses."
  - eggen1999 [73]. *Setting:* simulation.
    - "testing algorithms using KLI-based item selection performed better than or as well as those using FI-based item selection"
  - owen1975 [2]. As in M7.
- **Verdict:** moderate. The method is the established classification CAT. The numbers are simulation-tuned design choices.
- **Match:** yes. The comment's claim that the simulated Accurate mode places 926 of 928 testers is a simulation, not a field result.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M12. Quick check of skipped units; a unit missed twice is reopened

- **Where:** `services/unit-placement.js:1-20`, `services/unit-placement.js:244-301` (`markCheck`, with `misses >= 2` at line 290).
- **What:**
  - Each unit placed past gets one check question, unless placement already saw that unit answered right.
  - A unit whose check is missed twice is reopened, and the learner restarts at the earliest reopened unit.
- **Evidence:**
  - kalyuga2005 [80]. *Setting:* adaptive e-learning.
    - "instruction was dynamically tailored to changing levels of expertise using rapid tests of knowledge combined with measures of cognitive load"
  - kalyuga2004 [69]. *Setting:* rapid assessment.
    - "a highly significant correlation (up to.92) between performance on the rapid measure and traditional measures of knowledge"
  - lewis1990 [18].

        - "average test lengths can be reduced by half without sacrificing classification accuracy"
    - pan2023pretest [62].
        - "The evidence to date indicates that prequestioning and pretesting can often enhance learning"
- **Verdict:** design choice.
- **Match:** yes. Gap-only paths replace this with the unit belief (M13).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-C. Gap-only learning paths (`services/learner-path.js`, `services/gap-scan.js`)

01-M13. One evidence model per unit; nothing is skipped without the unit's own evidence

- **Where:** `services/learner-path.js:1-20`, `services/learner-path.js:30-58` (knobs), `services/learner-path.js:60-110` (`unitBelief`).
- **What:**
    - The model folds answers into a Bayes belief that the unit is known, using these knobs:

Knob	Value
Prior for a unit placed past (HOLE_PRIOR)	0.15
Prior at or after the placed unit (ISLAND_PRIOR)	0.65
SLIP	0.15
LUCK	0.15 + 1/choices
Known at	0.965 or above
Learn at	0.6 or below

- The model is asymmetric on purpose: a false skip is "the worst mistake", so undecided units are worked.

- **Evidence:**
    - vandesande2013 [81]. *Setting:* analysis of BKT.
        - "use the individual student's performance at each opportunity to apply a skill to update the conditional probability that the student has learned that skill"
    - baker2008 [71]. *Setting:* Cognitive Tutor.
        - "The model uses four parameters per skill, fit using student performance data, to relate performance to learning."
    - lewis1990 [18]. Asymmetric losses in Bayesian mastery decisions.
        - "A theoretical framework for mastery testing based on item response theory and Bayesian decision theory is described."

    - corbett1995 [82]. *Setting:* knowledge tracing in the ACT Programming Tutor. Title-level only.
- **Verdict:** strong for the method (Bayesian evidence and decision thresholds). The knob values are design choices.
- **Match:** yes.
    - Gap: the SLIP and LUCK knobs are global, not fitted per unit.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M14. The path steps over known units and records why

- **Where:** `services/mastery.js:221-273` (`advance`, `skippedCheckpointPatch`), `services/learner-path.js:1-20`.
- **What:**
    - Units that are known or done are stepped over. Their checkpoints are recorded as cleared, with their source: placement, quick check, test-out or gap scan.
    - A unit that is still unchecked is worked like `learn`.
- **Evidence:**
    - kalyuga2007 [84]. *Setting:* review, learner-tailored instruction.
        - "This paper reviews recent empirical findings associated with the expertise reversal effect"
    - kalyuga2003 [6].
        - "Instructional techniques that are highly effective with inexperienced learners can lose their effectiveness"
    - steenbergenhu2016 [85]. *Setting:* second-order meta-analysis, K-12.
        - "accelerated students significantly outperformed their nonaccelerated same-age peers (g = 0.70)"
    - kulik1984accel [86]. *Setting:* meta-analysis of 26 controlled studies, K-12.
        - "examination performance of accelerates surpassed by nearly one grade level the performance of nonaccelerates of equivalent age and intelligence"
    - major2021 [87]. *Setting:* meta-analysis, low- and middle-income countries.
        - "more personalised approaches which adapt or adjust to learners' level led to significantly greater impact (an effect size of 0.35)"
- **Verdict:** strong. Not re-teaching known material, and letting proven learners move on, both have converging meta-analytic support.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M15. Earlier-grade gap scan, offered and never forced

- **Where:** `services/gap-scan.js:1-40`.
- **What:**

    - The scan looks at units below the recommended grade: where placement misses map down, and every unit of the grade just below.
    - Every candidate unit gets one question before any gets a second, likeliest gaps first.
    - A unit is decided only after one of its own bank questions. Units still undecided when the questions or the clock run out are `unchecked`, and are worked.
- **Evidence:**
    - vanlehn1990 [88]. *Setting:* children's arithmetic.
        - "As children acquire arithmetic skills, they often develop small, local misconceptions that cause systematic errors."
    - brown1978 [68]. *Setting:* procedural bugs in arithmetic.
        - "explaining why a student is making a mistake as opposed to simply identifying the mistake"
    - kalyuga2005 [80]. Rapid diagnostic assessment.
        - "instruction was dynamically tailored to changing levels of expertise using rapid tests of knowledge"
    - koedinger2023 [89]. *Setting:* 27 datasets, 1.3 million observations.
        - "typically increasing by about 0.1 log odds or 2.5% in accuracy per opportunity"
- **Verdict:** design choice. Local gaps below the frontier are well documented. Fixing them by targeted practice fits koedinger2023 [89]'s finding that students differ mainly in initial performance, not in learning rate. No study sets the scan's order or the optional offer.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-D. The learning path and level layout (`config/curriculum.js`, `config/level-layout.js`)

01-M16. Level rhythm: lesson every 10, Master every 25, Grand Master every 100, Graduation at the end

- **Where:**
    - `config/level-layout.js:64` (`DEFAULT_RULE`) and `config/level-layout.js:267-285` (`ruleKind`).
    - `config/curriculum.js:468-470`.
    - Level 1 is a lesson.
- **What:**
    - On a grade without a stated layout, levels 10, 20, 30 and so on are lessons, 25, 50 and 75 are Master quizzes, 100 and 200 are Grand Master quizzes, and the last level is the Graduation exam.
    - Grades with stated layouts (the English, ACT and Spanish sheets) put levels where the sheet says.
- **Evidence:**

    - bangertdrowns1991freq [90]. *Setting:* meta-analysis of classroom testing frequency.
        - "students who took at least one test during a 15-week term scored about one half of a standard deviation higher"
    - yang2021 [21]. *Setting:* meta-analysis of classroom quizzing.
        - "overall testing (quizzing) raises student academic achievement to a medium extent (g = 0.499)"
    - kulik1990 [4]. *Setting:* meta-analysis of 108 mastery learning studies.
        - "mastery learning programs have positive effects on the examination performance of students"
    - khanna2013 [91]. *Setting:* college introductory psychology.
        - "Students completing courses with cumulative finals retained more than students who took noncumulative finals."
- **Verdict:** moderate.
    - Frequent checkpoints, mastery gating and a cumulative final are each supported.
    - The 10/25/100 spacing is a design choice.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M17. Practice levels between lessons

- **Where:**
    - `config/level-layout.js:267-285`, where levels that are not lessons or checkpoints are `practice`.
    - `config/curriculum.js:866-876` (`practiceUntil`).
    - `services/lessons.js:193-203`.
- **What:**
    - The nine or so levels after a lesson practise that lesson's topic until the next lesson.
    - The focus skill stays the same across the run. Review and weak-skill items are mixed in only by the session planner (M41-M43).
- **Evidence:**
    - rohrer2007 [92]. *Setting:* college mathematics.
        - "the practice problems are systematically shuffled so that each practice set includes a variety of problems drawn from many previous lessons"
    - rohrer2020 [13]. *Setting:* RCT in 7th-grade mathematics classes.
        - "the interleaved group outscored the blocked group, 61% versus 38%, d = 0.83"
    - taylor2010 [93]. *Setting:* children, mathematics.
        - "interleaving rather than blocking practice of different skills... usually improves subsequent test performance"
    - brunmair2019 [14]. *Setting:* meta-analysis.

        - "interleaving can effectively foster inductive learning... should be used with caution in certain conditions, especially for expository texts and words"
- **Verdict:** mixed.
    - Massed practice right after a lesson aids acquisition.
    - The interleaving studies show that mixing in earlier topics beats blocked runs on delayed tests, in mathematics especially. Brunmair limits this for expository text and words.
- **Match:** partly. Within a level, about 30% review and every third focus item come from other skills, so practice is partly interleaved. The run as a whole is still blocked on one topic. See P1.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M18. Grade-band difficulty, linear inside the band

- **Where:** `config/curriculum.js:749-787` (`difficultyForLevel`, `gradeBand`).
- **What:**
    - A grade gets a slice of the 1-10 scale, centred on its position in the track and at least 2 points wide (`GRADE_BAND_WIDTH`).
    - Difficulty climbs linearly from the first level to the Final.
    - The code itself calls this "Linear is the honest choice for v1: item statistics in Phase 8 replace it".
- **Evidence:**
    - wilson2019 [94]. *Setting:* theory and simulation.
        - "the optimal error rate for training is around 15.87% or, conversely, that the optimal training accuracy is about 85%"
    - wauters2010 [55].
        - "adaptive item sequencing by matching the difficulty of the items to the learner's knowledge level"
    - koedinger2023 [89].
        - "typically increasing by about 0.1 log odds or 2.5% in accuracy per opportunity"
    - metcalfe2005 [95]. *Setting:* study-time allocation (region of proximal learning); not evidence for a linear difficulty progression.
- **Verdict:** design choice.
    - A steady, slow increase in challenge fits the learning-rate data.
    - The linear map from level to item difficulty is uncalibrated.
- **Match:** partly. See P4: calibrate from real p-values, as the code already plans.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 01-P4, difficulty from live responses (needs data).

01-M19. Each level's focus skill, and the cumulative "introduced" set

- **Where:** `config/curriculum.js:789-850` (`levelPlan`).
- **What:**
  - The focus skill is the skill of the topic taught at that level.
  - `introduced` holds every skill met so far. Review, checkpoints and level-up all draw on that set, "which is what keeps level 200 from forgetting level 3".
- **Evidence:**
  - mayfield2002 [96]. *Setting:* college students, mathematics.
    - "On the final test, the cumulative group outscored the other groups on application and problem-solving items."
  - rohrer2007 [92].
    - "each practice set includes a variety of problems drawn from many previous lessons"
  - lawrence2013 [97]. *Setting:* introductory psychology.
    - "However, low-scoring students remembered more of the course material when they took cumulative exams throughout the semester."
  - szpunar2007 [98]. *Setting:* lab.
    - "subjects benefited more from initial testing when they also knew they would need to remember the information for ..."
- **Verdict:** strong. Cumulative practice and cumulative assessment converge.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M20. Retention levels (retrieval only)

- **Where:**
  - `services/block-levels.js:64-70` (`isRetention`) and `services/block-levels.js:81-91`.
  - `services/block-levels.js:765-791`, where `SHAPE.retention` is 3-6 blocks and 3-6 checks, and the prompt says "Every block makes the learner produce or recognize from memory; nothing re-teaches at length".
- **What:** practice levels the sheet labels "(retention)" are built as pure retrieval practice on material already taught, aimed at the learner's struggles.
- **Evidence:**
  - roediger2006 [1]. *Setting:* college, prose passages.
    - "Taking a memory test not only assesses what one knows, but also enhances later retention"
  - adesope2017 [11]. *Setting:* meta-analysis.
    - "practice tests are more beneficial for learning than restudying and all other comparison conditions"
  - agarwal2021 [99]. *Setting:* review of applied research in schools.
    - "the majority of which (57%) revealed medium or large benefits from retrieval practice"
  - rowland2014 [10]. *Setting:* meta-analysis.

        - "initial recall tests yielding larger testing benefits than recognition tests"
- **Verdict:** strong.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M21. Checkpoint lengths 15, 25 and 40; pass marks 75%, 75% and 90%

- **Where:** `config/curriculum.js:595-607`.
- **What:**
    - A Master quiz is 15 items, a Grand Master quiz 25 and the Graduation exam 40.
    - The pass marks are 75%, 75% and 90%. Because accuracy must reach the mark, 15 items need 12 right (80%), 25 need 19 and, before 25 October 2026, 40 needed 36 (the Final now needs 34/40; see Status).
- **Evidence:**
    - kulik1990 [4].
        - "mastery learning programs have positive effects on the examination performance of students"
    - guskey2007 [60].
        - "Bloom believed that all students could be helped to reach a high criterion of learning"
    - keller1968 [100]. *Setting:* the unit-mastery personalized system of instruction. Title-level only.
    - kulik1979 [101]. *Setting:* meta-analysis of PSI outcome studies. Title-level only.
    - babcock2012 [51].
        - "longer CATs yielded more accurate trait estimation" (a simulation; it supports the general principle that longer tests measure more accurately, not the checkpoint lengths themselves)
- **Verdict:** moderate. Mastery gating at a high criterion is supported. The exact marks and lengths are design choices.
- **Match:** partly. This section worked out the binomial pass rates:

Checkpoint	Learner who is truly at	Chance of passing
Master (15 items)	70%	0.30
Master (15 items)	80%	0.65
Master (15 items)	90%	0.94
Final (40 items)	90%	0.63

- So the Final fails about a third of learners who really are at the mark.
- Fixed-length checkpoints carry known classification error that nothing reports. See P3.

- **Status: Built on 25 October 2026** (switch `CHECKPOINT_SPRT`, from 01-P3). Pass decisions use a likelihood-ratio rule. Master 12/15 and Grand 19/25 are unchanged; the Final needs 34/40 (was 36/40). An exam ended early passes only on decisive evidence. The banners and the path page now show the mark each decision really needs (master 80%, grand

76%, Final 85%) through `passMarkFor` and `passMarksFor` in `config/curriculum.js`, used in `routes/learn.js`; the owner's base keeps 75%, 75% and 90%. **Open:** 01-P3, extending a near-mark checkpoint by a few questions (needs the quiz client to accept extra questions).

01-M22. Written checkpoints in English

- **Where:**
  - `config/curriculum.js:512-571` (`ESSAY_TRACKS`, `ESSAY_PASS_MARK 0.75`, `ESSAY_MIN_WORDS 80`, quick-check lengths 5 and 8).
  - `services/sessions.js:1279-1330` (`areasFromEssay`: a rubric criterion under 3/4 is an area to work on).
- **What:** English Master and Grand Master checkpoints, or the sheet's Essay Draft and Essay Revision rows, are written work marked on a rubric. Under-scoring criteria become drill areas.
- **Evidence:**
  - graham2015 [102]. *Setting:* meta-analysis, grades 1-8.
    - "feedback to students about writing from adults, peers, self, and computers statistically enhanced writing quality"
  - panadero2013 [103]. *Setting:* review of 21 rubric studies.
    - "Findings indicate that rubrics may have the potential to influence students learning positively"
  - jonsson2007 [104]. *Setting:* review of rubric reliability and validity. Title-level only.
  - kane2013 [105]. *Setting:* validity theory.
    - "To validate an interpretation or use of test scores is to evaluate the plausibility of the claims based on the scores."
- **Verdict:** moderate.
- **Match:** yes.
  - Gap: the AI rubric marking has no reported agreement with human markers, which the rubric-reliability literature asks for.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M23. The Challenge Exam: a 25-level clean run, then five bands of seven, each band perfect

- **Where:**
  - `config/curriculum.js:630-702` (`SKIP_STREAK_LEVELS 25`, `SKIP_MARK 0.9`, 5 bands × 7 questions).
  - `services/mastery.js:407-487`.
- **What:**
  - After 25 levels each cleared at 90% or better, a 35-question exam on the next five ten-level marks is offered.

    - The learner lands on the first band that is not 7/7. Missing the first band skips nothing.
- **Evidence:**
    - steenbergenhu2016 [85].
        - "acceleration appeared to have a positive, moderate, and statistically significant impact on students' academic achievement (g = 0.42)"
    - kulik1984accel [86].
        - "examination performance of accelerates surpassed by nearly one grade level the performance of nonaccelerates"
    - lewis1990 [18].
        - "shorter tests for individuals who have clearly mastered (or clearly not mastered) a given subject"
    - kalyuga2003 [6].
        - "can lose their effectiveness and even have negative consequences when used with more experienced learners" (supporting expertise-sensitive instruction rather than level skipping)
- **Verdict:** moderate. Letting proven learners skip ahead is well supported. The "perfect band" rule is a design choice.
- **Match:** partly.
    - A learner with a true 90% accuracy clears a 7/7 band only 48% of the time, and 70% at 95%. Clearing all five bands at 95% happens 17% of the time.
    - The rule is far stricter than its stated intent. See P3: a Bayesian band decision such as M13's.
- **Status: Built on 25 October 2026** (switch `CHALLENGE_SPRT`, from 01-P3). A gate band of 16 holds at 14 right, bands of 6 above it allow one miss, 40 questions (was 35). The owner keeps five bands of seven.

01-E. Lesson structure (`services/lessons.js`, `public/js/session.js`)

01-M24. The order of a lesson: topic, worked example, a harder example, summary, then "Check what stuck"

- **Where:**
    - `services/lessons.js:153-224` (`buildFallback` sections).
    - `public/js/session.js:283-297`, which renders intro → sections → worked example → common mistakes → "In short" → a "Check what stuck (N questions)" button.
- **What:** every lesson states its topic and what it covers, shows a worked example, shows one a step harder, summarises, then asks 3 check questions (2 or more for AI lessons).
- **Evidence:**
    - sweller1985 [7]. *Setting:* high-school algebra.
        - "schema acquisition is retarded by conventional problem-solving search techniques"

    - renkl2003 [106]. *Setting:* theory review.
        - "learning from worked-out examples, in comparison to problem solving, is very effective during the initial stages of cognitive skill acquisition"
    - atkinson2000 [9]. *Setting:* review.
        - "Worked examples are instructional devices that provide an expert's problem solution for a learner to study."
    - mayer1996 [107]. *Setting:* college, science summaries.
        - "recalled these steps and solved transfer problems as well as or better than students who received the full text along with the summary"
    - roediger2006 [1].
        - "Taking a memory test not only assesses what one knows, but also enhances later retention"
- **Verdict:** strong. Worked examples first, then retrieval, is one of the most replicated sequences.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 01-P2, faded worked examples and self-explanation prompts (changes the lesson format; not built).

01-M25. Examples at difficulty d−1 then d; checks at d and d±0.5

- **Where:** `services/lessons.js:118-151`.
- **What:**
    - Demo 1 sits at d−1 and demo 2 at d, both from the topic's family where one exists.
    - One on-topic check sits at d, and two from the rest of the skill sit at d−0.5 and d+0.5.
    - There is no fading: examples are complete, then checks are full problems.
- **Evidence:**
    - renkl2002 [108]. *Setting:* one field and two lab experiments.
        - "the fading procedure fosters learning, at least when near transfer performance is considered"
    - atkinson2003 [109]. *Setting:* two experiments, probability problems.
        - "this combination produced medium to large effects on near and far transfer without requiring additional time on task"
    - salden2010 [110]. *Setting:* Cognitive Tutor geometry.
        - "Both experiments provide evidence of improved learning results from adaptive fading of worked examples"
    - vanmerrienboer2003 [111]. *Setting:* design framework.
        - "we present a framework for scaffolding practice and just-in-time information presentation, aiming to control cognitive load effectively"
    - pollock2002 [112]. *Setting:* simple-to-complex sequencing of complex information. Title-level only.

- **Verdict:** moderate. Moving from easier to harder is supported. Fading worked steps, with self-explanation prompts, beats example-problem pairs.
- **Match:** partly. The easy-to-hard step matches. Fading and self-explanation prompts are missing. See P2.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 01-P2, faded worked examples (not built).

01-M26. Checks never repeat the examples, or anything asked recently

- **Where:**
  - `services/lessons.js:75-90` (`LESSON_HISTORY 400`) and `services/lessons.js:135-150`.
  - `services/sessions.js:109-125` (`HISTORY_WINDOW 600`) and `services/sessions.js:81-107`, which tries up to 10 regenerations for an unseen item.
- **What:**
  - A check is never the example shown above it.
  - Practice avoids the last 600 items the learner was asked, on the same skills; lessons avoid the last 400.
- **Evidence:**
  - pan2018 [113]. *Setting:* meta-analysis of transfer.
    - "testing can yield transferrable learning as measured relative to a nontesting reexposure control condition (d = 0.40"
  - butler2017 [114]. *Setting:* four experiments, college.
    - "All four experiments showed that variability during retrieval practice produced superior transfer of knowledge to new examples."
  - lievens2007 [585]. *Setting:* 941 candidates retaking a cognitive ability test in a high-stakes selection setting.
    - "the factor underlying the retest scores was less saturated with g and more associated with memory"
    - Scope: Shows that retest scores were less saturated with g and more associated with memory than initial scores, in a high-stakes cognitive ability test. It did not separate identical from varied items or test a no-repeat window.
- **Verdict:** strong. Repeating an item measures memory for that item, and varied items transfer better.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M27. The lesson says where it sits: topic N of M, which levels practise it, and the next lesson

- **Where:** `services/lessons.js:153-203` (`onward`, `about`, `intro`).

- **What:** the lesson opens with its topic and the creator's "covers" sentence, and ends by naming the practice levels and the next topic.
- **Evidence:**
  - ausubel1960 [117]. *Setting:* college students, an unfamiliar passage. Ausubel's stated hypothesis, which his study supported:
    - "can be facilitated by the advance introduction of relevant subsuming concepts (organizers)"
  - luiten1980 [118]. *Setting:* meta-analysis of 135 studies.
    - "Advance organizers were shown to have a facilitative effect on both learning and retention."
  - hattan2023 [120]. *Setting:* systematic review of 54 articles on prior-knowledge activation in reading. It replaces the unverified scorecard citation, whose title belongs to a 2015 paper by the same first author.
    - "the prompting of students' existing knowledge was beneficial during and after reading as well"
- **Verdict:** moderate. Advance organizers have meta-analytic support with small effects. Stating the objective is good practice.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M28. A lesson level is cleared by working through it, not by its check score

- **Where:** `services/sessions.js:1134-1150` (`resolveBlock`), `services/mastery.js:292-297` (`completeLesson`).
- **What:**
  - Finishing the lesson advances the learner whatever the checks scored.
  - The comment states that a failed check "should send the learner back over the lesson, not stall the path". No code sends them back.
- **Evidence:**
  - kapur2016 [121]. *Setting:* theory review.
    - "Conditions that maximize performance in the initial learning may not maximize learning in the longer term."
  - pan2023pretest [62].
    - "can benefit learning if there is an opportunity to study the correct answers afterwards"
  - pashler2005 [122]. *Setting:* college, foreign vocabulary.
    - "Supplying the correct answer after an incorrect response not only improved performance during the initial learning session--it also increased final retention by 494%."
  - hattie2007 [123]. *Setting:* review.
    - "Feedback is one of the most powerful influences on learning and achievement, but this impact can be either positive or negative."

- **Verdict:** design choice. A check that does not gate is defensible, because retrieval and feedback teach even when the answer is wrong. The literature says a miss should be followed by corrective study.
- **Match:** partly.
    - Feedback after each check exists.
    - Nothing targets re-study after missed lesson checks, though the next practice levels do pick up weak skills. See P7.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M29. The lesson writer is told the topic, its neighbours and the last 4 lessons

- **Where:** `services/ai.js:531-555`, `services/lessons.js:256-293` (the last 4 lessons, each summary cut to 200 characters).
- **What:**
    - The writer gets: "Teach this topic. Do not teach the whole skill, and do not substitute a neighbouring topic."
    - It also gets: "Do not teach any of those again. Take the next step beyond them".
- **Evidence:**
    - kalyuga2003 [6]. Redundant instruction harms more expert learners.
        - "can lose their effectiveness and even have negative consequences when used with more experienced learners"
    - liu2024lost [124]. *Setting:* LLM long-context benchmarks.
        - "significantly degrades when models must access relevant information in the middle of long contexts"
    - levy2024 [125]. *Setting:* LLM reasoning versus input length.
        - "a notable degradation in LLMs' reasoning performance at much shorter input lengths than their technical maximum"
    - shi2023 [126]. *Setting:* LLM distraction.
        - "model performance is dramatically decreased when irrelevant information is included"
- **Verdict:** design choice. Keeping the context short and relevant fits the LLM findings. Not re-teaching fits expertise reversal.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M30. The AI lesson writer's rules

- **Where:** `services/ai.js:556-585` (lesson system prompt).
- **What:** the prompt sets these rules.
    - Teach one idea.
    - Everything must be correct, so do the arithmetic.
    - The worked example shows every step with its reasoning.

    - Common mistakes are written as a fixed "Wrong: … Why: … Right: …" string, never as a rule.
    - Checks must be answerable from the lesson, with exactly one correct answer and four choices for multiple choice.
    - Plain text only. Never mention AI.
- **Evidence:**
    - harp1998 [127]. *Setting:* college, science passages.
        - "students who read expository passages with seductive details... recalled significantly fewer main ideas"
    - sundararajan2020 [128]. *Setting:* meta-analysis.
        - "learners exposed to interesting and irrelevant information, known as seductive details, do not perform as much as those who learned without seductive details"
    - schroeder2022 [129]. *Setting:* meta-analysis of refutation text.
        - "refutation text is associated with a positive, moderate effect (g = 0.41"
    - mclaren2015 [130]. *Setting:* middle school, decimals, web tutor.
        - "the erroneous examples group performed significantly better on a delayed test, given a week after the initial post-test (d = .33, for gain scores)"
    - tippett2010 [131]. *Setting:* science education review.
        - "These structures often contain misconceptions—inaccurate or incomplete information—that are highly resistant to change"
- **Verdict:** moderate. Coherence (one idea, no seductive detail), refutation of misconceptions and erroneous examples marked as wrong all have meta-analytic or delayed-test support.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M31. Validating lessons and items, including the key gate

- **Where:**
    - `services/ai.js:269-310` (`validateItem`).
    - `services/ai.js:606-640` (`validateLesson`).
    - `services/ai.js:592-604`: the key gate removes a multiple-choice check whose rationale describes another option and that a blind check does not confirm. A lesson left with fewer than 2 checks is refused.
- **What:** the validators enforce these rules.
    - Sections need 20 or more characters.
    - A worked example needs 2 or more steps and an answer.
    - Multiple choice has exactly 4 distinct choices, and the answer is one of them.
    - An answer longer than 6 characters must not appear inside the prompt.
    - No "NaN" or "undefined".
    - A rationale needs 10 or more characters.

- **Evidence:**
  - haladyna2002 [132]. *Setting:* review of 27 textbooks and 27 studies.
    - "A taxonomy of 31 multiple-choice item-writing guidelines was validated through a logical process"
  - little2012 [133]. *Setting:* college.
    - "multiple-choice tests could trigger productive retrieval processes-provided the alternatives were made plausible enough"
  - wang2022sc [135]. *Setting:* LLM reasoning benchmarks.
    - "selects the most consistent answer by marginalizing out the sampled reasoning paths"
  - rodriguez2005 [136]. *Setting:* meta-analysis of 80 years of research.
    - "More 3-option items can be administered than 4- or 5-option items per testing time while improving content coverage"
- **Verdict:** moderate. Machine checks for well-formed items follow item-writing guidelines. An independent re-solve (the key gate) follows self-consistency.
- **Match:** yes, with one note. Insisting on exactly 4 options is stricter than the evidence needs, since 3 options are psychometrically as good. That is harmless but costs writing tokens. See P8.
- **Status: Built on 25 October 2026** (switch `LEVEL_SELF_CHECK`, from 01-P10). Without the verifier key, a block level is served only after a self-check and a blind re-solve by a judge from another model family. A safety rule, so it applies to the owner too. **Open:** 01-P8, 3-option practice items (the gate is in `services/ai.js`; not built).

01-M32. How a block lesson is laid out

- **Where:** `services/block-levels.js:765-791` (`SHAPE`, `DEFAULT_LAYOUT`, `LAYOUT_SYSTEM`).
- **What:**
  - A lesson has 4-10 blocks and 2-5 checks. The default is rule card → text → steps → common error → cheat sheet.
  - The prompt sets this order: activate prior knowledge, teach in short chunks, make the learner retrieve and apply, confront the likely misconception, and end with a summary or memory anchor.
  - Coverage points are listed first. Prerequisites come before dependents, concrete before abstract, and practice right after its idea. No tag may be used more than twice.
- **Evidence:**
  - rey2019 [137]. *Setting:* meta-analysis of segmenting.
    - "reveals a significant segmenting effect with small to medium effects for retention and transfer performance"
  - noetel2022 [138]. *Setting:* meta-meta-analysis of multimedia design.
    - "We found 11 design principles that demonstrated significant, positive, meta-analytic effects on learning"
  - luiten1980 [118].

        - "Advance organizers were shown to have a facilitative effect on both learning and retention."
    - schroeder2022 [129].
        - "refutation text is associated with a positive, moderate effect (g = 0.41"
    - yang2021 [21].
        - "overall testing (quizzing) raises student academic achievement to a medium extent (g = 0.499)"
- **Verdict:** strong for each component.
- **Match:** yes. These are prompt instructions. Whether a given layout follows them is not checked by code beyond the schema, the tag limit and the coverage list.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M33. Checks may not be answerable from the text verbatim; they test understanding on new material

- **Where:**
    - `services/block-levels.js:896-907`: "A check must not be answerable from a block's text verbatim".
    - `services/block-checkpoints.js:68-80`: "Test understanding, not recognition: make the learner use, identify, correct, compare or justify on new material".
- **Evidence:**
    - anderson1972 [139]. *Setting:* measurement review.
        - "to propose the distinctions necessary to determine whether a person has comprehended an instructional communication"
    - agarwal2019 [140]. *Setting:* classroom and lab experiments.
        - "Critically, higher order and mixed quizzes improved higher order test performance, but fact quizzes did not."
    - pan2019clicker [141]. *Setting:* college biology.
        - "Benefits of jargon quizzing did not consistently generalize, however, to exam questions that assessed conceptual knowledge"
    - butler2017 [114].
        - "variability during retrieval practice produced superior transfer of knowledge to new examples"
- **Verdict:** moderate.
- **Match:** partly.
    - For multiple choice, the server only refuses an answer longer than 6 characters that appears in the prompt (`services/block-levels.js:1080-1083`).
    - Nothing checks overlap between a check and the block text. Enforcement rests on the prompt and the verifier.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M34. Fact research and the verification gate

- **Where:** `services/block-levels.js:1215-1260` (research and verify prompts), `services/block-levels.js:1346-1365` (`gateOn`, `servable`, `verifyOn`), `services/block-levels.js:1679-1690` (`tooMuchRemoved`).
- **What:**
  - Before writing, a researcher builds a fact sheet with source quotes, and marks each fact established, contested, opinion or theoretical.
  - After writing, a verifier checks every claim one by one. False and unsupported claims are removed. Contested claims stated as fact are reworded.
  - Under the gate, only a verified plan is served. A plan that loses more than a third of its blocks is written again.
  - The code notes that the in-house GLM route caught only 2 of 13 known errors, so the gate needs the owner's key.
- **Evidence:**
  - ji2023 [142]. *Setting:* survey of NLG hallucination.
    - "deep learning based generation is prone to hallucinate unintended text, which degrades the system performance"
  - huang2025 [143]. *Setting:* survey of LLM hallucination.
    - "LLMs are prone to hallucination, generating plausible yet nonfactual content."
  - dhuliawala2024 [144]. *Setting:* LLM experiments.
    - "we show COVE decreases hallucinations across a variety of tasks"
  - min2023 [145]. *Setting:* LLM biographies.
    - "breaks a generation into a series of atomic facts and computes the percentage of atomic facts supported by a reliable knowledge source"
  - lewis2020 [146]. *Setting:* knowledge-intensive question answering.
    - "their ability to access and precisely manipulate knowledge is still limited"
  - bastani2025 [38]. *Setting:* about 1,000 high-school maths students.
    - "48% improvement in grades for GPT Base... 17% reduction in grades for GPT Base"
- **Verdict:** strong. LLMs produce factual errors, and claim-level verification against retrieved sources reduces them.
- **Match:** yes when the gate is on.
  - Gap: with no verifier key configured, `gateOn()` is false and plans are served unverified.
  - Given 2 of 13 errors caught on the fallback route, that is the main residual risk in this section. See P10.
- **Status: Built on 25 October 2026** (switch `LEVEL_SELF_CHECK`, from 01-P10). Without the verifier key, unverified AI levels and model-written lessons are no longer served; the bank lesson is served instead.

01-M35. A 10% holdout gets the shared plan

- **Where:** `services/block-levels.js:29-35`, `services/block-levels.js:1909-1913`, `services/block-levels.js:2226-2230`.
- **What:** among tailored learners, 10% are served the untailored shared plan, so the effect of tailoring can be measured.
- **Evidence:**
  - kohavi2009 [147]. *Setting:* online controlled experiments in industry.
    - "Controlled experiments embody the best scientific design for establishing a causal relationship between changes and their influence on user-observable behavior."
  - doroudi2017 [79].
    - "existing methods for doing so can overestimate the performance of new teaching methods"
  - cheung2016 [148]. *Setting:* 645 education studies.
    - "Effect sizes were roughly twice as large for published articles, small-scale trials, and experimenter-made measures"
  - lortieforgues2019 [149]. *Setting:* 141 large educational RCTs.
    - "These sat within relatively large confidence intervals (mean width = 0.30 SDs), which meant that the results were often uninformative"
- **Verdict:** strong as method.
- **Match:** yes. lortieforgues2019 [149] found that even large educational trials often gave wide, uninformative confidence intervals (mean width 0.30 SDs), so the holdout comparison needs enough learners before it is read.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-F. Checkpoints (`services/block-checkpoints.js`, `services/sessions.js`)

01-M36. A fixed blueprint; tailoring changes only how questions are presented

- **Where:** `services/block-checkpoints.js:1-28`, `services/block-checkpoints.js:57-80`.
- **What:**
  - The skill and difficulty of each slot come from the course and are the same for everyone.
  - Tailoring picks the format and the style: plain, scenario, interest, data, error-spot or compare.
  - No hints, no easier wording and no fewer steps: "the style changes the surface, never the amount of thinking".
- **Evidence:**
  - aera2014 [150]. The *Standards for Educational and Psychological Testing* (AERA, APA, NCME, 2014). Title-level only.
  - messick1995 [151]. *Setting:* validity theory.
    - "The new unified concept of validity interrelates these issues as fundamental aspects of a more comprehensive theory of construct validity"

    - kane2013 [105].
        - "To validate an interpretation or use of test scores is to evaluate the plausibility of the claims based on the scores."
    - walkington2013 [152]. *Setting:* 9th-grade algebra in Cognitive Tutor.
        - "students in the personalization condition solved problems faster and more accurately within the modified unit"
- **Verdict:** moderate.
    - Standardising what is measured is the core fairness requirement.
    - walkington2013 [152] shows interest contexts can change performance, which suggests an "interest" style may affect performance even when the underlying problems are matched.
- **Match:** partly. Whether each style changes difficulty is asserted in the prompt but never measured. That is a validity question the item statistics could answer (P4).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M37. Checkpoint sampling covers every introduced skill, on a difficulty ramp

- **Where:** `services/sessions.js:179-192` (`planCheckpoint`).
- **What:**
    - Skills are drawn round-robin from `introduced`, or from all of the grade's skills for the Final.
    - Difficulty ramps across ±1 around the level, or ±2 for the Final.
    - The fill rule asks for the lessons taught to be covered "in proportion — none left out" (`services/block-checkpoints.js:68-80`).
- **Evidence:**
    - khanna2013 [91].
        - "Students completing courses with cumulative finals retained more than students who took noncumulative finals."
    - szpunar2007 [98].
        - "subjects benefited more from initial testing when they also knew they would need to remember the information for ..."
    - chan2006 [153]. *Setting:* college lab.
        - "testing enhanced recall 24 hr later for the initially nontested material."
    - porter2002 [119].
        - "tools for measuring the content of instruction, the content of instructional materials, and the alignment between these"
- **Verdict:** strong for cumulative, content-aligned checkpoints.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M38. Checkpoint formats and styles are varied

- **Where:** `services/block-checkpoints.js:35` (`STYLES`), `services/block-checkpoints.js:57-80`.
- **What:**
  - Formats mix typed answers and choices. Typed answers are favoured for learners who recognise more than they recall.
  - The six styles are spread across the quiz, and the multiple-choice answer position is varied.
- **Evidence:**
  - butler2017 [114].
    - "variability during retrieval practice produced superior transfer of knowledge to new examples"
  - mcdaniel2007 [154]. *Setting:* college course quizzes.
    - "short answer quizzes produced more robust benefits than multiple choice quizzes"
  - mcdermott2014 [155]. *Setting:* middle and high school classes.
    - "short-answer and multiple-choice classroom quizzing could enhance retention of information on classroom exams"
  - haladyna2002 [132]. Key position is among the 31 guidelines.
    - "A taxonomy of 31 multiple-choice item-writing guidelines was validated through a logical process"
  - mclaren2015 [130]. Supports the error-spot style.
    - "the erroneous examples group performed significantly better on a delayed test"
- **Verdict:** moderate.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M39. Every retry uses freshly written questions, with no cooldown

- **Where:** `config/curriculum.js:595-605`, `services/sessions.js:194-251` (`examPromptsSeen`, `REGEN_WAIT_MS 20000`, `prefetchRetry`).
- **What:**
  - A retry is the same level and topics with new questions, and every exam prompt already seen at that level is avoided.
  - The old cooldown was removed, because regenerating takes away the grinding it guarded against.
- **Evidence:**
  - hausknecht2007 [115].
    - "effects were larger when practice was accompanied by test coaching and when identical forms were used"
  - guskey2007 [60]. Bloom's model of varying instruction and learning time to match students' needs.

        - "Bloom labeled the strategy to accomplish this instructional variation and differentiation mastery learning."
    - calamia2012 [586]. *Setting:* meta-analysis of repeated neuropsychological testing.
        - "practice effects, defined as score increases due to factors such as memory for specific test items, learned strategies, or test sophistication"
        - Scope: Defines practice effects as including memory for specific items, and reports that the use of alternate forms was associated with the size of the change in many cases. It did not test freshly written questions on a retry.
    - lievens2007 [585]. *Setting:* 941 candidates retaking a cognitive ability test in a high-stakes selection setting.
        - "retesting induced both measurement and predictive bias"
        - Scope: Shows that retesting the same measure biases scores. It did not test freshly written questions.
- **Verdict:** strong that parallel (fresh) forms avoid the inflation from repeated identical forms.
- **Match:** yes.
    - Gap: Bloom's model puts correctives *before* the second test. Here the learner may retry at once.
    - That is the owner's product choice. The evidence favours nudging toward "Help me learn" after a low score (P7).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-G. Session planning (`services/sessions.js` planLesson, `services/learner.js` sessionParams, `content/pool.js`)

01-M40. Session length: 6 to 18 items from the profile

- **Where:** `services/learner.js:274-305`.
- **What:**
    - Length is `6 + 6·density + 6·pacing`, clamped to 5-18.
    - The scorecard can stretch it only through the insertion contract's calibrated gate.
- **Evidence:**
    - lyle2020 [157]. *Setting:* college precalculus.
        - "although some evidence suggested that the effect of amount of practice was less robust than the effect of spacing."
    - rawson2011 [24]. *Setting:* 533 college students.
        - "Across 3 experiments, 533 students learned conceptual material via retrieval practice with restudy."
    - debeer2014 [53].
        - "A decrease in examinee effort during the PISA reading assessment was found consistently across countries"
    - rey2019 [137].

        - "a significant segmenting effect with small to medium effects for retention and transfer performance"
- **Verdict:** design choice.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M41. About 30% due review, first, at d−1

- **Where:** `services/sessions.js:131-141`, `services/sessions.js:164-167`.
- **What:**
    - Up to `round(0.3 × count)` due reviews from introduced skills, at the level's difficulty minus 1.
    - They are placed before new material: "retrieval before new material".
- **Evidence:**
    - pastotter2014 [158]. *Setting:* review of lab studies. Title-level phrase: "Retrieval practice enhances new learning: the forward effect of testing".
    - lindsey2014 [27]. *Setting:* 8th-grade Spanish course.
        - "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)"
    - latimier2021 [159]. *Setting:* meta-analysis of 29 studies.
        - "We carried out meta-analyses on 29 studies investigating the benefit of spacing out retrieval practice episodes on final retention."
    - roediger2011 [160]. *Setting:* review.
        - "retrieval practice is actually a powerful mnemonic enhancer, often producing large gains in long-term retention relative to repeated studying"
- **Verdict:** strong.
- **Match:** yes. The full retention schedule is in `RETENTION-ENFORCER-EVIDENCE.md`.
- **Status: Built on 25 October 2026** (switch `CURVE_REVIEWS`, from 02-P4). Where the grade has an objective model, the 30% review share comes from the forgetting model's review queue. Otherwise SM-2 decides, as before.

01-M42. Support and stretch shares

- **Where:** `services/learner.js:284-299`, `services/sessions.js:143-162`.
- **What:**
    - Support is 0.1 of the session when challenge appetite is above 0.6, otherwise 0.2. Support items sit at d−1.5 on the weakest introduced skill.
    - Stretch is `0.1 + 0.2·appetite` at d+1.5 on the focus skill.
    - The scorecard can move each share, clamped to 0-0.35 for support and 0.05-0.35 for stretch.
- **Evidence:**
    - wilson2019 [94].
        - "the optimal training accuracy is about 85%"

    - tetzlaff2025 [61] (supports adaptive assistance by prior knowledge).
        - "Low prior knowledge learners learn better from high-assistance instruction (d = 0.505); high prior knowledge learners learn better from low-assistance instruction (d = −0.428)."
    - belland2017 [161]. *Setting:* meta-analysis, computer-based scaffolding in STEM.
        - "computer-based scaffolding showed a consistently positive ($\bar{g}$ = 0.46) effect on cognitive outcomes"
    - vandepol2010 [162]. *Setting:* review of scaffolding characteristics, not item-difficulty spacing.
        - "contingency, fading, and transfer of responsibility are discerned in this review as the three key characteristics of scaffolding"
- **Verdict:** moderate for the principle: easier items for the weak, harder ones for the ready.
- **Match:** partly. The shares are set from a profile, not from the learner's realised accuracy, so nothing checks that focus items land near the 80-85% band. See P9.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 01-P9, tuning the shares to realised accuracy (needs data).

01-M43. Interleaving within the focus items, and the order of roles

- **Where:** `services/sessions.js:153-167`.
- **What:**
    - Every third focus item comes from the next weakest skill.
    - The final order is review, then support, then focus, then stretch. The body is not shuffled.
- **Evidence:**
    - rohrer2015 [12]. *Setting:* 7th-grade mathematics.
        - "Cohen’s ds 0.42 and 0.79, respectively"
    - rohrer2020 [13].
        - "the interleaved group outscored the blocked group, 61% versus 38%, d = 0.83"
    - firth2021 [163]. *Setting:* systematic review.
        - "Memory... showed an interleaving benefit with effect sizes (Hedges’ g) of up to 0.65"
    - brunmair2019 [14].
        - "should be used with caution in certain conditions, especially for expository texts and words"
    - tauber2013 [164]. *Setting:* learner preferences.
        - "Participants overwhelmingly preferred to block their study, even though interleaving is normatively better for learning."
- **Verdict:** strong for interleaving in mathematics.
- **Match:** partly.
    - The comment says "an interleaved body", but the body is grouped by role: all support items, then the focus run, then all stretch items.

    - Only one focus item in three switches skill. See P1.
- **Status: Built on 25 October 2026** (switch `INTERLEAVE_SESSIONS`, from 01-P1). From the second practice level after a lesson the body is spread across skills; from the third, one focus item in two is an earlier topic. Spanish and Reading are unchanged.

01-M44. Question families, shape rotation and no-repeat history

- **Where:**
    - `content/pool.js:1-30` (families).
    - `content/pool.js:94-135` (`pickByDifficulty`: spent shapes are filtered out).
    - `content/pool.js:356-400` (`SHAPE_ROTATION_LIMIT 12`: the base generator and a shape generator alternate).
    - `services/sessions.js:81-125`.
- **What:**
    - One template with slot banks yields hundreds of questions. The code notes a 94% repeat rate on fixed lists before families.
    - A question's *shape* is tracked, so a spent shape is not served again with new nouns.
- **Evidence:**
    - butler2017 [114].
        - "variability during retrieval practice produced superior transfer of knowledge to new examples"
    - paas1994 [165]. *Setting:* geometry training, variability.
        - "Four computer-based training strategies for geometrical problem solving... were studied with regard to their effects on training performance, transfer performance, and cognitive load"
    - pan2018 [113].
        - "testing can yield transferrable learning"
- **Verdict:** moderate.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M45. Typed answers preferred for recall-leaning learners

- **Where:** `services/learner.js:287` (`retrieval_style > 0.55`), `services/sessions.js:81-107`.
- **What:** for those learners, the generator tries again for a format other than multiple choice.
- **Evidence:**
    - rowland2014 [10].
        - "initial recall tests yielding larger testing benefits than recognition tests"
    - mcdaniel2007 [154].

        - "short answer quizzes produced more robust benefits than multiple choice quizzes"
    - kang2007 [166]. *Setting:* college, no feedback.
        - "having had an intervening MC test led to better performance than an intervening SA test"
    - smith2014 [116].
        - "Practising retrieval in all format conditions enhanced retention, relative to a study-only control condition"
    - little2012 [133].
        - "multiple-choice tests could trigger productive retrieval processes-provided the alternatives were made plausible enough"
- **Verdict:** mixed.
    - Recall formats often help more.
    - With no feedback, multiple choice beat short answer (kang2007 [166]), and format differences are often small (smith2014 [116]).
- **Match:** yes. The preference is soft, and feedback is always given.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M46. One recall question at the end of each level, from what was missed

- **Where:** `services/sessions.js:1095-1126` (`blockRetentionItem`).
- **What:** one item on a skill missed in the level (or any skill, if none was missed), at the level's average difficulty minus 0.5.
- **Evidence:**
    - pashler2005 [122].
        - "Supplying the correct answer after an incorrect response ... also increased final retention by 494%."
    - pyc2007 [167]. *Setting:* college lab.
        - "it was particularly effective for learning initially incorrect items."
    - vaughn2011 [168]. *Setting:* college lab.
        - "increasing the criterion level enhanced associative memory"
    - metcalfe2011 [169] (especially for high-confidence errors). *Setting:* adults, general information.
        - "errors that are committed with high confidence are easier to correct than low-confidence errors"
- **Verdict:** moderate.
- **Match:** yes. It comes minutes after the miss: a short lag, which the spacing evidence (M47) rates below a delayed retest.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M47. The end-of-session recall set

- **Where:** `services/sessions.js:1510-1558` (`retentionPlan`).

- **What:**
  - One question per level covered, oldest level first, taken from what was missed there.
  - Topped up to 3 from the weakest remaining skills, at d−0.5.
- **Evidence:**
  - carpenter2009hist [170]. *Setting:* 8th-grade US history.
    - "Facts reviewed through testing were retained significantly better than facts reviewed through re-studying"
  - agarwal2021 [99].
    - "the majority of which (57%) revealed medium or large benefits from retrieval practice"
  - cepeda2006 [22]. *Setting:* 839 assessments in 317 experiments.
    - "This review found 839 assessments of distributed practice in 317 experiments located in 184 articles."
  - cepeda2008 [23].
    - "the optimal gap declined from about 20 to 40% of a 1-week test delay"
- **Verdict:** moderate.
- **Match:** partly. Within-session recall is a short gap. Its value is diagnostic plus a small retention gain. The larger gain needs a delay of a day or more, which the retention enforcers and due reviews supply. See P5.
- **Status: Built on 25 October 2026** (switch `ENFORCER_DEFER_SITTING_RETENTION`, from 02-P3). While a study session is running, the end-of-sitting recall set is not asked; its missed skills go to the session enforcer instead.

01-M48. Next-session difficulty moves ±0.5

- **Where:** `services/sessions.js:1500-1504` (rule-based analysis fallback).
- **What:** next session's difficulty goes up 0.5 at 85% accuracy or better, down 0.5 at 50% or worse, and is otherwise unchanged.
- **Evidence:**
  - wilson2019 [94].
    - "the optimal training accuracy is about 85%"
  - wauters2010 [55].
    - "adaptive item sequencing by matching the difficulty of the items to the learner's knowledge level"
  - mettler2016 [171]. *Setting:* adults, adaptive spacing.
    - "it may be better optimized by adaptive adjustments of spacing to learners' ongoing performance"
  - major2021 [87].
    - "more personalised approaches which adapt or adjust to learners' level led to significantly greater impact"
- **Verdict:** design choice.
- **Match:** yes.

- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 01-P9, closing the loop on challenge (needs data).

01-H. Mastery and advancement (`services/mastery.js`)

01-M49. Mastery is an exponential moving average, weighted by difficulty

- **Where:** `services/mastery.js:52-92`.
- **What:**
  - The learning rate is 0.35 for the first 3 attempts, then 0.15.
  - It is multiplied by `1 + 0.15·max(0, gap)`: credit for beating harder items, extra penalty for missing easier ones.
  - The rate is capped at 0.6.
  - An IRT/Elo estimate is stored beside it but does not drive decisions.
- **Evidence:**
  - pelanek2017mastery [172]. *Setting:* simulated and real data.
    - "We argue that a simple exponential moving average method is a suitable technique for mastery criterion"
  - pelanek2017mastery [172], same paper.
    - "the choice of data sources used for mastery decision and setting of thresholds are more important than the choice of a learner modeling technique"
  - pelanek2016 [77]. Elo in adaptive education. Title-level only.
  - koedinger2023 [89].
    - "typically increasing by about 0.1 log odds or 2.5% in accuracy per opportunity"
  - vandesande2013 [81].
    - "has the functional form of an exponential"
- **Verdict:** moderate.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M50. Coach-assisted answers count for less

- **Where:**
  - `services/mastery.js:87-90`: a correct answer the coach walked them to moves mastery half as far.
  - `services/mastery.js:675-680`: it earns 60% of the XP.
  - `services/coverage.js:103`: it carries 0.25 weight as evidence.
- **Evidence:**
  - aleven2003 [173]. *Setting:* review of help seeking.

        - "although effective help-seeking behavior in ILEs is related to better learning outcomes, learners are not using help facilities effectively"
    - aleven2016help [34]. *Setting:* ITS help-seeking research.
        - "feedback on help seeking helped students to use on-demand help more deliberately... but not to achieve better learning outcomes"
    - baker2008 [71]. Contextual guess estimates are what this discount approximates.
        - "using machine learning to make contextual estimations of the probability that a student has guessed or slipped"
    - wood1976 [174]. *Setting:* tutoring of young children. Title-level only.
- **Verdict:** design choice.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M51. Review intervals and ease

- **Where:** `services/mastery.js:15-18`, `services/mastery.js:94-102`.
- **What:**
    - Correct streaks step through 1, 3, 7, 16, 35, 70 and 140 days, scaled by ease/2.5.
    - Ease is clamped to 1.3-3.0, rising 0.05 per right answer and falling 0.15 per miss.
    - A miss returns in half a day. A hint halves the interval.
- **Evidence:**
    - kang2014 [175]. *Setting:* lab.
        - "However, the expanding schedule yielded much higher average recallability over the whole training period."
    - cepeda2008 [23].
        - "the optimal gap declined from about 20 to 40% of a 1-week test delay to about 5 to 10% of a 1-year test delay"
    - settles2016 [25]. *Setting:* Duolingo logs.
        - "reducing error by 45%+ compared to several baselines at predicting student recall rates"
    - tabibian2019 [26]. *Setting:* online flashcard users.
        - "current spaced repetition algorithms are simple rule-based heuristics with a few hard-coded parameters"
- **Verdict:** moderate.
- **Match:** yes as a heuristic. Fitted models beat hand-set intervals. `RETENTION-ENFORCER-EVIDENCE.md` covers the schedule in full.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M52. Due reviews: weakest and most overdue first

- **Where:** `services/mastery.js:143-151`.
- **What:** among skills that are due, the lowest mastery goes first, with ties broken by the earliest due date.

- **Evidence:**
  - lindsey2014 [27].
    - "personalized review yielded a 16.5% boost in course retention"
  - eglington2020 [28]. *Setting:* experiment plus simulation.
    - "up to 40% more items recalled in conditions where practice was scheduled optimally according to the model"
  - pyc2007 [167].
    - "it was particularly effective for learning initially incorrect items."
  - mettler2016 [171].
    - "it may be better optimized by adaptive adjustments of spacing to learners' ongoing performance"
- **Verdict:** moderate.
- **Match:** yes.
- **Status: Built on 25 October 2026** (switch `CURVE_REVIEWS`, from 02-P4). Where an objective model exists, due reviews come from the forgetting model; SM-2 remains the fallback.

01-M53. The level-up gate: average 0.8 or better, weakest 0.6 or better, 12 or more items

- **Where:** `services/mastery.js:20-24`, `services/mastery.js:167-199`, `services/mastery.js:275-290`.
- **What:**
  - Practice levels advance only when all three conditions hold over the introduced skills.
  - Lessons advance on completion (M28). Checkpoints advance on a pass (M21).
- **Evidence:**
  - bloom1984 [5]. *Setting:* review of the mastery learning and tutoring studies.
    - "Mastery Learning. Students learn the subject matter in a class with about 30 students per teacher."
  - kulik1990 [4].
    - "mastery learning programs have positive effects on the examination performance of students"
  - pelanek2017mastery [172].
    - "setting of thresholds are more important than the choice of a learner modeling technique"
  - doroudi2019fair [83].
    - "substantially more equitable than giving all students the same amount of practice"
- **Verdict:** moderate. Mastery gating is supported. The 0.8, 0.6 and 12 values are design choices.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-I. Grade-fit (`config/grade-skills.js`)

01-M54. Which skills fit which grade, and re-tagging by rule

- **Where:** `config/grade-skills.js:1-60` (FIT for math), `config/grade-skills.js:106-200` (ENGLISH, READING, SPANISH and BASKETBALL FIT, RULES and MODE), `config/grade-skills.js:205-290` (`remap`).
- **What:**
  - Each skill has a range of grade positions where it is neither above nor far below the grade. A skill outside the range is never served.
  - Lesson titles are re-tagged by keyword rules that must also fit the grade. The grade's skill list is rebuilt from its rows.
- **Evidence:**
  - ccss_8_ee [176]. Standard. Slope is addressed in grade 8, as the FIT comment cites.
    - "Graph proportional relationships, interpreting the unit rate as the slope of the graph."
  - ccss_3_nbt [177]. Standard. Rounding is grade 3.
    - "Use place value understanding to round whole numbers to the nearest 10 or 100."
  - porter2002 [119].
    - "The validity of data produced through use of these tools is found to be quite good."
  - anderson1972 [139]. Tests must state what they count.
    - "the primitive first requirement for a system of measurement, namely that there is a clear and consistent definition of the things being counted"
- **Verdict:** strong. Content is aligned to an official standard.
- **Match:** yes for math, English and Reading, where the comments cite the standards. IQ, ACT, chess and piano have no FIT, by design. Cross-reference `GRADE-FIT-STANDARDS.md` when it lands.
- **Status: Built on 25 October 2026** (switch `GRADE_SKILLS`, from the standards check). Ranges were checked against CCSS and the CEFR; 10 of 90 were changed.

01-M55. Difficulty caps per grade and skill, and the school grade passed to the bank

- **Where:** `config/grade-skills.js:293-341` (`DIFFICULTY_RANGE`), `config/grade-skills.js:343-394` (`capDifficulty`, `SCHOOL`, `fitEntry`).
- **What:**
  - Inside a skill, each grade gets a difficulty range. For example, Kindergarten `facts_within_20` is capped at 1-2.19, because sums to 20 are grade 1.
  - Every plan entry carries the US school grade, so the bank serves that grade's standard.

    - The planners, lessons, retention items and checkpoints all apply this (`services/sessions.js:170-177`, `services/lessons.js:130-133`).
- **Evidence:**
    - ccss_k_oa [178]. Standard.
        - "add and subtract within 10, e.g., by using objects or drawings to represent the problem."
    - ccss_k_cc [179]. Standard.
        - "Write numbers from 0 to 20."
    - ccss_3_nbt [177] (URL as in M54).
        - "Use place value understanding to round whole numbers to the nearest 10 or 100."
    - porter2002 [119].
        - "the alignment between these"
- **Verdict:** strong (standards-based).
- **Match:** yes.
- **Status: Built on 25 October 2026** (switch `GRADE_SKILLS`, from the standards check). Every math, English and Reading draw from preschool to grade 12 carries its school grade; reviews, drills and enforcer questions now also go through the cap.

01-J. Objective coverage (`services/coverage.js`)

01-M56. Objectives come from each lesson's own "covers" sentence

- **Where:** `services/coverage.js:1-20`, `services/coverage.js:135-240` (objective building and `levelMap`).
- **What:**
    - Each teaching lesson is one objective: its title, the author's "covers" sentence, its skill and the levels that teach and practise it.
    - Answers at a shared level (an entry check, a summary or a checkpoint) count for several objectives at once.
- **Evidence:**
    - porter2002 [119].
        - "This article describes tools for measuring the content of instruction"
    - anderson1972 [139].
        - "to propose the distinctions necessary to determine whether a person has comprehended an instructional communication"
    - messick1995 [151].
        - "a more comprehensive theory of construct validity"
    - kane2013 [105].
        - "to evaluate the plausibility of the claims based on the scores"
- **Verdict:** moderate.
- **Match:** yes.

- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M57. Knowledge tracing per objective, and what counts as proof

- **Where:** `services/coverage.js:20-35`, `services/coverage.js:84-105`, `services/coverage.js:360-370`.
- **What:**
  - Each objective is traced with BKT using these parameters:

Parameter	Value
PRIOR	0.2
LEARN	0.15
SLIP	0.1
Guess (choices)	1/n + 0.1
Guess (typed)	0.05 + 0.05

- Proof needs a belief of 0.95 or more, at least 3 units of unaided evidence, 3 right answers, and 2 right answers to questions written for that objective.
- An answer shared by k objectives counts 1/k. A coach-led answer counts 0.25. Rapid guesses are not evidence.

- **Evidence:**
  - corbett1995 [82]. Title-level only.
  - vandesande2013 [81].
    - "update the conditional probability that the student has learned that skill"
  - baker2008 [71].
    - "The model uses four parameters per skill, fit using student performance data"
  - wise2006 [35]. *Setting:* effort-moderated IRT.
    - "rapid guessing (i.e., reflecting low examinee effort)"
  - wise2017 [59].
    - "very short response time—termed rapid guessing—has been shown to indicate disengaged test taking"
- **Verdict:** strong for the method. The parameters are design choices: they are set by hand, not fitted from performance data.
- **Match:** yes. The "aligned answers" requirement is a sound guard against proving the skill rather than the objective.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M58. Retained or lapsed: a delayed, unaided check

- **Where:** `services/coverage.js:36-40`, `services/coverage.js:96-101`, `services/coverage.js:455-495`.
- **What:**

    - An objective is retained when it was proven, and then an unaided answer at least 1 day later came out right often enough: R of 0.9 or more, starting from a prior of 0.7.
    - It lapses when R falls to 0.4 or below. L is then capped at 0.5, and it must be proven again.
- **Evidence:**
    - cepeda2006 [22].
        - "This review found 839 assessments of distributed practice in 317 experiments"
    - janes2020 [180]. *Setting:* college biopsychology.
        - "successive relearning boosted students' learning of course content by at least 10% (with d s ranging from 0.54 to 1.10)"
    - higham2022 [181]. *Setting:* college course.
        - "Recall Of Course Material At The End Of The Semester Was Better For Relearning Compared With Restudying."
    - rawson2018 [182]. *Setting:* college.
        - "the advantage of successive relearning over single-session learning (i.e., relearning potency) was substantial"
- **Verdict:** moderate.
- **Match:** partly. A 1-day delay is the minimum that counts as a memory check. cepeda2008 [23] shows that the optimal gap generally lengthens for longer retention delays, so "retained" is a fairly weak claim. See P5.
- **Status: Built on 25 October 2026** (switch `RETAIN_REAL_DELAY`, from 01-P5). "Retained" needs a correct unaided check at least 7 days after proof (was 1).

01-M59. The coverage gate holds graduation until every objective is retained

- **Where:** `services/coverage.js:40-48`, `services/coverage.js:107-112` (`COVERAGE_GATE` defaults to `enforce` since 2026-10-24), `services/mastery.js:300-312`.
- **What:** passing the Final does not graduate a learner while objectives are unproven or lapsed. The record notes the hold.
- **Evidence:**
    - bloom1984 [5]. Title-level phrase only.
    - guskey2007 [60].
        - "Bloom believed that all students could be helped to reach a high criterion of learning"
    - kulik1990 [4].
        - "mastery learning programs have positive effects on the examination performance of students"
    - guskey1988 [183]. *Setting:* meta-analysis of 46 studies of group-based applications of mastery learning.
        - "Results show that such applications yield consistently positive effects on both cognitive and affective student learning outcomes"

    - keller1968 [100]. Unit-mastery requirement. Title-level only.
- **Verdict:** moderate.
- **Match:** yes.
- **Status: Built on 25 October 2026** (switch `RETAIN_REAL_DELAY`, from 01-P5). The gate now waits on the 7-day retained check. It never holds the owner.

01-K. Remediation after a failed checkpoint (`services/remediation.js`)

01-M60. Two doors: retry now, or "Help me learn"

- **Where:** `services/remediation.js:1-17`, `services/sessions.js:1220-1292` (`resolveCheckpoint`).
- **What:** below the mark, the learner stays on the checkpoint and chooses between a fresh retry and taught remediation. The retry stays available throughout.
- **Evidence:**
    - guskey2007 [60].
        - "Bloom argued that educators and teachers must increase variation in instructional approaches and learning time"
    - kulik1990 [4].
        - "mastery learning programs have positive effects on the examination performance of students"
    - aleven2003 [173]. Learners under-use help.
        - "learners are not using help facilities effectively"
- **Verdict:** moderate.
- **Match:** partly. Learners pick the faster door and under-use help, so an unguided choice tends toward a retry without the correctives. See P7.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 01-P7, "Help me learn" as the default door (UI; not built).

01-M61. Diagnosis from the actual wrong answers

- **Where:** `services/remediation.js:40-87` (`areasFrom`), `services/remediation.js:89-104` (`fallbackTeaching`).
- **What:**
    - Missed items are grouped by skill, worst accuracy first. Each group keeps the prompt, the right answer, the learner's answer and the error tag, plus the most common tag.
    - The tutor teaches from those answers: "they factored correctly and then dropped the sign".
- **Evidence:**
    - brown1978 [68].
        - "explaining why a student is making a mistake as opposed to simply identifying the mistake"

    - vanlehn1990 [88].
        - "they often develop small, local misconceptions that cause systematic errors."
    - vanderkleij2015 [15]. *Setting:* meta-analysis of computer-based feedback.
        - "elaborated feedback... produced larger effect sizes (0.49) than feedback regarding the correctness of the answer (KR; 0.05)"
    - butler2013 [17]. *Setting:* experimental study.
        - "explanation feedback produced superior performance on the new inference questions."
- **Verdict:** strong. Error-specific, elaborated feedback beats correctness-only feedback.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M62. Five drill questions per area; an area holds at the pass mark

- **Where:** `services/remediation.js:22-25` (`DRILL_COUNT 5`), `services/remediation.js:301-317` (`recordDrill`, where `held` means accuracy at or above the mark).
- **What:** holding every area does not advance the learner. Only passing the quiz does.
- **Evidence:**
    - rawson2011 [24].
        - "Across 3 experiments, 533 students learned conceptual material via retrieval practice with restudy."
    - vaughn2011 [168].
        - "increasing the criterion level enhanced associative memory"
    - janes2020 [180].
        - "successive relearning boosted students' learning of course content by at least 10%"
    - pyc2007 [167].
        - "particularly effective for learning initially incorrect items."
- **Verdict:** moderate. Practice to a criterion is supported. Five items at 75% is a reasoned size.
- **Match:** yes.
    - Gap: one round of 5 at 75% means 4 of 5 right, which is a coarse signal (compare the binomial note in M21).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M63. Diagnosis starts as soon as the quiz is lost

- **Where:** `services/remediation.js:106-170` (`startEarly`).
- **What:** once the best possible score has fallen below the mark, the diagnosis starts while the learner is still answering, so the teaching is ready when they click.
- **Evidence:**

    - shute2008 [16]. *Setting:* review of formative feedback.
        - "formative feedback should be nonevaluative, supportive, timely, and specific"
    - kulik1988 [184]. *Setting:* meta-analysis of 53 studies.
        - "Applied studies using actual classroom quizzes and real learning materials have usually found immediate feedback to be more effective than delayed."
    - butler2007 [185]. *Setting:* college.
        - "delayed feedback led to superior final test performance relative to immediate feedback."
    - metcalfe2009 [186]. *Setting:* college students, vocabulary.
        - "there was no difference between the delayed and immediate feedback conditions when the lag to test was controlled"
- **Verdict:** mixed. Classroom studies favour timely feedback. Lab studies often favour a delay.
- **Match:** yes. This is mainly a latency fix that keeps the learner from waiting, with no learning claim of its own.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01-M64. After a failed Final, remediation is planned from the whole record

- **Where:** `services/remediation.js:27-38` (`recordPlanOn`, default on since 2026-10-24).
- **What:** after a missed Final, the plan says what to learn again (objectives unproven, or proven and since lapsed) and what to retrieve (kept, but fading), not just which skills the Final missed.
- **Evidence:**
    - rawson2018 [182].
        - "the advantage of successive relearning over single-session learning (i.e., relearning potency) was substantial"
    - higham2022 [181].
        - "Recall Of Course Material At The End Of The Semester Was Better For Relearning Compared With Restudying."
    - janes2020 [180].
        - "successive relearning boosted students' learning of course content by at least 10%"
    - lindsey2014 [27].
        - "personalized review yielded a 16.5% boost in course retention"
- **Verdict:** moderate.
- **Match:** yes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

01 proposals (each backed by 3 or more sources; status on each)

Each proposal gives its status, its sources, the expected effect, the risk and the files it would touch.

01-P1. Interleave the session body

- **Status: BUILT (flag `INTERLEAVE_SESSIONS`, default on; 2026-10-25, implementer B).**
  - Review still comes first. The first practice level after a lesson stays blocked.
  - From the second, the body is spread evenly across skills (`sessions.spreadBySkill`). From the third, one focus item in two is on an earlier topic.
  - Spanish, Reading and English Reading are unchanged, and the owner gets the base product.
  - Sitting repeats did not rise: 20.6% vs 22.0% at level 13, and 19.3% vs 19.8% at level 44 (10 runs each, on vs off).
  - Tests: `scripts/curriculum-evidence-test.js`.
- **What:** shuffle support, focus and stretch items across skills within a level, with review still first. On the practice runs after a lesson, raise the share of focus items drawn from earlier topics, for example from 1 in 3 to 1 in 2, after the first one or two practice levels.
- **Sources:** rohrer2015 [12], rohrer2020 [13], taylor2010 [93], firth2021 [163], brunmair2019 [14].
- **Expected effect:** test gains in mathematics of d≈0.42 a day later and d≈0.79-0.83 about a month later (rohrer2015 [12], rohrer2020 [13]). The effect may be less reliable for vocabulary and expository material (brunmair2019 [14]).
- **Risk:** sessions feel harder, and learners prefer blocking (tauber2013 [164]). Keep a short block right after a lesson, then interleave. Leave the Spanish and reading text tracks unshuffled, or test them first.
- **Files:** `services/sessions.js` (`planLesson`, lines 153-167); `services/adaptation.js` if its plan builds the order.

01-P2. Fade worked examples, and add self-explanation prompts

- **Status: SKIPPED (2026-10-25).** This would change the lesson format.
  - It needs a new lesson schema and prompt in `services/ai.js`, and a graded blank step in the lesson render in `public*/js/session.js`.
  - Both are outside the implementer's files.
  - The bank fallback alone cannot fade a step without that input.
- **What:**
  - Turn "worked example, then one step harder" into complete example → example with its last step blank → full problem, fading backward.
  - Add a one-line "which rule justifies this step?" prompt to each worked step.
- **Sources:** renkl2002 [108], atkinson2003 [109], salden2010 [110], renkl2003 [106], bisra2018 [187] (g = .55 for inducing self-explanation).
- **Expected effect:** better near transfer from fading (renkl2002 [108]), and medium to large near and far transfer with prompts and fading together, without extra time (atkinson2003 [109]).

- **Risk:** longer lessons. The blank step needs a graded input. Expertise reversal suggests the blank-step support may be less effective, or even counterproductive, for advanced learners (kalyuga2003 [6]).
- **Files:**
  - `services/lessons.js` (`buildFallback`).
  - `services/ai.js` (lesson system prompt, `LESSON_SCHEMA`, `validateLesson`).
  - `public/js/session.js` (lesson render), plus its `public-v2/` and `public-redesign/` copies.
  - `services/block-levels.js` (`LAYOUT_SYSTEM`, `Basic_Steps` usage).

01-P3. Report and control classification error at checkpoints and the Challenge Exam

- **Status: BUILT (flags `CHECKPOINT_SPRT` and `CHALLENGE_SPRT`, default on; 2026-10-25, implementer B). Not built: extending a near-mark checkpoint by a few questions, which needs the quiz client to accept extra questions.**
- **The decision** (`config/curriculum.js` `passDecision`, `skipVerdicts`):
  - Each answer is weighed by Wald's log likelihood ratio between "should pass" and "should not":
    - Master and Grand: 0.85 vs 0.65.
    - Final: 0.90 vs 0.75.
    - Challenge: 0.90 vs 0.70.
  - Error rates: $\alpha = 0.05$ (0.10 for the Challenge) and $\beta = 0.10$.
  - At full length the test reduces to "right ≥ needed". `needed` is where the ratio crosses even: 0.760 of the answers for Master and Grand, 0.834 for the Final, and 0.814 for the Challenge.
  - An exam ended early, or with skipped questions, passes only on decisive evidence (the upper SPRT boundary). 3 of 3 no longer passes a Master.
- **Checkpoints, exact binomial P(pass) by true accuracy:**

	.65	**.70**	**.75**	**.80**	**.85**	**.90**	**.95**
Master 15, 12 needed (as before)	.173	.297	.461	.648	.823	.944	.995
Grand 25, 19 needed (as before)	.173	.341	.561	.780	.930	.991	1.000
Final 40, **34 needed**	.004	.024	.096	.286	.607	**.900**	.997
Final 40, 36 needed (old)	.000	.003	.016	.076	.263	.629	.952

The Master and Grand bands sit at the owner's 75% mark. A 70% learner, just under it, still passes a Master 30% of the time. Only the extension (not built) or a longer quiz can cut that.

- **Challenge Exam:**
  - A gate band of 16 holds at 14 right. Each band of 6 above it holds with at most one miss.

    - The evidence over all bands so far must also favour a pass. This is a guard; with these sizes it never binds.
    - 40 questions (was 35).
    - Gate alone, exact: P(hold) is .789 at 90% (7/7 was .478) and .099 at 70% (was .082).
    - Simulated, 200,000 exams per row:

true accuracy	P(skips anything) new / old	P(all five bands) new / old
.70	.100 / .083	.003 / .000
.80	.352 / .209	.065 / .000
.90	.790 / .479	.488 / .026
.95	.957 / .698	.838 / .168
.95 falling .05 a band	bands 1-5 reached .957 .848 .659 .431 .229 / old .698 .333 .105 .022 .003	

- The pages show "N of M needed" and "held". The rule is stored on each exam (`pass_rule`), so an open exam keeps the old rule.
- The owner keeps five bands of seven, each perfect.

- **Why not a sequential (stop-early) Challenge:** separating 90% from 70% needs about 12-15 answers per decision. The exam's items are all sent up front, and the answer route returns no band state. A curtailed SPRT at 7-16 per band gave P(.9) .80 and P(.7) .12, with a mean of 12.4 questions a band (`scratchpad/sprt.js`).
- **Tests:** `scripts/curriculum-evidence-test.js`. `self-test.js` pins the off state.
- **What:**
    - Use a sequential Bayes or SPRT decision for the Challenge Exam bands instead of "7/7 or fail". Today a learner at 90% clears a band 48% of the time.
    - Use the same decision to extend a checkpoint by a few items when the score sits within the error band of the mark, rather than failing a learner at the true mark about a third of the time, as happens on the Final.
- **Sources:** lewis1990 [18], spray1996 [19], eggen1999 [73], babcock2012 [51].
- **Expected effect:** the same classification accuracy with fewer items (lewis1990 [18]: test length halved), and fewer false fails for learners at the mark.
- **Risk:** variable-length quizzes look unequal to learners. Keep a minimum length and full skill coverage, and explain the extension.
- **Files:**
    - `config/curriculum.js` (`PASS_MARKS`, `CHECKPOINT_LENGTH`, `skipTargetFromBands`).
    - `services/sessions.js` (`planCheckpoint`, `resolveCheckpoint`, `planSkipExam`, `scoreSkipExam`).
    - `services/block-checkpoints.js` (write extra slots).

01-P4. Calibrate difficulty and placement items from live responses

- **Status: SKIPPED, needs data (2026-10-25).** Online Elo or IRT calibration needs answer volume per item, which the app does not have yet.
- **What:**
  - Fit item difficulty online, by Elo or IRT, from real answers. Use it to replace the linear `difficultyForLevel`, the judge-calibrated placement points and the 0.75-of-unit default.
  - Use the same statistics to check that checkpoint styles are difficulty-neutral (M36).
- **Sources:** klinkenberg2011 [78], pelanek2016 [77], wauters2010 [55].
- **Expected effect:** levels pitched nearer the target success rate. Placement error measured on real learners instead of simulated ones.
- **Risk:** this needs volume per item. Cold-start items keep the judged values. Shrink toward the judged value until enough answers arrive.
- **Files:**
  - `config/curriculum.js` (`difficultyForLevel`, `gradeBand`).
  - `services/adaptive-placement.js`, `services/unit-placement.js`.
  - `content/unit-placement-calibration/*.json`, `services/scorecard/irt.js`.

01-P5. Make "kept" mean a real delay

- **Status: BUILT, the "retained" half (flag `RETAIN_REAL_DELAY`, default on; 2026-10-25, implementer B).**
  - `RETAIN_SPAN_DAYS` defaults to 7 (was 1): an objective counts as retained only on a correct unaided check at least 7 days after its proof. Memory checks still start a day after.
  - Implementer C gated `adaptation.reviewQueue` so that, once retention is believed, a proven objective is not re-listed before day 7.
  - 21 days was not used, because graduation waits on "retained" (COVERAGE_GATE=enforce) and consolidation is capped at 21 days.
  - Not built: moving part of the end-of-session recall to the next session. It overlaps the retention section's P3 (the end-of-sitting set), which implementer A owns.
  - Tests: `coverage-test.js`.
- **What:**
  - Move part of the end-of-session recall to the start of the next session (one day or more).
  - Raise `RETAIN_DAYS` from 1 to a schedule that grows with time since proof, for example 1 day, then 7, then 21 or more, before an objective counts as retained at graduation.
- **Sources:** cepeda2006 [22], cepeda2008 [23], rawson2013power [188], carpenter2009hist [170].
- **Expected effect:** larger long-term retention gains than same-session recall. "Retained" becomes an honest claim.

- **Risk:** graduation takes longer. Coordinate with the retention enforcers (`RETENTION-ENFORCER-EVIDENCE.md`) so items are not double-scheduled.
- **Files:** `services/sessions.js` (`retentionPlan`, `createSession`); `services/coverage.js` (`RETAIN_DAYS`, `RETAIN_SPAN_DAYS`); `services/mastery.js` (`dueReviews`).

01-P6. Placement per-skill gaps on a belief, not a raw percentage

- **Status: BUILT (flag `PLACEMENT_SKILL_BELIEF`, default on; 2026-10-25, implementer B).**
  - `placement.score` uses unitBelief's Bayes update from an even prior: a strength needs P(known) ≥ 0.9, a gap P(known) ≤ 0.15, and anything between is a new "unsure" list.
  - One miss is no longer a gap, and two rights are no longer a strength.
  - An unsure skill is seen again the next day.
  - The owner keeps the percentage rule.
  - `routes/placement.js` does not display "unsure" yet.
  - Tests: `scripts/curriculum-evidence-test.js`.
- **What:** replace "under 60% is a gap" with a Bayesian belief using the same machinery as `unitBelief`, with a third outcome, "unsure", that the first sessions resolve.
- **Sources:** parsons2019 [66], morris1983 [63], vandesande2013 [81], baker2008 [71].
- **Expected effect:** fewer false gaps and false strengths seeded from 1-3 items.
- **Risk:** low. It changes only seeding and display.
- **Files:** `services/placement.js` (`score`, `seedSkillStates`), and the placement results display in `routes/placement.js`.

01-P7. After a missed check or a failed checkpoint, nudge toward correctives before the retry

- **Status: SKIPPED (2026-10-25).** Both halves are UI: making "Help me learn" the default door, and re-showing the worked example with the missed check explained. Both live in `public*/js/session.js`, outside the implementer's files, and no server-only version reaches the learner.
- **What:**
  - When a lesson's checks are mostly missed, queue a short targeted re-study, the lesson's worked example re-shown with the missed check explained, before the next practice level.
  - After a failed checkpoint, make "Help me learn" the default door. The retry stays one tap away.
- **Sources:** guskey2007 [60], bloom1984 [5], kulik1990 [4], pashler2005 [122], vanderkleij2015 [15], aleven2003 [173].
- **Expected effect:** the corrective step that defines Bloom's mastery learning gets used, while learners often fail to use optional help effectively (aleven2003 [173]).
- **Risk:** it slows learners who missed by carelessness. Make it skippable and short.

- **Files:** `services/sessions.js` (`resolveBlock`, `resolveCheckpoint`); `services/lessons.js`; `services/remediation.js`; `public/js/session.js`.

01-P8. Allow 3-option multiple choice where a fourth plausible distractor does not exist

- **Status: SKIPPED for practice and block items (2026-10-25).** The gate is `ai.validateItem` in `services/ai.js`, which is outside the implementer's files. Objective items now allow 3 options (03-scorecard-adaptation P11, flag `OBJ_ITEMS_3_OPTIONS`).
- **What:** accept 3 distinct options in `validateItem` and `whyRefused` for practice items, and keep 4 for checkpoints if the owner wants.
- **Sources:** rodriguez2005 [136], haladyna2002 [132]. (Two after the 27 October 2026 audit: little2012 [133] tests whether plausible distractors trigger retrieval, not how many options an item needs.)
- **Expected effect:** fewer implausible filler distractors, fewer repair calls, and more items per minute (rodriguez2005 [136]).
- **Risk:** guessing rises from 25% to 33%, so the coverage and unit-belief guess terms must read the real option count. They already take `choices`.
- **Files:** `services/ai.js` (`validateItem`); `services/block-levels.js` (`whyRefused`, `FILL_SYSTEM`); `services/block-checkpoints.js` (`FILL_SYSTEM`).

01-P9. Close the loop on challenge

- **Status: SKIPPED, needs data (2026-10-25).** Tuning the shares to realised accuracy by role needs logged real sessions.
- **What:** log realised accuracy by role (support, focus, stretch). Move `stretchShare` and `supportShare`, and the next-session difficulty delta, so that focus items run near 80-85%.
- **Sources:** wilson2019 [94], wauters2010 [55], klinkenberg2011 [78] (adaptive difficulty, with items sampled at a mean success probability of 0.75; it does not set the 80–85% target), tetzlaff2025 [61].
- **Expected effect:** time spent where learning per item is highest.
- **Risk:** feedback loops can oscillate. Clamp the step size, as `sessionParams` already clamps the shares.
- **Files:** `services/learner.js` (`sessionParams`); `services/adaptation.js`; `services/sessions.js` (`ruleBasedAnalysis`).

01-P10. Never serve an unverified AI level when the verifier key is missing

- **Status: BUILT (flag `LEVEL_SELF_CHECK`, default on; 2026-10-25, implementer B).**
  - With no `FACT_CHECK_ANTHROPIC_KEY`, a block level is served only after a self-check (`block-levels.selfCheck`):
    - a claim-by-claim read on the app's own route, with repairs;
    - a blind re-solve of every multiple-choice check by a judge from another family (`key-gate.blindResolve`; NVIDIA first, never xKiro or

UnoRouter). A disagreeing check is dropped; if no judge answers, the level fails.
    - Levels verified earlier under the key (Claude) are still served.
    - Everything else falls back to the ordinary session and the bank lesson.
    - Tailored checkpoints keep only the questions the blind judge confirms, and the bank fills the rest.
    - Model-written ordinary lessons (`lessons.js`) are not served without the key.
    - It is a safety rule, so it applies to the owner too.
    - With the key set, nothing changes.
    - Tests: `block-levels-test.js` and `key-gate-test.js`.
    - In the browser on a no-key test copy: an unverified English base plan went to the bank lesson, and a Sonnet-verified chess level was served.
- **What:** when `gateOn()` is false, run a cheaper independent check rather than serving the plan unverified:
    - a self-consistency re-solve of every graded check (the key gate already does this for lessons);
    - a Chain-of-Verification pass on factual claims, on a free model allowed by the owner's model rules;
    - otherwise fall back to the bank lesson.
- **Sources:** dhuliawala2024 [144], manakul2023 [189], wang2022sc [135], min2023 [145], ji2023 [142], bastani2025 [38].
- **Expected effect:** fewer factual and answer-key errors reaching learners. The code records the fallback route catching 2 of 13 known errors.
- **Risk:** latency and free-tier limits. The bank fallback guarantees nothing breaks.
- **Files:** `services/block-levels.js` (`gateOn`, `servable`, `verifyOn`, `verifyAndRepair`); `services/key-gate.js`; `services/providers`.

Section 02. Retention and review

This section covers:

- the retention enforcers and their calendar;
- the daily review and skill-level spaced review;
- consolidation before the Final;
- retrieval inside lessons ("Check what stuck", the per-level and end-of-sitting retention questions);
- feedback timing;
- interleaving across sessions;
- sleep and day boundaries;
- the forgetting model that gives every recall estimate here.

Design choices, stated plainly. No study sets any of these numbers. Each was chosen inside the direction the evidence gives, and each needs real-learner data:

- the 3/4/3/4 semi-weekly, weekly on day 7, bi-weekly on days 12, 26, 40, ..., and the monthly "pocket" day (owner rules);
- full under 10%, applied as 0.095;
- the scope shares (9.5 / 9.5 / 6 / 5 / 4 / 3%) and caps (8 / 10 / 12 / 15 / 15 / 20 min);
- the tier multipliers (0.35 / 0.65), the 20-minute daily cap, and the 0.5 weight for absorbed scopes;
- the score weights, the 0.35 value of a second try, the fragile and mature thresholds, and the lags;
- the 10-minute review, the 90% / 30-day / 21-day consolidation targets, and the 12-hour and 1-day gates.

In numbers: 50 mechanisms; 50 with 3 or more quoted sources; 5 strong, 25 moderate, 5 mixed, 15 design choice; 123 distinct sources cited or listed; 6 mechanisms built or changed on 25 October 2026.

02 at a glance

Mechanism	Verdict	Quoted sources	Status
02-M1 End-of-session enforcer is retrieval with feedback, not review	strong	6	existing
02-M2 End-of-day enforcer, which also takes in the day's unchosen session enforcers	moderate	4	existing
02-M3 Tier menu (quick, medium, full, none) with a preselected default	moderate	5	existing
02-M4 Full tier capped at 9.5% of covered active time; quick and medium at 0.35 and 0.65 of full	design choice	4	existing
02-M5 Per-scope budget shares and caps	design choice	4	existing
02-M6 20-minute daily cap across all enforcers	design choice	3	existing

Mechanism	Verdict	Quoted sources	Status
02-M7 Active-time clock that sizes every budget	design choice	4	existing
02-M8 Semi-weekly enforcer on strict 3/4/3/4-day gaps from the anchor	design choice	4	existing
02-M9 Weekly enforcer every 7 days, absorbing that day's semi-weekly	design choice	4	existing
02-M10 Bi-weekly enforcer on days 12, 26, 40, .	design choice	3	existing
02-M11 Monthly enforcer on a free "pocket" day within 3 days of day 30m	design choice	3	existing
02-M12 Absorption into one combined enforcer, with budget = largest part + 0.5 x the others	design choice	4	existing
02-M13 Cumulative windows, and questions only from the covered window	moderate	5	existing
02-M14 Retention calendar page (done, none, missed, due, nothing to cover)	moderate	4	existing
02-M15 Unchosen and unfinished enforcers are offered again; validity 24 h (session) and 48 h (day); "none" keeps items due	design choice	4	existing
02-M16 Selection by predicted forgetting: need = 1 - p_H at the scope's horizon	strong	5	existing
02-M17 Desirable-difficulty window d(p_now)	moderate	4	built on 25 October 2026
02-M18 Score flags: last answer wrong +0.30, confident error +0.25, right but slow or unsure +0.20, no delayed retrieval yet +0.15	moderate	5	built on 25 October 2026
02-M19 Defaults with no record: p_now 0.5, p_H 0.4; skill-only units p_H = 0.85 x mastery	design choice	3	existing
02-M20 Soft eligibility: minimum lag (5 min session, 30 min day, 1 day scheduled) and "kept in the last 20 hours" go last	moderate	4	existing
02-M21 Greedy fill by value per planned second, with a cost model per format	moderate	3	existing
02-M22 Allocation: one first try per topic; a second, varied try only for fragile topics on medium and full, worth 0.35 of a first	moderate	6	existing
02-M23 Drop-out retries: quick 1 attempt, medium 2, full until right (at most 3), on a different question at least 3 items later	moderate	4	existing
02-M24 Interleave by skill inside the sitting	mixed	4	existing
02-M25 Session scope asks the oldest-exposed topics first	moderate	3	existing
02-M26 Close on a question the curve expects right (p_now >= 0.8)	design choice	4	existing
02-M27 A new question instance each time; a varied second try; never the learner's own check as "varied"	moderate	5	existing

Mechanism	Verdict	Quoted sources	Status
02-M28 A "kept" answer is first try, unaided, not rapid (< 1200 ms) and machine-graded; only answers at a lag of 1 day or more are memory checks	moderate	4	built on 25 October 2026
02-M29 Low stakes: enforcer and review sittings are previews that move no level; "none" is never penalised	moderate	4	existing
02-M30 Daily review queue: memory checks first, then fading objectives, then the least durable; lapsed objectives to re-teach	moderate	4	existing
02-M31 Review target: predicted recall 0.90 now (ADAPT_REVIEW_AT) and 0.90 at 30 days (ADAPT_DURABLE_AT, ADAPT_DURABLE_DAYS)	moderate	4	existing
02-M32 Rested gate: a review waits at least 12 h after the last answer; a memory check comes at least 1 day after proof	moderate	3	existing
02-M33 Review budget: 10 minutes at 0.75 min per question, one question per objective	design choice	3	existing
02-M34 Daily review folded into the enforcers: it keeps only older work no enforcer will reach, and skips objectives kept in an enforcer in the last 20 h	design choice	3	existing
02-M35 Typed recall preferred in the review: "one right typed recall is a whole memory check, a multiple-choice one is half"	mixed	4	existing
02-M36 Lapsed objectives are re-taught as a lesson, at most 3 per session, 6 minutes each	moderate	3	existing
02-M37 Skill-level spaced review (mastery.dueReviews): an SM-2-style ladder	mixed	5	built on 25 October 2026
02-M38 Due reviews go first in a lesson, one difficulty step below the level	moderate	4	built on 25 October 2026
02-M39 Interleaving across sessions: earlier skills come back inside lessons; the focus skill alternates with the next weakest every third item	moderate	4	existing
02-M40 Final held for review-only days until 90% of objectives are durable at 30 days, for at most 21 days (ADAPT_CONSOLIDATE)	design choice	4	existing
02-M41 "Check what stuck": read the lesson, then answer its checks	strong	5	existing
02-M42 Per-level retention question after each level closes	moderate	4	existing
02-M43 End-of-sitting retention set: at most 3 questions, one per level from its misses, oldest level first, half a step easier	moderate	4	built on 25 October 2026
02-M44 Feedback timing: after each answer by default; held to the end of the level only when the learner's style answers ask for it; always immediate in enforcers, reviews and retention items	mixed	7	existing

Mechanism	Verdict	Quoted sources	Status
02-M45 Feedback content: the answer and why, plus "A second look" recap of misses at the end of an enforcer	strong	5	existing
02-M46 Day boundaries: local calendar days; scheduled scopes need a 1-day lag; the day enforcer closes the day; no bedtime gate	moderate	6	existing
02-M47 Half-life regression with a population fit, a guessing floor and capped counts	moderate	4	existing
02-M48 Curve shape: exponential by default; a power-law tail available but off	mixed	3	existing; open proposal
02-M49 Only spaced successes lengthen the half-life	strong	5	existing
02-M50 Personal forgetting multiplier from the learner's own memory checks	moderate	4	existing

02-A. Retention enforcers: what they are and what they cost

02-M1. End-of-session enforcer is retrieval with feedback, not review

- **Where:** services/enforcers.js:701-729 (onSessionEnded), :562-612 (plan), :640 (immediateFeedback: true)
- **What it does:** When a study session ends, the app offers a short sitting of questions on what was worked in that session. Every question is answered unaided and gets feedback at once.
- **Evidence:**
  - **roediger2006** [1]: "on the delayed tests, prior testing produced substantially greater retention than studying" (college students, prose passages, lab)
  - **karpicke2008** [20]: "Repeated studying after learning had no effect on delayed recall, but repeated testing produced a large positive effect." (college students, foreign-vocabulary pairs, lab)
  - **rowland2014** [10]: "initial recall tests yielding larger testing benefits than recognition tests" (meta-analysis of testing vs restudy, mostly adult lab studies)
  - **adesope2017** [11]: "practice tests are more beneficial for learning than restudying and all other comparison conditions" (meta-analysis, school to adult learners)
  - **yang2021** [21]: "overall testing (quizzing) raises student academic achievement to a medium extent (g = 0.499)" (meta-analysis of 222 classroom studies)
  - **lyle2011** [190]: "Exam scores were significantly and substantially higher in a section of the course taught with the PUREMEM procedure" (undergraduate course, end-of-lecture retrieval)
- **Verdict: strong**.
- **Code vs evidence:** Matches. Retrieval rather than rereading, with feedback, at the end of a learning episode.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M2. End-of-day enforcer, which also takes in the day's unchosen session enforcers

- **Where:** services/enforcers.js:704-716; services/retention-params.js:50 (day: 9.5%, 10-min cap, horizon 3.5 d, min lag 30 min)
- **What it does:** "End day session" offers one enforcer for the whole day. Session enforcers still waiting that day are merged into it, so the learner gets one sitting, not several. A topic must be at least 30 minutes past its last exposure to be first in line.
- **Evidence:**
    - **karpicke2011spaced** [191]: "Repeated retrieval with long intervals between each test produced a 200% improvement in long-term retention" (college students, foreign-vocabulary pairs, lab)
    - **karpicke2007** [192]: "Delaying the first test improved long-term retention, regardless of how the repeated tests were spaced." (college students, word pairs, lab)
    - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
    - **cepeda2009** [193]: "confirm the importance of cumulative reviews to promote retention over meaningful time periods" (laboratory studies, foreign vocabulary, facts, and names of visual objects, test delays up to 6 months)
- **Verdict: moderate**. A second retrieval later the same day is spaced retrieval at a short lag. The evidence supports spacing between tests, and delaying the first test. No study tests "30 minutes" or "end of day" as such.
- **Code vs evidence:** Matches the direction. The 30-minute lag and the 3.5-day horizon are design choices.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M3. Tier menu (quick, medium, full, none) with a preselected default

- **Where:** services/enforcers.js:386-425 (options, defaultTier); services/retention-params.js:67; public/js/enforcer.js:79
- **What it does:** The learner picks quick, medium, full or none. Medium is preselected. Full is preselected when 2 or more topics were last answered wrong, or when mean recall at the horizon is under 0.75. Quick is preselected for a session under 15 minutes, or after 3 quick choices in a row. "None" is always there.
- **Evidence:**
    - **jachimowicz2019** [194]: "a considerable influence of defaults (d = 0.68, 95% confidence interval = 0.53-0.83)" (meta-analysis of default studies, many domains (not learning apps))
    - **johnson2003** [195]: "finding large effects that significantly increase donation rates" (national organ-donor consent policies)
    - **patall2008** [46]: "providing choice enhanced intrinsic motivation, effort, task performance, and perceived competence" (meta-analysis, children and adults)
    - **ryan2000** [47]: "three innate psychological needs--competence, autonomy, and relatedness" (theory review)

    - **trumbo2016** [196]: "students were much more likely to take advantage of quizzing if it was a required course component" (large, diverse college course)
- **Verdict: moderate**. Defaults and choice both have meta-analytic support, but in other domains. Optional quizzes are taken much less than required ones, which is why a default is needed. The thresholds (2 wrong, 0.75, 15 min, 3 in a row) are design choices.
- **Code vs evidence:** Matches. The thresholds are unvalidated, and the default-tier experiment in RETENTION-ENFORCER-EVIDENCE.md section 12.3 is the test.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M4. Full tier capped at 9.5% of covered active time; quick and medium at 0.35 and 0.65 of full

- **Where:** services/retention-params.js:27 (FULL_CAP 0.095), :37 (multipliers), :190-196 (budgetMs); services/enforcers.js:370-386 (a tier that cannot fit one question is not offered)
- **What it does:** Budget = tier multiplier x min(scope share x covered active time, scope cap), never above 0.095 x covered time. The owner's rule is "full under 10%"; 0.095 leaves a margin for slow answers.
- **Evidence:**
    - **rohrer2007pashler** [197]: "immediate further study of the same material is an inefficient use of time" (review of spacing and overlearning studies)
    - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
    - **rohrer2006** [198]: "long-term retention was boosted by distributed practise and unaffected by overlearning" (college students, maths problems)
    - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school Spanish course, time-matched)
- **Verdict: design choice**. The principle is supported: short spaced retrieval buys more retention per minute than extra study now, and relearning sessions cost few trials. No study sets 10%, 9.5%, 0.35 or 0.65.
- **Code vs evidence:** Matches the principle. The numbers are owner rules and assumptions.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M5. Per-scope budget shares and caps

- **Where:** services/retention-params.js:48-55
- **What it does:** Full shares and caps are: session 9.5% / 8 min, day 9.5% / 10 min, semi-weekly 6% / 12 min, weekly 5% / 15 min, bi-weekly 4% / 15 min, monthly 3% / 20 min. Shares fall as the window grows, because a wider window holds more study time.
- **Evidence:**

    - **cepeda2009** [193]: "confirm the importance of cumulative reviews to promote retention over meaningful time periods" (adults, verbal materials, gaps and tests up to months)
    - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, conceptual material)
    - **lyle2020** [157]: "Across-semester retention benefited exclusively from increasing spacing." (college precalculus course)
    - **sievertsen2016** [199]: "for every hour later in the day, test performance decreases by 0.9% of an SD" (Danish national tests, children aged 8-15)
- **Verdict: design choice**. The evidence supports several cumulative spaced reviews, spacing over amount, and short sittings (fatigue builds over the day). It does not set any of these percentages or minute caps.
- **Code vs evidence:** No study to match. Tune them from telemetry, as RETENTION-ENFORCER-EVIDENCE.md section 12.1 describes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M6. 20-minute daily cap across all enforcers

- **Where:** services/retention-params.js:47; services/enforcers.js:352-367 (spentToday, budgetFor)
- **What it does:** Each day, enforcer time is capped at 20 minutes in total. A later enforcer gets only what is left of the cap.
- **Evidence:**
    - **sievertsen2016** [199]: "for every hour later in the day, test performance decreases by 0.9% of an SD" (Danish national tests, children aged 8-15)
    - **rohrer2007pashler** [197]: "immediate further study of the same material is an inefficient use of time" (review of spacing and overlearning studies)
    - **lyle2020** [157]: "Across-semester retention benefited exclusively from increasing spacing." (college precalculus course)
- **Verdict: design choice**. The sources support the direction: fatigue, low returns from more same-day practice, spacing over amount. They do not support the number 20.
- **Code vs evidence:** Fine as a guard. Log how often the cap binds.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M7. Active-time clock that sizes every budget

- **Where:** services/study-sessions.js:9-23 (rules), :33-37, :87 (idle_cap_ms 5 min), :111-117 (auto-end); services/retention-params.js:46
- **What it does:** Covered time is active time, not wall-clock time. It counts heartbeats every 30 seconds and every answer. A gap counts at most 5 minutes. A running session ends itself after 30 idle minutes; a paused one ends at the local day change or after 12 hours.
- **Evidence:**

    - **kovanovic2016** [200]: "our findings indicate that time-on-task estimation methods play an important role in shaping the final study results" (two university course log datasets (online and blended))
    - **nguyen2020tot** [201]: "what can be considered as outliers in time-on-task depends on the learning pattern of each student" (university course log data (LAK 2020, peer-reviewed proceedings))
    - **rotelli2022** [202]: "data pre-processing decisions have an impact on the analysis' outcomes and can lead to inaccurate predictions" (real-world student learning-activity data (LAK 2022, peer-reviewed proceedings))
    - **goldhammer2014** [203]: "the relation may be negative if working more fluently, and thus faster, reflects higher skill level" (1,020 adults aged 16-65, PIAAC field test)
- **Verdict: design choice**. Learning-analytics studies show that how time-on-task is estimated, and how idle gaps are trimmed, changes results. They argue for per-student and per-task thresholds rather than one fixed cut-off. No study sets 5 or 30 minutes for this.
- **Code vs evidence:** Partly. One fixed 5-minute idle cap is the simple method those papers criticise; per-learner trimming is possible later. The risk is small, because budgets only scale with this time.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-B. The calendar

02-M8. Semi-weekly enforcer on strict 3/4/3/4-day gaps from the anchor

- **Where:** services/retention-schedule.js:44 (isSemiweekly: day % 7 is 3 or 0), :11-12
- **What it does:** Semi-weekly enforcers fall on days 3, 7, 10, 14, 17, ... from the learner's first study day. Each covers the days since the last one.
- **Evidence:**
    - **cepeda2006** [22]: "the ISI producing maximal retention increased as retention interval increased" (meta-analysis, 839 assessments of verbal recall)
    - **cepeda2008** [23]: "the optimal gap declined from about 20 to 40% of a 1-week test delay to about 5 to 10% of a 1-year test delay" (over 1,350 adults online, facts, tests up to 1 year)
    - **kim2019bigdata** [204]: "the optimal amount of spacing between repeated retrieval events increased as the retention interval increased" (10,514 adults, naturally occurring workplace training data)
    - **cepeda2009** [193]: "An optimal gap improved final recall by up to 150%." (adults, verbal materials, gaps and tests up to months)
- **Verdict: design choice**. The best gap grows with the retention interval: about 20-40% of a 1-week horizon. A 3-4-day gap is 43-57% of a week, a little long, which is the side the curve penalises least: past the optimum, accuracy declines only gradually (cepeda2009 [193]). The exact 3/4/3/4 pattern is the owner's rule.
- **Code vs evidence:** Roughly on the evidence ridge. No study picks 3/4/3/4.

- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M9. Weekly enforcer every 7 days, absorbing that day's semi-weekly

- **Where:** services/retention-schedule.js:45, :98-109 (typesOn, absorbed)
- **What it does:** A weekly enforcer falls on days 7, 14, 21, ... On those days the weekly is the one enforcer, and it covers the week, the semi-week and the day.
- **Evidence:**
  - **cepeda2008** [23]: "the optimal gap declined from about 20 to 40% of a 1-week test delay to about 5 to 10%" (over 1,350 adults online, facts, tests up to 1 year)
  - **latimier2021** [159]: "a strong benefit of spaced retrieval practice in comparison with massed retrieval practice (g = 0.74)" (meta-analysis of spaced retrieval studies)
  - **kang2014** [175]: "the expanding schedule yielded much higher average recallability over the whole training period" (adults, long-term retrieval practice)
  - **karpicke2011spaced** [191]: "Repeated retrieval with long intervals between each test produced a 200% improvement in long-term retention" (college students, foreign-vocabulary pairs, lab)
- **Verdict: design choice**. Spaced retrieval beats massed retrieval (g = 0.74), and long gaps between tests beat short ones. A 7-day gap is about 23% of a 1-month horizon, inside the ridge.
- **Code vs evidence:** Matches the direction. The 7-day period is an owner rule.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M10. Bi-weekly enforcer on days 12, 26, 40, .

- **Where:** services/retention-schedule.js:46 (d % 14 === 12), :16-18
- **What it does:** Bi-weekly enforcers fall every 14 days, offset so there is a free day between them and the nearest semi-weekly and weekly. Each covers the 14 days ending on its day.
- **Evidence:**
  - **cepeda2009** [193]: "confirm the importance of cumulative reviews to promote retention over meaningful time periods" (adults, verbal materials, gaps and tests up to months)
  - **karpicke2011spaced** [191]: "there was no evidence that a particular relative spacing schedule (expanding, equal, or contracting) was inherently superior to another" (college students, foreign-vocabulary pairs, lab)
  - **latimier2021** [159]: "no significant difference between expanding and uniform spacing schedules of retrieval practice (g = 0.034)" (meta-analysis of spaced retrieval studies)
- **Verdict: design choice**. What matters is absolute spacing and cumulative review. Expanding, equal or contracting relative schedules do not differ reliably. So the offset that puts it on day 12 rather than day 14 has no evidential cost, and no evidence for it either.

- **Code vs evidence:** Neutral. The offset is purely an owner rule.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M11. Monthly enforcer on a free "pocket" day within 3 days of day 30m

- **Where:** services/retention-schedule.js:73-96 (monthlyDay, isMonthly), :19-26
- **What it does:** The monthly picks the day within +/-3 of each 30th day that has no other enforcer and whose nearest neighbour is the easiest one. It covers the days since the last monthly, 24 to 36 days.
- **Evidence:**
  - **cepeda2008** [23]: "the optimal gap declined from about 20 to 40% of a 1-week test delay to about 5 to 10%" (over 1,350 adults online, facts, tests up to 1 year)
  - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
  - **kim2019bigdata** [204]: "the optimal amount of spacing between repeated retrieval events increased as the retention interval increased" (10,514 adults, naturally occurring workplace training data)
- **Verdict: design choice**. A 30-day gap is about 8% of a 1-year horizon, inside the 5-10% band. Successive relearning recommends widely spaced relearning sessions. The pocket-day rule is the owner's.
- **Code vs evidence:** Matches the long-horizon ridge. The pocket-day rule is not evidence-based.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M12. Absorption into one combined enforcer, with budget = largest part + 0.5 x the others

- **Where:** services/retention-schedule.js:28-33, :182-188 (combinedWindow); services/retention-params.js:205-218 (combinedBudgetMs), :78 (others_weight 0.5); services/enforcers.js:330-341
- **What it does:** When enforcers share a day, one plan covers the union of their windows, with no per-scope quota. Its full budget is the largest part's budget plus 0.5 x the rest, capped at the largest cap and at 9.5% of the union's time.
- **Evidence:**
  - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
  - **vaughn2016** [205]: "to a higher initial criterion during the first session do not persist across relearning sessions" (college students, word pairs, lab)
  - **rawson2013relearn** [206]: "The benefits of initial spacing were substantial prior to relearning but were significantly attenuated after relearning." (college students, lab)

    - **lyle2020** [157]: "Across-semester retention benefited exclusively from increasing spacing." (college precalculus course)
- **Verdict: design choice**. Evidence supports the direction: one correct retrieval per topic per relearning session is the efficient dose, and extra same-sitting tries fade once relearning follows. The weight 0.5 is an assumption.
- **Code vs evidence:** Matches the direction. The 0.5 is untested.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M13. Cumulative windows, and questions only from the covered window

- **Where:** services/retention-schedule.js:168-180 (windowFor); services/enforcers.js:89-112 (windowItems)
- **What it does:** Each larger scope covers a larger cumulative window: since the last one of its kind, or 14 days for the bi-weekly. No question comes from outside that window; older work stays with the daily review.
- **Evidence:**
    - **mayfield2002** [96]: "the cumulative group outscored the other groups on application and problem-solving items" (college students, maths rules)
    - **khanna2013** [91]: "Students completing courses with cumulative finals retained more than students who took noncumulative finals." (college psychology courses)
    - **lawrence2013** [97]: "low-scoring students remembered more of the course material when they took cumulative exams throughout the semester" (introductory psychology course)
    - **szpunar2007** [98]: "subjects benefited more from initial testing when they also knew they would need to remember the information" (college students, word lists, lab)
    - **cepeda2009** [193]: "confirm the importance of cumulative reviews to promote retention over meaningful time periods" (adults, verbal materials, gaps and tests up to months)
- **Verdict: moderate**. Cumulative practice and cumulative exams improve retention, most for weaker students. Expecting a cumulative test raises what earlier testing does.
- **Code vs evidence:** Matches. The "only from the window" rule is the owner's. Coverage of older work depends on the daily review (mechanism 34).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M14. Retention calendar page (done, none, missed, due, nothing to cover)

- **Where:** services/enforcers.js:1077-1135 (calendar); routes/retention.js:133; public/js/calendar.js
- **What it does:** A month view shows each day's scheduled enforcer, what it absorbs, and how it went.
- **Evidence:**
    - **harkin2016** [207]: "promoted goal attainment (d+ = 0.40, 95% CI [0.32, 0.48])" (meta-analysis, 138 studies, N = 19,951, mostly health goals)

    - **wissman2012** [208]: "do not typically implement or understand the benefits of practising with longer lags (timing of practice)" (survey of 374 undergraduates)
    - **kornell2009cram** [209]: "spacing was more effective than massing for 90% of the participants, yet after the first study session, 72% of the participants believed" (college students, vocabulary flashcards)
    - **bjork2013** [210]: "people often have a faulty mental model of how they learn and remember" (review)
- **Verdict: moderate**. Learners do not schedule long lags themselves, and believe massing works. Monitoring progress toward a goal raises attainment, d+ = 0.40, but mostly in health studies. Physically recording progress made the effect larger.
- **Code vs evidence:** Matches. It is a record, not a goal, so the harkin2016 [207] effect is an analogy.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M15. Unchosen and unfinished enforcers are offered again; validity 24 h (session) and 48 h (day); "none" keeps items due

- **Where:** services/enforcers.js:735-747 (expire), :808-822 (offers), :828 (reopen), :966-990 (decline, skip); services/retention-params.js:68
- **What it does:** An enforcer left unchosen or half done is offered once on return. Unanswered items are re-scored, and answered ones are never re-asked. Choosing "none" marks nothing as reviewed, so the items come back.
- **Evidence:**
    - **trumbo2016** [196]: "students were much more likely to take advantage of quizzing if it was a required course component" (large, diverse college course)
    - **kornell2009spacing** [211]: "being allowed to drop flashcards had small but consistently negative effects on learning" (college students, flashcards)
    - **kornell2009son** [212]: "self-testing was motivated by a desire to diagnose learning rather than a desire to improve learning" (college students)
    - **jachimowicz2019** [194]: "a considerable influence of defaults (d = 0.68, 95% confidence interval = 0.53-0.83)" (meta-analysis of default studies, many domains (not learning apps))
- **Verdict: design choice**. Optional practice is skipped. Letting learners drop items they feel they know hurt learning. Defaults steer choices. The validity windows are assumptions.
- **Code vs evidence:** Matches the direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-C. Which questions an enforcer asks

02-M16. Selection by predicted forgetting: need = 1 - p_H at the scope's horizon

- **Where:** services/enforcers.js:246-300 (score), :264-266 (p_now, p_H from adaptation.predictedRecall); services/retention-params.js:48-55 (horizon_days 1 / 3.5 / 7 / 14 / 30 / 60)
- **What it does:** Each topic is scored by how much it is predicted to have faded at the scope's horizon. That is recall predicted 1 day out for a session enforcer, and up to 60 days out for a monthly one.
- **Evidence:**
  - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school Spanish course, time-matched)
  - **tabibian2019** [26]: "subject to a cost on the reviewing frequency, the optimal reviewing schedule is given by the recall probability itself" (theory plus a natural experiment on Duolingo data)
  - **tabibian2019** [26]: "learners who follow a reviewing schedule determined by our algorithm memorize more effectively than learners who follow alternative schedules" (natural experiment on Duolingo data)
  - **mettler2016** [171]: "adaptive scheduling outperformed fixed conditions at immediate and delayed tests of retention" (adults, lab)
  - **pavlik2008** [213]: "The optimized condition showed significant benefits with large effect sizes for both improved recall and recall latency." (college students, Japanese-English vocabulary)
  - **ye2022** [214]: "Experimental results have shown a 12.6% performance improvement over the state-of-the-art methods." (220 million student memory-behavior logs from an online language-learning app)
- **Verdict: strong**. Converging lab, classroom and app-scale studies favour model-based selection by predicted recall over fixed or generic schedules. Tabibian's optimal reviewing rate is driven by recall probability itself, which is what the need term does.
- **Code vs evidence:** Matches. The horizons are assumptions inside the evidence's direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M17. Desirable-difficulty window d(p_now)

- **Where:** services/enforcers.js:276-279; services/retention-params.js:61 (window lo 0.5, hi 0.9, below 0.6, above 0.3, too_easy 0.95)
- **What it does:** The need is multiplied by d: 1 for 0.50 <= p_now <= 0.90; 0.65 for 0.90 < p_now <= 0.95; 0.6 below 0.50; 0.3 above 0.95.
- **Evidence:**
  - **eglington2020** [28]: "the OET 86 and 94 conditions having significantly higher recall performance than the OET 40" (adults on Mechanical Turk, Japanese-English pairs, 22-min practice, test 3 days later (full text PMC7567101))
  - **eglington2020** [28]: "There was not a significant difference between OETs 94, 86, and 70" (adults on Mechanical Turk (full text PMC7567101))

    - **pyc2009** [215]: title-level only: "Does greater difficulty correctly recalling information lead to higher levels of memory?" (college students, lab)
    - **fazio2019** [216]: "the benefits are largest when the initial retrieval practice is effortful but successful" (review of retrieval studies in children)
    - **kornell2009** [217]: "Unsuccessful retrieval attempts enhanced learning with both types of materials." (college students, lab)
    - **wilson2019** [94]: "the optimal training accuracy is about 85%" (computational model and simulations, not learners)
- **Verdict: moderate**. With time costs counted, practising at a predicted recall of 0.86 or 0.94 beat 0.40 and 0.98; 0.70 did not differ from 0.86 or 0.94. Failed retrievals still teach when feedback follows. So down-weighting $p < 0.5$ and $p > 0.95$ is supported. Eglington is one within-session vocabulary study of adults, with a 3-day test.
- **Code vs evidence:** Gap. The 0.65 step for 0.90-0.95 penalises the region Eglington found efficient; see proposal P1. The params file does not list the 0.65 step at all.
- **Status: Built on 25 October 2026** (switch `ENFORCER_DIFFICULTY_WINDOW_95`, from 02-P1). d(p_now) = 1 from 0.50 to 0.95; the 0.65 step for 0.90-0.95 is gone.

02-M18. Score flags: last answer wrong +0.30, confident error +0.25, right but slow or unsure +0.20, no delayed retrieval yet +0.15

- **Where:** services/enforcers.js:280-288; services/retention-params.js:62
- **What it does:** Topics the learner got wrong, got wrong while sure, got right only slowly or unsurely, or has never retrieved after a delay move up the queue.
- **Evidence:**
    - **pashler2005** [122]: "Supplying the correct answer after an incorrect response not only improved performance during the initial learning session--it also increased final retention by 494%." (college students, foreign-word pairs, lab)
    - **butterfield2001** [218]: "highly confident errors were the most likely to be corrected in a subsequent retest" (adults, general-knowledge questions)
    - **metcalfe2017** [219]: "errorful learning followed by corrective feedback is beneficial to learning" (review)
    - **butler2008** [220]: "feedback also doubled the retention of correct low-confidence responses, relative to providing no feedback" (college students, general-knowledge multiple choice)
    - **karpicke2007** [192]: "Delaying the first test improved long-term retention, regardless of how the repeated tests were spaced." (college students, word pairs, lab)
- **Verdict: moderate**. Corrections matter most after errors. Confident errors are hypercorrected. Feedback doubles retention of low-confidence correct answers. A delayed first retrieval helps. The weights are assumptions.
- **Code vs evidence:** Matches the direction. Two flags in the params file, "left out last time" and "prerequisite", are not built (see P5). Neither is the "within 14 days" limit on last-wrong.

- **Status: Built on 25 October 2026** (switch `ENFORCER_LEFT_OUT_FLAG`, from 02-P5). A topic left out of the latest enforcers gets +0.10 x min(1, times / 2). A skill deferred from an end-of-sitting set gets +0.20 (switch `ENFORCER_DEFER_SITTING_RETENTION`).

02-M19. Defaults with no record: p_now 0.5, p_H 0.4; skill-only units p_H = 0.85 x mastery

- **Where:** services/enforcers.js:257-258, :271-275
- **What it does:** A topic with no objective record is treated as a coin-flip. A topic known only at skill level uses its mastery as p_now, and 0.85 of that at the horizon.
- **Evidence:**
  - **eglington2023** [221]: "learner models assuming equal learning rates for students will inevitably lead to systematic errors that result in suboptimal pedagogical decisions" (simulation plus six large educational datasets)
  - **zerr2018** [222]: "quicker learners were also more durable learners (i.e., exhibited better retention across a delay)" (adults, N = 281 and N = 92)
  - **sense2016** [223]: "We demonstrate that they are stable over time but not across materials." (adults, repeated lab sessions)
- **Verdict: design choice**. Learners differ in learning rate and forgetting, and population or rule-of-thumb parameters cause systematic errors. That argues for replacing these defaults with the learner's own data quickly. It does not set 0.5, 0.4 or 0.85.
- **Code vs evidence:** Fallback only. Labelled as an assumption in RETENTION-ENFORCER-EVIDENCE.md section 11.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M20. Soft eligibility: minimum lag (5 min session, 30 min day, 1 day scheduled) and "kept in the last 20 hours" go last

- **Where:** services/enforcers.js:226-233 (recentlyKept), :296-299; services/retention-params.js:49-54, :63
- **What it does:** A topic still inside its minimum lag, or answered right in another enforcer in the last 20 hours, is not dropped. It sorts after the eligible ones, so a short window still has questions.
- **Evidence:**
  - **karpicke2007** [192]: "Delaying the first test improved long-term retention, regardless of how the repeated tests were spaced." (college students, word pairs, lab)
  - **karpicke2011spaced** [191]: "Repeated retrieval with long intervals between each test produced a 200% improvement in long-term retention" (college students, foreign-vocabulary pairs, lab)
  - **cepeda2009** [193]: "An optimal gap improved final recall by up to 150%." (adults, verbal materials, gaps and tests up to months)

    - **pyc2007** [167]: "the efficiency of the various schedules depended critically on the interval between presentations of an item" (college students, Swahili-English pairs, lab)
- **Verdict: moderate**. Delaying the first test, and leaving long intervals between repeated tests, improve retention. Gaps that are too short cost more than gaps a little too long. The lag values and the 20 hours are assumptions.
- **Code vs evidence:** Matches the direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M21. Greedy fill by value per planned second, with a cost model per format

- **Where:** services/enforcers.js:132-163 (costModel, costOf), :312-341 (fillOrder, slots), :583-597; services/retention-params.js:38-45
- **What it does:** Slots are ranked by score per planned second and taken until the budget is spent. Planned time is the format time plus expected feedback, and the expected retry on medium and full. Format times are MC 20 s, cloze 25 s, typed 35 s and problem 60 s; feedback is 3 s right, 25 s wrong. After 20 answers in a format, the learner's own median replaces the default.
- **Evidence:**
    - **pavlik2008** [213]: "the item is at the spacing interval where long-term gain per unit of practice time is maximal" (college students, Japanese-English vocabulary)
    - **eglington2020** [28]: "with up to 40% more items recalled in conditions where practice was scheduled optimally according to the model" (adults on Mechanical Turk, Japanese-English pairs, 22-min practice, test 3 days later)
    - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school Spanish course, time-matched)
- **Verdict: moderate**. Optimal schedulers that count the time cost of failures beat conventional ones. The time constants are assumptions.
- **Code vs evidence:** Matches. Time per format is calibrated from the learner.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M22. Allocation: one first try per topic; a second, varied try only for fragile topics on medium and full, worth 0.35 of a first

- **Where:** services/enforcers.js:290-295 (fragile, mature), :330-341 (slots); services/retention-params.js:72-79
- **What it does:** Fragile means the last answer was wrong, p_now is under 0.50, or a third or more of 2+ answers were wrong. Mature means 3+ delayed checks, p_now >= 0.90 and the last answer right; mature topics never get a second try. Session enforcers stay at one try per topic.
- **Evidence:**

    - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
    - **vaughn2016** [205]: "substantial benefits of learning to a higher initial criterion during the first session do not persist across relearning sessions" (college students, word pairs, lab)
    - **pyc2007** [167]: "it was particularly effective for learning initially incorrect items" (college students, Swahili-English pairs, lab)
    - **vaughn2013** [224]: "equating retrieval success during practice did not eliminate normative item difficulty differences" (college students, word pairs, lab)
    - **butler2017** [114]: "variability during retrieval practice produced superior transfer of knowledge to new examples" (college students, four experiments)
    - **rawson2020end** [225]: "final test performance showed a significant but only small advantage of successive relearning over single-session learning" (college students, maths procedures)
- **Verdict: moderate**. One correct per relearning session is the efficient dose. Extra tries belong to items that missed. Hard items stay hard at equal recalls. A varied second try transfers better. For maths procedures, relearning beats one session by only d = 0.28. The 0.35 value, the cap of 2 and the thresholds are assumptions.
- **Code vs evidence:** Matches the direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M23. Drop-out retries: quick 1 attempt, medium 2, full until right (at most 3), on a different question at least 3 items later

- **Where:** services/retention-params.js:57-58; services/enforcers.js:841-885 (afterAnswer)
- **What it does:** After a miss, while the tier allows and the budget holds, the topic comes back on a different question at least 3 questions later. It is never the skill bank's off-grade question.
- **Evidence:**
    - **pyc2007** [167]: "it was particularly effective for learning initially incorrect items" (college students, Swahili-English pairs, lab)
    - **higham2022** [181]: "Recall Of Course Material At The End Of The Semester Was Better For Relearning Compared With Restudying." (undergraduate psychology course)
    - **karpicke2011spaced** [191]: "Repeated retrieval with long intervals between each test produced a 200% improvement in long-term retention" (college students, foreign-vocabulary pairs, lab)
    - **pashler2005** [122]: "Supplying the correct answer after an incorrect response not only improved performance during the initial learning session--it also increased final retention by 494%." (college students, foreign-word pairs, lab)
- **Verdict: moderate**. Drop-out schedules match fixed ones in fewer trials and help most with initially wrong items. Relearning to criterion beat restudy in a course. Spacing the repeats inside a sitting helps.

- **Code vs evidence:** Matches. "3 items later" is an assumption.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M24. Interleave by skill inside the sitting

- **Where:** services/adaptation.js:414-421 (interleave: round-robin by skill); services/enforcers.js:606; services/review-session.js (via reviewQueue)
- **What it does:** Neighbouring questions differ in skill where possible, keeping the order within each skill.
- **Evidence:**
  - **rohrer2015** [12]: "interleaved practice produced higher scores on both the immediate and delayed tests (Cohen's ds 0.42 and 0.79" (grade-7 maths classes)
  - **rohrer2020** [13]: "the interleaved group outscored the blocked group, 61% versus 38%, d = 0.83" (grade-7 maths classes, randomized trial)
  - **brunmair2019** [14]: "A multilevel meta-analysis revealed a moderate overall interleaving effect (Hedges' g = 0.42)." (meta-analysis of interleaving)
  - **brunmair2019** [14]: "An advantage of blocking compared with interleaving was found for studies based on words (g = -0.39)." (meta-analysis of interleaving)
  - **firth2021** [163]: "transfer (as tested by presenting novel items from a learned category) a benefit" (systematic review, concept learning)
- **Verdict: mixed**. Strong for maths and concepts (classroom RCT d = 0.83), but blocking beat interleaving for word materials (g = -0.39). Round-robin by skill does not deliberately pair look-alike words, which is the safe reading for vocabulary.
- **Code vs evidence:** Matches the safe reading. Confusable pairing for maths and concepts is not built.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M25. Session scope asks the oldest-exposed topics first

- **Where:** services/enforcers.js:605
- **What it does:** In an end-of-session enforcer, first tries go oldest last-exposure first, so the start of the session gets the longest lag.
- **Evidence:**
  - **karpicke2007** [192]: "Delaying the first test improved long-term retention, regardless of how the repeated tests were spaced." (college students, word pairs, lab)
  - **karpicke2011spaced** [191]: "Repeated retrieval with long intervals between each test produced a 200% improvement in long-term retention" (college students, foreign-vocabulary pairs, lab)
  - **pyc2007** [167]: "the efficiency of the various schedules depended critically on the interval between presentations of an item" (college students, Swahili-English pairs, lab)

- **Verdict: moderate**. A longer delay before the first retrieval improves long-term retention, and an item's efficiency depends critically on the interval between its presentations.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M26. Close on a question the curve expects right (p_now >= 0.8)

- **Where:** services/enforcers.js:610-611; services/retention-params.js:64 (close_easy_at 0.8)
- **What it does:** The last question of a sitting is one the learner is likely to get right.
- **Evidence:**
  - **finn2010** [226]: "An extremely effortful study episode extended by a more moderate interval was preferred to a shorter, unextended interval" (college students, effortful vocabulary study, lab)
  - **finn2016** [227]: "a challenging test episode extended by a more moderate interval was preferred to an unextended interval" (adults, difficult maths tests, five experiments)
  - **kahneman1993** [228]: "such evaluations are often dominated by the discomfort at the worst and at the final moments of episodes" (adults, cold-pressor task)
  - **fredrickson1993** [229]: "Retrospective evaluations appear to be determined by a weighted average of "snapshots" of the actual affective experience" (adults, film clips; foundational for peak-end effects but does not specify ending in this quote)
- **Verdict: design choice**. An episode that ends on an easier stretch is remembered as less unpleasant and is chosen again, including effortful study and maths tests. This supports coming back, not retention itself. No study sets 0.8.
- **Code vs evidence:** Matches the purpose: willingness to take the next enforcer. It costs one slot of ordering, not extra time.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M27. A new question instance each time; a varied second try; never the learner's own check as "varied"

- **Where:** services/enforcers.js:438-481 (ownCheck, questionFor)
- **What it does:** An enforcer uses a question written for the objective first, then the bank's question on the skill, then, as a last resort, the learner's own window check. A second try must be a new question on the same objective, or there is none.
- **Evidence:**
  - **butler2017** [114]: "variability during retrieval practice produced superior transfer of knowledge to new examples" (college students, four experiments)
  - **pan2018** [113]: "testing can yield transferrable learning as measured relative to a nontesting reexposure control condition (d = 0.40" (meta-analysis of transfer of testing)

    - **agarwal2019** [140]: "higher order and mixed quizzes improved higher order test performance, but fact quizzes did not" (college students)
    - **yeo2019** [230]: "repeated testing was as effective as repeated studying with nonidentical learning problems" (college students, maths procedure, 1-week test)
    - **pan2019clicker** [141]: "Benefits of jargon quizzing did not consistently generalize, however, to exam questions that assessed conceptual knowledge" (college biology course)
- **Verdict: moderate**. Varied retrieval transfers better, testing transfers (d = 0.40), and fact-only quizzes do not reach higher-order tests. Only multiple-choice objective items exist today, so "varied" is a new example, not a new format.
- **Code vs evidence:** Partly. Varied formats (explain, worked-example completion) are not built.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-D. What counts as kept, and stakes

02-M28. A "kept" answer is first try, unaided, not rapid (< 1200 ms) and machine-graded; only answers at a lag of 1 day or more are memory checks

- **Where:** services/coverage.js:96 (RETAIN_DAYS 1), :102 (OBJ_RAPID_MS 1200), :439, :459-460; services/retention-params.js:49-54 (spaced flag per scope)
- **What it does:** Answers faster than 1.2 s are rapid guesses and are not evidence. Session and day enforcers feed in-sitting evidence. Only answers a day or more after the last practice count as memory checks.
- **Evidence:**
    - **wise2005** [74]: "unmotivated examinees will answer too quickly (i.e., before they have time to read and fully consider the item)" (low-stakes computer-based tests)
    - **wise2017** [59]: "because rapid guesses do not contribute to measurement, it makes little sense to include them in scoring" (review of rapid guessing in computer-based tests)
    - **wise2006** [35]: "rapid guessing (i.e., reflecting low examinee effort)" (low-stakes computer-based tests)
    - **dunlosky2012** [231]: title-level only: "Overconfidence produces underachievement: Inaccurate self evaluations undermine students' learning and retention" (college students)
    - **karpicke2008** [20]: "Repeated studying after learning had no effect on delayed recall, but repeated testing produced a large positive effect." (college students, foreign-vocabulary pairs, lab)
- **Verdict: moderate**. Rapid guesses reflect disengagement, not knowledge, and should not be scored. Self-assessment is overconfident, so nothing self-scored counts. Delayed tests are what show retention. The 1200 ms threshold is fixed, while the measurement literature often sets thresholds per item.
- **Code vs evidence:** Matches. A per-item rapid threshold would be more exact, but it is outside this section.

- **Status: Built on 25 October 2026** (switch `RAPID_ITEM_RELATIVE`, from 03-P8). For an item with 30 or more timed answers, "rapid" is under max(1.2 s, 0.10 x the item's median).

02-M29. Low stakes: enforcer and review sittings are previews that move no level; "none" is never penalised

- **Where:** services/enforcers.js:634-637 (kind enforcer, preview: true); services/review-session.js:13-16, :201
- **What it does:** Nothing in an enforcer or daily review clears, advances or holds a level. There are no grades or penalties.
- **Evidence:**
  - **khanna2015** [232]: "Students taking ungraded pop quizzes outperformed students taking graded pop quizzes and students taking no quizzes on the final exam." (introductory psychology course)
  - **agarwal2014anx** [233]: "72% reported that retrieval practice made them less nervous for unit tests and exams" (1,408 middle and high school students (self-report))
  - **smith2016stress** [234]: "those who learned by retrieval practice were immune to the deleterious effects of stress" (adults, lab stress induction)
  - **higham2022** [181]: "Recall Of Course Material At The End Of The Semester Was Better For Relearning Compared With Restudying." (undergraduate psychology course)
- **Verdict: moderate**. Ungraded quizzes beat graded ones in a course, and students report less test anxiety with retrieval practice (self-report). Retrieval-practised memories hold up under stress.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-E. The daily review and spaced review

02-M30. Daily review queue: memory checks first, then fading objectives, then the least durable; lapsed objectives to re-teach

- **Where:** services/adaptation.js:288-327 (reviewQueue), :724-780 (dailyReview); services/review-session.js:87-131 (planFor)
- **What it does:** Memory checks the gate is waiting on come first, then objectives with predicted recall under 0.90, then those predicted under 0.90 at 30 days. Each gets one question, interleaved by skill.
- **Evidence:**
  - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school Spanish course, time-matched)

    - **mettler2016** [171]: "In both experiments, adaptive scheduling outperformed fixed conditions at immediate and delayed tests of retention." (adults, lab)
    - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
    - **cepeda2009** [193]: "confirm the importance of cumulative reviews to promote retention over meaningful time periods" (adults, verbal materials, gaps and tests up to months)
- **Verdict: moderate**. Personalised, model-based review beat massed and generic schedules. Successive relearning and cumulative reviews keep material. The ordering by reason is a design choice.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M31. Review target: predicted recall 0.90 now (ADAPT_REVIEW_AT) and 0.90 at 30 days (ADAPT_DURABLE_AT, ADAPT_DURABLE_DAYS)

- **Where:** services/adaptation.js:60, :63-64, :226-239 (durability)
- **What it does:** An objective is reviewed when its predicted recall falls under 0.90. It counts as durable when its predicted recall 30 days on is at least 0.90 of fresh.
- **Evidence:**
    - **eglington2020** [28]: "the OET 86 and 94 conditions having significantly higher recall performance than the OET 40" (adults on Mechanical Turk, Japanese-English pairs, 22-min practice, test 3 days later (full text PMC7567101))
    - **eglington2020** [28]: "There was not a significant difference between OETs 94, 86, and 70" (adults on Mechanical Turk (full text PMC7567101))
    - **wozniak1995** [235]: "A universal formula for computing inter-repetition intervals in paired-associate learning has been determined for the knowledge retention level of 95%." (the authors' own repetition data)
    - **wilson2019** [94]: "the optimal training accuracy is about 85%" (computational model and simulations, not learners)
    - **tabibian2019** [26]: "subject to a cost on the reviewing frequency, the optimal reviewing schedule is given by the recall probability itself" (theory plus a natural experiment on Duolingo data)
- **Verdict: moderate**. Practising at 0.86-0.94 predicted recall was most efficient, and 0.70 did not differ. SuperMemo-type schedules target 95% retention (the author's own claim). 85% is a model result. The exact 0.90 is a design choice that sits inside the efficient band.
- **Code vs evidence:** Matches. Keep 0.90 (see P1).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M32. Rested gate: a review waits at least 12 h after the last answer; a memory check comes at least 1 day after proof

- **Where:** services/adaptation.js:306-310 (RETAIN_DAYS x 0.5); services/coverage.js:96, :459-460
- **What it does:** No objective is reviewed within 12 hours of its last answer. Proof becomes a memory check only after a day.
- **Evidence:**
  - **bell2014** [236]: "a 12-hour spacing gap that includes sleep promotes long-term memory retention similar to the 24-hour gap" (adults, lab)
  - **karpicke2007** [192]: "Delaying the first test improved long-term retention, regardless of how the repeated tests were spaced." (college students, word pairs, lab)
  - **cepeda2006** [22]: "the ISI producing maximal retention increased as retention interval increased" (meta-analysis, 839 assessments of verbal recall)
- **Verdict: moderate**. A 12-hour gap that includes sleep retained as well as 24 hours, and delaying the first retrieval helps. The 12 h and 1 d values are design choices.
- **Code vs evidence:** Matches. The 12-hour gate is the overnight case only when the 12 h include sleep.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M33. Review budget: 10 minutes at 0.75 min per question, one question per objective

- **Where:** services/adaptation.js:78-79 (REVIEW_MINUTES, ITEM_MINUTES); services/review-session.js:113-119
- **What it does:** The daily review holds as many single questions as fit in 10 minutes, at least one when any is due. The rest wait for the next day.
- **Evidence:**
  - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
  - **rohrer2007pashler** [197]: "immediate further study of the same material is an inefficient use of time" (review of spacing and overlearning studies)
  - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school Spanish course, time-matched)
- **Verdict: design choice**. One correct per relearning session, short spaced sittings, and time-matched gains support the shape. They do not support 10 minutes or 0.75 min.
- **Code vs evidence:** Matches the shape.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M34. Daily review folded into the enforcers: it keeps only older work no enforcer will reach, and skips objectives kept in an enforcer in the last 20 h

- **Where:** services/review-session.js:68-131 (folded, planFor); services/enforcers.js:1151 (reachStart)

- **What it does:** While enforcers are on, the daily review asks only objectives last answered before the earliest window still to come. It also skips anything answered right in an enforcer in the last 20 hours. During consolidation it is the whole review.
- **Evidence:**
  - **cepeda2009** [193]: "An optimal gap improved final recall by up to 150%." (laboratory studies of foreign vocabulary, facts, and visual-object names, with test delays up to 6 months)
  - **karpicke2011spaced** [191]: "Repeated retrieval with long intervals between each test produced a 200% improvement in long-term retention" (laboratory experiment on repeated retrieval and long-term retention)
  - **rohrer2006** [198]: "long-term retention was boosted by distributed practise and unaffected by overlearning" (college students, maths problems)
- **Verdict: design choice**. This avoids two retrievals of the same objective too close together, which spend time for little gain. The split itself is an owner rule.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M35. Typed recall preferred in the review: "one right typed recall is a whole memory check, a multiple-choice one is half"

- **Where:** services/review-session.js:53-66 (questionFor -> generateWithPreference(entry, true))
- **What it does:** Where the skill has typed items, the review asks typed rather than multiple choice. Coverage weighs a typed right answer as a full check.
- **Evidence:**
  - **rowland2014** [10]: "initial recall tests yielding larger testing benefits than recognition tests" (meta-analysis of testing vs restudy)
  - **kang2007** [166]: "With feedback provided, having had an intervening SA test led to the best performance on the final test" (students in laboratory experiments)
  - **smith2014** [116]: "Practising retrieval in all format conditions enhanced retention, relative to a study-only control condition" (college students, prose, lab)
  - **little2012** [133]: "properly constructed multiple-choice tests can indeed trigger productive retrieval processes" (participants in two experiments)
- **Verdict: mixed**. Recall tests give larger testing benefits than recognition, and short answer with feedback was best. But Smith and Karpicke found all formats enhanced retention, and well-built multiple choice triggers productive retrieval. As evidence, a typed answer is less guessable, which is the stronger reason for full weight.
- **Code vs evidence:** Matches on direction. The "half" weight is a design choice.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M36. Lapsed objectives are re-taught as a lesson, at most 3 per session, 6 minutes each

- **Where:** services/adaptation.js:71 (RETEACH 3), :80 (RETEACH_MINUTES 6), :300-305
- **What it does:** An objective the record says is forgotten is not asked as a single question. It is taught again from its lesson, soon.
- **Evidence:**
  - **arnold2013** [237]: "unsuccessful retrieval attempts enhance the effectiveness of subsequent restudy, demonstrating that tests do potentiate subsequent learning" (college students, lab)
  - **kornell2009** [217]: "Unsuccessful retrieval attempts enhanced learning with both types of materials." (college students, lab)
  - **hays2013** [238]: "Failed tests followed by delayed feedback, by contrast, did not produce such a benefit" (college students, word pairs, lab)
- **Verdict: moderate**. A failed retrieval followed by restudy or feedback potentiates learning, and only when the restudy comes soon. Relearning is faster than first learning. The caps are assumptions.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M37. Skill-level spaced review (mastery.dueReviews): an SM-2-style ladder

- **Where:** services/mastery.js:15-18 (ease 2.5 in 1.3-3.0; intervals 1, 3, 7, 16, 35, 70, 140 d), :94-103, :144-151 (dueReviews); services/sessions.js:135
- **What it does:** A right answer moves the skill one step up the ladder, scaled by ease (+0.05 right, -0.15 wrong); a coached answer halves the step; a miss is due in 0.5 d. Due skills, weakest then most overdue, fill up to 30% of a lesson.
- **Evidence:**
  - **kang2014** [175]: "the expanding schedule yielded much higher average recallability over the whole training period" (adults, long-term retrieval practice)
  - **latimier2021** [159]: "no significant difference between expanding and uniform spacing schedules of retrieval practice (g = 0.034)" (meta-analysis of spaced retrieval studies)
  - **wozniak1995** [235]: "A universal formula for computing inter-repetition intervals in paired-associate learning has been determined for the knowledge retention level of 95%." (the authors' own repetition data)
  - **mettler2016** [171]: "adaptive scheduling outperformed fixed conditions at immediate and delayed tests of retention" (adults, lab)
  - **tabibian2019** [26]: "learners who follow a reviewing schedule determined by our algorithm memorize more effectively than learners who follow alternative schedules" (natural experiment on Duolingo data)
- **Verdict: mixed**. Expanding intervals are fine: no final-test difference from uniform, and higher recall along the way. But adaptive, model-based schedules beat fixed heuristic ladders. The app already has a better model (adaptation.predictedRecall) used everywhere else.

- **Code vs evidence:** Gap. Two schedulers: lessons use the fixed ladder, enforcers and the review use the fitted curve. See P4.
- **Status: Built on 25 October 2026** (switch `CURVE_REVIEWS`, from 02-P4). Where an objective model exists, lesson reviews come from the fitted curve; the SM-2 ladder is the fallback.

02-M38. Due reviews go first in a lesson, one difficulty step below the level

- **Where:** services/sessions.js:135-141, :162-166 (planLesson)
- **What it does:** A lesson opens with its due review items at difficulty -1, then the new material.
- **Evidence:**
  - **pastotter2014** [158]: "recall testing of previously studied information can enhance learning of subsequently presented new information" (review of lab and applied studies)
  - **chan2018** [240]: "testing reliably potentiates the future learning of new materials by increasing correct recall or by reducing erroneous intrusions" (meta-analysis of test-potentiated new learning)
  - **yang2018fwd** [241]: "interim testing of studied information facilitates learning and retrieval of new information—the forward testing effect" (review)
  - **szpunar2008** [242]: "interpolating tests during the study sequence insulates against this negative influence" (college students, word lists, lab)
- **Verdict: moderate**. Retrieving studied material before new material improves learning of the new material (the forward testing effect, meta-analysed). The easier step serves "effortful but successful", and its size is a design choice.
- **Code vs evidence:** Matches.
- **Status: Built on 25 October 2026** (switch `CURVE_REVIEWS`, from 02-P4). The reviews that go first are now chosen by the forgetting model where one exists.

02-M39. Interleaving across sessions: earlier skills come back inside lessons; the focus skill alternates with the next weakest every third item

- **Where:** services/sessions.js:135, :155-160
- **What it does:** Past skills return as due reviews, and the body mixes the weakest introduced skill into every third focus slot.
- **Evidence:**
  - **rohrer2015** [12]: "interleaved practice produced higher scores on both the immediate and delayed tests (Cohen’s ds 0.42 and 0.79" (grade-7 maths classes)
  - **rohrer2020** [13]: "the interleaved group outscored the blocked group, 61% versus 38%, d = 0.83" (grade-7 maths classes, randomized trial)
  - **emeny2021** [243]: "spaced practice produced higher test scores than did massed practice" (grade-7 maths classes)
  - **brunmair2019** [14]: "A multilevel meta-analysis revealed a moderate overall interleaving effect (Hedges’ g = 0.42)." (meta-analysis of interleaving)

- **Verdict: moderate**. Interleaved maths assignments, which also space each skill across sessions, beat blocked ones in classrooms. Spaced maths practice beat massed.
- **Code vs evidence:** Partly. The body is two-thirds focus skill, mild interleaving, which suits a first exposure. Reviews are blocked at the start rather than mixed in.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-F. Consolidation before the Final

02-M40. Final held for review-only days until 90% of objectives are durable at 30 days, for at most 21 days (ADAPT_CONSOLIDATE)

- **Where:** services/consolidation.js:1-23, :59-90; services/adaptation.js:81-82, :783-803
- **What it does:** On first reaching the Graduation exam, the Final waits. The learner does daily reviews of at most 10 minutes until 90% of the grade's objectives are predicted at least 0.90 recalled 30 days on, or 21 days pass. The next grade can be started meanwhile.
- **Evidence:**
  - **rawson2011** [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals" (533 college students, key-term definitions)
  - **khanna2013** [91]: "Students completing courses with cumulative finals retained more than students who took noncumulative finals." (college psychology courses)
  - **cepeda2008** [23]: "the optimal gap declined from about 20 to 40% of a 1-week test delay to about 5 to 10%" (over 1,350 adults online, facts, tests up to 1 year)
  - **rohrer2006** [198]: "long-term retention was boosted by distributed practise and unaffected by overlearning" (college students, maths problems)
  - **kapler2015** [244]: title-level only: "Spacing in a simulated undergraduate classroom: Long-term benefits for factual and higher-level learning" (simulated undergraduate classroom)
- **Verdict: design choice**. Relearning in spaced sessions before a test, cumulative finals, and gaps sized to the retention interval all support the shape. The 90%, 30 days and 21 days are design choices. The simulation ablation in the code comment (95.6% vs 80.0%) is the app's own model, not evidence.
- **Code vs evidence:** Matches the shape. It costs calendar days, not study time.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-G. Retrieval inside lessons and sittings

02-M41. "Check what stuck": read the lesson, then answer its checks

- **Where:** public-v2/js/session.js:264-290 (lesson page, start-checks button); also public/js/session.js, public-redesign/js/session.js
- **What it does:** After the lesson text, worked example and common mistakes, one button starts that block's questions.

- **Evidence:**
  - **roediger2006** [1]: "on the delayed tests, prior testing produced substantially greater retention than studying" (college students, prose passages, lab)
  - **karpicke2008** [20]: "Repeated studying after learning had no effect on delayed recall, but repeated testing produced a large positive effect." (college students, foreign-vocabulary pairs, lab)
  - **smith2014** [116]: "Practising retrieval in all format conditions enhanced retention, relative to a study-only control condition" (college students, prose, lab)
  - **agarwal2021** [99]: "the majority of which (57%) revealed medium or large benefits from retrieval practice" (classroom experiments, total n = 5,374, school and college)
  - **yang2021** [21]: "overall testing (quizzing) raises student academic achievement to a medium extent (g = 0.499)" (meta-analysis of 222 classroom studies)
- **Verdict: strong**. Study followed by retrieval beats more study, in the lab and in classrooms.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M42. Per-level retention question after each level closes

- **Where:** services/sessions.js:1096-1124 (blockRetentionItem), :1186; routes/learn.js:134 (retention items always reveal feedback)
- **What it does:** When a level is resolved, one recall question is added from what wobbled in it (a random miss, else a random item), half a step easier.
- **Evidence:**
  - **szpunar2008** [242]: "interpolating tests during the study sequence insulates against this negative influence" (college students, word lists, lab)
  - **szpunar2013** [245]: "interpolating online lectures with memory tests can help students sustain attention to lecture content" (college students, online lecture video)
  - **lyle2011** [190]: "Exam scores were significantly and substantially higher in a section of the course taught with the PUREMEM procedure" (undergraduate course, end-of-lecture retrieval)
  - **pashler2005** [122]: "Supplying the correct answer after an incorrect response not only improved performance during the initial learning session--it also increased final retention by 494%." (college students, foreign-word pairs, lab)
- **Verdict: moderate**. Interpolated tests protect against interference and hold attention, and end-of-lesson retrieval raised exam scores. Feedback on a miss is what makes it teach.
- **Code vs evidence:** Matches. The pick is random within misses rather than scored.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M43. End-of-sitting retention set: at most 3 questions, one per level from its misses, oldest level first, half a step easier

- **Where:** services/sessions.js:1510-1561 (retentionPlan), :1703; routes/learn.js:676-720
- **What it does:** When a sitting finishes, the analyst's own retention questions or bank items on the wobbly skills are added. They are excluded from accuracy, and feedback always shows.
- **Evidence:**
  - **lyle2011** [190]: "Exam scores were significantly and substantially higher in a section of the course taught with the PUREMEM procedure" (undergraduate course, end-of-lecture retrieval)
  - **karpicke2007** [192]: "Delaying the first test improved long-term retention, regardless of how the repeated tests were spaced." (college students, word pairs, lab)
  - **rohrer2006** [198]: "long-term retention was boosted by distributed practise and unaffected by overlearning" (college students, maths problems)
  - **vaughn2016** [205]: "to a higher initial criterion during the first session do not persist across relearning sessions" (college students, word pairs, lab)
- **Verdict: moderate**. End-of-lesson retrieval helps, but more practice right away adds little once spaced relearning follows. The first test does better when delayed.
- **Code vs evidence:** Gap. There is no lag rule, and when enforcers are on it duplicates what the session enforcer asks when the study session ends. See P3.
- **Status: Built on 25 October 2026** (switch `ENFORCER_DEFER_SITTING_RETENTION`, from 02-P3). Not asked while a study session is running; the session enforcer covers it.

02-H. Feedback

02-M44. Feedback timing: after each answer by default; held to the end of the level only when the learner's style answers ask for it; always immediate in enforcers, reviews and retention items

- **Where:** services/learner.js:279, :308 (immediateFeedback: timing < 0.55); routes/learn.js:85-93 (holdsFeedbackByChoice), :134; services/enforcers.js:640; services/review-session.js:201
- **What it does:** The style quiz's feedback-timing dimension can switch a lesson to end-of-level feedback. Checkpoints hold feedback for their own reasons.
- **Evidence:**
  - **kulik1988** [184]: "Applied studies using actual classroom quizzes and real learning materials have usually found immediate feedback to be more effective than delayed." (review of feedback-timing studies)
  - **hays2013** [238]: "Failed tests followed by delayed feedback, by contrast, did not produce such a benefit" (college students, word pairs, lab)
  - **butler2007** [185]: "delayed feedback led to superior final test performance relative to immediate feedback" (college students, multiple-choice, lab)
  - **metcalfe2009** [186]: "there was no difference between the delayed and immediate feedback conditions when the lag to test was controlled" (grade-6 children and college students, vocabulary)

    - **mullet2014** [246]: "When students received delayed feedback, they performed better on subsequent course exams that contained new problems" (engineering course)
    - **corral2021** [247]: "no posttest differences were observed between the feedback conditions in Experiments 1 and 2" (adults, complex concept learning)
    - **pashler2008** [31]: "we found virtually no evidence for the interaction pattern mentioned above" (review of learning-styles (meshing) evidence)
- **Verdict: mixed**. Applied studies favour immediate feedback, and after a failed retrieval delayed feedback lost the benefit. But lab multiple-choice and some course studies favour delayed feedback, and with the lag controlled there was no difference. Matching to a stated preference has no support in the learning-styles evidence.
- **Code vs evidence:** Partly. Retention sittings are immediate plus a delayed recap, the safe combination. Holding lesson feedback by preference rests on preference, not evidence. It is not proposed for change because the evidence on timing is mixed.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M45. Feedback content: the answer and why, plus "A second look" recap of misses at the end of an enforcer

- **Where:** public/js/enforcer.js:246-257 (missed), :287; services/sessions.js:1085-1094 (answer, rationale, feedback)
- **What it does:** A wrong answer shows the correct answer and its rationale. At the end of the enforcer the misses are shown again, with "comes back in the next enforcers".
- **Evidence:**
    - **pashler2005** [122]: "Supplying the correct answer after an incorrect response not only improved performance during the initial learning session--it also increased final retention by 494%." (college students, foreign-word pairs, lab)
    - **vanderkleij2015** [15]: "elaborated feedback (EF; e.g., providing an explanation) produced larger effect sizes (0.49) than feedback regarding the correctness of the answer (KR; 0.05)" (meta-analysis, computer-based learning)
    - **butler2013** [17]: "explanation feedback produced superior performance on the new inference questions" (college students, lab)
    - **butler2008mc** [248]: "both immediate and delayed feedback increased the proportion of correct responses and reduced, rather than eliminated, the proportion of intrusions" (college students, multiple choice)
    - **smith2010kimball** [249]: "delaying feedback strengthens initially correct responses due to the spacing of encoding opportunities" (college students, lab; it found little effect of the delay on correcting errors)
- **Verdict: strong**. The answer after an error raised retention 494%. Elaborated feedback (0.49) beats right/wrong (0.05). Explanations help inference. Feedback reduces, though it does not remove, multiple choice's lure risk. A delayed second look adds spacing.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-I. Sleep and day boundaries

02-M46. Day boundaries: local calendar days; scheduled scopes need a 1-day lag; the day enforcer closes the day; no bedtime gate

- **Where:** services/retention-schedule.js:190-226 (localDate, dayIndex); services/retention-params.js:51-54 (min_lag 24 h); services/study-sessions.js:111-117 (paused session ends at the day change)
- **What it does:** A "day" is the learner's local date from their anchor. Every scheduled enforcer counts only answers at least a day old as spaced, so at least one night lies between learning and the first scheduled retrieval.
- **Evidence:**
  - **berres2021** [250]: "we found a moderate overall sleep benefit in episodic memory (g = 0.44)" (meta-analysis, 823 effect sizes; g = 0.28 after bias correction)
  - **gais2006** [251]: "declarative memory is enhanced when sleep follows within a few hours of learning, independent of time of day" (lab, declarative (vocabulary) learning)
  - **mazza2016** [239]: "interleaving sleep between learning sessions not only reduced the amount of practice needed by half" (adults, foreign-word pairs, 6-month follow-up)
  - **bell2014** [236]: "a 12-hour spacing gap that includes sleep promotes long-term memory retention similar to the 24-hour gap" (adults, lab)
  - **bauml2014** [252]: "sleep benefited recall of restudied items but left recall of retrieved items unaffected" (adults, lab)
  - **wilhelm2008** [253]: "Beneficial effects of sleep on retention of declarative memories were comparable in children and adults." (children and adults, lab)
- **Verdict: moderate**. Sleep after learning helps, g = 0.44 (0.28 after bias correction), and children benefit like adults. Sleep between sessions halves the relearning needed. But sleep helps restudied items more than retrieved ones, so retrieval just before bed is not a special booster.
- **Code vs evidence:** Matches, with one gap in the docs. RETENTION-ENFORCER-EVIDENCE.md section 3.8 says "the app suggests doing the day enforcer in the evening", but no such suggestion exists in the code or UI. That line should read "not built". Given bauml2014 [252], sleep is not shown to provide an extra recall boost after retrieval practice.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-J. The forgetting model and its recall estimates

02-M47. Half-life regression with a population fit, a guessing floor and capped counts

- **Where:** services/scorecard/forgetting.js:1-35 (P = g + (pk - g) x 2^(-dt/h); log2 h = a + b nc + c nw + m), :88-105 (predict); DEFAULT_CURVE {a 1, b 0.3, c -0.1}; CAP 15

- **What it does:** Recall is predicted from the days since last practice, the counts of earlier right and wrong answers, the ability model's fresh expectation, and a guess floor. The curve is fitted to all learners, with a default until there is enough data.
- **Evidence:**
  - **settles2016** [25]: "reducing error by 45%+ compared to several baselines at predicting student recall rates" (Duolingo practice logs (error in predicting recall, not a learning gain))
  - **ye2022** [214]: "Experimental results have shown a 12.6% performance improvement over the state-of-the-art methods." (220 million review logs from an online language-learning app)
  - **pavlik2005** [254]: "The relative benefit of spacing increased with increased practice and with longer retention intervals." (college students, Japanese-English vocabulary)
  - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school foreign-language course, time-matched)
- **Verdict: moderate**. Half-life regression predicts recall well at scale, and model-based schedulers outperform. An activation model with power-law decay explains spacing-by-practice interactions. The default parameters are assumptions until fitted.
- **Code vs evidence:** Matches. Its calibration on real memory tests is not yet known.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M48. Curve shape: exponential by default; a power-law tail available but off

- **Where:** services/scorecard/forgetting.js:36-58 (FEATURES basic, SHAPE exp, SHAPES list)
- **What it does:** Forgetting is read as 2^(-dt/h). The power-law option (SCORECARD_FORGET_SHAPE=fit) is off, because on the simulation's own exponential memory it over-predicted.
- **Evidence:**
  - **wixted1991** [255]: "the course of forgetting was best described by a simple power function of time" (adults (word recall, face recognition) and pigeons)
  - **rubin1996** [256]: "The best fits were to the logarithmic function, the power function, the exponential in the square root of time" (210 published retention data sets)
  - **murre2015** [257]: "the Ebbinghaus forgetting curve has indeed been replicated" (one adult, 70 hours of list learning (replication))
  - **averell2011** [258]: title-level only: "The form of the forgetting curve and the fate of memories" (re-analysis of forgetting data)
- **Verdict: mixed**. Pooled human forgetting data are fit better by power-like functions, and an exponential under-predicts long-interval recall. The simulation that rejected the power tail assumes exponential memory, so it cannot test this. Long-horizon p_H is probably pessimistic, which means more reviews: the safe direction.
- **Code vs evidence:** Gap. The choice was made on simulated data. See P2.

- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 02-P2, choosing the curve shape on real memory tests (needs data).

02-M49. Only spaced successes lengthen the half-life

- **Where:** services/adaptation.js:196-209 (predictedRecall reads memory checks, not same-day answers)
- **What it does:** Without the fitted curve, the half-life grows only with right memory checks (plus the proof) and shrinks with missed ones. Same-day practice answers do not count, after they made a freshly drilled objective "look good for ten days".
- **Evidence:**
  - **cepeda2006** [22]: "the ISI producing maximal retention increased as retention interval increased" (meta-analysis, 839 assessments of verbal recall)
  - **rohrer2006** [198]: "long-term retention was boosted by distributed practise and unaffected by overlearning" (college students, maths problems)
  - **karpicke2011spaced** [191]: "Repeated retrieval with long intervals between each test produced a 200% improvement in long-term retention" (lab study)
  - **donovan1999** [259]: "an overall mean weighted effect size of 0.46, indicating that individuals in spaced practice conditions performed significantly" (better than those in massed practice; meta-analysis, 63 studies of practice distribution)
  - **kornell2009cram** [209]: "spacing was more effective than massing for 90% of the participants, yet after the first study session, 72% of the participants believed" (participants, GRE-type word-pair flashcards)
- **Verdict: strong**. Spaced practice beats massed practice across meta-analyses, and overlearning in one sitting adds little long-term.
- **Code vs evidence:** Matches. The fitted "basic" curve still counts all answers alike, and the spaced refit is a data step.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

02-M50. Personal forgetting multiplier from the learner's own memory checks

- **Where:** services/adaptation.js:239-279 (memoryAdjust: Beta posterior, >= 10 checks; -1 doubling if P(fail rate > 0.25) >= 0.95; +0.5 if >= 20 checks and P(fail rate < 0.05) >= 0.95)
- **What it does:** The curve is shortened for a learner who clearly forgets faster than expected, and lengthened (by less) for one who clearly keeps more. The bar is 95% posterior certainty.
- **Evidence:**
  - **sense2016** [223]: "We demonstrate that they are stable over time but not across materials." (adults, repeated lab sessions)
  - **zerr2018** [222]: "quicker learners were also more durable learners (i.e., exhibited better retention across a delay)" (adults, N = 281 and N = 92)

    - **eglington2023** [221]: "learner models assuming equal learning rates for students will inevitably lead to systematic errors that result in suboptimal pedagogical decisions" (simulation plus six large educational datasets)
    - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school Spanish course, time-matched)
- **Verdict: moderate**. Individual forgetting rates are stable over time and differ between people. Models that assume equal rates make systematic errors. Personalised review beat generic review. The asymmetric steps and the 10/20/0.95 bars are design choices; the asymmetry errs toward more review.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

The eglington2020 [28] question: should the enforcer's selection weights shift?

What the study found. Eglington & Pavlik (2020): adults on Mechanical Turk practised Japanese-English word pairs for 22 minutes and were tested 3 days later.

- In each condition the scheduler practised "whichever item the model estimated was closest to (but less than) the recall probability threshold" (full text).
- Thresholds of 0.86 and 0.94 gave significantly better final recall than 0.40 and 0.98. That held with the time cost of failures counted: failed retrievals take longer, because of feedback.
- **0.70 did not differ significantly from 0.86 or 0.94.**
- 0.86 was numerically highest.

What the code did before 25 October 2026 (`services/enforcers.js:276-279`, `services/retention-params.js:61`; the 0.90-0.95 row was changed that day, see 02-M17 Status):

p_now	d now	Broadly consistent with eglington2020 [28]?
< 0.50	0.6	yes (0.40 was worse)
0.50 to 0.90	1.0	yes (0.70 and 0.86 equal-best)
0.90 to 0.95	**0.65**	**no: 0.94 was among the best**
> 0.95	0.3	yes (0.98 was worse)

The need term (1 - p_H) is not in question. The theory in tabibian2019 [26] puts the optimal reviewing rate on the recall probability itself, while its large natural experiment found that the resulting algorithm outperformed heuristic schedules.

Verdict: shift one weight. Do not re-centre the window.

02-P1. Raise d(p_now) for 0.90 < p_now <= 0.95 from 0.65 to 1.0

BUILT (flag `ENFORCER_DIFFICULTY_WINDOW_95`, on by default; 2026-10-25, implementer A). Steps 1, 3 and 4 were done as written. Step 2 differs: the brief requires an off switch, so the step branch stays, but only behind the switch. It now lives in `retention-params.js` `difficultyWeight`, and `score` calls that. With the switch on, d is 1 from 0.50 to 0.95 and no step applies. `=off` restores

0.90 and the 0.65 step. Tests: `scripts/retention-enforcers-test.js`, section "P1". The scorecard-sim A/B suggested under "How to check it" was not run.

- **The exact change:**
  1. `services/retention-params.js:61`: set `window: { lo: 0.5, hi: 0.95, below: 0.6, above: 0.3, too_easy: 0.95 }` (hi 0.90 → 0.95).
  2. `services/enforcers.js:279`: `else if (pNow > dw.hi) d = (1 + dw.above) / 2;` can then never fire. Delete it, so the code says what it does.
  3. `curriculum-build/retention-enforcer-params.json` `selection.step_2_score.difficulty_window`: set it to "d(p_now) = 1.0 if 0.50 <= p_now <= 0.95; 0.6 if p_now < 0.50; 0.3 if p_now > 0.95". Add eglington2020 [28] to its `basis` as general support for adaptive, item-level scheduling. (`fromResearch` does not read the window today, so this is documentation unless that is added too.)
  4. `curriculum-build/RETENTION-ENFORCER-EVIDENCE.md`: update section 5, step 2, and the section 11 line on the 0.65 step.
- **What stays the same, and why:**
  - **below = 0.6**: practice aimed at 0.40 was worse.
  - **above = 0.3 for p_now > 0.95**: 0.98 was worse, through overpractice and short spacing.
  - **need = 1 - p_H**: this operationalizes the principle that review priority should track memory strength (tabibian2019 [26]).
  - **ADAPT_REVIEW_AT = 0.90**: it sits inside 0.86-0.94.
  - **The 0.50-0.90 plateau**: 0.70 did not differ from 0.86 or 0.94. Peaking the weight at 0.86-0.94 would push down topics in 0.50-0.85. In an enforcer those topics have fallen *past* the efficient point, and Eglington's own rule practises items under the threshold first.
- **Not recommended yet: a closest-to-threshold rule.** This means replacing d with a rule like "closest under 0.90", Khajah's heuristic (b). Test it in `curriculum-build/scorecard-sim` before considering it, because it changes behaviour on backlogs. Eglington and Khajah tested it on fresh learning or in simulation, not on a cross-day backlog.
- **If the first-try accuracy band is ever built** (`global.target_first_try_accuracy.band`, not built today): use [0.70, 0.90] rather than [0.70, 0.85]. 0.70-0.94 was one efficient region in eglington2020 [28], while 0.85 comes from a model of learning algorithms, not a band tested with learners (wilson2019 [94]).
- **Evidence:**
  - **eglington2020** [28]: "the OET 86 and 94 conditions having significantly higher recall performance than the OET 40" (adults on Mechanical Turk, Japanese-English pairs, 22-min practice, test 3 days later (full text PMC7567101))
  - **eglington2020** [28]: "There was not a significant difference between OETs 94, 86, and 70" (adults on Mechanical Turk (full text PMC7567101))
  - **eglington2020** [28]: "a simulated student practiced whichever item the model estimated was closest to (but less than) the recall probability threshold" (simulation and experiment design (full text PMC7567101))

    - **pavlik2008** [213]: "the item is at the spacing interval where long-term gain per unit of practice time is maximal" (college students, Japanese-English vocabulary)
    - **khajah2014** [260]: "review material whose predicted memory strength is closest to a particular threshold" (simulations with two memory models (ACT-R, MCM), semester-long review)
    - **tabibian2019** [26]: "subject to a cost on the reviewing frequency, the optimal reviewing schedule is given by the recall probability itself" (theory plus a natural experiment on Duolingo data)
    - **wilson2019** [94]: "the optimal training accuracy is about 85%" (computational model and simulations, not learners)
- **Expected effect: small.** A worked example at a weekly horizon:
    - Topic A, p_now 0.93 and p_H 0.70: its score rises from 0.30 × 0.65 = 0.195 to 0.30.
    - Topic B, p_now 0.80 and p_H 0.55: 0.45 before and after.
    - Flagged misses keep their +0.30 and stay on top.
    - The change mostly stops near-threshold topics, including freshly learned ones in session and day enforcers (the fallback curve tops out at 0.95), from losing to topics with the same need but a different p_now.
- **Risk: low.**
    - The evidence is one within-session study: adults, vocabulary, a 3-day test. Enforcers are cross-day and mix materials.
    - Near-threshold items are cheap (3 s feedback when right), so the budget cost is small.
- **How to check it:**
    - A scorecard-sim A/B (hi 0.90 vs 0.95, four seeds), with 30-day recall as the outcome.
    - Then live calibration: predicted p against the next memory check, as in RETENTION-ENFORCER-EVIDENCE.md section 12.2.

Other proposals (3 or more sources each; status marked on each)

02-P2. Choose the forgetting-curve shape on real memory tests, not on the simulation

SKIPPED: needs data. It needs real learners' memory tests, enough for MIN_FIT_TESTS, and we do not have them yet (2026-10-25).

- **What.** Once the population fit has enough memory tests (MIN_FIT_TESTS), fit both `SCORECARD_FORGET_SHAPE=exp` and `=fit` (the power-law tail) on real learners' memory tests. Compare held-out log-likelihood and 30-day calibration, and adopt the power tail only if it wins. The current rejection used the simulation's exponential memory, which cannot show a power tail.
- **Evidence:**
    - **wixted1991** [255]: "the course of forgetting was best described by a simple power function of time" (adults (word recall, face recognition) and pigeons)

    - **rubin1996** [256]: "The best fits were to the logarithmic function, the power function, the exponential in the square root of time" (210 published retention data sets)
    - **murre2015** [257]: "the Ebbinghaus forgetting curve has indeed been replicated" (one adult, 70 hours of list learning (replication))
    - **averell2011** [258]: title-level only: "The form of the forgetting curve and the fate of memories" (re-analysis of forgetting data)
- **Expected effect.** Better long-horizon p_H, so bi-weekly and monthly enforcers spend less time on topics the exponential wrongly calls faded.
- **Risk.** A power tail that over-predicts recall means under-review, the unsafe direction. Keep the exponential when the likelihoods tie.
- **Files.** `services/scorecard/forgetting.js:56` (SHAPE default) and `fitCurve`; a new comparison script in `scripts/`. Nothing in the enforcers changes.

02-P3. Stop the end-of-sitting retention set from duplicating the session enforcer

BUILT (flag `ENFORCER_DEFER_SITTING_RETENTION`, on by default; 2026-10-25, implementer A).

- **When it applies.** `enforcers.deferSittingRetention` is called from `routes/learn.js` (the finish route). It applies when enforcers are on, the account is not the owner's base version, the sitting is not an enforcer, and a study session is open (running or paused) that holds at least one of the sitting's answers. Then no end-of-sitting retention items are created, and the response carries `retention_deferred: true`. The per-level question is unchanged.
- **How the skills reach the enforcer.** The set's skills that were *missed* in the sitting are stored on the study session (`retention_deferred`). The session or day enforcer then gives those topics +0.20 (`flags.deferred_retention`; new params key `wobbled_in_deferred_sitting_set`), but only when neither "last answer wrong" nor "slow or unsure" already fired. The weight is an assumption: the same as "right but unsure".
- **Not changed:** `services/sessions.js` (`retentionPlan`) belongs to implementer B; nothing there was needed.
- **Tests:** `scripts/retention-enforcers-test.js`, section "P3".
- **Browser check:** on a test copy, the recap had no Retention check card, and the session enforcer flagged the missed skill's topic.
- **What.** When `RETENTION_ENFORCERS` is on and a study session is running, skip the at-most-3 end-of-sitting retention questions (`routes/learn.js:676-720`). Pass their skills to the session enforcer as flagged instead. That enforcer asks the same window's topics when the study session ends, with a 5-minute lag and oldest first. Keep the set when no study session is running, and keep the per-level question.
- **Evidence:**
    - **karpicke2007** [192]: "Delaying the first test improved long-term retention, regardless of how the repeated tests were spaced." (college students, word pairs, lab)

    - **rohrer2006** [198]: "long-term retention was boosted by distributed practise and unaffected by overlearning" (college students, maths problems)
    - **rohrer2005over** [261]: "The overlearners recalled far more than the low learners at the 1-week test, but this difference decreased dramatically thereafter." (college students, facts, tests to 9 weeks)
    - **vaughn2016** [205]: "to a higher initial criterion during the first session do not persist across relearning sessions" (college students, word pairs, lab)
    - **rohrer2007pashler** [197]: "immediate further study of the same material is an inefficient use of time" (review of spacing and overlearning studies)
- **Expected effect.** About 1-2 minutes saved per sitting with no expected loss of retention. One effortful delayed retrieval replaces two quick ones close together.
- **Risk.** The analyst's tailored retention questions are no longer shown in that case. Learners who skip the session enforcer lose this check, though the day and scheduled enforcers still cover it.
- **Files.** `routes/learn.js:676-720`, `services/sessions.js:1510-1561` (retentionPlan), `services/study-sessions.js` (is a session running), `services/enforcers.js:186-224` (unitsFrom, to take the flag).

02-P4. Drive lesson reviews from the fitted curve, not the SM-2 ladder

Status: BUILT (flag `CURVE_REVIEWS`, default on; 2026-10-25, implementer B).

- Where the grade has an objective model, planLesson's review share (up to 30%) comes from `adaptation.reviewQueue`'s memory checks, fading and least durable objectives (`sessions.curveReviews`). These are mapped to introduced skills, at most two per skill.
- Each review item is tagged with its objective (`objective_id`, `review_reason`) and uses a question written for that objective when one is left, so the answer is that objective's memory check.
- With no model, `mastery.dueReviews` (SM-2) decides, as before. An empty queue means nothing is due.
- The owner keeps SM-2.
- About 1 ms per plan on the live data.
- Tests: `scripts/curriculum-evidence-test.js`.
- **What.** In `planLesson`, where an objective model exists, pick the due-review share (up to 30%) from `adaptation.reviewQueue`'s fading objectives, mapped to their skills. Keep `mastery.dueReviews` as the fallback, so one scheduler decides what is due everywhere.
- **Evidence:**
    - **mettler2016** [171]: "adaptive scheduling outperformed fixed conditions at immediate and delayed tests of retention" (adults, lab)
    - **tabibian2019** [26]: "learners who follow a reviewing schedule determined by our algorithm memorize more effectively than learners who follow alternative schedules" (natural experiment on Duolingo data)
    - **lindsey2014** [27]: "personalized review yielded a 16.5% boost in course retention over current educational practice (massed study)" (middle-school foreign-language course, time-matched)

    - **pavlik2008** [213]: "The optimized condition showed significant benefits with large effect sizes for both improved recall and recall latency." (college students, Japanese-English vocabulary)
    - **ye2022** [214]: "Experimental results have shown a 12.6% performance improvement over the state-of-the-art methods." (220 million review logs from a vocabulary app)
- **Expected effect.** Modest: no second schedule that ignores what enforcers just retrieved, and review aimed by predicted recall.
- **Risk.** Lesson composition changes. Objective-to-skill mapping is coarse. Test it on the simulator first.
- **Files.** `services/sessions.js:131-141`, `services/mastery.js:144-151`, `services/adaptation.js:288-327`.

02-P5. Build the "left out last time" rotation flag (low priority)

BUILT (flag `ENFORCER_LEFT_OUT_FLAG`, on by default; 2026-10-25, implementer A).

- **Why it was built.** It was built on anderson1994 [262] and murayama2014 [263]. A third source, cepeda2009 [193], was removed in the 27 October 2026 audit because it concerns spacing of material that is practised, not rotating topics, so this change now rests on two sources. chan2006 [153] is counter-evidence, so the weight stays the params file's small +0.10. The risk is very small.
- **How it works.** `plan()` stores the eligible topics it did not ask on the enforcer row (`left_out`). `leftOutCounts` counts, for each topic, the run of the learner's latest planned enforcers that left it out; a plan that asked it ends the run. `score` adds 0.10 × min(1, times / 2). `fromResearch` now parses the params string. The owner's base version stores and reads nothing.
- **Tests:** `scripts/retention-enforcers-test.js`, section "P5".
- **What.** The params file defines +0.10 × min(1, times left out / 2), but neither `services/enforcers.js:280-288` nor `fromResearch` reads it. Store the eligible-but-not-chosen keys on each enforcer row in `plan()` and add the flag.
- **Evidence:**
    - **anderson1994** [262]: "Category-cued recall of unpracticed members of practiced categories was impaired on a delayed test." (category-exemplar lists)
    - **murayama2014** [263]: "Retrieving a subset of items can cause the forgetting of other items" (meta-analysis of retrieval-induced forgetting)
- **Counter-evidence:**
    - **chan2006** [153]: "testing enhanced recall 24 hr later for the initially nontested material" (prose)
    - Retrieval-induced forgetting is not universal, which is why this is low priority.
- **Expected effect.** Small: fewer topics left out twice in a row in tight budgets.
- **Risk.** Very small.
- **Files.** `services/enforcers.js:280-288`, `:562-597`; `services/retention-params.js:62`, `:118-125`.

Considered, not proposed

- **Holding lesson feedback by preference** (mechanism 44). The timing evidence is mixed, so no change is justified either way.
- **An evening prompt for the day enforcer.** Sleep favours restudied over retrieved items bauml2014 [252]. The evidence document should mark this as not built (mechanism 46).
- **Putting confusable maths concepts side by side.** Interleaved mathematics practice is supported (brunmair2019 [14]; rohrer2015 [12]), but needs confusability data the bank does not have.

Corrections to RETENTION-ENFORCER-EVIDENCE.md found here

- **Section 3.8:** "the app suggests doing the day enforcer in the evening" is not in the code or UI. Mark it "(not built)".
- **Section 5, step 2:** the 0.65 step is listed, but the params file's `difficulty_window` text omits it. P1 removes the step.
- **Section 5:** the "Left out last time" and "Prerequisite" flags are correctly marked not built, and the "within 14 days" limit on last-wrong is also not built. Both are already noted there.
- **Applied 2026-10-25 (implementer A).** Section 3.8 now marks the evening suggestion as not built. Section 5, step 2 states the new window. The params `difficulty_window` text now reads 1.0 up to 0.95, with `difficulty_window_basis` naming eglington2020. The left-out flag is marked built in section 5.

Section 03. Scorecard and adaptation

This section covers:

- the learner scorecard (`services/scorecard/*`) and its insertion contract (`contract.js`);
- how the scorecard drives teaching: `ADAPTIVE_TEACHING` and the directives (`services/adaptation.js`, `services/learner.js` sessionParams);
- the effort monitor (`ADAPT_EFFORT`, `services/effort-monitor.js`);
- the coach guards (`services/ai.js`, `services/assist.js`);
- the style quiz, the about-you survey and onboarding;
- owner base vs optimized (`config/owner.js`, `services/owner-skips.js`): product behaviour, no evidence needed;
- grading leniency (accents, number forms) and the grading injection check;
- objective items;
- the young-learner oral mode (`YOUNG_ORAL`);
- the lesson sandbox (`LESSON_SANDBOX`).

The forgetting curve, review timing and the daily review are section 02. Mechanisms 28 (typed recall) and 39 (the 1.2 s rule) overlap section 02's mechanisms 35 and 28; the verdicts agree.

Design choices, stated plainly. No study sets any of these numbers:

- edge 20, MIN_CONF 30, GATE_Z 2, the MIN_EVIDENCE counts, SIDE_P 0.98, MEAN_Z 0.88, VALIDITY 0.9;
- the 6 and 8 directive caps;
- the knob sizes (+/-0.05 to 0.25), the 0.15 base shares and the 0.05 to 0.35 clamps, item counts 4 to 18;
- the 1.2 s rapid line, the 0.3 x expected-time line, 2 reserves, the 5% / 95% / 30-answer rush flag;
- coached weight 0.25, XP x 0.6, attestation weight 0.25;
- questionnaire weight 0.6, the 0.75 cap on observed weight;
- the split-half 0.6 cut-off.

The 10 weak rules: what the new research says

#	Rule (where)	New evidence	Decision
1	Typed answers for a rushing learner (`adaptation.js:103`, `contract.js:66`, `learner.js:300`)	greving2018 [264] found short answer helped and multiple choice did not. smith2014 [116] found no advantage, kang2007 [166] found it depends on feedback, and little2012 [133] shows good MC works. No study tests format against rushing. The effort monitor (mechanism 41, wise2006cbt [265], wise2019 [266]) already deals with rushers directly.	**Remove** the `typed: true` knob, or run it only as an experiment arm (P1)
2	At most 6 directives (`translate.js:2`	he2024cello [267]: LLMs "often ignore semantic constraints". liu2024lost [124], levy2024 [125] and	**Keep** as a design choice. Make the coach cap match

#	Rule (where)	New evidence	Decision
	9; coach 8, `contract.js:73`)	shi2023 [126] show that length, position and irrelevant text cost accuracy. No paper tests a cap near 6.	(6), or test both caps on the app's own writer (P9)
3	The 85% challenge target (challenge_appetite; `sessions.js:1484-1502`)	wilson2019 [94] is a machine-learning result. The human studies point to high success: jansen2013 [32] (children: a higher pre-set success rate meant more practice and bigger gains), eglington2020 [28] (practice at 0.86-0.94 predicted recall beat 0.40 and 0.98), lomas2013 [268] (easier was more engaging). None supports a harder extension for learners who say they like a challenge.	**Soften** the high side: stretch +0.1 to +0.05, stretch clamp 0.35 to 0.25 (P2). Cite the human studies, not wilson2019 [94] alone
4	Brevity (`adaptation.js:97`, `contract.js:60`)	mayer1996 [107]: summaries alone did as well as or better than the full text. sundararajan2020 [128] and harp1998 [127]: seductive details hurt. rey2019 [137]: segmenting helps. kirschner2013 [269]: learners do not know best, so a stated wish for "full detail" is weak grounds for longer text.	**Keep** the high side. **Soften** the low side: lengthen only for an objective that is new to this learner (P3)
5a	Split-half r >= 0.6 (`SCORECARD-EVIDENCE-REPORT.md:545`)	koo2016 [270]: for ICCs, below 0.5 is poor and 0.5-0.75 moderate, judged on the 95% CI. cicchetti1994 [271] and lance2006 [272]: cut-offs are rules of thumb. Spearman-Brown correction (brown1910 [273], spearman1910 [274] title-level) is needed for a split-half; parsons2019 [66] recommends reporting split-half reliability.	**Soften and specify**: Spearman-Brown corrected, averaged over random splits, point >= 0.6 **and** lower 95% bound >= 0.5 (P7)
5b	Posterior >= 0.98 (`translate.js:59`)	newton2004 [275] and efron2008 [276]: thresholding posterior probabilities defines the expected false-discovery share, so a calibrated P >= 0.98 means at most about 2% of issued directives are on the wrong side. kass1995 [277]: posterior odds (here 49:1). berry2006 [278]: Bayesian thresholds on accruing data.	**Keep**. The basis is now principled, not only the simulation. Its calibration still needs the section 11 real-data check
6	Logit counting variance (`index.js:529-541`)	The code is the Haldane-Anscombe empirical logit, log((k+0.5)/(n-k+0.5)), with its standard variance 1/(k+0.5)+1/(n-k+0.5): haldane1956 [279], anscombe1956 [280] (title-level), gart1967 [281] (small-sample bias of these estimators), warton2011 [282] (logit over arcsine).	**Keep**. Now **strong**. Fix the report: it is a textbook estimator, not just a simulation choice
7	Hint dependency / help seeking (`stats.js:218-223`)	aleven2016help [34]: better help use did not raise learning. koedinger2007 [33]: the assistance dilemma is open. baker2004 [283]: gaming the system exploits feedback and help and is associated with reduced learning. Worked examples are better supported (renkl2014 [8], salden2010 [110]).	**Soften** the hint_dependency text to a worked example plus a try-first prompt (P5). **Keep** help_seeking
8	Low self-efficacy (`stats.js:264-266`)	usher2008 [284] and byarswinston2017 [285] (k = 61): mastery experience is the strongest source. bandura1981 [286]: proximal subgoals raised mastery and self-efficacy. schunk1983 [287]: ability feedback on progress raised both. talsma2018 [288]: performance drives later self-efficacy, which supports "make step one achievable".	**Keep**, and strengthen the text (P6). The measure ("sure" taps on right answers) stays a proxy; validity 0.8 is already applied

#	Rule (where)	New evidence	Decision
9	Misconception persistence (`stats.js:351-353`)	The fix is strong: schroeder2022 [129] (g = 0.41), tippett2010 [131], kendeou2018 [289]. The premise is mixed: shtulman2012 [290] says naive theories survive for years; payne1990 [291] says mal-rules are "very unstable" within a student; smith1994 [292].	**Keep**. Contrasting the error helps whoever gets it. Read the stat as "this tag recurs", not as a trait
10	Frustration, persistence, momentum	Frustration: dmello2012 [293] (unresolved frustration leads to boredom); camachomorles2021 [294] found frustration's link to performance near zero (rho = -0.02), anger -0.35. Persistence: botelho2019 [295], ventura2013ce [296]. Momentum: mace1988 [297] and lee2005 [298] are about compliance, not learning; aamodt1992 [299] (easy-first ordering, d =.11).	Frustration: **keep** the gentle ramp (low risk), but the "rapid after a miss" proxy is unvalidated. Persistence: **keep**, moderate. Momentum: **keep off** (it never fires); design choice

In numbers: 61 mechanisms; 57 with 3 or more quoted sources; 9 strong, 25 moderate, 7 mixed, 18 design choice, 2 product rule; 165 distinct sources cited or listed; 12 mechanisms built or changed on 25 October 2026.

03 at a glance

Mechanism	Verdict	Quoted sources	Status
03-M1 Score gate before a population exists	moderate	4	existing
03-M2 Minimum evidence of each kind	design choice	4	existing
03-M3 Standing gate: P(on that side) >= 0.98 and |z| >= 0.88	design choice	5	existing
03-M4 Behaviour is not trait: the posterior widened at validity 0.9	design choice	4	existing
03-M5 Memory-kind and secondary rates never steer	design choice	3	existing
03-M6 Empirical-Bayes standing among learners	strong	4	existing
03-M7 Logit counting variance for rates	strong	3	existing
03-M8 Split-half r >= 0.6 real-data check	design choice	4	built on 25 October 2026; open proposal
03-M9 Rank by impact x distance x certainty; cap at 6 (coach 8)	design choice	5	existing; open proposal
03-M10 Combined instructions for pairs that pull apart	design choice	3	existing
03-M11 Hand-written directive text; no learner numbers in model text	moderate	4	existing
03-M12 One versioned contract; it fails to the new-learner default; never a skip	design choice	3	existing
03-M13 Style quiz: 20 forced choices between concrete behaviours	moderate	3	existing
03-M14 The questionnaire is a weak prior with no measured weight	strong	4	existing

Mechanism	Verdict	Quoted sources	Status
03-M15 Stated preferences blended into session settings, never fully overruled	mixed	3	built on 25 October 2026
03-M16 About-you survey; interest examples only when interests are named	moderate	3	existing
03-M17 Guidance from the learner's own record on the objective	strong	4	existing
03-M18 A scorecard reading cannot take worked examples away from an unseen objective	moderate	3	existing
03-M19 Length knob	mixed	4	built on 25 October 2026
03-M20 Items per block: multipliers, 4 to 18, compressed at most 6	design choice	4	existing
03-M21 Stretch and support shares	mixed	4	built on 25 October 2026
03-M22 challenge_appetite directive text	mixed	3	built on 25 October 2026
03-M23 Frustration low: a gentler ramp and more support	mixed	3	existing
03-M24 Frustration measured as rapid answers right after a miss	design choice	4	existing
03-M25 Persistence low: payoff early, short sections, 20% fewer items	moderate	3	existing
03-M26 Persistence measured as stopping after 3+ misses in a row	moderate	4	existing
03-M27 Momentum: an early easy win	design choice	4	existing
03-M28 A rushing learner gets typed answers	mixed	5	built on 25 October 2026
03-M29 Rush directive text: the key sentence alone, one "slow down" warning	moderate	3	existing
03-M30 Hint dependency and help seeking directives	mixed	6	built on 25 October 2026
03-M31 Low self-efficacy: an achievable first step, named as achieved	moderate	6	built on 25 October 2026
03-M32 Misconception persistence: state, refute, replace	strong	6	existing
03-M33 Interest context knob	moderate	3	existing
03-M34 Recap knob (consolidation need)	strong	4	existing
03-M35 Practice ends on proof: compress or extend	strong	3	existing
03-M36 Blocked while new, interleaved once proven	moderate	3	existing
03-M37 Proof check with questions written for the objective	moderate	3	existing
03-M38 Session summary at 85% / 60%; difficulty_delta computed but unused	design choice	2	existing
03-M39 Answers under 1.2 s are not evidence	moderate	4	built on 25 October 2026

Mechanism	Verdict	Quoted sources	Status
03-M40 Relative rapid line: under 0.3 x expected time	design choice	4	existing
03-M41 "Too fast to count" on checkpoints and the Final	moderate	4	built on 25 October 2026
03-M42 Rush flag and the one-line slow-down note	moderate	4	existing
03-M43 The coach never gives the answer first; leak guard	moderate	4	built on 25 October 2026
03-M44 Coach turn rules: one question, the smallest next step	moderate	4	existing
03-M45 The coach stays on the question on screen	design choice	2	existing
03-M46 Coached right answers weigh 0.25; XP x 0.6	design choice	4	existing
03-M47 No assist on checkpoints and essays	moderate	3	existing
03-M48 Form leniency: case, accents, $, punctuation, "x ="	moderate	4	existing
03-M49 Number forms: 0.5 = 1/2, mixed numbers, strict thousands commas	moderate	3	existing
03-M50 Strict on exactly the form the prompt asks for	moderate	3	existing
03-M51 Spanish: optional subject pronoun, el/él and tu/tú kept apart	design choice	3	existing
03-M52 Grading injection check and fenced judge input	moderate	3	existing
03-M53 Objective-item code checks	strong	5	built on 25 October 2026
03-M54 Two blind judges from different model families agree on the key	design choice	3	existing
03-M55 Objective items served first; aligned right answers needed for proof	moderate	3	existing
03-M56 Spoken checkpoints for preschool and Kindergarten	moderate	4	existing
03-M57 A grown-up's "I heard it" weighs 0.25 and never proves alone	design choice	4	existing
03-M58 The recording stays on the device	moderate	3	existing
03-M59 Lessons in a sandboxed frame with a strict CSP	strong	4	existing
03-M60 Owner base vs optimized	product rule	0 (flagged)	product rule
03-M61 Owner skips	product rule	0 (flagged)	product rule

03-A. The scorecard gate and the insertion contract

03-M1. Score gate before a population exists

- **Where:** `services/scorecard/translate.js:25-33, 194-209`
- **What it does:** with no population yet, a stat steers only when all four hold:
  - its 0-100 score is at least 20 from 50;
  - its confidence is at least 30;
  - it is 2 standard errors from 50;
  - the trait reading still puts it on that side with P >= 0.98.
- **Evidence:**
  - gelman2012 [30] — "the problem of multiple comparisons can disappear entirely when viewed from a hierarchical Bayesian perspective" — education research methods.
  - efron1975 [29] — "This estimator is reviewed briefly in an empirical Bayes context." — statistics.
  - benjamini1995 [300] — "It calls for controlling the expected proportion of falsely rejected hypotheses — the false discovery rate." — statistics.
  - parsons2019 [66] — "subpar reliability has severe implications for measures' validity and interpretation" — cognitive measures.
- **Verdict:** moderate. The principle (shrink noisy readings, guard against many comparisons) is strong; the numbers are chosen.
- **Code vs evidence:** matches. The 2-SE rule is a frequentist stand-in, used only until the population gate (mechanism 3) takes over.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M2. Minimum evidence of each kind

- **Where:** `translate.js:34-49, 95-106`
- **What it does:** a stat needs enough data of its own kind before it may steer. Examples: 8 weighted updates, 30 rate chances (10 once calibrated), 15 contrast answers, 20 memory tests, 150 "reach" answers.
- **Evidence:**
  - hedge2018 [67] — "Experimental effects become well established - and thus those tasks become popular - when between-subject variability is low." — cognitive tasks.
  - rouder2019 [301] — "correlations among tasks that seemingly have processes in common are often low" — cognitive tasks; the paper asks whether trial-level noise and low individual variability cause this, which is the reason to ask for enough evidence before reading a difference.
  - parsons2019 [66] (above).
  - enkavi2019 [302] — "Accurate characterization of individual differences requires measuremen[t reliability]" — adults, self-regulation tasks.
- **Verdict:** design choice. The literature says thin behavioural data is unreliable for individuals; the counts come from the simulation.
- **Code vs evidence:** matches the direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M3. Standing gate: P(on that side) >= 0.98 and |z| >= 0.88

- **Where:** `translate.js:50-60, 186-193` (weak rule 5b)
- **What it does:** once there is a population, a directive needs both:
  - a posterior probability of at least 0.98 that the learner is on that side of the typical learner;
  - a best estimate in about the top or bottom fifth.
- **Evidence:**
  - newton2004 [275] — "Gene-specific posterior probabilities of differential expression form the basis of statistical inference; they define short gene lists and false discovery rates" — many simultaneous decisions (genomics).
  - efron2008 [276] — "with particular attention focused on Benjamini and Hochberg's False Discovery Rate method" — the empirical-Bayes two-groups model.
  - kass1995 [277] — "the Bayes factor, which is the posterior odds of the null hypothesis when the prior probability on the null is one-half" *(fetched during the review)* — statistics.
  - berry2006 [278] — "Accumulating results can be assessed at any time, including continually, with the possibility of modifying the design of the trial" — Bayesian adaptive trials.
  - gelman2012 [30] (above).
- **Verdict:** design choice, now with a principled basis. Taking only decisions whose posterior is at least 0.98 bounds the expected share of wrong-side directives at about 2%, provided the posterior is calibrated (newton2004 [275]; efron2008 [276]). 0.98 is posterior odds of 49:1.
- **Code vs evidence:** matches. The open question is calibration on real learners: the section 11 check (95% of directives still on the same side 30 levels later) stays required.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M4. Behaviour is not trait: the posterior widened at validity 0.9

- **Where:** `translate.js:61-71`
- **What it does:** reads the measured behaviour as if it correlated 0.9 with the trait behind it (0.8 for self-efficacy). This widens the posterior before the gate.
- **Evidence:**
  - schmidt1996 [303] — "This article examines 26 real-world case studies and explicates them on the basis of the principles of reliability theory" — measurement.
  - messick1995 [151] — "The new unified concept of validity interrelates these issues as fundamental aspects of a more comprehensive theory of construct validity" — measurement.
  - kane2013 [105] — "To validate an interpretation or use of test scores is to evaluate the plausibility of the claims based on the scores." — measurement.
  - hedge2018 [67] (above).

- **Verdict:** design choice. The direction is right (attenuation; validity is about the use); the 0.9 is set on the simulation.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M5. Memory-kind and secondary rates never steer

- **Where:** `translate.js:72-78, 171-179`
- **What it does:** two kinds of stat are shown but never acted on:
  - decay, retention and retrieval;
  - rates read inverted as a second stat.
- **Evidence:**
  - kane2013 [105] (above).
  - messick1995 [151] (above).
  - parsons2019 [66] (above).
- **Verdict:** design choice. A measure is acted on only once the use is validated; the simulation found these posteriors over-confident.
- **Code vs evidence:** matches. `SCORECARD_MEMORY_STEER=on` exists only for re-measuring.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M6. Empirical-Bayes standing among learners

- **Where:** `services/scorecard/index.js:180-215, 529-560`
- **What it does:** each learner's reading is shrunk toward the population, by its own noise against the population's true spread (method of moments).
- **Evidence:**
  - efron1975 [29] (above).
  - morris1983 [63] — "reviews the state of multiparameter shrinkage estimators with emphasis on the empirical Bayes viewpoint" — statistics.
  - raudenbush1985 [304] — "This strategy enables the analyst to (a) estimate the variance of the effect size parameters" — education meta-analysis.
  - dersimonian1986 [305] — "We discuss a random effects approach to combining evidence from a series of experiments comparing two treatments." — clinical-trial random-effects modeling of between-study heterogeneity, not learner-level shrinkage.
- **Verdict:** strong.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M7. Logit counting variance for rates

- **Where:** `index.js:529-541` (weak rule 6)

- **What it does:** a rate of k events in n chances becomes the logit of (k + 0.5)/(n + 1), with variance 1/(k + 0.5) + 1/(n - k + 0.5). That logit is exactly log((k+0.5)/(n-k+0.5)).
- **Evidence:**
  - haldane1956 [279] — title-level only — the +0.5 correction to the log ratio of frequencies.
  - anscombe1956 [280] — title-level only — the same correction for binomial responses.
  - gart1967 [281] — "The bias of several logit estimators and their corresponding variance estimators is investigated in small samples." — statistics.
  - warton2011 [282] — "The logit transformation is proposed as an alternative approach to address these issues." — analysis of proportions.
  - dersimonian1986 [305] (above): random-effects pooling of estimates from related studies.
- **Verdict:** strong. This is the standard empirical logit and its standard variance estimate, pooled the way meta-analysis pools log-odds.
- **Code vs evidence:** matches. The report (section 7) says the choice "rests on the simulation"; it should say it is the Haldane-Anscombe estimator, with gart1967 [281]'s small-sample bias caveat. The small-sample caveat is relevant to keeping a minimum-evidence threshold; the specific 10+ threshold is a design choice.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M8. Split-half r >= 0.6 real-data check

- **Where:** `curriculum-build/SCORECARD-EVIDENCE-REPORT.md:545` (weak rule 5a)
- **What it does:** this is a planned check, not code yet. Each learner's evidence is split by odd and even days, the two standings are correlated, and a stat that issues directives must reach r >= 0.6.
- **Evidence:**
  - koo2016 [270] — "values less than 0.5, between 0.5 and 0.75, between 0.75 and 0.9, and greater than 0.90 are indicative of poor, moderate, good, and excellent" — ICC reliability guideline (judged on the 95% CI).
  - cicchetti1994 [271] — "Criteria, guidelines, and simple rules of thumb are provided to assist the clinician faced with the challenge of choosing an appropriate test instrument" — assessment guideline.
  - lance2006 [272] — "trace four widely cited and reported cutoff criteria to their (alleged) original sources" — methods review.
  - parsons2019 [66] (above).
  - brown1910 [273] and spearman1910 [274] — title-level only — the Spearman-Brown step-up for split halves.
- **Verdict:** design choice.
- **Code vs evidence:** as written, it leaves out two things:
  - the Spearman-Brown correction: a raw half-data r understates full-data reliability;

    - the uncertainty: koo2016 [270] grades on the confidence interval, not the point estimate.

0.6 corrected is "moderate", which is reasonable for low-stakes, reversible directives. See P7.

- **Status: Built on 25 October 2026** (no switch: a check, from 03-P7). `services/scorecard/reliability.js` `splitHalf`: Spearman-Brown corrected, 500 random day splits, passes at point >= 0.6 and lower 95% bound >= 0.5. Running it needs real learners. **Open:** running the check on real learners (needs data: 20+ learners with 30+ levels).

03-M9. Rank by impact x distance x certainty; cap at 6 (coach 8)

- **Where:** `translate.js:27-29, 214-260`; `contract.js:71-73, 206` (weak rule 2)
- **What it does:**
    - ranks the confident directives so the most important comes first;
    - sends at most 6 to the lesson writer and 8 to the coach;
    - adds a certainty note (`translate.js:269-288`).
- **Evidence:**
    - liu2024lost [124] — "significantly degrades when models must access relevant information in the middle of long contexts" — LLMs.
    - levy2024 [125] — "a notable degradation in LLMs' reasoning performance at much shorter input lengths than their technical maximum" — LLMs.
    - shi2023 [126] — "model performance is dramatically decreased when irrelevant information is included" — LLMs (preprint).
    - he2024cello [267] — "LLMs often ignore semantic constraints from task descriptions, generate incorrect formats, violate length or sample count constraints" — LLM benchmark.
    - jaroslawicz2025 [306] — "even the best frontier models only achieve 68% accuracy at the max density of 500 instructions" — LLM benchmark (preprint).
- **Verdict:** design choice. Fewer, front-loaded instructions are the right direction; no study tests about 6.
- **Code vs evidence:** matches the direction. The coach's 8 has no stated reason to differ from 6 (P9).
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 03-P9, testing the directive caps on the app's own writer (needs data).

03-M10. Combined instructions for pairs that pull apart

- **Where:** `translate.js:108-141, 233-256`
- **What it does:** 14 hand-written texts replace two directives that would conflict (for example, show every step + be brief). A combined text ranks with the stronger part's priority.
- **Evidence:**
    - he2024cello [267] (above).

    - jiang2024follow [307] — "a Multi-level mechanism that incrementally adds a single constraint to the initial instruction at each increased level" — LLM benchmark (design only; it tests following added constraints, not resolving conflicting ones).
    - qin2024 [308] — "particularly in complex instruction-following" — LLM benchmark.
- **Verdict:** design choice. Multi-constraint prompts are where models fail; resolving a conflict before the model sees it removes one kind of complexity. No study tests this exact move.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M11. Hand-written directive text; no learner numbers in model text

- **Where:** `services/scorecard/stats.js:24-28`; `contract.js:32-40, 75-82`
- **What it does:**
    - the model gets fixed sentences only, never scores, the evidence log, survey answers, the name or the email;
    - prompt lines go only to providers allowed learner data.
- **Evidence:**
    - gdpr-art5-1c [309] — "adequate, relevant and limited to what is necessary in relation to the purposes for which they are processed" — EU regulation (from section 05).
    - slade2013 [310] — "The field of learning analytics has the potential to enable higher education institutions to increase their understanding of their students' learning needs" — ethics of learning analytics.
    - drachsler2016 [311] — "Privacy and analytics: it's a DELICATE issue a checklist for trusted learning analytics" — LA privacy checklist.
    - dwork2013 [312] — "the need increases for a robust, meaningful, and mathematically rigorous definition of privacy" — privacy theory.
- **Verdict:** moderate (data minimisation is a legal standard; the learning effect of fixed wording is untested).
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M12. One versioned contract; it fails to the new-learner default; never a skip

- **Where:** `contract.js:1-47, 180-232`
- **What it does:**
    - every consumer reads one object, so they cannot disagree;
    - on any error, the status is `fallback`: no directives and the default curve;
    - there is no field that could mark content as known: coverage stays on bank evidence.

- **Evidence:**
  - saltzer1975 [313] — "Section I describes desired functions, design principles, and examples of elementary protection and authentication mechanisms." — security design (fail-safe defaults is one of its principles).
  - kulik1990 [4] — "mastery learning programs have positive effects on the examination performance of students" — 108 evaluations.
  - bloom1984 [5] — "Mastery Learning. Students learn the subject matter in a class with about 30 students per teacher." — classroom studies.
- **Verdict:** design choice. Failing safe, and advancing only on demonstrated mastery, are both well founded; the contract applies them.
- **Code vs evidence:** matches. But `learner.effective()` reads the scorecard outside the contract (mechanism 15).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-B. Style quiz, survey and onboarding

03-M13. Style quiz: 20 forced choices between concrete behaviours

- **Where:** `public/js/style-quiz.js:1-3`; `services/scorecard/index.js:1408-1434`
- **What it does:** each item asks which of two concrete behaviours the learner does, instead of a 1-5 self-rating. Skipping is allowed.
- **Evidence:**
  - brown2011fc [314] — "Multidimensional forced-choice formats can significantly reduce the impact of numerous response biases typically associated with rating scales." — psychometrics.
  - cao2019 [315] — "FC scales were consistently found to be more faking-resistant when constructed with statements balanced in social desirability" — meta-analysis.
  - nisbett1977 [316] — "there may be little or no direct introspective access to higher order cognitive processes" — review.
- **Verdict:** moderate.
- **Code vs evidence:** mostly matches. brown2011fc [314] also warns that forced choices scored classically "produce ipsative data". Here each choice is read as one dimension's 0-1 value, which is classical scoring. That is acceptable only because the quiz is a weak prior (mechanism 14), but see mechanism 15.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M14. The questionnaire is a weak prior with no measured weight

- **Where:** `index.js:1408-1434`; `contract.js:33-34`
- **What it does:** quiz answers nudge the shown score (weight 0.6) but add no measured mass, so on its own the quiz can never make a directive confident.
- **Evidence:**

    - pashler2008 [31] — "the instructional method that proves most effective for students with one learning style is not the most effective method for students with a different learning style" (the test meshing fails) — review.
    - rogowsky2015 [317] — "failed to statistically support the meshing hypothesis either for verbal comprehension aptitude or learning based on mode of instruction" — adults.
    - kirschner2013 [269] — "The three legends can be seen as variations on one central theme, namely, that it is the learner who knows best" — review.
    - nisbett1977 [316] (above).
- **Verdict:** strong.
- **Code vs evidence:** matches in the scorecard.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M15. Stated preferences blended into session settings, never fully overruled

- **Where:** `services/learner.js:233-257, 274-306`
- **What it does:** `effective()` sets these from the quiz's stated value, blended with observed behaviour:
    - item count, stretch and support shares;
    - typed-answer preference, feedback timing, explanation depth.

Before 25 October 2026 the observed weight was capped at 0.75, so "a stated preference is never fully overruled" (the cap is now 1.0; see Status). Then, at scorecard confidence >= 30, the measured score is blended in. That is outside the calibrated gate.

- **Evidence:**
    - pashler2008 [31] (above).
    - rogowsky2015 [317] (above).
    - kirschner2013 [269] (above).
- **Verdict:** mixed. The evidence says self-reported preference is a poor basis for tailoring. This code keeps at least 25% stated weight forever, and lets an ungated reading move structure.
- **Code vs evidence: gap**. The contract's rule ("nothing the questionnaire nudged steers") holds for directives, but not for these session parameters. See P4.
- **Status: Built on 25 October 2026** (switch `SCORECARD_OBSERVED_PREFS`, from 03-P4). Observed weight is capped at 1.0 (was 0.75), and the scorecard term is blended in only when the contract's calibrated gate passes it. A shadow mode is available.

03-M16. About-you survey; interest examples only when interests are named

- **Where:** `services/learner.js:96-150`; `routes/profile.js:107-125`; `translate.js:213`; `contract.js:180-232` (`hasInterests`)
- **What it does:**
    - five free-text boxes (interests, studying, hobbies, sports, notes), all optional;
    - the interest directive is dropped when no interest was named;

    - the answers are never sent as numbers.
- **Evidence:**
    - walkington2013 [152] — "students in the personalization condition solved problems faster and more accurately within the modified unit" — 145 algebra students.
    - bernacki2018 [318] — "High school students in the personalization condition reported greater triggered situational interest in experimental units" — high school maths.
    - cordova1996 [319] — "examined the effects on the learning process of 3 complementary strategies—contextualization, personalization, and provision of choices" — elementary pupils.
- **Verdict:** moderate.
- **Code vs evidence:** matches. Interests are asked, not inferred, and are used only when given.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-C. How the scorecard drives teaching (ADAPTIVE_TEACHING, default on)

`adaptation.js:47-51`: `on` applies, `shadow` computes and records, `off` computes nothing. The owner account is never adapted (mechanism 60).

03-M17. Guidance from the learner's own record on the objective

- **Where:** `services/adaptation.js:445-458`
- **What it does:**
    - **worked examples first:** no answers yet, belief under 0.5 (`ADAPT_NOVICE_BELOW`), or lapsed;
    - **faded examples:** belief in between;
    - **problems first:** proven or retained (and the teaching text is one step shorter).
- **Evidence:**
    - kalyuga2003 [6] — "Instructional techniques that are highly effective with inexperienced learners can lose their effectiveness and even have negative consequences when used with more experienced learners." — review.
    - tetzlaff2025 [61] — "Low prior knowledge learners learn better from high-assistance instruction (d = 0.505)." — meta-analysis, 60 studies, 5,924 participants.
    - renkl2014 [8] — "Learning from examples is a very effective means of initial cognitive skill acquisition." — review.
    - salden2010 [110] — "Both experiments provide evidence of improved learning results from adaptive fading of worked examples" — tutored problem solving.
- **Verdict:** strong.

- **Code vs evidence:** matches. tetzlaff2025 [61] says the evidence is less clear for younger students and some fields of study, including humanities and language learning: worth a holdout check there.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M18. A scorecard reading cannot take worked examples away from an unseen objective

- **Where:** `adaptation.js:95-96, 466-476`
- **What it does:** prior_knowledge and scaffolding_fade move guidance one step either way. A move toward less guidance is refused on an objective with no answers yet, unless the move comes from prior_knowledge itself.
- **Evidence:**
  - tetzlaff2025 [61] (above).
  - kalyuga2007 [84] — "This paper reviews recent empirical findings associated with the expertise reversal effect, their interpretation within cognitive load theory, relations to ATI studies" — review.
  - salden2010 [110] (above).
- **Verdict:** moderate.
- **Code vs evidence:** matches. Expertise is specific to the task, so the objective's own record outranks a general trait.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M19. Length knob

- **Where:** `adaptation.js:97-98, 102, 477, 485-486, 582-586`; `contract.js:60` (weak rule 4)
- **What it does:**
  - **shorter teaching:** brevity high, fast pace, short attention runway, problems-first;
  - **fuller teaching:** brevity low ("wants full detail").
- **Evidence:**
  - mayer1996 [107] — "recalled these steps and solved transfer problems as well as or better than students who received the full text along with the summary or the full text alone" — college students, science lessons.
  - sundararajan2020 [128] — "learners exposed to interesting and irrelevant information, known as seductive details, do not perform as much as those who learned without seductive details" — meta-analysis.
  - harp1998 [127] — "students who read expository passages with seductive details... recalled significantly fewer main ideas" — college students.
  - kirschner2013 [269] (above).
- **Verdict:** mixed:
  - the high side (concise, nothing required dropped) is supported for learners in general;

    - the low side (longer because the learner asked) rests on a stated preference, which kirschner2013 [269] and mechanism 14 say is weak grounds.
- **Code vs evidence:** partly. The low side can lengthen text for a learner who already knows the objective, which expertise reversal argues against. See P3.
- **Status: Built on 25 October 2026** (switch `SCORECARD_LENGTHEN_NEW_ONLY`, from 03-P3). brevity_preference low makes teaching fuller only for an objective taught with worked examples.

03-M20. Items per block: multipliers, 4 to 18, compressed at most 6

- **Where:** `adaptation.js:71-73, 98-102, 526-529`
- **What it does:** confident sides scale the item count:
    - pace high +20%, low -20%;
    - persistence low -20%;
    - attention runway low -25%, high +15%.

The count is clamped to 4-18, and a proven objective gets at most `MIN_ITEMS + 2` = 6.

- **Evidence:**
    - rey2019 [137] (above).
    - bandura1981 [286] — "Under proximal subgoals, children progressed rapidly in self-directed learning, achieved substantial mastery of mathematical operations" — children with arithmetic deficits.
    - rohrer2006 [198]: "this extra effort had no effect on test scores 1 or 4 weeks later" (two experiments, 216 college students, one type of mathematics problem). Scope: Backs only the 'less practice once mastered' half: extra practice added nothing in this task. It did not study learners who stop early.
    - rohrer2005over [261]: "overlearning (and its concomitant demand for additional study time) is an inefficient strategy for learning material for meaningfully long periods of time" (two experiments, 218 college students, recall tested 1 to 9 weeks later). Scope: Backs only the 'less practice once mastered' half. It did not test adaptive block sizes.
- **Verdict:** design choice. Shorter units for those who stop early, and less practice once mastered, both follow the evidence; the percentages do not come from it.
- **Code vs evidence:** matches the direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M21. Stretch and support shares

- **Where:** `adaptation.js:99-100, 544-545`; `learner.js:284-299`; `sessions.js:143-160` (weak rule 3)
- **What it does:**
    - the base stretch share is 0.1 + 0.2 x the quiz's challenge appetite; support is 0.1 or 0.2;
    - challenge_appetite high adds 0.1 stretch; low moves 0.05 from stretch to support;

    - frustration low adds 0.1 support and takes 0.05 stretch;
    - stretch is clamped to 0.05-0.35 (0.05-0.25 since 25 October 2026; see Status) and support to 0-0.35;
    - stretch items run +1.5 difficulty, support items -1.5.
- **Evidence:**
    - jansen2013 [32] — "the higher the pre-set success rate, the more problems were attempted, and the larger the improvement in math performance" — 207 children, grades 3-6, six weeks.
    - eglington2020 [28] — "with up to 40% more items recalled in conditions where practice was scheduled optimally according to the model of practice" *(fetched during the review)* — adults, vocabulary; optimal predicted recall 0.86-0.94.
    - lomas2013 [268] — "subjects were more engaged and played longer when the game was easier, which seems to contradict the generality of the Inverted-U Hypothesis" — large online game experiments.
    - wilson2019 [94] — "the optimal error rate for training is around 15.87% or, conversely, that the optimal training accuracy is about 85%" — machine learners, not people.
- **Verdict:** mixed. The human studies favour a high success rate. Extra stretch for a self-declared challenge seeker is not supported, and a stretch share of 0.35 at +1.5 difficulty pushes success down.
- **Code vs evidence:** partly. Support for low appetite and low frustration matches. The high-side +0.1 and the 0.35 ceiling go further than the evidence. See P2.
- **Status: Built on 25 October 2026** (switch `SCORECARD_SOFT_STRETCH`, from 03-P2). challenge_appetite high adds +0.05 stretch (was +0.1); the stretch share is capped at 0.25 (was 0.35).

03-M22. challenge_appetite directive text

- **Where:** `services/scorecard/stats.js:158-160` (weak rule 3)
- **What it does:**
    - **high:** "a genuinely challenging worked example ... a harder 'try this' extension";
    - **low:** "comfortable difficulty ... build confidence before anything hard".
- **Evidence:**
    - jansen2013 [32] (above).
    - lomas2013 [268] (above).
    - lomas2017 [320] — "even moderate difficulty can reduce player engagement" — online game experiments.
- **Verdict:** mixed. The low side is supported (jansen2013 [32]). The high side is a stated preference with no human evidence that it helps learning.
- **Code vs evidence:** partly. An *optional* extension at the end, after the core, is low-risk. Making the worked example itself hard is not supported. P2 softens the wording.
- **Status: Built on 25 October 2026** (switch `SCORECARD_SOFT_STRETCH`, from 03-P2). New text: teach the core at the usual level, then end with one optional harder "try this" case.

03-M23. Frustration low: a gentler ramp and more support

- **Where:** `stats.js:215-217`; `adaptation.js:100`; combined text `translate.js:129-130` (weak rule 10)
- **What it does:**
  - the text says: "ramp difficulty gently and never leave a gap the learner has to bridge alone";
  - the structure gets +0.1 support and -0.05 stretch;
  - with high challenge appetite, the combined text asks for a hard case in small signposted steps.
- **Evidence:**
  - dmello2012 [293] — "failure to restore equilibrium as well as obstacles that block goals trigger frustration, which, if unresolved, will eventually lead to boredom" *(fetched during the review, ERIC)* — college students, 32-35 min computer tutoring.
  - camachomorles2021 [294] — "For frustration, the relation with performance was near zero ([rho] = - 0.02)." — meta-analysis; frustration 9 samples, N = 1,418 (anger rho = -0.35).
  - pekrun2006 [321] — "appraisals of control and values are central to the arousal of achievement emotions, including activity-related emotions such as enjoyment, frustration, and boredom" — theory review.
  - loderer2020 [322] — title-level only (search snippet: 186 studies of emotions in technology-based learning) — meta-analysis.
- **Verdict:** mixed:
  - control-value theory and dmello2012 [293] support restoring equilibrium when a learner is stuck.
  - camachomorles2021 [294] finds frustration itself barely related to performance.
- **Code vs evidence:** acceptable. The action (more support items, smaller steps) is low risk and fits the support literature. Do not claim frustration harms learning; say unresolved impasses do.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M24. Frustration measured as rapid answers right after a miss

- **Where:** `services/scorecard/signals.js:62-63, 215, 266-270`
- **What it does:** a contrast: the learner's rapid-answer rate just after a miss, against their own rapid rate otherwise. It needs a population before it can steer.
- **Evidence:**
  - wise2017 [59] — "very short response time—termed rapid guessing—has been shown to indicate disengaged test taking" — both high- and low-stakes contexts.
  - baker2004 [283] — "no other type of off-task behavior is associated nearly so strongly with reduced learning as "gaming the system"" — Cognitive Tutor classrooms.
  - dmello2012 [293] (above).

    - ocumpaugh2014 [323] — "The automated detection of these constructs allows EDM researchers to develop intervention strategies" — detectors must be validated on their own population.
- **Verdict:** design choice. A rapid answer after a miss is a plausible marker of disengagement; nothing has validated it as frustration.
- **Code vs evidence:** matches the caution (contrast form, population gate). It still needs validation against self-report before it is called frustration.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M25. Persistence low: payoff early, short sections, 20% fewer items

- **Where:** `stats.js:321-323`; `adaptation.js:101` (weak rule 10)
- **What it does:** the text asks to "put the payoff early and keep each section short enough to finish easily", and the structure cuts items per block by 20%.
- **Evidence:**
    - bandura1981 [286] (above): proximal subgoals beat distal goals for mastery and self-efficacy.
    - rey2019 [137] — meaningful, learner-paced segmentation of multimedia instruction modestly improves retention and transfer and reduces cognitive load (above).
    - botelho2019 [295] — "high persistence is not always productive for students, where additional practice does not help the student move toward a state of mastery" — students on an online maths platform.
- **Verdict:** moderate.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M26. Persistence measured as stopping after 3+ misses in a row

- **Where:** `signals.js:70, 714, 750-785`
- **What it does:** of the moments with three or more misses in a row and questions left, the share where the learner stopped (left the level, or no answer for 30 minutes).
- **Evidence:**
    - botelho2019 [295] (above): stopout and wheel-spinning as the low and unproductive-high ends of persistence.
    - ventura2013ce [296] — "Persistence was measured by recording the time spent on unsolved anagrams and riddles." — online performance task.
    - beck2013 [325] — "rather than a fixed number of practices, students continue to practice a skill until they have mastered it. However, an implicit assumption" — adaptive practice (snippet).
    - kai2018 [326] — "wheel-spinning, a type of unproductive persistence where students spend too much time struggling without achieving mastery of skills" — skill builders.

- **Verdict:** moderate. Behavioural measures of persistence are established. The 3-miss trigger and the 30-minute gap are chosen.
- **Code vs evidence:** matches. As the report notes, real quits have other causes.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M27. Momentum: an early easy win

- **Where:** `stats.js:126-128`; `signals.js:109, 461-462, 497` (weak rule 10)
- **What it does:**
  - the text asks for an easy success in the first section;
  - the measure compares the rest of a level after two opening misses with the same positions after a good start;
  - it needs a population, and it has never been confident.
- **Evidence:**
  - mace1988 [297] — "Results showed the antecedent high-probability command sequence increased compliance and decreased compliance latency and task duration." — five experiments on compliance, not learning.
  - lee2005 [298] — "Overall, high-p request sequences were found to be an effective method to increase compliance." — meta-analysis of applied studies.
  - aamodt1992 [299] — "Exams beginning with easier items and then moving to more difficult items are slightly easier than exams with randomly ordered items (d = .11)" — meta-analysis.
  - miller2018 [327] — "a subtle but substantial bias exists in a common measure of the conditional dependence of present outcomes on streaks of past outcomes" — streak measurement.
- **Verdict:** design choice. Behavioural momentum is real for compliance; easy-first ordering helps a little; no study ties a measured "momentum need" to learning.
- **Code vs evidence:** matches the caution. It cannot fire before a population, and miller2018 [327] is why it is read against the same positions. Keep it off until a holdout test.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M28. A rushing learner gets typed answers

- **Where:** `adaptation.js:103, 481, 546`; `contract.js:66`; `learner.js:300` (weak rule 1)
- **What it does:** rush_tendency high sets `prefer_type_in`, so items are served as typed recall instead of multiple choice.
- **Evidence:**
  - greving2018 [264] — "The results indicated no testing effect for multiple-choice testing." — university lecture, N = 92 (short answer did help).
  - rowland2014 [10] — "initial recall tests yielding larger testing benefits than recognition tests" — meta-analysis.
  - smith2014 [116] — "there were little or no advantages of answering short-answer or hybrid format questions over multiple-choice questions" — students, lab.

    - kang2007 [166] — "having had an intervening MC test led to better performance than an intervening SA test, but... With feedback provided, having had an intervening SA test led to the best" — students, lab.
    - little2012 [133] — "multiple-choice tests could trigger productive retrieval processes-provided the alternatives were made plausible enough" — students, lab.
- **Verdict:** mixed for typed recall in general. Not shown at all for rushers.
- **Code vs evidence:** does not match a tested claim. The effort monitor (mechanism 41), which has direct evidence, already handles rushing. See P1.
- **Status: Built on 25 October 2026** (switch `SCORECARD_DROP_RUSH_TYPED`, from 03-P1). A confident rush reading no longer sets typed answers; the effort monitor handles rushers.

03-M29. Rush directive text: the key sentence alone, one "slow down" warning

- **Where:** `stats.js:132-134`
- **What it does:** asks the writer to put the most important sentence on its own line and add one "slow down here" warning where a hasty reader goes wrong.
- **Evidence:**
    - wise2009 [328] — "items with more text or those occurring later in the test were related to increased rapid guessing" — low-stakes university tests.
    - wise2006cbt [265] — "an effort-monitoring CBT increased examinee effort and yielded more valid test scores than a conventional CBT" — university students.
    - wise2019 [266] — "after a proctor notification was triggered, test-taking engagement" (rose) during tests — K-12 test takers; whether formatting warnings in learning content work similarly is untested.
- **Verdict:** moderate. Less text to skim, plus a targeted warning, both follow from the rapid-guessing work; the exact layout is untested.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M30. Hint dependency and help seeking directives

- **Where:** `stats.js:218-223`; measured at `signals.js:51-52, 311-325` (weak rule 7)
- **What it does:**
    - **hint_dependency high:** "Build gentle hints into the explanation";
    - **help_seeking low:** "Anticipate the learner's questions and answer them in the text";
    - rapid answers when stuck count as frustration, not help seeking.
- **Evidence:**
    - aleven2016help [34] — "feedback on help seeking helped students to use on-demand help more deliberately... but not to achieve better learning outcomes" — ITS studies.
    - koedinger2007 [33] — "How should learning environments balance information or assistance "giving" and "withholding" to achieve optimal student learning?" — review of Cognitive Tutor experiments.

    - aleven2003 [173] — "although effective help-seeking behavior in ILEs is related to better learning outcomes, learners are not using help facilities effectively" — review.
    - ryan2001 [329] — "Why do some students avoid seeking help in the classroom when they need it?" — review of help avoidance.
    - baker2004 [283] — gaming the system can involve exploiting feedback and help to obtain correct answers, and is associated with lower post-test performance.
    - Also: razzaq2010 [330] and mclaren2016 [331] — title-level only (proactive vs on-demand hints; worked examples vs tutored problems).
- **Verdict:** mixed:
    - giving help-avoiders the answers to likely questions fits the help-avoidance work;
    - "build hints in" for heavy hint users is the open assistance dilemma: heavy hint use is partly gaming, and better help use did not raise learning.
- **Code vs evidence:** partly. P5 replaces the hint text with what has better support: a worked example of the exact step, then a try-first prompt (renkl2014 [8], salden2010 [110]).
- **Status: Built on 25 October 2026** (switch `SCORECARD_WORKED_HINT_TEXT`, from 03-P5). The hint_dependency text now opens with one fully worked example, then a try-first step.

03-M31. Low self-efficacy: an achievable first step, named as achieved

- **Where:** `stats.js:264-266`; measured at `signals.js:81, 297-304, 542, 692-697` (weak rule 8)
- **What it does:**
    - the text asks to "Make the first step clearly achievable and point out explicitly that it was achieved; show that the hard part is learnable";
    - the measure is the share of "sure" taps among right answers, at validity 0.8.
- **Evidence:**
    - usher2008 [284] — "Although mastery experience is typically the most influential source of self-efficacy" — review, school settings.
    - byarswinston2017 [285] — "PA showed by far the strongest unique association with self-efficacy beliefs." — meta-analysis, k = 61, N = 8,965 (PA = performance accomplishments).
    - bandura1981 [286] (above).
    - schunk1983 [287] — "Children given only ability feedback demonstrated the highest subtraction skill and self-efficacy" — children deficient in subtraction.
    - talsma2018 [288] — "Performance had a net positive effect on subsequent self-efficacy (β = 0.205...), significantly larger than the effect of self-efficacy on performance (β = 0.071" — meta-analysis.
    - Context: vandinther2011 [332] found that educational programmes based on social cognitive theory could enhance self-efficacy — review, higher education.

- **Verdict:** moderate for the action. The success-first mechanism is the best-supported source of self-efficacy, and the direction (performance to belief) supports it. The measure is a confidence proxy, not self-efficacy itself.
- **Code vs evidence:** matches. P6 adds what schunk1983 [287] and bandura1981 [286] show: feedback on progress.
- **Status: Built on 25 October 2026** (switch `SCORECARD_SUBGOAL_TEXT`, from 03-P6). Adds short sub-goals and specific feedback after the first one.

03-M32. Misconception persistence: state, refute, replace

- **Where:** `stats.js:351-353`; measured at `signals.js:56-58, 531` (weak rule 9)
- **What it does:**
  - the measure is the same tagged mistake again, out of misses on a skill with a past tagged mistake (per subject);
  - the text says to confront the recorded misconception: state it, show a case where it fails, give the correct idea.
- **Evidence:**
  - schroeder2022 [129] — "refutation text is associated with a positive, moderate effect (g = 0.41" — meta-analysis.
  - tippett2010 [131] — "These structures often contain misconceptions—inaccurate or incomplete information—that are highly resistant to change" — science education review.
  - kendeou2018 [289] — "the structure of the refutation texts facilitated the coactivation and integration of the explanation with the commonsense belief, resulting in knowledge revision" — readers, think-aloud and reading times.
  - shtulman2012 [290] — "naïve theories survive the acquisition of a mutually incompatible scientific theory, coexisting with that theory for many years to follow" — adults with many years of science education.
  - payne1990 [291] — "mal-rules are very unstable, students typically use mal-rules very irregularly" — algebra students, three schools.
  - smith1994 [292] — counter-view: the "flawed ideas that are strongly held" picture "overemphasizes the discontinuity" — theory.
  - Also: ecker2022 [333] — title-level only.
- **Verdict:** strong for the action (refutation). Mixed for the premise that persistence is a stable learner trait (payne1990 [291] vs shtulman2012 [290]).
- **Code vs evidence:** matches. Measuring recurrence of a *tagged* error per subject reads behaviour, not a global trait.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M33. Interest context knob

- **Where:** `adaptation.js:104, 482`; `translate.js:213`
- **What it does:** a confident interest_transfer high builds examples on the learner's stated interests, only when interests exist.

- **Evidence:**
  - walkington2013 [152] (above).
  - bernacki2018 [318] (above).
  - cordova1996 [319] (above).
- **Verdict:** moderate.
- **Code vs evidence:** matches. major2021 [87] found interest-linked personalisation had the smaller effect, so this stays a text change, not a structural one.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M34. Recap knob (consolidation need)

- **Where:** `adaptation.js:105, 483, 600`; `translate.js:127-128`
- **What it does:** the lesson ends with "a three-line recap they answer from memory".
- **Evidence:**
  - rowland2014 [10] (above).
  - yang2021 [21] — "overall testing (quizzing) raises student academic achievement to a medium extent (g = 0.499)" — classroom meta-analysis.
  - adesope2017 [11] — "practice tests are more beneficial for learning than restudying and all other comparison conditions" — meta-analysis.
  - bertsch2007 [334] — "The size of the generation effect across the 86 studies was .40" — meta-analysis.
- **Verdict:** strong for retrieval at the end (it helps everyone). The trigger (scroll-backs) is a design choice.
- **Code vs evidence:** matches. Rereading often added little or no benefit in experiments with educational texts (callender2009 [335]).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M35. Practice ends on proof: compress or extend

- **Where:** `adaptation.js:488-506`; invariant check `adaptation.js:605-629`
- **What it does:**
  - **compress:** a proven objective skips its remaining practice levels, but keeps its spaced-recall level;
  - **extend:** an unproven one whose own practice is used up gets more practice;
  - `checkWhat` refuses any plan that changes what the level covers.
- **Evidence:**
  - kulik1990 [4] (above).
  - bloom1984 [5] (above).
  - corbett1995 [82] — title-level only (knowledge tracing to mastery).
  - beck2013 [325] (above).
- **Verdict:** strong (mastery learning).
- **Code vs evidence:** matches.

- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M36. Blocked while new, interleaved once proven

- **Where:** `adaptation.js:532`; `adaptation.js:413-421`
- **What it does:** practice is blocked on a new objective under worked-example guidance, and interleaved with other skills otherwise.
- **Evidence:**
  - brunmair2019 [14] — "interleaving can effectively foster inductive learning... should be used with caution in certain conditions, especially for expository texts and words" — meta-analysis.
  - rohrer2015 [12] — "Cohen's ds 0.42 and 0.79, respectively" — 126 seventh graders.
  - firth2021 [163] — "Memory... showed an interleaving benefit with effect sizes (Hedges' g) of up to 0.65" — review.
- **Verdict:** moderate.
- **Code vs evidence:** partly matches brunmair2019 [14], which found interleaving less helpful in some conditions, especially for expository texts and words; blocking while an objective is new is a reasoned reading of that caution, not something the meta-analysis tested.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M37. Proof check with questions written for the objective

- **Where:** `adaptation.js:514-524`; `services/coverage.js:95`
- **What it does:** when belief is already above 0.95 but too few right answers came from questions written for this objective (needs 2), it asks a few of those instead of more practice.
- **Evidence:**
  - cohen1987 [336] — "When critical features of instructional stimuli match those of assessment, effect sizes routinely reach 1.2 to 3 sigma." — review of alignment studies.
  - martone2009 [337] — "alignment research is critical for ensuring the standards-assessment-instruction cycle facilitates student learning" — review.
  - kulik2016 [37] — "suggesting that alignment of test and instructional objectives is a critical determinant of evaluation results" *(fetched during the review)* — meta-analysis, 50 ITS evaluations.
- **Verdict:** moderate.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M38. Session summary at 85% / 60%; difficulty_delta computed but unused

- **Where:** `services/sessions.js:1484-1502` (weak rule 3)

- **What it does:** the rule-based end-of-session analysis says:
  - 85% or more: "the next one steps up";
  - 60% or more: "a few soft spots";
  - below that: "stays here".

It writes `difficulty_delta` of +0.5 (85% or more) or -0.5 (50% or less; a different cut from the 60% in the text above, as in the code), but nothing reads that value (only `next_session.note` is shown).

- **Evidence:**
  - wilson2019 [94] (supports an approximately 85% training-accuracy sweet spot in AI and animal-learning models).
  - jansen2013 [32] (above; it supports the general principle that experiencing success helps, not the exact cut-off).
- **Verdict:** design choice. 85% as "strong" fits the human studies, which favour high success. *After the 27 October 2026 audit this mechanism has 2 quoted sources, fewer than three: a source that did not support it was removed (see the audit section).*
- **Code vs evidence:** the summary matches. `difficulty_delta` is dead data: either wire it to the next session's difficulty or drop it, so no one later mistakes it for a live 85% target.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-D. The effort monitor (ADAPT_EFFORT, default on)

03-M39. Answers under 1.2 s are not evidence

- **Where:** `services/coverage.js:75, 102, 436-441`
- **What it does:** an answer under 1,200 ms is counted as rapid and never as evidence of knowing (`OBJ_RAPID_MS`).
- **Evidence:**
  - kong2007 [338] — "compared four methods for setting item response time thresholds to differentiate rapid-guessing behavior from solution behavior" — low-stakes tests.
  - wise2017 [59] (above).
  - soland2021 [339] — "we help provide guidance to developers and administrators of large-scale assessments on the tradeoffs involved in using a given method" — 728,923 students in grades 3-8.
  - wise2009 [328] (above): items with more text are associated with more rapid guessing.
- **Verdict:** moderate. Response-time thresholds are well established; the fixed value is a choice.
- **Code vs evidence:** partly. A fixed 1.2 s is generous on a short item and lenient on a long one; kong2007 [338] and soland2021 [339] compare threshold-setting methods. See P8. Section 02 mechanism 28 agrees.

- **Status: Built on 25 October 2026** (switch `RAPID_ITEM_RELATIVE`, from 03-P8). Items with 30 or more timed answers use max(1.2 s, 0.10 x the item's median).

03-M40. Relative rapid line: under 0.3 x expected time

- **Where:** `signals.js:266, 411, 581-583, 635`
- **What it does:** the scorecard's rush evidence is an answer under 0.3 x the item's expected time, or a lesson page read in under 0.2-0.25 of its expected time.
- **Evidence:**
  - kong2007 [338] (above).
  - soland2021 [339] (above).
  - wise2005 [74] — "unmotivated examinees will answer too quickly (i.e., before they have time to read and fully consider the item)" — low-stakes tests.
  - lee2014rt [340] — "A five-step procedure was proposed to identify rapid-guessing behavior" — NAEP study.
- **Verdict:** design choice. Item-relative thresholds are among the compared methods; 0.3 is chosen.
- **Code vs evidence:** matches the direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M41. "Too fast to count" on checkpoints and the Final

- **Where:** `services/effort-monitor.js:1-24, 109-131`; `adaptation.js:397-411`
- **What it does:**
  - an answer under 1.2 s on a scored sitting is set aside, unscored;
  - the learner is told why and gets a parallel question (same skill, difficulty and objective), at most 2 times;
  - the pass mark and the number of scored questions do not change;
  - retention items and the owner are exempt.
- **Evidence:**
  - wise2006cbt [265] (above).
  - wise2019 [266] (above).
  - wise2006 [35] — "rapid guessing (i.e., reflecting low examinee effort)" — the effort-moderated IRT model.
  - rios2017 [341] — "the true mean was underestimated by around 0.20 SDs if the total amount of careless responses exceeded 6.25%, 12.5%, and 12.5% for easy, moderately difficult, and difficult tests" — group-based low-stakes assessments.
- **Verdict:** moderate. Effort monitoring with warnings is directly tested; the swap for a parallel question is a reasoned extension.
- **Code vs evidence:** matches.
- **Status: Built on 25 October 2026** (switch `RAPID_ITEM_RELATIVE`, from 03-P8). The item-relative line also sets answers aside on checkpoints and the Final (at most 2 reserves). A set-aside answer never counts as wrong.

03-M42. Rush flag and the one-line slow-down note

- **Where:** `adaptation.js:65-69, 377-395`; `effort-monitor.js:133-147`
- **What it does:** outside scored sittings, a rapid answer from a learner whose record flags rushing gets "Read each question to the end". The flag needs a Beta posterior 95% sure that more than 5% of their answers are rapid, on at least 30 answers. No score changes.
- **Evidence:**
  - wise2006cbt [265] (above).
  - wise2019 [266] (above).
  - kluger1996 [342] — "FIs improved performance on average (d =.41) but that over 1/3 of the FIs decreased performance" — meta-analysis.
  - shute2008 [16] — "formative feedback should be nonevaluative, supportive, timely, and specific" — review.
- **Verdict:** moderate.
- **Code vs evidence:** matches. The note is task-focused, which is what kluger1996 [342] points to; its non-evaluative form is supported by shute2008 [16].
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-E. Coach guards

03-M43. The coach never gives the answer first; leak guard

- **Where:** `services/ai.js:1202-1205, 1320, 1415-1430`
- **What it does:**
  - before the learner has answered, the coach must not state the answer or say which choice is correct, even when asked or told a teacher allowed it;
  - `leakGuard` replaces any reply that contains the normalised key (unless the learner said it first) with a redirecting question.
- **Evidence:**
  - bastani2025 [38] — "48% improvement in grades for GPT Base... 17% reduction in grades for GPT Base" — high school maths; a tutor with guardrails (prompts designed to safeguard learning) largely mitigated the harm.
  - baker2004 [283] (above).
  - sinha2021 [605]: "Our results showed a significant, moderate effect in favor of PS-I" (meta-analysis of 53 studies with 166 comparisons of problem solving before instruction against instruction first). Scope: Shows that attempting problems before being taught helped on average. It found the opposite trend for second to fifth graders, and it did not test a tutor that withholds the answer to the question on screen.
  - kapur2014 [606]: "students who engaged in problem solving before being taught demonstrated significantly greater conceptual understanding" (two randomized controlled studies in mathematics). Scope: Supports letting learners attempt a

problem before being told. It did not study an AI coach or withholding the answer to a question on screen.

- **Verdict:** moderate.
- **Code vs evidence:** partly. The guard checks for the key's *text* only. A reply such as "it's option 2" or "the second choice" is not caught; the system prompt forbids it, but nothing checks it. See P10.
- **Status: Built on 25 October 2026** (switch `COACH_LEAK_CHOICE`, from 03-P10). The leak guard also catches a reply that names the key's choice number, letter or position before the learner does.

03-M44. Coach turn rules: one question, the smallest next step

- **Where:** `ai.js:1207-1218`; `services/assist.js:47-60`
- **What it does:**
  - ends each turn with one question that moves the learner one step;
  - exposes a flaw by asking, not telling;
  - changes approach when going in circles;
  - two or three sentences per turn.
- **Evidence:**
  - chi2001 [343] — "Surprisingly, students learned just as effectively even when tutors were suppressed from giving explanations and feedback." *(fetched during the review)* — one-to-one human tutoring.
  - graesser2004 [344] — "AutoTutor has produced learning gains of approximately.70 sigma for deep levels of comprehension." *(fetched during the review)* — Newtonian qualitative physics and computer literacy.
  - nye2014 [345] — "Research on early versions of AutoTutor documented the impact on deep learning by co-constructed explanations, feedback, conversational scaffolding" *(fetched during the review)* — review.
  - vanlehn2011 [36] — "Instead, it found that the effect size of human tutoring was much lower: d = 0.79." *(fetched during the review)* — review; intelligent tutoring systems about as effective.
- **Verdict:** moderate.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M45. The coach stays on the question on screen

- **Where:** `ai.js:1264-1298`
- **What it does:** the coach may never substitute a different problem, new numbers or an easier example; it asks about a part of this question instead.
- **Evidence:**
  - cohen1987 [336] (above).
  - kulik2016 [37] (above).

- **Verdict:** design choice. Help aligned to the step being assessed is the better-supported design; the ban itself came from an observed failure (drift). *After the 27 October 2026 audit this mechanism has 2 quoted sources, fewer than three: a source that did not support it was removed (see the audit section).*
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M46. Coached right answers weigh 0.25; XP x 0.6

- **Where:** `services/coverage.js:76, 103, 447`; `services/mastery.js:670-680`
- **What it does:**
  - a right answer reached with the coach counts 0.25 toward proof;
  - XP for it is 60% of the uncoached value: "enough to make working it out alone feel worth more, not enough to discourage asking for help".
- **Evidence:**
  - koedinger2007 [33] (above; reviews balancing assistance and withholding, not the evidential weight of assisted answers).
  - aleven2016help [34] — review of help seeking with intelligent tutoring systems; in its classroom study, feedback on help seeking made students use on-demand help more deliberately but did not improve learning outcomes.
  - feng2009 [346] — "we can do a reliably better job predicting student end-of-year exam scores by leveraging the interaction data" — ASSISTment system, middle-school math; the interaction data include the help a student needed.
  - beck2008help [347] — "help both promotes student long-term learning and provides additional scaffolding on the current problem" — Reading Tutor data; help can affect performance on the current problem.
- **Verdict:** design choice. That assisted answers are weaker evidence is standard in tutoring assessment; 0.25 and 0.6 are chosen.
- **Code vs evidence:** matches the direction.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M47. No assist on checkpoints and essays

- **Where:** `services/assist.js:1-10`
- **What it does:** checkpoints and essays are the test, so no chat help is offered there.
- **Evidence:**
  - haladyna2004 [348] — "One major threat is construct-irrelevant variance (CIV)." — measurement review.
  - messick1995 [151] (above).
  - kane2013 [105] (above).
- **Verdict:** moderate.
- **Code vs evidence:** matches. Help during a test would mix what the tutor knows into the score.

- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-F. Grading leniency and the injection check

03-M48. Form leniency: case, accents, $, punctuation, "x ="

- **Where:** `services/grading.js:1-8, 23-35, 116-130`
- **What it does:**
  - the rules forgive capitals, accents, "$", trailing punctuation, unicode dashes and quotes;
  - one leading "x =" is forgiven both ways;
  - a wrong value is never forgiven.
- **Evidence:**
  - haladyna2004 [348] (above).
  - messick1995 [151] (above).
  - downing2005 [349] — "flawed multiple-choice test items, which violate well established and evidence-based principles of effective item writing, disadvantage some medical students" — medical students; indirect support (flawed multiple-choice items, not answer-format leniency).
  - burrows2015 [350] — "Automatic short answer grading (ASAG) is the task of assessing short natural language responses to objective questions using computational methods." — review of 35 systems; background on automatic short-answer grading, not direct support for notation leniency.
- **Verdict:** moderate. Scoring the construct and not the notation is the validity principle; the exact list is the app's.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M49. Number forms: 0.5 = 1/2, mixed numbers, strict thousands commas

- **Where:** `grading.js:38-90, 278-288`
- **What it does:**
  - integers, decimals, fractions and mixed numbers are compared by value;
  - a comma is a thousands separator only in the exact 1,250 shape, so "12, 6, 4" is not read as 1264.
- **Evidence:**
  - ccss-4nf-c6 [351] — "For example, rewrite 0.62 as 62/100; describe a length as 0.62 meters; locate 0.62 on a number line diagram." — US standard.
  - ccss-4nf-a1 [351] same URL — "Use this principle to recognize and generate equivalent fractions." — US standard.
  - ccss-3nf-a3 [352] — "Express whole numbers as fractions, and recognize fractions that are equivalent to whole numbers." — US standard, grade 3 (a separate page from 4.NF).

    - haladyna2004 [348] (above).
- **Verdict:** moderate (the standards treat these forms as the same number).
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M50. Strict on exactly the form the prompt asks for

- **Where:** `grading.js:132-185, 254-277`
- **What it does:** when the prompt names a form (lowest terms, as a fraction or decimal, capitalisation, punctuation, accent), forgiveness is switched off in exactly that dimension. "24/32" for "lowest terms" is wrong.
- **Evidence:**
    - messick1995 [151] (above).
    - haladyna2004 [348] (above).
    - ccss-4nf-a1 [351] (above): generating an equivalent fraction is itself the skill.
- **Verdict:** moderate. When the form is the construct, forgiving it would itself be construct-irrelevant.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M51. Spanish: optional subject pronoun, el/él and tu/tú kept apart

- **Where:** `grading.js:19-20, 92-115`
- **What it does:**
    - a leading subject pronoun may be included or left out;
    - an unaccented "el" or "tu" counts as the pronoun only when the other side spells it with its accent, so "el libro" never accepts "libro".
- **Evidence:**
    - otheguy2007 [353] — "Subject personal pronouns are highly variable in Spanish but nearly obligatory in many contexts in English" — corpus of 142 speakers.
    - rae-tilde [354] — "Se llama tilde diacrítica al acento gráfico que permite diferenciar en la escritura ciertas palabras de igual forma, pero distinto valor" — Spanish language academy; its table pairs él (pronoun) with el (article) and tú with tu.
    - haladyna2004 [348] (above).
- **Verdict:** design choice.
- **Code vs evidence:** matches. Accents are forgiven elsewhere (mechanism 48), so the accent-only distinction el/él needs this special case, and has it.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M52. Grading injection check and fenced judge input

- **Where:** `services/typed-grading.js:26-33, 137-204`

- **What it does:**
  - an answer is marked incorrect before any model sees it when it talks to the grader: instructions, "the answer is correct", fake system lines, chat tokens, JSON verdicts, role-play (English and Spanish), unless the item's own text says the same;
  - every answer a judge sees is cleaned (invisible characters, NFKC), put on one line, capped at 600 characters, and fenced between random markers;
  - `GRADING_INJECTION_CHECK=off` skips the pre-check only.
- **Evidence:**
  - greshake2023 [355] — "LLM-Integrated Applications blur the line between data and instructions" — deployed LLM applications.
  - shi2024judge [356] — "JudgeDeceiver injects a carefully crafted sequence into an attacker-controlled candidate response" — LLM-as-a-judge.
  - owasp-llm01 [357] — "Separate and clearly denote untrusted content to limit its influence on user prompts." — industry standard.
- **Verdict:** moderate.
- **Code vs evidence:** matches the standard (fencing). shi2024judge [356] shows JudgeDeceiver outperforms manually crafted injection attacks and that tested defenses are insufficient. The stronger guard here is the consensus rule and "never accept on doubt" (typed-grading.js header), and the regex check is the first layer. Keep both.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-G. Objective items

03-M53. Objective-item code checks

- **Where:** `services/objective-items.js:43-97`
- **What it does:** an item is refused unless it has all of these:
  - exactly 4 distinct, non-empty options (3 or 4 since 25 October 2026; see Status);
  - no "all/none of the above";
  - a key that is one of the options and is not written in the question;
  - no length cue;
  - a rationale that points to the key;
  - at most 80 words;
  - no reference to a missing picture or audio;
  - no duplicate for the objective;
  - for maths, the math-claims checks.
- **Evidence:**
  - haladyna2002 [132] — "A taxonomy of 31 multiple-choice item-writing guidelines was validated through a logical process" — 27 textbooks, 27 studies.
  - tarrant2008 [359] — "Fewer examinees passed the standard scale than the total scale" — nursing students, 10 high-stakes papers.

    - downing2005 [349] (above).
    - gierl2017 [360] — "the task of creating distractors has received much less attention" — review.
    - rodriguez2005 [136] — "More 3-option items can be administered than 4- or 5-option items per testing time while improving content coverage, without detrimental effects" — meta-analysis.
- **Verdict:** strong (item-writing guidelines, and the harm of flawed items).
- **Code vs evidence:** matches. One difference: the evidence allows 3 options (rodriguez2005 [136]; aamodt1992 [299]: 3-choice items take less time), and the code requires exactly 4. Not a defect (P11, optional).
- **Status: Built on 25 October 2026** (switch `OBJ_ITEMS_3_OPTIONS`, from 03-P11). Objective items may have 3 or 4 options. Today's files are all 4-option, so nothing changes until a build writes 3-option items.

03-M54. Two blind judges from different model families agree on the key

- **Where:** `objective-items.js:8-13, 82-83`; `services/key-agreement.js`
- **What it does:** only items kept by the content factory are served: code checks, plus two blind judges from different model families agreeing on the key and suitability. Here the rationale must also point to the key.
- **Evidence:**
    - zheng2023 [358] (above): self-enhancement bias is a limitation of LLM-as-a-judge.
    - pardos2024 [361] — "ChatGPT-generated help failed quality checks on 32% of problems" — maths help.
    - gierl2017 [360] (above).
- **Verdict:** design choice.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M55. Objective items served first; aligned right answers needed for proof

- **Where:** `objective-items.js:1-20`; `coverage.js:66-67, 95`
- **What it does:** reviews and remediation ask questions written for one objective before bank questions, and a proof needs at least 2 right answers to such questions.
- **Evidence:**
    - cohen1987 [336] (above).
    - martone2009 [337] (above).
    - kulik2016 [37] (above).
- **Verdict:** moderate.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-H. The young-learner oral mode (YOUNG_ORAL, default on)

03-M56. Spoken checkpoints for preschool and Kindergarten

- **Where:** `services/essays.js:50-95`
- **What it does:** the sheet's "Explain Aloud" and "Oral Retell" checkpoints are no longer an 80-word typed essay:
  - a tap-to-answer quick check, marked by the app, decides the pass;
  - the child then says it aloud to a grown-up;
  - no model is called.
- **Evidence:**
  - ccss-wk2 [362] — "Use a combination of drawing, dictating, and writing to compose informative/explanatory texts" — US standard, Kindergarten.
  - puranik2009 [363] — "Preschoolers appear to progress along a continuum from scribbling to conventional spelling" — 372 preschoolers.
  - reed2011 [364] — "Retell was moderately correlated with standardized measures of reading comprehension" — 54 studies, K-12.
  - rittlejohnson2006 [365] — "self-explanation promoted transfer regardless of instructional condition" — children aged 8-11.
- **Verdict:** moderate.
- **Code vs evidence:** matches:
  - a typed 80-word essay asks for transcription that preschoolers do not have;
  - reed2011 [364] also says retell "has not yet demonstrated adequacy as a progress monitoring instrument", which supports treating retell as supplementary rather than as the sole progress-monitoring measure.

rittlejohnson2006 [365] is about older children, so the benefit of explaining aloud at age 4-5 is extrapolated.

- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M57. A grown-up's "I heard it" weighs 0.25 and never proves alone

- **Where:** `coverage.js:77, 104, 447`; `essays.js:63-66`
- **What it does:** an adult's confirmation is a light, unverified answer, and it can never prove an objective on its own.
- **Evidence:**
  - kane2013 [105] (above; general principle that a claim needs adequate supporting evidence).
  - messick1995 [151] (above; general principle that score meaning needs sufficient evidence).
  - zell2014 [587]: "the mean correlation between ability self-evaluations and performance outcomes across meta-analyses was moderate (M = .29, SD = .11)" (metasynthesis of 22 meta-analyses of self-evaluations against objective

performance). Scope: These are self-reports. Read across to a grown-up's unverified report, it shows only that unverified reports track performance loosely. It did not study an adult confirming a child's skill.
    - bernes2021 [588]: "Weak correlations were found between the rating scales and laboratory measures, indicating poor convergent validity for the behavior rating scale." (children with prenatal alcohol exposure, children with ADHD and controls; informant rating scales against laboratory measures). Scope: Shows that informant ratings can line up only weakly with directly measured performance in these groups. It did not study an adult confirming a skill, or any weight such as 0.25.
- **Verdict:** design choice.
- **Code vs evidence:** matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-M58. The recording stays on the device

- **Where:** `essays.js:58-62`; `public/js/say-aloud.js`
- **What it does:** the optional recording stays on the device: nothing is uploaded or kept.
- **Evidence:**
    - coppa-312.2 [366] — "A photograph, video, or audio file where such file contains a child 's image or voice" — US rule (from section 05).
    - gdpr-art5-1c [309] (above).
    - livingstone2019 [367] (no DOI) — "The relation between privacy and data online is becoming highly complex" — children's privacy (background only; it does not address voice recordings).
- **Verdict:** moderate.
- **Code vs evidence:** matches. A child's voice is personal information, so not collecting it is the minimal option.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-I. The lesson sandbox (LESSON_SANDBOX, default on)

03-M59. Lessons in a sandboxed frame with a strict CSP

- **Where:** `routes/base-lessons.js:180-207, 236-290`
- **What it does:**
    - lesson pages load in `<iframe sandbox="allow-scripts allow-forms">`, without allow-same-origin;
    - the response itself sends `Content-Security-Policy: sandbox ...`, so the page is opaque wherever it is opened;
    - scripts run only with a per-response nonce, or with hashed inline handlers;
    - `connect-src`, `form-action`, `frame-src` and `object-src` are `'none'`;
    - the page talks only by checked postMessage.

- **Evidence:**
  - whatwg-iframe-sandbox [368] — "The sandbox attribute, when specified, enables a set of extra restrictions on any content hosted by the iframe element." — W3C Working Draft. The same page warns that allow-scripts plus allow-same-origin together let a page break out.
  - w3c-csp3 [369] — "This document defines a mechanism by which web developers can control the resources which a particular page can fetch or execute" — web standard.
  - weichselbaum2016 [370] — "75.81% of distinct policies use script whitelists that allow attackers to bypass CSP" — 1,680,867 hosts.
  - saltzer1975 [313] (above): least privilege.
- **Verdict:** strong (standards plus measured CSP failures).
- **Code vs evidence:** matches:
  - nonces rather than host whitelists, as weichselbaum2016 [370] recommends;
  - no allow-same-origin, as the WHATWG warning requires.

The policy still lists `'self'` and the host in script-src. That is the one remaining whitelist entry, and safe only while the app serves no JSONP-style endpoint.

- **Status:** Existing (in the code before 25 October 2026; not changed that day).

03-J. Owner base vs optimized (product behaviour, no evidence needed)

03-M60. Owner base vs optimized

- **Where:** `config/owner.js:1-51`; `adaptation.js:652-658, 683-686`; `learner.js:239, 293`; `effort-monitor.js:38-45`
- **What it does:**
  - the owner's account gets the base product, unmodified: no adaptation plan, neutral session settings, no scorecard steering;
  - every other account gets every modification;
  - `OWNER_BASE=off` (or `ADAPT_OWNER=1`) makes the owner an ordinary learner;
  - safety and integrity rules never ask this.
- **Note:** a product rule from docs/PRODUCT-DECISIONS.md. One detail: `effort-monitor.js` exempts the owner through `isOwner` (skip rights), not `ownerBase`. So `OWNER_BASE=off` adapts the owner but still exempts them from "too fast to count". The owner.js comment, by contrast, says the too-fast rule never asks ownerBase. That is consistent, but worth knowing when testing as the owner.
- **Status:** Existing product rule (owner decision).
- **Flag:** product rule; no learning evidence is needed (0 sources).

03-M61. Owner skips

- **Where:** `services/owner-skips.js:1-16`

- **What it does:** the owner may skip a placement quiz, skip a whole grade, and stand on any level ahead or behind. A skipped grade graduates without the Final, the coverage gate, the skill gate or consolidation. Only `canSkip` routes reach it.
- **Note:** a product rule, no evidence needed.
- **Status:** Existing product rule (owner decision).
- **Flag:** product rule; no learning evidence is needed (0 sources).

03 proposals (each backed by 3 or more sources; status on each)

#	Change	Files and values	Sources	Expected effect	Risk
03-P1	**BUILT (flag `SCORECARD_DROP_RUSH_TYPED`, 2026-10-25).** Drop the typed-answer switch for rushers, or run it as an experiment arm	`services/adaptation.js:103` remove `rush_tendency: { high: { typed: true } }`. The same in `services/scorecard/contract.js:66`. `services/learner.js:300` then never sees `s.typed` from rush. Update `scripts/scorecard-contract-test.js` (KNOBS equality)	smith2014 [116], kang2007 [166], little2012 [133], greving2018 [264], rowland2014 [10]; wise2006cbt [265] and wise2019 [266] for the effort monitor that remains	Removes an untested format switch. Rushers are still handled by mechanisms 41-42	Low. If typed recall does help rushers, an experiment arm would show it
03-P2	**BUILT (flag `SCORECARD_SOFT_STRETCH`; knob, clamp and text).** Soften the challenge high side	`adaptation.js:99` `challenge_appetite.high.stretch` +0.1 to **+0.05**; `contract.js` KNOBS the same. Stretch clamp 0.35 to **0.25** at `adaptation.js:544` and `learner.js:299`. `stats.js:158-160` high text to "Teach the core at the usual level, then end with one optional harder 'try this' case"	jansen2013 [32], eglington2020 [28], lomas2013 [268], lomas2017 [320]	Higher success rates and more practice attempted (jansen2013 [32])	Strong learners may find levels easy. Problems-first on proven objectives (mechanism 17) still moves them on
03-P3	**BUILT (flag `SCORECARD_LENGTHEN_NEW_ONLY`).** Lengthen only for new material	`adaptation.js:477`: apply `length: +1` from brevity_preference low only when `g === 2` (worked examples). Keep the -1 side as is	mayer1996 [107], sundararajan2020 [128], kalyuga2003 [6], kirschner2013 [269]	Shorter text for learners who already know the objective	Low
03-P4	**BUILT (flag `SCORECARD_OBSERVED_PREFS` = on, shadow or off; live).** It was run in shadow on a	`learner.js:242`: cap `Math.min(0.75, ...)` to **1.0**, so after 7 analysed sessions the stated value has no weight.	pashler2008 [31], rogowsky2015 [317], kirschner2013	Session structure follows behaviour, as the	Medium. Changes session size and shares for existing

#	Change	Files and values	Sources	Expected effect	Risk
	copy of the live DB: 2 of 107 profiles change, both stretch 0.125 to 0.1. The recovery sim (`--quick`) gives the same numbers as before. That sim does not exercise `learner.effective`, so the evidence that nothing regressed comes from the scorecard tests and the shadow pass. Let measured behaviour fully replace stated preference, and gate the scorecard term	`learner.js:251-255`: replace the confidence >= 30 blend with `contract.forSession` sides only, or drop it	[269], nisbett1977 [316]	contract already promises	learners; run in shadow first
03-P5	**BUILT (flag `SCORECARD_WORKED_HINT_TEXT`).** A better-supported hint_dependency directive	`stats.js:218-220` high text to "Open with one fully worked example of exactly this step, then ask the learner to try the next step themselves before any hint"	koedinger2007 [33], renkl2014 [8], salden2010 [110], baker2004 [283]	Help up front without bottom-out hinting	Low. Text only
03-P6	**BUILT (flag `SCORECARD_SUBGOAL_TEXT`).** A stronger self_efficacy directive	`stats.js:264-266` low text: add "Split the task into short sub-goals, and after the first one say specifically what they did right and that it shows they can do this kind of step"	schunk1983 [287], bandura1981 [286], usher2008 [284], byarswinston2017 [285]	More mastery experience and specific progress feedback	Low. Avoid generic praise (mueller1998 [371])
03-P7	**BUILT (no flag: a check, not a behaviour).** `services/scorecard/reliability.js` `splitHalf`, plus the report wording in section 11. Running it on real learners needs data (20+ learners with 30+ levels). Specify the split-half check	`SCORECARD-EVIDENCE-REPORT.md:545`, and the check script when it is written: "Spearman-Brown-corrected split-half, averaged over 500 random splits (not only odd/even days), point >= 0.6 and lower 95% bound >= 0.5"	koo2016 [270], parsons2019 [66], brown1910 [273], spearman1910 [274] (title-level), cicchetti1994 [271]	A check that matches reliability practice	None (documentation, then a script)
03-P8	**BUILT (flag `RAPID_ITEM_RELATIVE`; knobs `OBJ_RAPID_REL` 0.10, `OBJ_RAPID_MIN_N` 30).** Applied in	`coverage.js:102, 439` and `adaptation.js:408-411` (sittingRule): rapid if `ms < max(1200, 0.10 x the item's`	kong2007 [338], soland2021 [339], wise2009 [328], wise2017 [59]	Catches skimming on long items that 1.2 s misses	More answers set aside on long items; the cap of 2

#	Change	Files and values	Sources	Expected effect	Risk
	coverage, in set-asides on scored sittings, and in the slow-down note when the caller passes the item. A set-aside answer never counts as wrong. An item-relative rapid line on long items	`typical time)` once an item has 30+ timed answers; 1.2 s until then			reserves still holds
03-P9	**SKIPPED: needs data.** No simulation in the repo can measure how well the writer and coach models follow 3, 6, 9 or 12 directives: the recovery sim stubs the model. Measuring it needs model runs and per-directive compliance scoring. COACH_MAX stays 8. Test, then match, the directive caps	Run the writer and coach models on 3 / 6 / 9 / 12 directives and score compliance per directive. Set `contract.js:73` COACH_MAX to whatever wins (6 unless the coach shows a reason)	liu2024lost [124], levy2024 [125], he2024cello [267], jaroslawicz2025 [306]	A measured cap instead of two unmeasured ones	Low. Evaluation cost only
03-P10	**BUILT (flag `COACH_LEAK_CHOICE`; non-streamed and streamed coach).** Close the leak guard's choice-number gap	`ai.js:1421-1430`: also redirect when, before the learner names it, the reply names the key's choice number, letter or position ("(2)", "option B", "the second one")	bastani2025 [38], baker2004 [283], aleven2016help [34]	The no-answer rule holds on multiple choice too	False redirects when an explanation mentions an option number innocently. Limit the check to the key's own index
03-P11 (optional). **BUILT (flag `OBJ_ITEMS_3_OPTIONS`, default on; 2026-10-25, implementer B):** `objective-items.problems` accepts 3	Allow 3-option objective items	`objective-items.js:68`: accept 3 or 4 choices; the build script may then drop the weakest distractor	rodriguez2005 [136],	More items per minute at no loss in quality	Low. Existing files stay valid

#	Change	Files and values	Sources	Expected effect	Risk
or 4 distinct options, and coverage already guesses 1/n. The build script's distractor drop is not done; today's files are all 4-option. Tests: `objective-items-test.js`.					

Report corrections (no code). DONE 2026-10-25 in SCORECARD-EVIDENCE-REPORT.md (sections 6.2, 6.3, 6.4, 6.6, 6.7, 6.10, 7 and 11; the new sources are added to section 13):

- SCORECARD-EVIDENCE-REPORT.md section 7: "Rates use logit counting variance ... rests on the simulation" should say it is the Haldane-Anscombe empirical logit with the Gart-Zweifel variance (haldane1956 [279]; anscombe1956 [280]; gart1967 [281]).
- The 0.98 has the posterior false-discovery justification (newton2004 [275]; efron2008 [276]).
- Section 6.7: list jansen2013 [32] and eglington2020 [28] next to wilson2019 [94], and drop "about 85% for gradient-based learners" as a basis for people.
- Section 6.6: add camachomorles2021 [294] on frustration and academic performance.
- Section 6.10: add payne1990 [291] and shtulman2012 [290] on the persistence premise.

Section 04. Skill tracks (chess, piano, basketball) and learning games

This section covers:

- the skill curricula: `curriculum-build/SKILL-CURRICULA.md` and the subjects' research files in `../curriculum-build/<subject>/research`;
- stage checks and the stage gate: `content/skill-stage-checks.js`, `services/skill-coverage.js`, `services/skill-record.js`, `services/skill-gate.js`, `SKILL_STAGE_GATE` and "two cold days";
- the practice dose: `services/skill-adaptation.js`, `services/skill-dose.js`;
- the chess sittings: `content/skill-chess-sittings.js`, `services/skill-chess-sittings.js`, `content/chess-trainer.js`;
- the piano and basketball sittings: `services/skill-piano-bb-sittings.js` and `content/skill-piano-bb-sittings.js`. These cover sight-reading, scale evenness, memory restarts, decision scenes, shooting by zone and box scores;
- piano score-following: `public/engine/piano/*`;
- basketball body tracking: `public/engine/basketball/*`, including `pose.js`;
- deliberate practice, the "film it, rate it" loop, cold first attempts and two-day lessons;
- learning games: `services/game-lab.js`, `services/game-payloads.js`, `GAME_HOLDS` and the game payloads.

Fact re-verification.

- `curriculum-build/FACT-REVERIFY-2026-10-25.md` is present and is cited here for:
  - the FIBA free-throw line (5.80 m, Art. 2.5.3);
  - the Tran & Silverberg aim cue, corrected to "toward the back of the ring";
  - the scale-spelling display bug: 59 of 102 battery items showed flats as sharps. The fix needs an app restart on 6767, which is the coordinator's job.

Code version read: the working tree on 2026-10-25. Line numbers are as of that read.

Notation.

- [R] marks a number the code itself calls provisional.
- "Pop." gives the population or setting each source studied.

In numbers: 58 mechanisms; 56 with 3 or more quoted sources; 11 strong, 19 moderate, 7 mixed, 21 design choice; 169 distinct sources cited or listed; 13 mechanisms built or changed on 25 October 2026.

04 at a glance

Mechanism	Verdict	Quoted sources	Status
04-M1 Cold-only sittings	strong	4	existing
04-M2 Two different days	mixed	7	built on 25 October 2026

Mechanism	Verdict	Quoted sources	Status
04-M3 A belief that cannot be gamed by lucky passes	design choice	3	existing
04-M4 Estimates (microphone, camera) count only when confirmed	strong	5	existing
04-M5 Self-report, film and imported ratings never prove	moderate	4	existing
04-M6 The retention ("kept") sitting	strong	4	existing
04-M7 Lapse after two cold misses	design choice	3	existing
04-M8 Stale proofs need a fresh cold pass at exit	design choice	3	existing
04-M9 The stage gate, per stage, with a default shadow list	design choice	3	built on 25 October 2026; open proposal
04-M10 Decision-scene sets need split-half ≥ 0.8 before they gate	moderate	3	built on 25 October 2026; open proposal
04-M11 Dose by distance from target, every unproven metric kept	moderate	4	built on 25 October 2026
04-M12 When to call a cold sitting	design choice	4	existing
04-M13 Maintenance share and the over-learning dose	mixed	4	existing
04-M14 Carry-over across stages is off by default	mixed	3	existing; open proposal
04-M15 Lesson share while knowledge is open	design choice	4	existing
04-M16 Session length, set length and items per set	design choice	4	built on 25 October 2026
04-M17 Answers held on the server; fixed-form sittings	design choice	3	existing
04-M18 Mastery cuts of 80-95%	moderate	4	existing
04-M19 Unannounced chess items	moderate	3	built on 25 October 2026
04-M20 Moves are played or typed, not picked from a menu	strong	3	existing
04-M21 The puzzle trainer: Glicko-2 rating and puzzle choice	mixed	4	built on 25 October 2026
04-M22 The puzzle-rating stage cut	design choice	2	existing
04-M23 Play-outs against perfect defence	design choice	2	existing
04-M24 Recorded rapid games against measured bots	moderate	4	existing
04-M25 CH0 foundations: star mazes at par, set-up from memory	design choice	3	existing
04-M26 Calculation: type the whole line	design choice	3	existing
04-M27 Score-informed scoring and the timing window	design choice	4	existing
04-M28 Scale evenness as a CV	moderate	4	built on 25 October 2026

Mechanism	Verdict	Quoted sources	Status
04-M29 Scales drawn at random	moderate	4	existing
04-M30 Sight-reading staircase with never-seen excerpts	moderate	4	existing
04-M31 Memory restarts from random bars	moderate	3	existing
04-M32 Dynamics from MIDI velocity	design choice	3	existing
04-M33 Legato pedalling from CC64	design choice	3	existing
04-M34 Musicality is never machine-graded	moderate	3	existing
04-M35 The microphone pipeline	strong	3	existing
04-M36 Learn mode waits; play mode runs at tempo	moderate	4	built on 25 October 2026
04-M37 Self-rate before the measurement	mixed	5	built on 25 October 2026; open proposal
04-M38 Ear items	moderate	4	existing
04-M39 Steady beat	design choice	3	existing
04-M40 A piece learned alone	design choice	3	existing
04-M41 Two-day lessons	mixed	5	built on 25 October 2026
04-M42 The nine-activity lesson ladder	strong	7	existing
04-M43 Beta-posterior make-rate gates	strong	4	existing
04-M44 Shooting by zone	design choice	4	existing
04-M45 Own baseline, never a ranking	moderate	4	existing
04-M46 Decision scenes	moderate	7	existing; open proposal
04-M47 Applied tasks: box scores and rules	design choice	3	existing
04-M48 Camera tracking on the device; the learner confirms	moderate	4	existing; open proposal
04-M49 Film it, rate it	moderate	5	existing; open proposal
04-M50 One external cue per drill	strong	4	built on 25 October 2026
04-M51 Court session: health check, warm-up first, pain and enjoyment logged	strong	6	built on 25 October 2026; open proposal
04-M52 Handle and react; eyes-up dribbling	design choice	4	existing
04-M53 Serving holds	moderate	4	existing
04-M54 Payload inspection, including the explanation on every question	strong	7	existing
04-M55 The explanation-key gate	design choice	3	existing
04-M56 evenOrder	mixed	4	existing
04-M57 Games practise what the last lesson taught	strong	3	existing
04-M58 Tailored games with a holdout arm; games as practice	moderate	4	existing

04-A. How a stage is proven (skill-coverage, skill-record, skill-gate)

04-M1. Cold-only sittings

- **Where:** `services/skill-coverage.js:66` (`COLD_HOURS: 20`), `:242-246` (`isCold`), `:256-258`; `services/skill-record.js:8-9, 198-210` (`coldAt`); `services/skill-dose.js:19-20`.
- **What it does:**
  - A sitting counts only when it is cold: the metric's first sitting that day, and at least 20 hours after the metric was last practised or sat.
  - Warm practice sets never prove anything. Each one makes the next sitting warm for 20 hours.
- **Evidence:**
  - soderstrom2015 [39]: "people often mistakenly interpret their performance during acquisition as a reliable guide to long-term learning". Pop.: review of motor and verbal learning.
  - kantak2012 [372]: "challenging practice conditions may impair practice performance, but enhance long-term retention of motor skills". Pop.: focused review of motor learning studies.
  - schmidt1992 [373]: "Manipulations that maximize performance during training can be detrimental in the long term". Pop.: motor and verbal training paradigms.
  - lee1988 [374]: "there is strong evidence that massed practice depresses performance and learning (when learning is assessed by absolute retention measures)". Pop.: review and meta-analysis of motor practice distribution.
  - adams1961 [375]: background only: review of warm-up decrement, showing performance after a rest differs from warmed-up performance (title-level; does not directly test cold-sitting validity).
- **Verdict: strong** for the principle that a cold, delayed sitting measures learning better than in-session performance. The 20-hour window is a design choice.
- **Code vs evidence:** Matches. The rule cannot be bypassed by repeated retakes: `sit()` marks every non-first sitting of a day as warm, so it never counts.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M2. Two different days

- **Where:** `services/skill-coverage.js:67` (`PROOF_DAYS: 2`), `:270-280`; the "two cold days" text in `services/skill-gate.js:208`.
- **What it does:** A metric is proven only by cold passes on at least two different UTC days, with the belief at or above 0.95 (M03).
- **Evidence:**
  - walker2002 [376]: "a night of sleep results in a 20% increase in motor speed without loss of accuracy". Pop.: adults, finger-tapping sequence.
  - simmons2006 [377]: "We found significant sleep-dependent improvements in performance accuracy in the retests that followed intervals of sleep". Pop.: adults learning a 12-note keyboard melody.

    - abushanab2013 [378]: "on a retention test 2 days later, participants were faster with melodies from the random-order condition". Pop.: 25 trained pianists. This shows that a 2-day-later retention test reverses the in-session ranking.
    - mazza2016 [239]: "interleaving sleep between learning sessions not only reduced the amount of practice needed by half but also ensured much better long-term retention". Pop.: 40 adults, vocabulary.
    - rickard2008 [379]: "Whereas sleep may play a role in protection from forgetting of motor skills, it does not result in performance enhancement." Pop.: adults, motor sequence, with confounds controlled.
    - wilhelm2008 [253]: "children showed smaller improvement in finger-tapping skill across retention sleep than wakefulness". Pop.: 15 children aged 6-8 and 15 adults.
    - brashers1996 [380]: "consolidation of a motor skill was disrupted when a second motor task was learned immediately after the first". Pop.: adults, reaching task.
- **Verdict: mixed**:
    - A delayed test is a better index of learning than performance during practice (soderstrom2015 [39], abushanab2013 [378]); holding it on a second day is a design choice.
    - Whether sleep *improves* a motor skill is contested (walker2002 [376] and simmons2006 [377] against rickard2008 [379]).
    - For children, the procedural sleep benefit may be absent (wilhelm2008 [253]).
- **Code vs evidence:** The rule is sound as a delayed-retention and anti-luck rule, and that is what the code implements. It makes no sleep-gain claim. Any learner-facing text that promises improvement "after a night's sleep" would overstate the evidence for under-12s; see proposal P9.
- **Status: Built on 25 October 2026** (switch `SKILL_TWO_DAY_TEXT`, from 04-P9). Stage pages say the two-day rule is a later test, not a promise of getting better overnight.

04-M3. A belief that cannot be gamed by lucky passes

- **Where:** `services/skill-coverage.js:68-71` (`PRIOR 0.5`, `PASS_IF_AT 0.85`, `LUCK 0.15`, `PROVEN_AT 0.95`), `:261-268`.
- **What it does:**
    - Each counting sitting updates P(at threshold) by Bayes' rule. A pass has likelihood 0.85 if the learner is at the bar and 0.15 if below it.
    - Two passes from the even prior reach 0.97. Each miss in between needs another pass.
    - So retaking until two lucky passes come cannot prove a metric (sequential testing).
- **Evidence:**
    - gelman2013 [381]: "Bayesian Data Analysis describes how to conceptualize, perform, and critique statistical analyses from a Bayesian perspective." Textbook (CRC).

    - brown2001 [382]: "we recommend the Wilson interval or the equal-tailed Jeffreys prior interval for small n". Pop.: statistical analysis of binomial intervals. Small-n Wald intervals can have erratic coverage; Wilson or Jeffreys intervals are recommended for small n.
    - hopkins2000 [383]: "Better reliability implies better precision of single measurements and better tracking of changes in measurements in research or practical settings." Pop.: sports-science measurement review.
    - wald1945 [384]: title-level only (sequential tests).
- **Verdict: design choice.** The method follows standard sequential Bayesian updating. The 0.85 and 0.15 likelihoods come from the staged gates' simulated error rates, not from learner data.
- **Code vs evidence:** Matches in form. Gap: the likelihoods should be re-estimated from real sittings once enough are recorded. A fixed 0.85 / 0.15 may be wrong for noisy metrics such as shooting.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M4. Estimates (microphone, camera) count only when confirmed

- **Where:**
    - `services/skill-coverage.js:29-32, 256-258`;
    - `services/skill-checks.js:713-721`: a microphone take counts as confirmed only if every part has two passing takes;
    - `content/skill-stage-checks.js:17-23`;
    - `services/skill-record.js:11-14`.
- **What it does:**
    - Exact sources count directly: board, MIDI, on-screen keys, taps and screen answers.
    - Microphone and camera results are estimates. They count only with a confirmation: a second passing take, or the learner confirming the count.
    - A source the metric does not list never counts.
- **Evidence:**
    - decheveigne2002 [385]: "Error rates are about three times lower than the best competing methods". Pop.: YIN F0 estimation on speech and music. Even the best single-pitch method has errors.
    - emiya2010 [386]: "It addresses the issue of overlapping overtones by modeling the spectral envelope of the overtones of each note". Pop.: multipitch piano recordings (MAPS). Chords are a known hard case.
    - kong2021 [387]: "Our system achieves a pedal onset F1 score of 91.86%, which is the first benchmark result on the MAESTRO dataset." Pop.: piano transcription benchmark. Even state-of-the-art audio is below MIDI's exactness.
    - washabaugh2022 [388]: "Pose estimation methods differed in their ability to measure kinematics." Pop.: healthy gait, open-source pose methods against motion capture.

    - stenum2021 [389]: "Mean absolute error of sagittal plane hip, knee and ankle angles between motion capture and OpenPose were 4.0°, 5.6° and 7.4°." Pop.: adults walking, 2D video.
    - benetos2013 [390]: title-level only (automatic music transcription challenges).
- **Verdict: strong.** Audio and video estimates carry known errors, so treating them as estimates is well supported.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M5. Self-report, film and imported ratings never prove

- **Where:** `services/skill-coverage.js:33-36`; `services/skill-gate.js:17-22`; `attest` list at `:146`.
- **What it does:** Metrics whose only sources are self-report, film for a person to review, or imported stats are "attest". They are recorded and shown, and never hold or prove a stage. Whether a person's sign-off should hold a stage is left to the owner.
- **Evidence:**
    - kruger1999 [391]: "participants scoring in the bottom quartile on tests of humor, grammar, and logic grossly overestimated their test performance and ability". Pop.: US undergraduates.
    - bergee1997 [392]: "self-evaluation correlated poorly with both faculty and peer evaluation". Pop.: college applied-music performers.
    - hewitt2001 [393]: "participants who listened to a model during self-evaluation improved more than those not listening to a model". Pop.: 82 junior-high instrumentalists. Self-evaluation helps learning when it is anchored, but that is not the same as proof.
    - andrade2019 [394]: "An argument is made for less research on consistency and summative self-assessment, and more on the cognitive and affective mechanisms of formative self-assessment." Pop.: review of 76 studies.
- **Verdict: moderate.** Self-assessment is useful formatively and unreliable as a summative measure.
- **Code vs evidence:** Matches. Self-ratings are used formatively (M37, M49) and never as proof.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M6. The retention ("kept") sitting

- **Where:** `services/skill-coverage.js:18, 288-293`; `services/skill-record.js:245`. Stage `retention_days` comes from the progression ("again after 7+ days").
- **What it does:** A stage that names a retention test needs a cold pass at least that many days after the proof. Only then does the metric move from `proven` to `retained`.
- **Evidence:**
    - arthur1998 [395]: "there is substantial skill loss with nonpractice or nonuse". Pop.: meta-analysis of 53 articles on skill decay.

    - roediger2006 [1]: "on the delayed tests, prior testing produced substantially greater retention than studying". Pop.: college students, prose.
    - cepeda2008 [23]: "The optimal gap increased as test delay increased." Pop.: 1,350+ adults, facts, delays up to a year.
    - soderstrom2015 [39] (as M01).
- **Verdict: strong.**
- **Code vs evidence:** Matches. The 7-day figure is the progression's own number.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M7. Lapse after two cold misses

- **Where:** `services/skill-coverage.js:69` (`LAPSE_MISSES: 2`), `:294-305`.
- **What it does:** After proof, two cold misses in a row, or a belief below 0.5, lapse the metric. Its proof days are cleared and it must be proven again.
- **Evidence:**
    - arthur1998 [395] (as M06).
    - hopkins2000 [383] (as M03). A single miss can reflect trial-to-trial measurement noise.
    - rawson2011 [24]: "practice recalling concepts to an initial criterion of 3 correct recalls and then to relearn them 3 times at widely spaced intervals". Pop.: 533 students. Relearning at widely spaced intervals improves long-term retention at relatively minimal practice cost.
- **Verdict: design choice.** Two misses is a reasoned noise allowance.
- **Code vs evidence:** Consistent with the evidence. No study fixes "two".
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M8. Stale proofs need a fresh cold pass at exit

- **Where:** `services/skill-coverage.js:75` (`STALE_DAYS: 7`), `:319-338`; `services/skill-record.js:246`; `services/skill-gate.js:7-9`.
- **What it does:** When a stage is left, any proven metric whose last counting pass is more than 7 days old is `stale`. It holds the stage until a fresh cold sitting passes, so the battery stands at exit.
- **Evidence:**
    - arthur1998 [395] (as M06).
    - kim2013 [396]: "we discuss how we might improve training by better spacing the iterations between training sessions to support proceduralisaion to improve skill retention". Pop.: theory of skill retention, with training implications.
    - cepeda2008 [23] (as M06).
    - rawson2013 [188]: title-level only (successive relearning).
- **Verdict: design choice.**
- **Code vs evidence:** Consistent. The 7-day cutoff has no study behind the exact number.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M9. The stage gate, per stage, with a default shadow list

- **Where:** `services/skill-gate.js:39-42` (`mode`, default `on`), `:54` (`DEFAULT_SHADOW`), `:56-69` (`stageMode`), `:160-178` (`atGraduation`), `:194-213` (path node).
- **What it does:**
  - A passed Final does not graduate a stage while a provable metric is unproven, but only on stages where every provable metric has a working check.
  - The stages listed below stay in shadow. In shadow, the gate's verdict is only written to the enrollment.
  - Stages in shadow and why:
    - CH2-CH5: games need an owner call; Maia-2.
    - PI5-10: library pieces and [R] tempos.
    - BB1-BB7: decision-scene reliability is not yet measured.
    - BB8-12: a person is needed.
- **Evidence:**
  - aera2014 [150] (AERA/APA/NCME *Standards for Educational and Psychological Testing*, 2014) (official, open access). This is the standard requiring validity and reliability evidence for an intended score use.
  - downing1997 [397]: "A primary type of validity evidence derives from the item development process and item responses." Pop.: high-stakes exam item development.
  - hopkins2000 [383] (as M03).
  - kohavi2009 [147]: "Controlled experiments embody the best scientific design for establishing a causal relationship between changes and their influence on user-observable behavior." Pop.: online controlled experiments. Shadow mode lets a gate's effect be measured before it is enforced.
- **Verdict: design choice.** It is a conservative rollout: no learner is held on a test without evidence it measures the skill.
- **Code vs evidence:** Matches.
- **Status: Built on 25 October 2026** (switch `SCENE_RELIABILITY_GATE`, from 04-P5). BB1-BB3 leave the hand-kept shadow list; a stage with scene sets leaves shadow only when every set reaches 0.8 on 100+ learners. Today BB1-BB7 all stay in shadow. **Open:** BB1-BB7 stay in shadow until real scene data reaches 0.8 (needs data).

04-M10. Decision-scene sets need split-half ≥ 0.8 before they gate

- **Where:** `services/skill-gate.js:49-50`; `content/skill-piano-bb-sittings.js:24-25` ("a set used as a gate needs split-half >= 0.8 after 100 learners").
- **What it does:** It keeps BB1-BB7 in shadow until the drawn-scene sets show adequate reliability.
- **Evidence:**
  - rosch2021 [398]: "the diagnostic results discriminate between performance levels". Pop.: 17 German U16 national-team players and 17 sports-class students, video 4v4 test.

    - cortina1993 [399]: "a demonstration of the effects of test length and dimensionality on alpha". Pop.: psychometric methods. Reliability depends on set length, which matters for 10-30 scene sets.
    - hopkins2000 [383] (as M03).
    - aera2014 [150] (AERA Standards) (as M09).
- **Verdict: moderate.** A reliability bar before high-stakes use is standard practice. 0.8 is a common convention.
- **Code vs evidence:** Gap. No code computes split-half yet: it appears only in comments. The bar is enforced by hand through the shadow list. See P5.
- **Status: Built on 25 October 2026** (switch `SCENE_RELIABILITY_GATE`, from 04-P5). Split-half (Spearman-Brown corrected) is now computed from learners' first scene sittings and gates automatically. **Open:** the gate opens only with 100+ learners per set (needs data).

04-B. How much and when to practise (skill-adaptation, skill-dose)

04-M11. Dose by distance from target, every unproven metric kept

- **Where:** `services/skill-adaptation.js:8-15, 183-194`; `checkPlan` (header lines 4-6).
- **What it does:**
    - The session's practice sets go to the stage's unproven metrics in proportion to how far each is from its target on the learner's recent warm sets.
    - Each unproven metric gets at least one set, so none is dropped.
- **Evidence:**
    - duke2009 [400]: "the strategies employed during practice were more determinative of performance quality at retention than was how much or how long the pianists practiced". Pop.: 17 advanced piano students, next-day retention.
    - ericsson1993 [40]: "Individual differences, even among elite performers, are closely related to assessed amounts of deliberate practice." Pop.: musicians.
    - guadagnoli2004 [41]: "motor tasks represent different challenges for performers of different abilities". Theory paper (challenge point).
    - platz2014 [401]: "Our results corroborate the central role of long-term (deliberate) practice for explaining expert performance in music." Pop.: meta-analysis, 13 studies, N = 788. Shows that accumulated practice time predicts musical achievement (r = 0.61); it does not test weakness targeting.
- **Verdict: moderate.** Targeting weaknesses is the core of deliberate practice. Proportional allocation is a reasoned rule.
- **Code vs evidence:** Matches.
- **Status: Built on 25 October 2026** (switch `SKILL_LOAD_CAPS`, from 04-P7). Basketball court work is trimmed to youth weekly caps (3/5/10/14 h) with rest days.

04-M12. When to call a cold sitting

- **Where:** `services/skill-adaptation.js:45-46` (`READY 3`, `CHECK_DAYS 5`), `:130-137`.

- **What it does:** A cold stage test is scheduled when the last 3 warm sets were all at target, or at least every 5 days. The second rule means a learner whose warm sets under-state them is not held.
- **Evidence:**
  - soderstrom2015 [39] (as M01). Warm performance is an unreliable guide, so it only *times* the test and never decides it.
  - abushanab2013 [378] (as M02). The ranking of practice conditions flipped by the 2-day test.
  - rawson2011 [24] (as M07). Criterion-then-relearn scheduling.
  - kornell2009 [217]: "Unsuccessful retrieval attempts enhanced learning with both types of materials." Pop.: verbal materials. A failed early test is not wasted, which supports the 5-day forced check.
- **Verdict: design choice.**
- **Code vs evidence:** Consistent.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M13. Maintenance share and the over-learning dose

- **Where:** `services/skill-adaptation.js:50-60` (`MAINT_SHARE -1`, i.e. "even"; `OVERLEARN 2`; `OVERLEARN_MIN 4`), `:96-107` (`overlearnDeficit`), `:157-170`.
- **What it does:**
  - While other metrics are unproven, a proven metric keeps the share of sets that even practice would give it.
  - It keeps that share until it has had twice the sets it took to prove it, with a minimum of 4.
- **Evidence:**
  - driskell1992 [402]: "overlearning produces a significant effect on retention of moderate overall magnitude". Pop.: meta-analysis, cognitive and physical tasks.
  - shibata2017 [403]: "overlearning in humans abruptly changes neurochemical processing, to hyperstabilize and protect trained perceptual learning from subsequent new learning". Pop.: adults, visual perceptual learning.
  - rohrer2005over [261]: "overlearning (and its concomitant demand for additional study time) is an inefficient strategy for learning material for meaningfully long periods of time". Pop.: 218 college students, facts and definitions.
  - arthur1998 [395] (as M06).
- **Verdict: mixed.** Over-learning helps retention in the medium term (driskell1992 [402]) but its advantage fades over weeks (rohrer2005over [261]). It may also block new learning immediately afterwards (shibata2017 [403]).
- **Code vs evidence:** Reasonable. The retention sitting (M06) and the stale rule (M08) catch the fade that rohrer2005over [261] describes. The 2× factor is a design choice.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M14. Carry-over across stages is off by default

- **Where:** `services/skill-adaptation.js:61-66` (`CARRY_DAYS 0`, `CARRY_MAX 2`).
- **What it does:** Once a stage is left, its metrics get no further maintenance sets. The comment explains that "the audit's truth model has no spacing benefit".
- **Evidence:**
  - lee1988 [374] (as M01). Massed practice depresses motor learning.
  - donovan1999 [259]: "individuals in spaced practice conditions performed significantly higher than those in massed practice conditions". Pop.: meta-analysis, 63 studies.
  - wiseheart2017 [404]: "No spacing effect was found, suggesting that the effect may not always be demonstrable for complex motor skills". Pop.: adult novices, 17-note piano sequence, lags of 0-15 minutes.
- **Verdict: mixed.** Spacing showed an overall advantage across tasks, moderated by task type and interval (lee1988 [374], donovan1999 [259]). The piano null result used minute-scale lags and a 5-minute test (wiseheart2017 [404]).
- **Code vs evidence:** Gap. The default rests on a simulation whose truth model has no spacing benefit. The day-scale motor literature points the other way. See P8: measure it rather than assume it.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 04-P8, carry-over maintenance as an experiment (needs data).

04-M15. Lesson share while knowledge is open

- **Where:** `services/skill-adaptation.js:118, 153` (`lessonShare = 0.3`); `services/skill-dose.js:103`.
- **What it does:** While the stage's knowledge gate is open, 30% of the session goes to its lessons. The rest is practice.
- **Evidence:**
  - charness2005 [405]: "serious study alone was the strongest predictor of chess skill in both samples". Pop.: two large samples of rated chess players.
  - macnamara2014ps [406]: "deliberate practice explained 26% of the variance in performance for games, 21% for music, 18% for sports". Pop.: meta-analysis across domains.
  - macnamara2016 [407]: "deliberate practice accounted for only 1% of the variance in performance among elite-level performers". Pop.: meta-analysis, sports.
  - renkl2003 [106]: "we propose a fading procedure in which problem-solving elements are successively integrated into example study". Review of example-based learning.
- **Verdict: design choice.** Study and practice both matter, and no study fixes the split.
- **Code vs evidence:** Consistent.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M16. Session length, set length and items per set

- **Where:** `services/skill-dose.js:36` (`ITEMS_PER_SET 5`), `:63-67` (session = midpoint of the stage's `session_minutes`, else 30); `services/skill-adaptation.js:47-48` (`SET_MIN 4`, `SIT_MIN 4`).
- **What it does:** It sets a short, age-banded session made of 4-minute sets.
- **Evidence:**
  - bergeron2015 [408]: "presented recommendations for developing healthy, resilient and capable youth athletes". IOC consensus.
  - difiori2018 [409]: "an overemphasis on competitive success in youth sports may limit the benefits of participation, and could increase the risk of injury, burnout, and disengagement". NBA/USA Basketball expert review.
  - lee1988 [374] (as M01).
  - donovan1999 [259] (as M14).
- **Verdict: design choice.** Short distributed blocks are supported, and the numbers are reasoned.
- **Code vs evidence:** Consistent. Age load caps are documented but not enforced in code (P7).
- **Status: Built on 25 October 2026** (switch `SKILL_LOAD_CAPS`, from 04-P7). Court sessions count toward the weekly cap; screen sets do not.

04-C. The stage checks: items, takes and counts (content/skill-stage-checks.js, skill-checks.js)

04-M17. Answers held on the server; fixed-form sittings

- **Where:** `content/skill-stage-checks.js:7-12`; `content/chess-trainer.js:6-12` ("The page never sees a solution before the puzzle is over").
- **What it does:** Item sittings are served and graded on the server, and keys never reach the page before the sitting ends. Chess positions are computed or taken from Stockfish-verified pools.
- **Evidence:**
  - aera2014 [150] (AERA Standards) (as M09), on test security and score validity.
  - downing1997 [397] (as M09).
  - haladyna2002 [132]: "A taxonomy of 31 multiple-choice item-writing guidelines was validated through a logical process". Classroom-assessment item writing.
  - itc2015security [410]: "The amount and severity of security threats have increased considerably over the past two decades, calling into question the validity of assessments". International Test Commission guidelines; they cover security while tests are developed and administered, including online delivery.
- **Verdict: design choice.** It follows assessment standards.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M18. Mastery cuts of 80-95%

- **Where:** `content/skill-stage-checks.js:38-117`. Examples: CH0 95% of squares; CH1 18/20 and 13/15; PI0 9/10; PI2 aural 9/10; CH2-3 tactics 80%; BB1 rules 8/10 [R].
- **What it does:** A stage's items must be mastered at a high bar, not just passed.
- **Evidence:**
  - kulik1990 [4]: "mastery learning programs have positive effects on the examination performance of students in colleges, high schools, and the upper grades in elementary schools". Meta-analysis of 108 evaluations.
  - bloom1984 [5]: "Formative tests (the same tests used with the conventional group) are given for feedback followed by corrective procedures and parallel formative tests". Mastery learning against tutoring.
  - guskey2007 [60]: "Students for whom the instructional methods and amount of time are appropriate learn well". Review of Bloom's mastery learning.
  - gobet2007 [411]: "The slower player needed 8 times as much practice to reach master level than the faster player." Pop.: 104 Argentinian chess players. This is why pace must follow measured mastery rather than time.
- **Verdict: moderate.** Mastery learning is supported. The exact cut points are design choices.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M19. Unannounced chess items

- **Where:**
  - `content/skill-stage-checks.js:44-55`: 10 of 15 positions mate in one and 5 do not; tactics have motifs not named;
  - `services/skill-checks.js:189-210`: the "no mate" option;
  - gaps noted in the code at `:50` and `:55`.
- **What it does:** The learner must decide whether a tactic exists, as in a game, rather than being told one does.
- **Evidence:**
  - rohrer2015 [12]: "Interleaved practice requires students to choose a strategy on the basis of the problem itself". Pop.: 126 seventh graders.
  - bilalic2008 [412]: "the first idea that comes to mind, triggered by familiar features of a problem, prevents a better solution being found". Pop.: expert chess players, eye tracking. Shows familiar cues can trigger a first solution that blocks a better one; it does not test no-tactic items.
  - burns2004 [414]: "up to 81% of variance in chess skill (measured by rating) was accounted for by how players performed". Pop.: blitz tournament archives. Blitz results track rating closely, which suggests fast pattern recognition accounts for much of skill (correlational).
- **Verdict: moderate.**
- **Code vs evidence:**
  - The mate-or-no-mate sitting is fully unannounced.

    - Partial gap before 25 October 2026 (closed that day; see Status): the CH2 and CH3 tactics sets contained no "nothing to find" positions. The code says so at `:50` and `:55`, and every position in them has a winning move. See P12.
    - vandermaas2005 [413] also supports using response times, which are not used for chess (P12 note).
- **Status: Built on 25 October 2026** (switch `CHESS_NOTHING_TO_FIND`, from 04-P12). CH2 has 5 of 30 and CH3 6 of 30 quiet positions; every position offers "Nothing wins".

04-M20. Moves are played or typed, not picked from a menu

- **Where:** `content/skill-stage-checks.js:34` (`MOVE_PROMPT`); the item kind `move` in `services/skill-checks.js`.
- **What it does:** Chess answers are produced on the board (recall), not recognised in a list.
- **Evidence:**
    - rowland2014 [10]: "initial recall tests yielding larger testing benefits than recognition tests". Meta-analysis.
    - vandermaas2005 [413] (as M19). Choose-a-move tasks are a validated measure of chess skill.
    - roediger2006 [1] (as M06).
- **Verdict: strong.**
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M21. The puzzle trainer: Glicko-2 rating and puzzle choice

- **Where:**
    - `content/chess-trainer.js:33` (`PUZZLE_RD 80`), `:46` (provisional when RD > 110), `:57-66` (`choose`: nearest to the target rating plus up to 120 points of random jitter, excluding the last 30 puzzles seen), `:89` (target = the learner's rating);
    - `content/chess-rating.js:13` (`TAU 0.5`), `:16` (start 1000 / 350 / 0.06).
- **What it does:** It rates the learner and serves puzzles near their own rating. Before 25 October 2026 that meant an expected success of about 50%; practice puzzles now aim at a 0.75 expected score (see Status).
- **Evidence:**
    - glickman1999 [415]: "improves over the commonly used algorithm of Elo by incorporating the variability in parameter estimates". Pop.: chess and tennis ratings.
    - glickman2001 [416]: "cognitive ability not only changes over time, but also can change abruptly". Glicko-2 volatility.
    - eggen2006 [417]: "results in each individual answering about 50% of the items correctly". Pop.: computerised adaptive testing simulations. This is optimal for *measurement*, not necessarily for practice.
    - guadagnoli2004 [41] (as M11).
    - pelanek2016 [77]: title-level only (Elo in adaptive educational systems).

    - klinkenberg2011 [78]: title-level only (Math Garden adaptive practice).
- **Verdict: mixed.** The rating model is well founded. The success target is contested: about 50% maximises information (eggen2006 [417]), while practice systems often aim higher for motivation (the challenge-point framework).
- **Code vs evidence:** Gap. `SKILL-CURRICULA.md:246` says the chess target is "expected score 0.75 (0.6 for calculation)". The code serves nearest-rating puzzles, which gives about 0.5. See P1.
- **Status: Built on 25 October 2026** (switch `CHESS_PRACTICE_TARGET`, from 04-P1). Practice puzzles aim at a 0.75 expected score (191 points below the learner); stage sittings stay at the rating.

04-M22. The puzzle-rating stage cut

- **Where:** `content/skill-stage-checks.js:60-65` (CH4 1590, CH5 1790, RD ≤ 80, ≥ 10 rated puzzles [R]).
- **What it does:** It passes a stage on a rating that is high enough and certain enough.
- **Evidence:**
    - glickman1999 [415] (as M21). RD is the rating's uncertainty.
    - vandermaas2005 [413] (as M19). Puzzle-type tasks measure chess skill reliably.
- **Verdict: design choice.** The cuts are [R]. *After the 27 October 2026 audit this mechanism has 2 quoted sources, fewer than three: a source that did not support it was removed (see the audit section).*
- **Code vs evidence:** Consistent. The doc notes that puzzle ratings run above game ratings, which is why games are a separate metric (M24).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M23. Play-outs against perfect defence

- **Where:** `content/skill-chess-sittings.js:33-50`. KQK: 5 positions, mate within 6 moves of the fastest. KRK: within 8, inside the 50-move rule. KPK: win 4 and hold 2. Lucena and Philidor are played against Stockfish. `services/skill-chess-sittings.js:9-16`.
- **What it does:** It proves technique against the strongest possible defence, not against a lenient bot.
- **Evidence:**
    - ash1990 [418]: "the advantages of segmentation methods in part-task training". Pop.: adults, keyboard task. Endgames are taught as isolated part-tasks (the chess grades' "Backward Chaining" and "Play Against Perfect Defence" activities).
    - pinder2011 [419]: "to ensure functionality and action fidelity in training and learning environments". Representative learning design.
    - thompson1986 [420] and bourzutschky2005 [421]: title-level only (endgame analysis).

- **Verdict: design choice.** *After the 27 October 2026 audit this mechanism has 2 quoted sources, fewer than three: a source that did not support it was removed (see the audit section).*
- **Code vs evidence:** Consistent. The slack values are [R].
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M24. Recorded rapid games against measured bots

- **Where:** `content/skill-chess-sittings.js:26` (15+10), `:35-52` (12 games per band; median hung pieces ≤ 1; score 0.4 or 0.5; performance rating 1250 or 1450). `services/skill-chess-sittings.js:17-23`: a blunder is a move that loses 20% or more of the winning chances.
- **What it does:** It measures play in real games, scored by the engine.
- **Evidence:**
  - regan2011 [422]: "formulas for representing playing strength at chess by the quality of moves played, rather than by the results of games". Pop.: large sets of human tournament games.
  - mcilroy2020 [423]: "predicts human moves at a much higher accuracy than existing engines". Maia, human-like bots.
  - charness2005 [405] (as M15). Tournament play is a weaker *practice* predictor than study. That is a training finding; here games are used as a *measure*.
  - burns2004 [414] (as M19).
  - guid2006 [424]: title-level only.
- **Verdict: moderate.**
- **Code vs evidence:** Consistent. Twelve games carry a large standard error, as SKILL-CURRICULA §1.4 notes. That is why the stages stay in shadow (M09).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M25. CH0 foundations: star mazes at par, set-up from memory

- **Where:** `content/skill-chess-sittings.js:29-32`. 4 mazes, one per piece, 3 of 4 at the breadth-first-search par. All 32 pieces placed correctly.
- **What it does:** It proves that the beginner knows how each piece moves and what the board looks like.
- **Evidence:**
  - gobet1996 [425]: "chess players, like experts in other recall tasks, use long-term memory retrieval structures". Pop.: class A players to grandmasters.
  - chase1973 [426]: title-level only.
  - gobet1996b [589]: "strong players generally maintain some superiority over weak players even with random positions" (review of recall experiments with strong and weak chess players). Scope: Shows that board recall differs by skill even for random positions. It did not test beginners' knowledge of piece moves or board set-up.

    - astudillo2026 [590]: "novices performed better reconstructing real (median 4.70 correct pieces) than random positions (median 3.65 correct pieces)" (51 amateur players with 6 to 24 months of training; cross-sectional). Scope: Shows that novices reconstruct real positions better than random ones, and links real-position reconstruction to rating. It did not test knowing how each piece moves or setting up the full board.
- **Verdict: design choice.**
- **Code vs evidence:** Consistent.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M26. Calculation: type the whole line

- **Where:** `content/skill-chess-sittings.js:45-46` (CH4, 7 of 10).
- **What it does:** The learner types their moves and every reply without moving the pieces.
- **Evidence:**
    - vandermaas2005 [413] (as M19). Includes a predict-a-move task.
    - gobet1996 [425] (as M25; supports internal chess-position memory structures, not move calculation).
    - debruin2007 [427]: title-level only (prediction plus self-explanation in chess endgames).
    - chabris2003 [591]: "the frequency and magnitude of blunders did not differ in rapid versus blindfold play" (23 grandmasters in classical, rapid and rapid blindfold games). Scope: Shows that grandmasters' blunder frequency and size did not differ between rapid and blindfold play, despite the extra visualisation load. It did not test typing whole lines as a training method.
- **Verdict: design choice.**
- **Code vs evidence:** Consistent.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-D. Piano: score-following and the sittings (public/engine/piano, skill-piano-bb-sittings)

04-M27. Score-informed scoring and the timing window

- **Where:** `public/engine/piano/piano-core.js:186-215` (alignment; tolerance `max(80 ms, 15% of a beat)`; rushing and dragging flags).
- **What it does:**
    - Played notes are aligned to the target.
    - Note accuracy and the share of on-beat notes are scored separately.
    - At 80 bpm the window is about ±112 ms.
- **Evidence:**
    - repp2005 [428]: "SMS is controlled jointly by two error correction processes (phase correction and period correction)". Review of tapping.

    - repp2013 [429]: "Sensorimotor synchronization (SMS) is the coordination of rhythmic movement with an external rhythm". Review.
    - friberg1995 [430]: "the relative jnd constant at 2.5% of the tone interonsets above 240 ms". Pop.: 30 listeners. Listeners hear far finer deviations than the window allows.
    - dannenberg2006 [431]: "By relating musical sound to musical notation, these systems generate tireless, expressive musical accompaniment to follow". Score following.
- **Verdict: design choice.**
- **Code vs evidence:** Consistent. The window is deliberately lenient for learners, and the evenness CV (M28) measures fine timing.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M28. Scale evenness as a CV

- **Where:** `public/engine/piano/piano-core.js:217-229`: SD over mean of beat-normalised inter-onset gaps, needing at least 3 gaps. The cuts are PI3/PI4 0.12 (`content/skill-stage-checks.js:91-97`), PI5-7 0.09 and PI8-10 0.06 (`content/skill-piano-bb-sittings.js:46, 52`), all [R].
- **What it does:** It measures scale evenness with a number.
- **Evidence:**
    - jabusch2004 [432]: "Scale Analysis is an effective and precise tool for quantification of focal dystonia in pianists and provides fine resolution". Pop.: 16 professional pianists, MIDI scale analysis.
    - vanvugt2013 [433]: "even non-expressive playing of scales reveals consistent, partially effector-unspecific, but inaudible inter-individual differences". Pop.: professional pianists.
    - jabusch2009 [434]: "Expert pianists with an average daily practice time of 3.75h or more showed an improvement of performance in this selected motor skill." Pop.: 19 pianists over 27 months.
    - simmons2006 [377] (as M02): "temporal evenness improved in the absence of practice 24 hours after training" (same abstract).
- **Verdict: moderate.** The measure is validated. The cut points are not normed.
- **Code vs evidence:** Consistent. The cuts need norms from learners.
- **Status: Built on 25 October 2026** (no switch: a fact re-verification fix). Scale parts show the scale's own note spelling; 59 of 102 items showed a wrong name before. Grading is unchanged.

04-M29. Scales drawn at random

- **Where:** `content/skill-stage-checks.js:91-97`: PI3 and PI4 "Two of this grade's scales, chosen at random". The same holds for PI5-10 (`content/skill-piano-bb-sittings.js:46, 52`).

- **What it does:** The test samples the battery randomly, so learners cannot over-practise one scale.
- **Evidence:**
  - shea1979 [42]: "retention was greater following high interference (random) acquisition than after low interference (blocked) acquisition". Pop.: adults, three motor tasks.
  - brady2004 [435]: "Significant differences were also obtained between the effect sizes for adults (.50) and those for younger learners (.10)." Meta-analysis, 61 studies.
  - abushanab2013 [378] (as M02).
  - wulf2002 [436]: "situations with low processing demands benefit from practice conditions that increase the load and challenge the performer". Review, complex skills.
  - magill1990 [437]: title-level only.
- **Verdict: moderate.** Random practice is robust in adults and weaker in children (brady2004 [435]).
- **Code vs evidence:** Matches for testing. Practice order for young children should ramp from blocked to random, as documented in SKILL-CURRICULA §5. The skill dose does not show that ramp, but it is low priority.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M30. Sight-reading staircase with never-seen excerpts

- **Where:**
  - `content/piano-sight-reading.js:29-38`: SR1-SR7 levels, `STAGE_MINIMUM`, and a note that SR8-10 are listed but not generated;
  - `services/skill-piano-bb-sittings.js:59-62`: the excerpt hash;
  - `:104-108`: up to 30 tries for an unseen excerpt;
  - `content/skill-piano-bb-sittings.js:38-45`: 2 excerpts, each played once, 80% of the notes and 80% on the beat.
- **What it does:** Each sight-reading sitting is a new excerpt at the stage's level, drawn in notation.
- **Evidence:**
  - mishra2014a [438]: "treatments categorized as “Aural Training,” “Controlled Reading,” “Creative Activities,” and “Singing/Solfege” significantly and positively affecting sightreading". Meta-analysis of 92 studies.
  - mishra2014b [439]: "improvisational skills, ear-training ability, technical ability, and music knowledge correlating most closely with sight-reading". Meta-analysis.
  - kopiez2006 [440]: "when sight reading stimuli is easy, general pianistic expertise is sufficient to be able to excel". Pop.: 52 piano students, 5 difficulty levels.
  - mcpherson1997 [441]: "Performing a repertoire of rehearsed music was found to be influenced most by an ability to sight-read". Pop.: 101 high-school wind players.

    - lehmann1996 [442]: title-level only.
- **Verdict: moderate.**
- **Code vs evidence:** Matches for PI2-PI7. Gap: SR8-SR10 are not generated (`piano-sight-reading.js:23`), so PI8-10 have no sight-reading sitting.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M31. Memory restarts from random bars

- **Where:** `services/skill-piano-bb-sittings.js:147-162`; `content/skill-piano-bb-sittings.js:48-49` (3 restarts of 2 bars each, nothing shown).
- **What it does:** It proves a memorised piece can be restarted anywhere, not just played from the top.
- **Evidence:**
    - chaffin2002 [443]: "the pianist used the formal structure as a retrieval scheme and performance cues as retrieval cues". Concert pianist case study.
    - williamon2002 [444]: "This article explores the use of structure in the encoding and retrieval of music and its relation to level of skill." Pop.: 22 pianists at four skill levels.
    - chaffin2010 [445]: "The cellist organized her practice around the musical structure, developed interpretation before working on technique and practised memory retrieval at each stage." Cellist case study.
- **Verdict: moderate.** Converging expert studies, mostly case studies.
- **Code vs evidence:** Partial. The bars are random, but the expert work points to *structural* cue points: section starts and performance cues. Restarting at phrase or section boundaries would match the evidence more closely. This is a minor point, not a proposal.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M32. Dynamics from MIDI velocity

- **Where:** `public/engine/piano/piano-core.js:231-237` (MIDI only); `content/skill-piano-bb-sittings.js:50-51` (median f velocity at least 15 above median p, out of 127, plus 90% of the notes).
- **What it does:** It checks soft and loud contrast only where the input can measure it.
- **Evidence:**
    - jabusch2004 [432] (as M28). Its method analysed "Key velocities and timing parameters" from MIDI.
    - goebl2001 [446]: "an emphasized voice (the melody) tends to be played not only louder than the other voices, but also about 30 ms earlier". Pop.: 22 skilled pianists on a computer-monitored grand.
    - lerch2020 [447]: "Different performances of the same composition may vary in terms of performance parameters such as timing or dynamics". Review of music performance analysis.
- **Verdict: design choice.** The 15-velocity margin has no study behind it.
- **Code vs evidence:** Consistent.

- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M33. Legato pedalling from CC64

- **Where:** `content/skill-piano-bb-sittings.js:54-55` (90% of chord changes lifted within 0.15 s, 80% of chords right); `public/engine/piano/piano-core.js:1044`.
- **What it does:** It grades pedal changes against chord onsets, on MIDI only.
- **Evidence:**
  - repp1996 [448]: "Pedal timing emerges as having a complex pattern that is sensitive to local and global tempo changes in varying degrees". Pop.: two pianists.
  - lerch2020 [447] (general background on performance-timing analysis; as M32).
  - liang2018 [592]: "The onset and offset times of each pedaling technique were computed using signal processing algorithms." (piano pedaling measured with a sensor beside the sound; signal-processing and machine-learning study). Scope: Supports that sustain-pedal onset and offset can be measured and classified by technique. It did not use MIDI CC64 or grade pedalling against chord onsets.
- **Verdict: design choice.**
- **Code vs evidence:** Consistent. Repp's tempo dependence suggests the 150 ms window may need to scale with tempo.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M34. Musicality is never machine-graded

- **Where:** `public/engine/piano/piano-core.js:250` (`notMeasured: ['phrasing', 'tone', 'character', ...]`); `services/skill-gate.js:17-22` (recitals are attest-only).
- **What it does:** The app reports what it cannot measure and never scores interpretation.
- **Evidence:**
  - lerch2020 [447] (as M32). Performance assessment is still a research field.
  - bergee1997 [392]: "Ranging from .23 to .93, total score faculty interjudge reliability was mixed." Even expert judges disagree.
  - vanvugt2013 [433] (as M28). Timing signatures can be inaudible, so measurable is not the same as musical.
- **Verdict: moderate.**
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M35. The microphone pipeline

- **Where:**
  - `public/engine/piano/piano-core.js:45-108`: YIN with minHz 55, threshold 0.12 and minRms 0.01;
  - `:472-500`: the chord-ghost check and spectrum verifier;

    - `public/engine/piano/piano-song.js:18-22, 96-99`: the per-device latency measurement and "told which notes are due";
    - `content/skill-stage-checks.js:84-95`: lower mic note bars (0.93 or 0.95), and one hand at a time on a microphone.
- **What it does:** It hears single lines reliably. Chords are checked against the expected notes. Mic results are always labelled as estimates.
- **Evidence:**
    - decheveigne2002 [385] (as M04; supports YIN fundamental-frequency estimation for speech and music, not polyphonic chord detection).
    - emiya2010 [386] (as M04).
    - kong2021 [387] (as M04; supports piano-note transcription, not the broader claim that polyphonic audio is hard).
    - benetos2013 [390]: title-level only.
- **Verdict: strong.** Polyphonic audio is hard, and the design respects that.
- **Code vs evidence:** Matches. The YIN chord bug is documented in `SKILL-APPS-MECHANICS.md` §2.4.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M36. Learn mode waits; play mode runs at tempo

- **Where:** `public/engine/piano/piano-song.js:11-13` (learn waits at each note; play has a count-in, and the app plays the other hand at full, ¾ or ½ speed), `:190`; `piano-core.js:396-430` (stars: 60% notes and 40% timing; 3 stars from 0.9; capped at 2 in learn mode or below 0.75 speed).
- **What it does:** First meetings with a piece are guided and nearly errorless. Full credit needs independent play in time.
- **Evidence:**
    - maxwell2001 [449]: "the reduction of errors during learning encourages the use of implicit, unselective learning processes, which confer insusceptibility to performance breakdown under distraction". Pop.: adults, golf putting.
    - duke2009 [400] (as M11). The percentage of correctly performed trials was significantly related to next-day retention-test quality.
    - fitts1954 [450]: title-level only (speed-accuracy).
    - poolton2005 [593]: "The environment was constrained early in learning to minimize performance error." (golf putting over 400 trials; constrained early learning against explicit instruction). Scope: Supports minimising errors early in learning. It did not test note-by-note waiting, slow practice or independent play in time.
    - furuya2013 [594]: "Piano practice of sequential finger movements at a submaximal speed over days progressively increased the maximum speed of trained movements." (piano finger-movement practice over days, with retention and transfer tests). Scope: Supports practising below top speed before playing at tempo. It did not test waiting at each note or errorless learning.
- **Verdict: moderate.**

- **Code vs evidence:** Matches. Before 25 October 2026 the tempo ladder and hotspot loop documented in `SKILL-CURRICULA.md` §4 were not built (P3); both were built that day (see Status).
- **Status: Built on 25 October 2026** (switch `PIANO_TEMPO_LADDER`, from 04-P3). Warm parts start at 65% tempo (+10% after 2 clean takes, -10% after 2 misses); `PIANO_HOTSPOT_LOOP` loops the worst 1-2 bars at 65%; `PIANO_FEEDBACK_FADE` fades feedback after 2 clean takes. Cold takes stay at target tempo.

04-M37. Self-rate before the measurement

- **Where:** `public/engine/piano/piano-blocks.js:158-191` (`selfRate`; the report says whether the learner's rating matched the measurement); `services/block-levels.js:450`.
- **What it does:** The learner rates the take first, then sees notes, timing and evenness, and is told whether their own judgement matched.
- **Evidence:**
  - hewitt2001 [393] (as M05). Self-evaluation helped *with* a model.
  - kruger1999 [391] (as M05).
  - bergee1997 [392] (as M05): "self-evaluation correlated poorly with both faculty and peer evaluation".
  - andrade2019 [394] (as M05).
  - potdevin2018 [451]: "This study highlights the potential relevance of using VFB in fostering motor learning, motivation and self-assessment during a PE programme with young children." Pop.: novice children, gymnastics.
- **Verdict: mixed.** Calibration practice is useful, but self-evaluation alone is poorly accurate. The gain in hewitt2001 [393] came with a model to compare against.
- **Code vs evidence:** Partial. There is no model performance to listen to before rating. See P4, which covers piano too.
- **Status: Built on 25 October 2026** (switch `PIANO_FEEDBACK_FADE`, from 04-P2). After 2 clean warm takes, a take shows only what is outside the bar; "Show all the numbers" shows the rest. **Open:** 04-P4, anchor clips for piano self-rating (needs media).

04-M38. Ear items

- **Where:** `content/skill-stage-checks.js:74-75` (PI0 echo: 5 patterns, 4 of 5), `:82-83` (PI1: a 5-note tune within 3 tries), `:89-90` (PI2: major or minor, 9 of 10); `services/block-levels.js:465` (`Piano_EarEcho`).
- **What it does:** Playing by ear is a proved skill, not an optional extra.
- **Evidence:**
  - mcpherson2005 [452]: "the importance of helping students to develop a repertoire of task-appropriate strategies". Pop.: 157 children aged 7-9, tested on five skills including playing by ear.
  - mcpherson1997 [441]: "an ability to improvise was most markedly influenced by an ability to play by ear" (same abstract as M30's quote).

    - woody2010 [453]: "vernacular musicians required fewer trials than formal musicians". Pop.: 24 musicians, listen-then-perform cycles.
    - mishra2014b [439] (as M30). Ear training correlates with sight-reading.
- **Verdict: moderate.**
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M39. Steady beat

- **Where:** `content/skill-stage-checks.js:72-73` (16 beats at 80 bpm, 75% on the beat).
- **What it does:** It is the first rhythm gate for ages 3-6.
- **Evidence:**
    - drake2000 [454]: "ability to synchronize tapping and discriminate tempo (improved attunement)". Developmental study.
    - provasi2003 [455]: "children could slow down their tap rhythm when the auditory stimulation became slower". Pop.: children aged 2½-4.
    - repp2005 [428] (as M27), a general review of sensorimotor synchronization and tapping variability.
- **Verdict: design choice.** 75% is a lenient, age-appropriate bar.
- **Code vs evidence:** Consistent. Young children synchronise less precisely, and the wide window (M27) allows for that.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M40. A piece learned alone

- **Where:** `content/skill-piano-bb-sittings.js:56-57` (a new 16-bar SR7 piece, kept 14 days, played cold).
- **What it does:** It proves independent learning: the strategy, not just the piece.
- **Evidence:**
    - mcpherson2005 [452] (as M38).
    - hallam2012 [456]: "use of recordings for listening and feedback and use of the metronome". Pop.: 3,325 young instrumentalists. Practice strategies develop with expertise.
    - duke2009 [400] (as M11).
- **Verdict: design choice.**
- **Code vs evidence:** Consistent.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M41. Two-day lessons

- **Where:**
    - `curriculum-build/SKILL-CURRICULA.md:187` (piano: "The day-2 cold take is the learning score"), `:225-227` (basketball day A and day B);

    - in code, the cold-sitting rule (M01, M02) for stage metrics;
    - FACT-REVERIFY §2 (B05:99 "Judge Shooting by the Next-Day Cold Set", confirmed against Shamshiri et al. 2025).
- **What it does:** The lesson's own score comes from a cold attempt on the next day, not from the end of the practice session.
- **Evidence:**
    - walker2002 [376] (as M02; supports sleep-dependent motor-skill improvement after practice, not specifically a child day-2 scoring rule).
    - simmons2006 [377] (as M02).
    - abushanab2013 [378] (as M02).
    - mazza2016 [239] (as M02; indirect evidence from foreign-vocabulary learning for spacing sessions across sleep).
    - rickard2008 [379] (as M02).
- **Verdict: mixed**, for the same reason as M02. The day-2 test is right; the sleep-gain rationale is contested and weaker for children.
- **Code vs evidence:** Partial. Stage metrics are enforced this way. The per-lesson day-2 cold take is in the lesson design (the "Spaced Recall" activity, M42), but no code makes the lesson's score the day-2 take. This is noted, not proposed: it belongs to the lesson engine, which is outside this section.
- **Status: Built on 25 October 2026** (switch `SKILL_TWO_DAY_TEXT`, from 04-P9). The stage pages describe the two-day rule as a delayed test.

04-M42. The nine-activity lesson ladder

- **Where:**
    - `content/basketball-grades.js:43-51`: Watch and Copy, then Play It Slowly, Game With Help, Game on Your Own, New Way to Play, Play With a Friend, Oops Fix It, Free Recall and Mixed Review, Spaced Recall;
    - `content/piano-grades.js:43`: for example Hands Separately, Hands Together Slowly, Section at Target Tempo, Join the Seams, Critique Your Take;
    - `content/chess-grades.js:43`: for example Solve With Hints, Solve on Your Own, Backward Chaining, Play Against Perfect Defence.
- **What it does:** Each lesson moves through:

1. a model;
2. guided practice;
3. independent practice;
4. variation;
5. a peer;
6. an error-fixing task;
7. two retrieval activities.

- **Evidence:**

    - renkl2003 [106] (as M15). Fading from examples to problems.
    - kalyuga2003 [6]: "Instructional techniques that are highly effective with inexperienced learners can lose their effectiveness and even have negative consequences when used with more experienced learners." Review.
    - stemarie2012 [457]: "practitioners should first assess the observer's characteristics and the task characteristics for which any observation intervention is being created". Review of observation in sport.
    - wulf2010 [458]: "Observation of others, particularly when it is combined with physical practice, can make important contributions to learning." Review.
    - kornell2009 [217] (as M12). Errors and retrieval attempts.
    - rowland2014 [10] (as M20).
    - cepeda2008 [23] (as M06).
- **Verdict: strong** for each component.
- **Code vs evidence:** Matches. The ladder is fixed for every learner. Expertise reversal (kalyuga2003 [6]) supports adapting instruction for learners who show mastery, which the general adaptation layer handles (section 03).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-E. Basketball: counts, scenes, tracking and the court (skill-checks, skill-piano-bb-sittings, public/engine/basketball)

04-M43. Beta-posterior make-rate gates

- **Where:** `services/skill-checks.js:94-104` (`rateAtLeast`: exact P(p ≥ bar) under a uniform Beta(1,1) prior); `content/skill-stage-checks.js:111-124`. Examples:
    - BB1 lay-ups: P(rate ≥ 0.5) ≥ 0.9 with n ≥ 10 per hand;
    - BB2 form shots: ≥ 0.6 at confidence 0.8 with n ≥ 20;
    - BB3 free throws: ≥ 0.5 at 0.8 with n ≥ 20;
    - BB4-5 free throws: ≥ 0.65 at 0.8 with n ≥ 50 (`content/skill-piano-bb-sittings.js:73-74`).
- **What it does:** A shooting gate passes on the probability that the learner's true rate clears the bar, not on one lucky set.
- **Evidence:**
    - brown2001 [382] (for related small-sample binomial interval guidance).
    - gelman2013 [381] (as M03).
    - thall1994 [595]: "Sampling continues until E is shown to be either promising or not promising relative to S with high posterior probability" (phase IIB single-arm trials; binary responses monitored continuously). Scope: Supports deciding on the posterior probability that a binary success rate is promising. It did not use a Beta(1,1) prior or a learner gate.
    - tuyl2008 [596]: "Beta distributions with both parameters equal to 0, ½, or 1 are the usual choices for “noninformative” priors" (Bayesian estimation of a binomial proportion; two worked examples). Scope: Notes that Beta(1,1) is among the

usual non-informative priors for a binomial rate, and favours the Bayes–Laplace prior over Jeffreys in its sensitivity analysis. It did not study pass gates.

- **Verdict: strong**, as a method.
- **Code vs evidence:** Matches. The bars are [R]. FACT-REVERIFY §3.2 confirms the FIBA free-throw line: "It shall have its furthest edge 5.80 m from the inner edge of the endline" (fiba2026rules [460] (FIBA *Official Basketball Rules 2026 v1.1*), Art. 2.5.3, https://assets.fiba.basketball/image/upload/documents-corporate-fiba-official-rules-2026-v1-1.pdf). Brown recommends the Jeffreys prior for small n; the uniform prior used here differs very little at n ≥ 10.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M44. Shooting by zone

- **Where:** `services/skill-piano-bb-sittings.js:288-316`: per-zone Beta rule, an 80% interval shown, every zone needing `min_n` shots. `content/skill-piano-bb-sittings.js:71-76`:
  - BB4-5: 5 zones, ≥ 17 shots each, mid-range 45% or three 33% [R];
  - BB6-7: ≥ 25 each, 50% or 36% [R].
- **What it does:** It builds a shooting profile with honest uncertainty for each zone.
- **Evidence:**
  - brown2001 [382] (as M03).
  - hopkins2000 [383] (as M03).
  - gilovich1985 [462]: title-level only (people misread random shooting streaks).
  - shortridge2014 [597]: "We propose an empirical Bayesian smoothing rate estimate that uses a novel local spatial neighborhood tailored for basketball shooting." (2011–2012 NBA shot data). Scope: Supports smoothing noisy local shooting rates in basketball. It did not set a minimum number of shots per zone or show per-zone intervals.
  - wongtoi2023 [598]: "we quantify the shooting pattern via the field goal attempts and percentages over twelve non-overlapping regions around the front court." (NBA players, two seasons, twelve court regions). Scope: Supports describing shooting by court region. It did not study small-sample noise, minimum shots or intervals.
- **Verdict: design choice.** The method is sound and the marks are [R].
- **Code vs evidence:** Consistent. At n = 17 a zone's 80% interval is still wide (about ±15 points), so the gate is deliberately cautious. The code shows the interval, as the evidence suggests it should.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M45. Own baseline, never a ranking

- **Where:**
  - `services/skill-checks.js:774-806`: baseline and growth rules; the first count sets the baseline;

    - `content/skill-stage-checks.js:105-106`: BB0 catch and bounce, against the learner's own baseline;
    - `:125-126`: Mikan drill, 20% growth;
    - `SKILL-CURRICULA.md` §2 rule 7.
- **What it does:** Children's athletic numbers are compared only with their own past numbers.
- **Evidence:**
    - cobley2009 [463]: "Annual age-grouping is a common organizational strategy in sport. However, such a strategy appears to promote relative age effects (RAEs)." Meta-analysis of 38 studies.
    - lloyd2012 [464]: "PREVIOUS LONG-TERM ATHLETE DEVELOPMENT MODELS HAVE CLASSIFIED YOUTH-BASED TRAINING METHODOLOGIES IN RELATION TO CHRONOLOGIC AGE GROUPS, AN APPROACH THAT HAS DISTINCT LIMITATIONS". Youth physical development model.
    - difiori2018 [409] (as M16).
    - crane2015 [465]: "five major areas emerged: lack of enjoyment, perceptions of competence, social pressures, competing priorities and physical factors (maturation and injuries)". Systematic review of 43 studies on youth sport dropout.
- **Verdict: moderate.**
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M46. Decision scenes

- **Where:**
    - `content/skill-piano-bb-sittings.js:60-70`: BB1 2v1, 8 of 10; BB2 3v3, 12 of 15; BB3 4v4, each shown twice, 15 of 18, response times recorded but not gated; BB4-7 at 80%;
    - `content/basketball-scenes.js:26`: "Defence and live decisions are never machine-graded";
    - `:161-170`: the key is re-derived from the players' positions, for example the shot-clock rules in late-clock scenes.
- **What it does:** It trains and tests reads, meaning which option is right, on drawn plays paused at the moment of decision.
- **Evidence:**
    - rosch2021 [398] (as M10): "response time discriminated within the elite athletes' performance data but not between performance levels" (same abstract).
    - mann2007 [466]: "experts are better than nonexperts in picking up perceptual cues, as revealed by measures of response accuracy and response time". Meta-analysis, 42 studies.

    - zhu2024 [467]: "its transfer effect on real-game performance improvement (ES = 0.65) was inferior to laboratory performance improvement (ES = 1.51)". Meta-analysis, 22 studies, team sports.
    - broadbent2015 [468]: "The main aim of any training intervention should be the efficacy of retention and transfer of learning from training to field situations". Review.
    - hadlow2018 [469]: "Modified perceptual training (MPT) collectively describes on- or off-field sports training tasks that are specifically designed to develop visual and perceptual-cognitive skill." Framework.
    - abadrobles2020 [470]: "they did not show significant improvements in skill execution". Meta-analysis of game-based teaching, which did improve decision making.
    - harvey2014 [471]: "development of decision-making skills and effective decision-makers". Review of game-centred approaches.
- **Verdict: moderate.** Decision training transfers, but at about half the lab effect. Drawn top-down scenes are less representative than first-person video.
- **Code vs evidence:** Partial:
    - The scenes are drawn from above, not first-person (the code admits this gap at `:65`).
    - Response time is recorded and not used.
    - Reliability is not yet measured (M10).
    - See P5 and P10.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 04-P10, first-person scenes and response-time bands (needs media and data).

04-M47. Applied tasks: box scores and rules

- **Where:** `services/skill-piano-bb-sittings.js:240-258` (number keys with a tolerance of 0.051); `content/skill-piano-bb-sittings.js:77-79` (BB8-12: 8 stats items and 2 rules items, 8 of 10); `public/engine/basketball/tracking.js:489-490` (eFG% in the shot log); `content/skill-stage-checks.js:116-117` (BB1 rules).
- **What it does:** It checks the parts of advanced basketball knowledge that a machine can verify.
- **Evidence:**
    - kubatko2007 [461] (as M44). eFG%, TS%, possessions and ratings come from its framework.
    - fiba2026rules [460] (FIBA *Official Basketball Rules 2026 v1.1*) (official URL as M43). FACT-REVERIFY §2 confirmed the B06 rules items: the 14-second reset (Art. 29), the no-charge semicircle (radius 1.30 m to the outer edge of its line, Art. 2.5.7; 1.25 m to the inner edge) and the injury allowance of about 15 seconds (Art. 5.3).
    - rohrer2015 [12] (as M19). Mixed problem types.

    - bangert1991 [472]: "Feedback effects were found to vary with control for presearch availability, type of feedback, use of pretests, and type of instruction". Meta-analysis of 58 effect sizes.
- **Verdict: design choice.**
- **Code vs evidence:** Consistent. The scouting report and game plan are honestly left to a person (`:78`).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M48. Camera tracking on the device; the learner confirms

- **Where:**
    - `public/engine/basketball/pose.js:1-26`: MediaPipe Pose Landmarker lite, served locally, up to 2 bodies, detection confidence 0.5, GPU with a CPU fallback;
    - `tracking.js:100-130`: homography from the learner's floor taps, which refuses collinear taps;
    - `tracking.js:404-449`: the rim-crossing make or miss;
    - `camera-blocks.js:5-11, 33`: results are marked estimated, the learner fixes them, and there is a manual tap fallback;
    - `content/skill-stage-checks.js:20-23`: a camera count counts only once confirmed.
- **What it does:** The phone counts shots and places them on the court. The count is an estimate the learner checks, and no picture leaves the device.
- **Evidence:**
    - cao2021 [473]: "the first open-source realtime system for multi-person 2D pose detection, including body, foot, hand, and facial keypoints". OpenPose.
    - stenum2021 [389] (as M04). Joint-angle errors of 4.0-7.4°.
    - washabaugh2022 [388] (as M04).
    - hartley2004 [474]: "Techniques for solving this problem are taken from projective geometry and photogrammetry." Textbook (Cambridge UP), the basis of the floor homography.
- **Verdict: moderate.** Pose accuracy is known for gait. Make/miss from a phone is unvalidated, as `SKILL-CURRICULA.md` §1.3 notes.
- **Code vs evidence:** Matches the honesty rule. Validation is missing (P6).
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 04-P6, validating camera make/miss and pose per device (needs data).

04-M49. Film it, rate it

- **Where:** `public/engine/basketball/basketball-blocks.js:212-323`:
    - the clip stays on the device;
    - "Watch your clip at least twice: once for the whole, once for one detail";
    - a Good / OK / Needs work rubric;

    - "Your one focus for the next three sessions ... close in first, then from your normal spots, then with a defender; film again after".
- `services/block-levels.js:342-355`: the rubric has 2-4 criteria "in external-cue words".
- **What it does:** It runs a closed loop of film, review, rating, a single focus, practice, and filming again.
- **Evidence:**
    - potdevin2018 [451] (as M37).
    - modinger2022 [475]: "Video-based visual feedback seems to be effective to enhance motor learning in physical education and seems to be more effective than solely verbal feedback." Systematic review of 11 studies.
    - stemarie2012 [457] (as M42).
    - hewitt2001 [393] (as M05).
    - wulf2010 [458] (as M42).
- **Verdict: moderate.**
- **Code vs evidence:** Partial. The reviews' strongest effects come with expert models or anchors next to self-modelling (modinger2022 [475], hewitt2001 [393]). There is no model or anchor clip yet. BB3's "rated against anchor clips" in `SKILL-CURRICULA.md` §5 is not built. See P4. The progression from close-in to defender matches the constraints literature (M52).
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 04-P4, model and anchor clips for "film it" (needs media).

04-M50. One external cue per drill

- **Where:**
    - `services/block-levels.js:377`: each drill requires a `cue`;
    - `:409`: "One external cue";
    - `:351`: the rubric "in external-cue words";
    - `content/basketball-grades.js:6171`: the lesson "External Focus Cues: Aim at the Target";
    - FACT-REVERIFY §2: B07 corrected to "the back of the ring, not your elbow".
- **What it does:** Every drill names one cue about the movement's effect, never the body.
- **Evidence:**
    - chua2021 [476]: "an external focus is superior to an internal focus whether considering tests of motor performance or learning". Meta-analysis, 73 and 40 studies.
    - wulf2013 [477]: "an external focus (i.e., on the movement effect) enhances motor performance and learning relative to an internal focus (i.e., on body movements)". Review of 15 years.
    - wulf2016 [478]: "We speculate that enhanced expectancies and an external focus propel performers' cognitive and motor systems in productive "forward" directions". OPTIMAL theory.

    - tran2008 [479]: "should aim the ball towards the back of the ring, and should launch the ball at 52 degrees". Simulated free throws.
- **Verdict: strong.**
- **Code vs evidence:** Matches. FACT-REVERIFY fixed the two lines that had the aim point wrong.
- **Status: Built on 25 October 2026** (no switch: a fact re-verification fix). The B03 aim cue now says "toward the back of the ring", and B07's says "the back of the ring, not your elbow".

04-M51. Court session: health check, warm-up first, pain and enjoyment logged

- **Where:** `services/block-levels.js:368-380` ("a health check, the warm-up always first, timed drills with one cue each, then effort, pain and enjoyment, logged"); `content/skill-stage-checks.js:109-110` (BB0 safety: 4 of 4).
- **What it does:** It puts safety and enjoyment into every court session.
- **Evidence:**
    - soligard2008 [480]: "This indicates that a structured warm-up programme can prevent injuries in young female football players." Cluster RCT with 1,892 players aged 13-17.
    - longo2012 [481]: "The FIFA 11+ warm-up program is effective in reducing the rates of injuries in elite male basketball players." RCT.
    - emery2015 [482]: "There is evidence for the effectiveness of neuromuscular training strategies in the reduction of injury in numerous team sports." Meta-analysis, youth.
    - difiori2014 [483]: "Such an excessive focus on early intensive training and competition at young ages rather than skill development can lead to overuse injury and burnout." AMSSM position statement.
    - crane2015 [465] (as M45). Lack of enjoyment drives dropout.
    - bergeron2015 [408] (as M16).
- **Verdict: strong.**
- **Code vs evidence:** Partial. The warm-up content is free text, up to 6 items of 20-300 s each. The trials used *structured neuromuscular* programmes such as the 11+. See P11.
- **Status: Built on 25 October 2026** (switch `SKILL_LOAD_CAPS`, from 04-P7). Rest days have no court work; court sets are refused once the week is used. **Open:** 04-P11, a structured warm-up template (exercise content not sourced).

04-M52. Handle and react; eyes-up dribbling

- **Where:** `services/block-levels.js:397-410` (a health check first; targets wait 3 s at the steady pace and 1.8 s at the quick pace; touch them with the free hand while dribbling); `content/skill-stage-checks.js:113-114` (call out a held-up number while dribbling; weak hand at least 70% of the strong hand).
- **What it does:** It practises ball handling under a perceptual load, as in games.
- **Evidence:**

    - mann2007 [466] (as M46).
    - renshaw2010 [484]: "physical educators could use task, performer and environmental constraints to channel acquisition of movement skills and decision-making behaviours in learners". Constraints-led approach.
    - pinder2011 [419] (as M23).
    - broadbent2015 [468] (as M46).
- **Verdict: design choice.** The principle is supported. The timings and the 0.7 ratio are reasoned values.
- **Code vs evidence:** Consistent.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-F. Learning games (game-lab, game-payloads, GAME_HOLDS)

04-M53. Serving holds

- **Where:** `services/game-payloads.js:43-83`; `content/game-holds.json` (70 levels, the English Proficient sets the 2026-10-23 audit found wrong). A held level is not served, not regenerated and not answered, and opens as an ordinary session. It is released only by a `verified_replacement` that two blind judges kept. `GAME_HOLDS=off` lifts the holds.
- **What it does:** It keeps questions known to be faulty away from learners.
- **Evidence:**
    - downing2005 [349]: "flawed multiple-choice test items, which violate well established and evidence-based principles of effective item writing, disadvantage some medical students". Pop.: medical students, 4 exams.
    - tarrant2008 [359]: "47.3% of all items were flawed." Pop.: nursing exams, 10 papers.
    - haladyna2002 [132] (as M17).
    - downing1997 [397] (as M09).
- **Verdict: moderate.**
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M54. Payload inspection, including the explanation on every question

- **Where:** `services/game-payloads.js:160-198`. Each question must have:
    - a prompt and an explanation;
    - exactly the expected number of choices, with no duplicates;
    - an answer index in range;
    - no markup;
    - a difficulty from 1 to 10;
    - no answer identifiable by its length.
- Across the set, no answer slot may hold more than 75% of keys once 6 or more are answered. `services/game-lab.js:14-19` rotates answer slots.

- **What it does:** It rejects sets with item-writing flaws and makes every answer carry a one-sentence why.
- **Evidence:**
  - haladyna2002 [132] (as M17).
  - haladyna1989 [485]: "for nearly one half the rules, no research was found". This is a caution that some item rules are convention.
  - attali2003 [486]: "questions with middle correct answers are easier and less discriminating than questions with extreme correct answers". Pop.: test takers and test makers.
  - millman1965 [487]: title-level only (test-wiseness, length cues).
  - hattie2007 [123]: "the type of feedback and the way it is given can be differentially effective". Review.
  - wisniewski2020 [488]: "feedback has higher impact on cognitive and motor skills outcomes than on motivational and behavioral outcomes". Meta-analysis, 435 studies.
  - shute2008 [16]: "formative feedback should be nonevaluative, supportive, timely, and specific". Review.
  - bangert1991 [472] (as M47).
- **Verdict: strong.** Item-writing guidance and feedback both rest on reviews and meta-analyses.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M55. The explanation-key gate

- **Where:** `services/game-payloads.js:509-519` (`KEY_GATE=on`). A question whose explanation describes another option is re-keyed blind. If the blind check does not confirm it, it is rewritten once and checked again. If it still fails, the set is not stored as clean.
- **What it does:** It catches mismatches between key and explanation.
- **Evidence:**
  - downing1997 [397] (as M09). Quality-assurance review of items.
  - haladyna2002 [132] (as M17).
  - tarrant2008 [359] (as M53).
- **Verdict: design choice.**
- **Code vs evidence:** Consistent.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M56. evenOrder

- **Where:** `services/game-payloads.js:202-244`: questions are grouped by difficulty, each group is shuffled, and groups are dealt in proportion with a seeded order. Tailored plans may instead set a "confidence first then climb" curve (`:607`).

- **What it does:** It makes every stretch of a game mix the set's difficulties, so no sitting is all easy or all hard.
- **Evidence:**
  - leary1985 [489]: "The literature has produced evidence of context effects, but has not demonstrated that the effects are so strong as to invalidate test theory". Review of item order.
  - eggen2006 [417]: computerized adaptive testing can select easier or more difficult item sets.
  - hamari2016 [490] and engeser2008 [491]: title-level only (challenge and flow).
  - hambleton1974 [599]: "item order had an effect on the stress generated during the test" (mathematics test, difficult-to-easy against easy-to-difficult order, high and low test-anxious examinees). Scope: Shows that item order changes performance and stress. It did not test mixed difficulty in a game.
  - wang2021 [600]: "the difficulty of item group and the administration order of item groups of different difficulty levels had significant influences on answering performance." (381 fifth graders, counterbalanced experiment, science achievement test). Scope: Shows that item-group difficulty and the order of difficulty groups affected test performance in this sample. It did not test mixed difficulty in a game or motivation.
- **Verdict: mixed.** Order effects are small but real. Evidence on rising versus mixed difficulty for motivation is not settled.
- **Code vs evidence:** Consistent. It is a reasonable default with a tailored override.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M57. Games practise what the last lesson taught

- **Where:** `services/game-payloads.js:543-558` (`taughtFor`: "the questions practise these points, spread across all of them").
- **What it does:** A level's game is retrieval practice on the lesson it follows.
- **Evidence:**
  - roediger2006 [1] (as M06).
  - agarwal2021 [99]: "the majority of which (57%) revealed medium or large benefits from retrieval practice". Review of 50 classroom experiments.
  - rowland2014 [10] (as M20).
  - biggs1996 [492]: title-level only (constructive alignment).
- **Verdict: strong.**
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04-M58. Tailored games with a holdout arm; games as practice

- **Where:** `services/game-payloads.js:651-676`. The tailored set is planned from the profile, and `armFor` returns `holdout` for a control group, which gets the shared set. `services/game-lab.js` serves the preview lab.

- **What it does:** It personalises games and measures whether personalising helps.
- **Evidence:**
  - wouters2013 [493]: "serious games were found to be more effective in terms of learning (d= 0.29, p <.01) and retention (d = 0.36, p <.01)". Meta-analysis, k = 77.
  - clark2016 [494]: "the results highlight the affordances of games for learning as well as the key role of design beyond medium". Meta-analysis, K-16.
  - sailer2020 [43]: "the effect of gamification on cognitive learning outcomes was stable in a subsplit analysis of studies employing high methodological rigor". Meta-analysis of gamification (game elements in instruction), distinct from dedicated learning games.
  - kohavi2009 [147] (as M09).
- **Verdict: moderate.** Games help modestly. Whether tailoring helps is exactly what the holdout tests.
- **Code vs evidence:** Matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

04 proposals (each backed by 3 or more sources; status on each)

The proposals below don't change any stage rule. They are listed in priority order. **Status 2026-10-25 (implementer E):** the "Status" column and each proposal's first bullet. Flags, tests and rollback: `curriculum-build/LIVE-CHANGES-2026-10-24.md`, section "Skill tracks: P1, P2, P3, P5, P7, P9, P12 (implementer E)".

#	Proposal	Priority	Status
04-P1	Serve puzzles at the documented 0.75 expected score in practice	high	BUILT (flag `CHESS_PRACTICE_TARGET`)
04-P2	Reduced, bandwidth and on-request feedback on piano takes	medium	BUILT (flag `PIANO_FEEDBACK_FADE`)
04-P3	Build the tempo ladder and the hotspot loop	medium	BUILT (flags `PIANO_TEMPO_LADDER`, `PIANO_HOTSPOT_LOOP`)
04-P4	Model and anchor clips for film-it and self-rating	high	SKIPPED: needs media
04-P5	Compute split-half reliability of scene sets and gate on it automatically	high	BUILT (flag `SCENE_RELIABILITY_GATE`); BB1-BB7 stay in shadow until real data meets 0.8
04-P6	Validate camera make/miss and pose per device before trusting counts	medium	SKIPPED: needs data
04-P7	Enforce youth load caps and rest days in the dose	high	BUILT (flag `SKILL_LOAD_CAPS`)
04-P8	Run carry-over maintenance as an experiment instead of leaving it off	medium	SKIPPED: needs data
04-P9	Describe the two-day rule as a delayed test, not a sleep gain, for children	low (text)	BUILT (flag `SKILL_TWO_DAY_TEXT`)

#	Proposal	Priority	Status
04-P10	First-person decision scenes and response-time bands	low (cost)	SKIPPED: needs media and data
04-P11	A structured neuromuscular warm-up template	medium	SKIPPED: exercise content not sourced
04-P12	Add "nothing to find" positions to the CH2/CH3 tactics sets	medium	BUILT (flag `CHESS_NOTHING_TO_FIND`)

04-P1. Serve practice puzzles at the 0.75 expected score the curriculum names

- **Status: BUILT (flag `CHESS_PRACTICE_TARGET`, default on; `=0.6` or any 0.5-0.95 sets another score).** `content/chess-trainer.js` `practiceTarget`: rated puzzles aim 400·log10(E/(1-E)) points below the learner (191 for 0.75), by the Elo expectation. The rating update is unchanged. While a stage check's rated-puzzle sitting is open (CH4/CH5 `tactics_rating`), and for the owner, the target stays at the rating. Tests: `scripts/chess-trainer-test.js` (8 checks). The 0.6 calculation score has no puzzle set to apply to yet (CH4 calculation is `not_wired`).
- **Gap:** M21. `SKILL-CURRICULA.md:246` says 0.75, but `content/chess-trainer.js:89` targets the learner's own rating, which gives about 0.5.
- **Sources:**
  - eggen2006 [417]: "results in each individual answering about 50% of the items correctly". 50% is the *measurement* optimum, and the paper shows how to shift the target easier.
  - guadagnoli2004 [41] (as M11).
  - wulf2016 [478] (as M50). Expectancy of success.
  - klinkenberg2011 [78] reports arithmetic practice items sampled at a mean success probability of 0.75, so tasks were challenging but not too difficult; this is not direct evidence about chess or persistence.
  - pelanek2016 [77]: title-level only.
- **Expected effect:** More successful practice and better persistence for weaker learners. The rated mode keeps its information.
- **Risk:** The rating converges more slowly, because easier items carry less information. Mitigation: shift only in practice mode, and keep stage-check sittings at the nearest rating.
- **Files:** `content/chess-trainer.js` (`choose`, line 89: target = rating − offset from `chess-rating.solveChance`), `content/chess-rating.js` (inverse of `solveChance`).

04-P2. Reduced, bandwidth and on-request feedback on piano takes

- **Status: BUILT (flag `PIANO_FEEDBACK_FADE`, default on), on warm practice takes.** After 2 clean warm takes in a row on a metric (`practice_streak`), a take's result flags only what is outside the bar (`piano-core.outsideBar`); "Show all the numbers" shows the rest on request (`public/js/skill-kit.js`). A miss resets the streak, so full feedback returns. Cold stage takes always get the full result. Not built: the fade inside lesson blocks (`piano-blocks.js` has no server flag to read, and its schema is in `services/block-`

levels.js). Tests: scripts/piano-items-test.js, scripts/skills-r3-piano-bb-test.js.

- **Gap:** M36 and M37. SKILL-CURRICULA.md §4 documents fading (real-time, then end-of-take, then bandwidth, then self-rate first). The piano engine always shows full real-time highlights in learn mode and a full report after play.
- **Sources:**
  - sherwood1988 [495]: "giving KR about a relatively large bandwidth enhances movement consistency". Pop.: adults, timed elbow flexion.
  - wulf2010 [458]: "including feedback and model demonstrations controlled by the learner, has been found to be more effective than externally controlled practice conditions".
  - wulf2002 [436] (as M29). Complex skills (high processing load) benefit from conditions that reduce the load to manageable levels.
  - salmoni1984 [496] discusses the guidance hypothesis (knowledge of results as guidance); it does not by itself establish benefits of reduced feedback frequency.
  - winstein1990 [497]: title-level only (reduced feedback frequency).
- **Expected effect:** Better retention on day-2 takes. Less dependence on the falling-note display.
- **Risk:** Beginners get frustrated. Mitigation: fade only after 2 clean takes, and keep it learner-requestable.
- **Files:** public/engine/piano/piano-song.js (highlight modes), public/engine/piano/piano-blocks.js (report: a bandwidth filter on offBeatNotes and wrong), services/block-levels.js (Piano_PlayAlong schema: a feedback field).

04-P3. Build the tempo ladder and the hotspot loop

- **Status: BUILT (flags PIANO_TEMPO_LADDER and PIANO_HOTSPOT_LOOP, default on), on warm practice sets.** piano-core.js has the pure rules: tempoLadder (start at 65%, +10% after 2 clean takes, -10% after 2 misses, between 60% and 100%), barErrors, hotspot (the worst bar plus its worse neighbour), sliceBars. services/skill-dose.js applies them. A warm set below the target tempo, or a loop set, never counts as "at target", so warm sets only call a cold test at full speed. A library piece that misses gets its worst 1-2 bars looped at 65% in the next set, and the set after rejoins the whole piece. Cold stage takes stay at the target tempo. Tests: scripts/piano-items-test.js (13 checks), scripts/skills-r3-piano-bb-test.js (13 checks).
- **Gap:** M36. Documented in SKILL-CURRICULA.md §4 ("Start at 60-70% of target. Step up 10% after 2 clean takes"; "The worst 1-2 bars loop at 65%") but not in code. The speed control is a fixed set: ½, ¾ or full.
- **Sources:**
  - duke2009 [400] (as M11). Error handling and the share of correct trials predict retention.
  - ash1990 [418] (as M23). Segmentation beats whole-task practice.
  - chaffin2010 [445] (as M31): "The cellist repeatedly took the piece apart section-by-section and then re-integrated the sections".

    - fitts1954 [450]: title-level only.
- **Expected effect:** More correct trials per minute on the hard bars, and a better next-day take.
- **Risk:** Learners stay at low tempo. Mitigation: the ladder only steps up, and the cold take stays at target.
- **Files:** `public/engine/piano/piano-song.js` (speed state, a loop of the bars with the most `offBeatNotes`/`wrong`), `public/engine/piano/piano-core.js` (per-bar error summary), `services/skill-dose.js` (warm-set tempo).

04-P4. Model and anchor clips for film-it and for piano self-rating

- **Status: SKIPPED: needs media.** No model or anchor clips exist, and making them needs licensing and consent, or drawn animations. Revisit when clips are produced.
- **Gap:** M37 and M49. Self-rating has no reference. BB3's "rated against anchor clips" is not built.
- **Sources:**
    - hewitt2001 [393] (as M05). Junior high instrumentalists who used a model during self-evaluation improved more in several performance areas than those self-evaluating without a model.
    - modinger2022 [475] (as M49). Covers expert modelling and self-modelling.
    - stemarie2012 [457] (as M42).
    - bergee1997 [392] (as M05). Undergraduate musicians' self-ratings correlated poorly with faculty and peer ratings.
- **Expected effect:** Better calibrated self-ratings and larger technique gains from the film loop.
- **Risk:**
    - Licensing and consent for basketball clips. Drawn or animated anchors would avoid this.
    - Storage, since clips stay on the device.
- **Files:** `public/engine/basketball/basketball-blocks.js` (`renderFilm`: an anchor pane before rating), `services/block-levels.js` (the `Basketball_Film` and `Piano_PlayAlong` schemas: `anchor` and `model` fields), `public/engine/piano/piano-blocks.js` (`selfRate`: play the model first).

04-P5. Compute split-half reliability of scene sets and gate on it automatically

- **Status: BUILT (flag `SCENE_RELIABILITY_GATE`, default on).** `services/skill-checks.js` now keeps each scene item's mark on the sitting (`detail.marks`). `services/skill-scene-reliability.js` computes the odd/even split-half from each learner's first scene sitting, steps it up with Spearman-Brown, and asks for at least 0.8 on at least 100 learners. In `services/skill-gate.js`, BB1-BB3 leave the hand-kept shadow list. Any stage with scene sets stays in shadow until every set meets the bar, so **BB1-BB7 stay in shadow today** (BB4-5 and BB6-7 also stay listed for their [R] zone marks). Report: `node`

`services/skill-scene-reliability.js`. Tests: `scripts/skills-r3-piano-bb-test.js` (8 checks).

- **Gap:** M10. The 0.8 bar exists only in comments, and BB1-BB7 are held in shadow by a hand-kept list.
- **Sources:**
  - cortina1993 [399] (as M10).
  - hopkins2000 [383] (as M03).
  - aera2014 [150] (AERA Standards).
- **Expected effect:** The basketball stages leave shadow when, and only when, their sets prove reliable.
- **Risk:** Before 100 learners the estimates are noisy. Require the 100 learners the progression names.
- **Files:** new `scripts/scene-reliability.js` (split-half with the Spearman-Brown correction over `skill_records` sittings), `services/skill-gate.js` (`forcedShadow` reads the computed result), `content/basketball-scenes.js` (scene ids stable across sittings).

04-P6. Validate camera make/miss and pose against hand logs, per device

- **Status: SKIPPED: needs data.** It needs paired hand-tapped and camera counts from real learners on each device.
- **Gap:** M48. Make/miss is unvalidated. Every camera count needs the learner to confirm it.
- **Sources:**
  - washabaugh2022 [388] (as M04).
  - stenum2021 [389] (as M04).
  - hopkins2000 [383] (as M03): agreement and typical error.
- **Expected effect:** A measured error rate for each device and set-up. With good agreement, confirmation becomes a light check rather than a recount.
- **Risk:** Privacy. Only counts are compared, never frames, which keeps the current rule.
- **Files:** `public/engine/basketball/camera-blocks.js` (paired hand-tap and camera log), `public/engine/basketball/tracking.js` (agreement statistics), `services/skill-checks.js` (store agreement with the count).

04-P7. Enforce youth load caps and rest days in the dose

- **Status: BUILT (flag `SKILL_LOAD_CAPS`, default on), basketball only.** `services/skill-dose.js` applies the NBA and USA Basketball caps by age, taken from the age screen's month and year of birth: 3/5/10/14 hours a week and 2/2/1/1 rest days for ages 7-8 (and younger), 9-11, 12-14 and high school. Adults, and accounts with no age answer, get no cap. The week is the 7 UTC days ending today. Court minutes are the logged sets plus the cold sittings; screen sets (scenes, rules) do not count. Over the week's limit, court work is cut: spare sets first, then practice down to one set each, then sittings. On a rest day there is no court work. Court sets (`/log`) and court counts (`skill-checks.count`) are refused on a rest day or once the week is used. The plan page and the path say why. **Hook for the**

lesson engine: court sessions in `services/sessions.js` can add their minutes with `skillDose.logLoad(userId, 'basketball', minutes)`. Tests: `scripts/skills-r3-dose-test.js` (11 checks).

- **Gap:** M16 and M51. `SKILL-CURRICULA.md` §5 lists caps by age: 3 / 5 / 10 / 14 hours a week and 2 / 2 / 1 / 1 rest days. No service enforces them. `skill-dose.js` sizes a session but does not cap the week.
- **Sources:**
  - difiori2014 [483] (as M51).
  - difiori2018 [409] (as M16).
  - bergeron2015 [408] (general youth-athlete development guidance; not evidence for the specific caps).
- **Expected effect:** Lower risk of overuse injury and burnout.
- **Risk:** Some motivated learners are held back. That is intended. A parent can see the cap.
- **Files:** `services/skill-dose.js` (weekly minutes from `skill_records` and court logs), `services/skill-adaptation.js` (`planSession` takes `weekLeft`), `content/skill-progressions/basketball.json` (caps by age band).

04-P8. Run carry-over maintenance as an experiment

- **Status: SKIPPED: needs data.** An arm only means something with real learners' retention sittings. The arm assignment would also go in `services/personalize.js`, which is outside this section's files.
- **Gap:** M14. `CARRY_DAYS 0` rests on a simulation without spacing. The motor-spacing literature points the other way.
- **Sources:**
  - lee1988 [374] (as M01).
  - donovan1999 [259] (as M14).
  - arthur1998 [395] (as M06).
  - kim2013 [396] (as M08).
  - kohavi2009 [147] (as M09).
- **Expected effect:** If spacing helps here, better retention of earlier stages' skills when retested.
- **Risk:** Spends session minutes, as the round-3 README warns. A small arm of `CARRY_DAYS 7`, `CARRY_MAX 1` keeps the cost small.
- **Files:** `services/skill-adaptation.js` (knobs per arm), `services/personalize.js` (`armFor` experiment), `services/skill-dose.js`.

04-P9. Describe the two-day rule as a delayed test, not a sleep gain, for children

- **Status: BUILT (flag `SKILL_TWO_DAY_TEXT`, default on).** A scan found no learner-facing skill text that promises a sleep gain. `content/skill-learner-text.js` `TWO_DAY_RULE` ("A check on a later day ... is a later test, not a promise that you get better overnight") is served by `skill-gate.left` and `skill-checks.overview`, and shown on the stage-check and stage-left pages. `SKILL-CURRICULA.md` §1.2 now gives the delayed-test reason. A

guard test fails if any skill learner text promises a sleep gain: `scripts/skills-r3-gate-test.js` (5 checks).

- **Gap:** M02 and M41. The rule is right. Any wording that promises sleep improvement overstates the evidence for children.
- **Sources:**
  - wilhelm2008 [253] (as M02).
  - rickard2008 [379] (as M02; indirect motor-sequence evidence in its studied population).
  - soderstrom2015 [39] (as M01).
- **Expected effect:** Honest learner and parent text. No change to behaviour.
- **Risk:** None.
- **Files:** `content/skill-learner-text.js`, `curriculum-build/SKILL-CURRICULA.md` §1.2.

04-P10. First-person decision scenes and response-time bands

- **Status: SKIPPED: needs media and data.** First-person scenes need a renderer or clips. Response-time bands need each learner's own band, which needs real sittings. Response times are already recorded (BB3).
- **Gap:** M46. The scenes are top-down, and response time is recorded but not used.
- **Sources:**
  - rosch2021 [398] (as M10). A first-person video test.
  - zhu2024 [467] (as M46). Transfer is limited, and representative stimuli help.
  - broadbent2015 [468] (as M46).
  - hadlow2018 [469] (as M46).
- **Expected effect:** Better transfer of reads to games. RT adds diagnostic information.
- **Risk:** Production cost. Clips need consent or licensing. Rendered first-person scenes avoid both.
- **Files:** `content/basketball-scenes.js` (a camera-view renderer), `public/engine/basketball/basketball-blocks.js` (`Basketball_Play`), `content/skill-piano-bb-sittings.js` (an RT rule once the learner's band is known).

04-P11. A structured neuromuscular warm-up template

- **Status: SKIPPED: the exercise content is not sourced.** The three sources show that structured programmes such as the 11+ prevent injuries. They do not give the exercises themselves (the FIFA 11+ manual's parts 1-3), or a version for under-10s (11+ Kids). Neither is among the verified sources. A template a child follows needs that content checked first. The `Basketball_Session` schema is also in `services/block-levels.js`, which is outside this section's files.
- **Gap:** M51. The warm-up is free text. The trials used structured programmes.
- **Sources:**
  - soligard2008 [480] (as M51).
  - longo2012 [481] (as M51). In basketball.

    - emery2015 [482] (as M51).
- **Expected effect:** Fewer injuries, especially lower-limb and overuse injuries.
- **Risk:** Session time, about 10-15 minutes. It needs an age-appropriate version for under-10s.
- **Files:** `services/block-levels.js` (`Basketball_Session`: a `warmup_template` id), `public/engine/basketball/basketball-blocks.js`, and the basketball lesson content.

04-P12. Add "nothing to find" positions to the CH2/CH3 tactics sets

- **Status: BUILT (flag `CHESS_NOTHING_TO_FIND`, default on).** Stockfish 19 is installed next to the repo (`local-tools/chess-engines`, used by `services/chess-engine.js`). `scripts/chess-quiet-pool-build.js` wrote `content/chess-quiet-pool.js`. Each position uses a real tactics-pool position's pieces on random squares, so piece counts give nothing away. White is to move with a capture to consider. A position is kept only when Stockfish (1M nodes, MultiPV 3) finds no mate and no move ending more than 80 cp better than the material. CH2 has 5 of 30 quiet, CH3 has 6 of 30. Every position offers "Nothing wins", which is right only on a quiet one. The owner's sittings have none. Tests: `scripts/skills-r3-checks-test.js` (8 checks). Response times are not used yet.
- **Gap:** M19. The code's own `gap` notes at `content/skill-stage-checks.js:50, 55`.
- **Sources:**
    - rohrer2015 [12] (as M19).
    - bilalic2008 [412] (as M19).
    - vandermaas2005 [413] (as M19). Its quote on response latency also supports recording solve times.
- **Expected effect:** Tactics sittings measure *spotting* as well as *executing*, closer to game play.
- **Risk:** Quiet positions need Stockfish checks that the best move wins nothing, which is a build job.
- **Files:** `content/chess-tactics-pool.js` (quiet positions), `services/skill-checks.js` (`buildChessTactics`: mix in a "no tactic" option as in `buildChessMateOrNot`), `content/skill-stage-checks.js`.

Section 05. Engagement, UX and wellbeing

This section covers: streaks, XP and rewards; the study-session strip; break and length guidance; feedback wording and rationales; hints and the coach; question formats and `preferTypeIn`; accessibility; mobile layout; dashboard, progress and calendar; reminders; privacy and consent for minors; owner skips; error and empty states.

Honest limits, up front:

- **No source found tests a streak counter against learning outcomes.** Streak evidence is indirect (habit formation, loss aversion, goal-gradient). It is marked mixed or design choice, never strong.
- **XP and gamification evidence is mixed.** Meta-analyses find small to moderate positive effects (Sailer 2019 g = .49 cognitive; Zeng 2024 g = 0.782) but those effects are unstable on motivation and behaviour and depend on context. The overjustification literature conflicts (Deci 1999 vs Cameron 1994). See M3.
- Evidence on UI-layout choices (touch targets, single call to action, bottom tab bar) is mostly HCI lab work with adults, or standards. Children are under-studied.
- The ICO code and COPPA are law and regulator guidance, not learning evidence. They set what is allowed; they do not show what teaches.

Secondary summaries read for context only (not quoted, not counted): Hanus and Fox 2015 is reported by secondary reviews to have found less intrinsic motivation, satisfaction and lower exam scores in a gamified course of 80 students; Bai 2020 is reported at g = 0.504 with no moderator difference by element type. Treat both as unconfirmed until the abstracts are read.

In numbers: 50 mechanisms; 50 with 3 or more quoted sources; 6 strong, 8 moderate, 8 mixed, 11 design choice, 15 standard, 2 gap; 141 distinct sources cited or listed; 11 mechanisms built or changed on 25 October 2026.

05 at a glance

Mechanism	Verdict	Quoted sources	Status
05-M1 Day streak counter	mixed	5	built on 25 October 2026; open proposal
05-M2 Zero streak and zero XP are not shown	design choice	3	existing
05-M3 XP formula	mixed	8	existing
05-M4 Always-on header stats	mixed	4	built on 25 October 2026
05-M5 Coach discount on XP (60%)	design choice	3	existing; open proposal
05-M6 Segmented progress bar and counter	moderate	4	existing
05-M7 Dashboard single "next step" block	design choice	3	existing
05-M8 Recent sessions table, "Due for review" list, scale per subject	mixed	4	existing

Mechanism	Verdict	Quoted sources	Status
05-M9 The learning-style quiz sets session length, formats and explanation depth	mixed	5	built on 25 October 2026
05-M10 Start, pause and end strip	design choice	5	existing
05-M11 Active-time accounting	design choice	3	existing
05-M12 Session length: 6 to 18 items, stage minutes	mixed	4	existing
05-M13 Break guidance: none in the code	gap	4	built on 25 October 2026
05-M14 Retention-enforcer prompt in the strip	strong	3	existing
05-M15 Verdict wording	strong	5	existing
05-M16 Praise and tone	strong	4	existing
05-M17 A rationale on every item	strong	4	existing
05-M18 Held versus immediate feedback	mixed	3	existing
05-M19 Optional "How sure are you?" row	moderate	4	existing
05-M20 Multiple choice versus type-in	strong	7	existing
05-M21 `preferTypeIn`	mixed	6	built on 25 October 2026; open proposal
05-M22 Neutral "Checking your answer..." state	moderate	4	existing
05-M23 The coach: asks, never tells, before the answer	strong	5	existing; open proposal
05-M24 Help is withheld on scored sittings	design choice	4	existing
05-M25 "Ask about this" after the answer	moderate	3	existing
05-M26 "Report a problem."	design choice	3	built on 25 October 2026
05-M27 "Too fast to count" note and re-served question	moderate	3	built on 25 October 2026
05-M28 Optional "Give Input" form that never gates progress	design choice	4	existing
05-M29 Touch targets	standard	3	built on 25 October 2026
05-M30 Colour and contrast tokens	standard	3	existing
05-M31 Right and wrong are never colour-only	standard	3	existing
05-M32 Reduced motion	standard	3	existing
05-M33 Announcements, toasts and focus	standard	3	built on 25 October 2026
05-M34 Skip link and focus ring	standard	3	existing
05-M35 Mobile layout: bottom tab bar, fixed Next bar, safe areas	moderate	3	existing
05-M36 Kids audience variant	design choice	3	existing

Mechanism	Verdict	Quoted sources	Status
05-M37 Retention calendar	moderate	3	built on 25 October 2026
05-M38 Reminders: none exist	gap	4	built on 25 October 2026; open proposal
05-M39 Error states	standard	3	existing
05-M40 Empty states	design choice	3	existing
05-M41 Owner-only skips	design choice	3	existing
05-M42 Neutral age screen	standard	4	existing
05-M43 Country age-of-consent table	standard	4	existing
05-M44 Email-plus parental consent with delayed confirmation	standard	4	existing
05-M45 Camera and microphone off for unconsented child accounts	standard	4	existing
05-M46 On-device only camera and mic	standard	3	existing
05-M47 Retention timers	standard	3	existing
05-M48 Parent review, download, revoke and delete	standard	3	existing
05-M49 Pinned notice version	standard	3	existing
05-M50 Trust cues: AI disclosure and coach limits	moderate	3	existing

05-A. Streaks, XP and rewards

05-M1. Day streak counter

- **Where:** `services/sessions.js:1604-1613` (finishing a session), `services/review-session.js:256-263` (finishing the daily review), shown at `public/js/app.js:329-343` (header pill) and `public/js/dashboard.js:345-349` (up to 7 tick marks).
- **What it does:** a day counts when a session or review ends on a new date; yesterday's date extends the run, any gap resets to 1. No freeze, no grace day.
- **Evidence:** lally2009 [45] (one missed day "did not materially affect" habit formation, and automaticity took 18 to 254 days); kahneman1979 [498] (losses weigh more than gains, which is the pull a streak leans on); kivetz2006 [499] (effort rises near a reward goal); shortt2021 [500] (Duolingo literature studies the app design, not learning outcomes); ico-aadc-5 [50] (reward loops are allowed only if the child can disengage "without feeling pressurised").
    - lally2009 [45]: "Missing one opportunity to perform the behaviour did not materially affect the habit formation process."; "ranged from 18 to 254 days" (96 adult volunteers, 12 weeks, daily behaviour)
    - kahneman1979 [498]: "is generally steeper for losses than for gains" (Lab choice problems)

    - kivetz2006 [499]: "participants in a real café reward program purchase coffee more frequently the closer they are to earning a free coffee" (Café customers, online song raters (adults))
    - shortt2021 [500]: "The focus on app design marks an emphasis on the creation of tools rather than the process and outcomes of language learning" (Systematic review, 35 Duolingo studies)
    - ico-aadc-5 [50]: "make it easy for them to disengage without feeling pressurised or disadvantaged if they do so"; "Do not use children's personal data in ways that have been shown to be detrimental to their wellbeing" (UK statutory code, services likely to be used by children)
- **Verdict: mixed**. Habit evidence favours regular practice and shows one miss is harmless; loss aversion explains why streaks pull; no study found links a streak counter to learning, and the ICO treats reward loops as a wellbeing risk for children.
- **Code vs evidence:** partial match. Gaps: (1) there is no forgiveness, although lally2009 [45] says one missed opportunity did not materially affect habit formation; (2) the dates are UTC (`toISOString().slice(0,10)` at `sessions.js:1605` and `review-session.js:257`) while study sessions use the learner's own time zone (`services/study-sessions.js:today/zoneFor`), so an evening learner far from UTC can lose or keep a streak wrongly; (3) enforcer-only or lesson-only days do not touch the streak. See proposal P2.
- **Status: Built on 25 October 2026** (switch `STREAK_LOCAL_GRACE`, from 05-P2). The streak uses the learner's local day and forgives one missed day in 7. The owner keeps the old rule. **Open:** the rest of 05-P2, counting a study session or review toward the streak (not built).

05-M2. Zero streak and zero XP are not shown

- **Where:** `public/js/app.js:329-343`.
- **What it does:** the header pills render only when the value is above zero; the dashboard streak card likewise (`dashboard.js:345`).
- **Evidence for the principle:** kluger1996 [342] (feedback that draws attention to the self rather than the task can lower performance); hattie2007 [123] (feedback effects can be positive or negative); henderlong2002 [501] (praise helps when it promotes controllable attributions and attainable standards).
    - kluger1996 [342]: "FIs improved performance on average (d =.41)" (Meta-analysis, 23,663 observations)
    - hattie2007 [123]: "Feedback is one of the most powerful influences on learning and achievement, but this impact can be either positive or negative." (Review)
    - henderlong2002 [501]: "Provided that praise is perceived as sincere, it is particularly beneficial to motivation when it encourages performance attributions to controllable causes" (Review of children's praise research)
- **Verdict: design choice** (no study of showing versus hiding a zero). Code matches the principle: nothing shames a new learner.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M3. XP formula

- **Where:** `services/mastery.js:675-683` (`xpForItems`): a wrong answer earns 2; a right answer earns 10 plus the item difficulty above 5; 60% of that if the coach was used. Awarded at `sessions.js:1603`, `review-session.js:245`, `enforcers.js:922`; shown in outcome cards (`session.js:1214`, `review.js:254`, `skip.js:202`) and the recent-sessions table (`dashboard.js:321-329`).
- **What it does:** XP is a running total with no spending, ranks or leaderboards. It is informational, not a prize.
- **Evidence:** sailer2020 [43] (small positive effects overall, "less stable" for motivation and behaviour); zeng2024 [502] (g = 0.782 on academic performance, 22 studies); deci1999 [44] (expected tangible, engagement-, completion- and performance-contingent rewards undermined intrinsic motivation; positive feedback enhanced it); cameron1994 [503] ("overall, reward does not decrease intrinsic motivation"; the only negative effect is expected tangible rewards given for just doing a task); lepper1973 [504] (children given an expected reward later showed less interest); eisenberger1996 [505] (harm occurs only "under highly restricted, easily avoidable conditions"); sailer2017 [506] (effects depend on the specific element); hamari2014 [507] (depends on context and user).
    - sailer2020 [43]: "significant small effects of gamification on cognitive (g = .49, 95% CI [0.30, 0.69], k = 19, N = 1686)"; "effects on motivational and behavioral outcomes were less stable" (Meta-analysis of gamified vs non-gamified learning)
    - zeng2024 [502]: "the results revealed a moderately positive effect of gamification on student academic performance (Hedges's g = 0.782, p < 0.05)" (Meta-analysis, 22 experimental studies, 2008-2023)
    - deci1999 [44]: "engagement-contingent, completion-contingent, and performance-contingent rewards significantly undermined free-choice intrinsic motivation"; "Positive feedback enhanced both free-choice behavior (d = 0.33) and self-reported interest (d = 0.31)." (Meta-analysis, 128 experiments, children and college students)
    - cameron1994 [503]: "Results indicate that, overall, reward does not decrease intrinsic motivation."; "The only negative effect appears when expected tangible rewards are given to individuals simply for doing a task." (Meta-analysis, 96 experiments)
    - lepper1973 [504]: "subjects in the expected-award condition would show less subsequent intrinsic interest in the target activity" (Field experiment with children)
    - eisenberger1996 [505]: "detrimental effects of reward occur under highly restricted, easily avoidable conditions" (Review of 25 years of reward research)
    - sailer2017 [506]: "gamification is not effective per se, but that specific game design elements have specific psychological effects" (Randomised online-simulation study)
    - hamari2014 [507]: "the effects are greatly dependent on the context in which the gamification is being implemented, as well as on the users using it" (Literature review of empirical gamification studies)

- **Verdict: mixed**, and deliberately so. The overjustification meta-analyses disagree (deci1999 [44] vs cameron1994 [503]). Gamification meta-analyses are positive but heterogeneous, and the cognitive-outcome effect is the stable one. XP here is non-tangible and not expected for mere participation in the way a prize would be, which sits on the safer side of both reviews, but no study tests this exact XP design.
- **Code vs evidence:** partial. Two details lean toward the risk side: "+2 for a wrong answer" may raise a related concern about contingent rewards (deci1999 [44] flags engagement-contingent rewards); and XP is shown on every page header (M4).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M4. Always-on header stats

- **Where:** `public/js/app.js:341-344`.
- **What it does:** streak and XP appear in the header of every signed-in page.
- **Evidence:** sailer2017 [506] (elements have specific effects; badges, leaderboards and performance graphs raised competence need satisfaction); przybylski2010 [508] (engagement is sustained by competence, autonomy and relatedness, not by points per se); ryan2000 [47] (controlling contexts undermine motivation); ico-aadc-5 [50] (reward loops under scrutiny for children).
    - sailer2017 [506]: "gamification is not effective per se, but that specific game design elements have specific psychological effects" (Randomised online-simulation study)
    - przybylski2010 [508]: "both the appeal and well-being effects of video games are based in their potential to satisfy basic psychological needs for competence, autonomy, and relatedness" (Theory and review (self-determination theory))
    - ryan2000 [47]: "three innate psychological needs--competence, autonomy, and relatedness--which when satisfied yield enhanced self-motivation and mental health" (Theory review)
    - ico-aadc-5 [50]: "make it easy for them to disengage without feeling pressurised or disadvantaged if they do so"; "Do not use children's personal data in ways that have been shown to be detrimental to their wellbeing" (UK statutory code, services likely to be used by children)
- **Verdict: mixed / design choice**. Competence feedback is supported; permanence on every screen is not tested. No way for the learner to hide it (autonomy: patall2008 [46]). See proposal P5.
- **Status: Built on 25 October 2026** (switch `UX_HIDE_STATS`, from 05-P5). Learners (or parents) can hide the streak and XP; they are still counted.

05-M5. Coach discount on XP (60%)

- **Where:** `services/mastery.js:679`.
- **What it does:** a correct answer after using the coach earns 0.6 of the base XP, "enough to make working it out alone feel worth more, not enough to discourage asking for help" (code comment).

- **Evidence:** bastani2025 [38] (unguarded help lowered later unaided grades by 17%; guardrails mitigated this); aleven2003 [173] (learners do not use help facilities effectively); roll2011 [607].
    - bastani2025 [38]: "students actually perform worse than those who never had access (17% reduction in grades for GPT Base)"; "These negative learning effects are largely mitigated by the safeguards in GPT Tutor." (Nearly 1,000 high-school maths students)
    - aleven2003 [173]: "although effective help-seeking behavior in ILEs is related to better learning outcomes, learners are not using help facilities effectively" (Selective review)
    - roll2011 [607]: "the improved help-seeking skills transferred to learning new domain-level content" (two studies, 58 and 67 students, with a geometry cognitive tutor). Scope: Shows that feedback on help-seeking errors improved help-seeking and the gain carried over. It did not test an XP discount or whether help replaces trying.
- **Verdict: design choice** (the 0.6 has no study). The principle, that help should not be a free substitute for trying, is supported. Gap: nothing shows the learner why the number is lower, which matters for informational (not controlling) feedback.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 05-P7, a visible "why XP is lower with help" line, with the hint ladder (not built).

05-M6. Segmented progress bar and counter

- **Where:** `public/js/session.js:202-225`, markup at `session.js:146-162`.
- **What it does:** one segment per question, outlined where you are, filled when answered, in the right/wrong colour only when feedback is not being held (`session.js:214-217`); "n / N" counter.
- **Evidence:** kivetz2006 [499] (effort rises as the goal nears); nunes2006 [510] (a display that starts with progress already shown raises persistence); cheema2011 [511] (goal visualization helps only near the goal); myers1985 [512] (people prefer progress indicators).
    - kivetz2006 [499]: "participants in a real café reward program purchase coffee more frequently the closer they are to earning a free coffee" (Café customers, online song raters (adults))
    - nunes2006 [510]: "people provided with artificial advancement toward a goal exhibit greater persistence toward reaching the goal" (Consumer field experiments (adults))
    - cheema2011 [511]: "the beneficial effects of visualization exist only when people are close to the goal" (Consumers, swimmers, lab tasks)
    - myers1985 [512]: "people prefer to have progress indicators" (Formal experiment (adults))
- **Verdict: moderate** (consumer and HCI evidence; none from learner questions). Code matches, and the hold on verdict colours during exams and held-feedback sittings protects the assessment (design choice).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M7. Dashboard single "next step" block

- **Where:** `public/js/dashboard.js:176-212` (`leadBlock`) and `nextStep` (`dashboard.js:62-90`).
- **What it does:** one large primary action for the most urgent subject; other subjects in rows beneath.
- **Evidence for the principle (proximal, visible goals and choice management):** locke2002 [513] (35 years of goal-setting research, abstract-level only); kivetz2006 [499]; patall2008 [46] (choice helps, with limits).
    - locke2002 [513]: "The authors summarize 35 years of empirical research on goal-setting theory." (Review (abstract-level support only))
    - kivetz2006 [499]: "participants in a real café reward program purchase coffee more frequently the closer they are to earning a free coffee" (Café customers, online song raters (adults))
    - patall2008 [46]: "providing choice enhanced intrinsic motivation, effort, task performance, and perceived competence" (Meta-analysis of choice studies)
- **Verdict: design choice**. No study of a single call to action on a learning dashboard was found.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M8. Recent sessions table, "Due for review" list, scale per subject

- **Where:** `dashboard.js:318-333`, `dashboard.js:336-344`, `rowScale` (app.js scale code at `app.js:573-611`).
- **What it does:** shows score, XP and time for recent sittings; the weakest due items with a mastery percentage.
- **Evidence:** bodily2017 [514] (93 articles on student-facing dashboards); matcha2020 [515] (dashboards "rarely grounded in learning theory"); sedrakyan2020 [516] (title-level only); kluger1996 [342] and hattie2007 [123] (information about performance can help or harm).
    - bodily2017 [514]: "Out of the 945 articles retrieved from databases and journals, 93 articles were included in the analysis." (Review of student-facing dashboards)
    - matcha2020 [515]: "existing LADs are rarely grounded in learning theory" (Systematic review of dashboard studies)
    - sedrakyan2020 [516]: proposes dashboard feedback grounded in learning regulation and learner goals.
    - kluger1996 [342]: "FIs improved performance on average (d =.41)" (Meta-analysis, 23,663 observations)
    - hattie2007 [123]: "Feedback is one of the most powerful influences on learning and achievement, but this impact can be either positive or negative." (Review)
- **Verdict: mixed**. Showing learners their own record is plausible but the dashboard literature itself is weak. Code matches the modest version (no ranks, no comparison with peers, which Hanus and Fox report as a risk, unconfirmed).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M9. The learning-style quiz sets session length, formats and explanation depth

- **Where:** `public/js/dashboard.js:354` (copy), `public/js/style-quiz.js:87`, `services/learner.js:276-310`.
- **What it does:** twenty choices become a profile; `sessionParams` turns it into item count 6-18, `preferTypeIn`, explanation depth and feedback timing.
- **Evidence against matching:** pashler2008 [31] (the meshing hypothesis is the claim tested); rogowsky2015 [317] (no relationship between preference and comprehension in adults); rogowsky2020 [517] (same in fifth graders); willingham2015 [518] ("scientific support for these theories is lacking"); cuevas2015 [519] (sounder studies refute the hypothesis).
    - pashler2008 [31]: "the meshing hypothesis, according to which instruction is best provided in a format that matches the preferences of the learner" (Review)
    - rogowsky2015 [317]: "Results failed to show a statistically significant relationship between learning style preference (auditory, visual word) and learning aptitude" (College-educated adults)
    - rogowsky2020 [517]: "Consistent with earlier findings with adults, results failed to find a significant relationship between auditory or visual learning style preference and comprehension." (Fifth graders)
    - willingham2015 [518]: "scientific support for these theories is lacking" (Review)
    - cuevas2015 [519]: "the more methodologically sound studies have tended to refute the hypothesis" (Review of post-2009 studies)
- **Verdict: mixed**. Modality matching is not supported. Using preference to set quantity (how many items) or comfort (explanation depth) is a harmless design choice, but the dashboard copy "sets ... the mix of question formats" invites the learning-styles reading. Gap: reword, and see M20 for what evidence says about formats.
- **Status: Built on 25 October 2026** (switch `UX_STYLE_COPY`, from 05-P6). The copy now says the quiz sets pace and comfort, not how you are taught.

05-B. Study sessions, length and breaks

05-M10. Start, pause and end strip

- **Where:** `public/js/study.js` (whole file; `draw` at 78-136, `act` at 138-180), `services/study-sessions.js`.
- **What it does:** a one-line strip under the header: "Start session"; then a clock in minutes, Pause or Resume, and End; End asks "End session" or "End day session" with "Keep going" and Escape to cancel; focus moves to the first button; every action is announced (`App.announce`).
- **Evidence:** biwer2023 [520] (87 students: systematic, pre-planned breaks gave mood and efficiency benefits over self-regulated ones); cepeda2006 [22] (distributed practice, 839 assessments); szpunar2013 [245] (structured testing within a session sustains attention); wcag-4.1.3 [521] and wcag-2.4.7 [522] (announcements and visible focus).

    - biwer2023 [520]: "Taking pre-determined, systematic breaks during a study session had mood benefits and appeared to have efficiency benefits" (87 Dutch university students, one day)
    - cepeda2006 [22]: "This review found 839 assessments of distributed practice in 317 experiments located in 184 articles." (Quantitative synthesis)
    - szpunar2013 [245]: "the simple act of interpolating online lectures with memory tests can help students sustain attention to lecture content" (Two experiments, online lectures)
    - wcag-4.1.3 [521]: "status messages can be programmatically determined through role or properties such that they can be presented to the user by assistive technologies without receiving focus." (W3C Level AA)
    - wcag-2.4.7 [522]: "Any keyboard operable user interface has a mode of operation where the keyboard focus indicator is visible." (W3C Level AA)
- **Verdict: design choice** for the strip itself (self-timing has no direct trial here); the underlying bounded, structured session is supported (**moderate**). Accessibility behaviour matches the standards.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M11. Active-time accounting

- **Where:** `public/js/study.js:16-17,45-48,193-206`; `services/study-sessions.js:7-21,33-37,70-76`.
- **What it does:** heartbeat every 30 s only while the page is visible and the learner has acted in the last 4 minutes; a gap between two signs of activity counts at most 5 minutes; a running session with 30 minutes of nothing ends at its last activity; a paused one ends after 12 hours or at the day change.
- **Evidence for treating idle time as non-study:** wise2005 [74] (fast, unengaged responding is detectable from timing); wisedemars2005 [523] (low motivation lowers performance and distorts scores).
    - wise2005 [74]: "unmotivated examinees will answer too quickly (i.e., before they have time to read and fully consider the item)" (Computer-based test examinees; about fast answers, so treating idle time as non-study is a related extension, not something it tested)
    - wisedemars2005 [523]: "low student motivation is associated with a substantial decrease in test performance" (Low-stakes assessment review)
    - baker2008 [509]: "present evidence on which student behaviors, motivations, and emotions are associated with the choice to game the system" (Students in interactive learning environments)
- **Verdict: design choice** for the numbers (4 min, 5 min, 30 min, 12 h); principle **moderate**. Gap: the thresholds are untested on children, who may read slowly.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M12. Session length: 6 to 18 items, stage minutes

- **Where:** `services/learner.js:283,303` (`length = 6 + density*6 + pacing*6`, clamped 5-18; scorecard may scale it), `services/adaptation.js:546-575` (plan bounds), `services/skill-dose.js:63-67` (stage `session_minutes` midpoint, default 30).
- **Evidence:** ackerman2009 [52] (fatigue rises with time-on-task, yet performance in longer tests did not fall); sievertsen2016 [199] (performance falls 0.9% of an SD per hour of the school day; a 20-30 minute break recovers 1.7%); borgonovi2016 [54] (endurance on a long low-stakes test differs across students); cepeda2006 [22] (spacing).
    - ackerman2009 [52]: "subjective fatigue increased with increasing time-on-task. However, mean performance increased in the longer test length conditions" (239 first-year university students)
    - sievertsen2016 [199]: "for every hour later in the day, test performance decreases by 0.9% of an SD"; "a 20- to 30-minute break improves average test performance by 1.7% of an SD" (All Danish public-school children, 2009-2013)
    - borgonovi2016 [54]: "academic endurance, defined as students' ability to maintain their baseline rates of successful test completion for the duration of a low-stakes cognitive test" (International comparison of students)
    - cepeda2006 [22]: "This review found 839 assessments of distributed practice in 317 experiments located in 184 articles." (Quantitative synthesis)
- **Verdict: mixed**. Fatigue is real and breaks help at school scale; evidence that a 6-18 item session is optimal does not exist. Code is a reasoned range (**design choice**).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M13. Break guidance: none in the code

- Searched `public/js`, `routes`, `services` for break prompts: none. The only break-like behaviour is idle auto-pause (M11) and the enforcer after End.
- **Evidence that guidance would help:** albulescu2022 [524] (micro-breaks raise vigor d = .36 and cut fatigue d = .35, but the performance effect was "non-significant"); biwer2023 [520]; blasche2018 [525]; sievertsen2016 [199].
    - albulescu2022 [524]: "statistically significant but small effects of micro-breaks in boosting vigor (d = .36, p < .001; k = 9, n = 913)"; "a non-significant effect on increasing overall performance (d = .16, p = .116; k = 15, n = 1132)" (Meta-analysis, 22 samples, N = 2335, mostly working adults and students)
    - biwer2023 [520]: "Taking pre-determined, systematic breaks during a study session had mood benefits and appeared to have efficiency benefits" (87 Dutch university students, one day)
    - blasche2018 [525]: "the unstructured rest break led to an increase in vigor" (66 students, 4-hour university classes)
    - sievertsen2016 [199]: "for every hour later in the day, test performance decreases by 0.9% of an SD"; "a 20- to 30-minute break improves average test performance by 1.7% of an SD" (All Danish public-school children, 2009-2013)
- **Verdict: gap, not a mechanism**; evidence is **moderate for wellbeing, weak for performance**. See proposal P1.

- **Status: Built on 25 October 2026** (switch `BREAK_NUDGE`, from 05-P1). After 25 active minutes without a pause (`BREAK_NUDGE_MIN`), one quiet, dismissible break line. Learners can turn it off; not shown to the owner.

05-M14. Retention-enforcer prompt in the strip

- **Where:** `study.js:62-75,108-115`: "Start a session to time your study. At the end, a short retention enforcer locks it in."
- **Evidence:** roediger2006 [1], rowland2014 [10] (retrieval beats restudy; recall tests larger than recognition). These support the practice; the enforcer's design is covered in the retention section.
    - roediger2006 [1]: "Testing is a powerful means of improving learning, not just assessing it." (College students, prose passages)
    - rowland2014 [10]: "initial recall tests yielding larger testing benefits than recognition tests" (Meta-analysis of testing vs restudy)
    - cepeda2006 [22]: "This review found 839 assessments of distributed practice in 317 experiments located in 184 articles." (Quantitative synthesis)
- **Verdict: strong** for the practice. Wording gap: "locks it in" overclaims; the evidence says retrieval improves later retention, not that anything is locked.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-C. Feedback, rationales, hints and the coach

05-M15. Verdict wording

- **Where:** `public/js/player.js:439-455`: "Right." or "Not quite. The answer is X.", then an optional feedback line and the rationale; correct and wrong marks are drawn plus announced (`player.js:466-480,493`).
- **Evidence:** pashler2005 [122] (supplying the correct answer after an error raised final retention by 494%; feedback after correct answers made little difference); shute2008 [16] (feedback should be "nonevaluative, supportive, timely, and specific"); hattie2007 [123]; wisniewski2020 [488] (d = 0.48); metcalfe2017 [219] (errorful learning with corrective feedback is beneficial).
    - pashler2005 [122]: "Supplying the correct answer after an incorrect response not only improved performance during the initial learning session"; "feedback after correct responses made little difference either immediately or at a delay" (N = 258 adults, word pairs, one-week test)
    - shute2008 [16]: "formative feedback should be nonevaluative, supportive, timely, and specific" (Review)
    - hattie2007 [123]: "Feedback is one of the most powerful influences on learning and achievement, but this impact can be either positive or negative." (Review)
    - wisniewski2020 [488]: "medium effect (d = 0.48) of feedback on student learning" (Meta-analysis)

    - metcalfe2017 [219]: "errorful learning followed by corrective feedback is beneficial to learning"; "Errors committed with high confidence are corrected more readily than low-confidence errors." (Review of laboratory and classroom studies)
- **Verdict: strong**. Code matches: neutral wording, no trait praise, correct answer shown after an error (M17 covers rationale).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M16. Praise and tone

- No praise strings in `player.js` verdict; `services/ai.js:1202-1225` coach prompt says "Warm and direct, never patronizing".
- **Evidence:** mueller1998 [371] (praise for intelligence worsened motivation and persistence after failure versus praise for effort); henderlong2002 [501] (sincere, controllable-cause praise helps); skipper2012 [601]; morris2014 [602].
    - mueller1998 [371]: "praise for intelligence had more negative consequences for students' achievement motivation than praise for effort" (Six studies, fifth graders)
    - henderlong2002 [501]: "Provided that praise is perceived as sincere, it is particularly beneficial to motivation when it encourages performance attributions to controllable causes" (Review of children's praise research)
    - skipper2012 [601]: "After one failure, participants who received person praise reacted most negatively on all dependent measures." (British children aged 9 to 11 and university students; imagined school scenarios, two studies). Scope: Shows that person praise worsens the response to a later failure. It did not test warmth or sincerity.
    - morris2014 [602]: "Ambiguous praise, similar to verbal effort praise, produced higher persistence and self-evaluations, and fewer fixations on error after failure compared to verbal trait praise." (95 children aged 5 to 6, praise assigned at random, measured after failure). Scope: Shows that ambiguous praise, like effort praise, led to better persistence than trait praise after failure. It did not test warmth or sincerity.
- **Verdict: strong** that trait praise backfires; code matches (none used).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M17. A rationale on every item

- **Where:** `public/js/player.js:445` (rationale div, inside the verdict at 439-455); `routes/learn.js:119,535` (rationale stored on the item and returned with the result); fallback text at `learn.js:966`.
- **Evidence:** butler2013 [17] (explanation feedback beat correct-answer feedback on new inference questions); vanderkleij2015 [15] (elaborated feedback effect larger than correctness-only, 0.49 against 0.05); little2012 [133] (multiple-choice tests teach more when learners can retrieve why wrong alternatives are wrong); wisniewski2020 [488].
    - butler2013 [17]: "explanation feedback produced superior performance on the new inference questions" (College students, prose passages, test two days later)

    - vanderkleij2015 [15]: "than feedback regarding the correctness of the answer (KR; 0.05)" (Meta-analysis of computer-based feedback)
    - little2012 [133]: "properly constructed multiple-choice tests can indeed trigger productive retrieval processes" (Two experiments, undergraduates)
    - wisniewski2020 [488]: "medium effect (d = 0.48) of feedback on student learning" (Meta-analysis)
- **Verdict: strong** for explanation feedback on transfer. Code matches for the display. Gap: rationale quality is generated; the evidence applies only if the explanation is correct (covered by the fact-check sections, not here).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M18. Held versus immediate feedback

- **Where:** `services/learner.js:308` (`immediateFeedback: timing < 0.55`), `public/js/session.js:80,188-191,214-217,1036-1063` (held answers shown together at the end), `routes/learn.js:92`.
- **What it does:** most sittings show the verdict at once; learners whose profile prefers delay (and every checkpoint or exam) see everything at the end.
- **Evidence:** kulik1988 [184] (classroom studies mostly favour immediate feedback); butler2007 [185] (delayed feedback gave better final test performance); mullet2014 [246] (delay can promote transfer, although students prefer immediate).
    - kulik1988 [184]: "Applied studies using actual classroom quizzes and real learning materials have usually found immediate feedback to be more effective than delayed." (Review of timing studies)
    - butler2007 [185]: "delayed feedback led to superior final test performance relative to immediate feedback" (Adults, multiple-choice tests on prose)
    - mullet2014 [246]: "a growing body of research has challenged this assumption by demonstrating that delaying feedback can facilitate learning" (College students)
- **Verdict: mixed**; sources conflict by setting (lab versus classroom, retention versus transfer). Held feedback on scored sittings is an assessment-integrity choice, not a learning claim. Gap: for practice, delay is set by learner preference (the thing mullet2014 [246] says students get wrong), not by evidence.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M19. Optional "How sure are you?" row

- **Where:** `public/js/player.js:79-95`, enabled by `session.js:760` (`!isExam() && index % 3 === 1`), sent with the answer at `routes/learn.js:464`.
- **What it does:** three chips (Sure, Think so, Guessing) on every third non-exam question; optional; a second tap clears it. It feeds the scorecard calibration stat.
- **Evidence:** butler2008 [220] (feedback corrects low-confidence correct answers, improving retention); metcalfe2017 [219] (high-confidence errors are corrected more readily); stankov2012 [526] (confidence combined with accuracy gives a bias score, the

idea behind the calibration stat); dunlosky2012 [231] (more accurate self-judgments went with better retention two days later).
    - butler2008 [220]: "when correct responses are made with low confidence, feedback serves to correct this initial metacognitive error, enhancing retention" (Undergraduates, general-knowledge questions)
    - metcalfe2017 [219]: "errorful learning followed by corrective feedback is beneficial to learning"; "Errors committed with high confidence are corrected more readily than low-confidence errors." (Review of laboratory and classroom studies)
    - stankov2012 [526]: "can be combined with accuracy information to obtain bias scores that may be useful for group comparisons and for identification of misconceptions" (15-year-olds in Singapore, mathematics (N = 1940) and English (N = 1786))
    - dunlosky2012 [231]: "greater accuracy was associated with higher levels of retention" (College students, key-term definitions, two studies, final test two days later)
- **Verdict: moderate**. Code matches (optional, spaced, not on exams). The chips are 36 px high (M28).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M20. Multiple choice versus type-in

- **Where:** `public/js/player.js:100-170` (both renderers; type-in and short answer differ only by box size), formats chosen by the item generators.
- **Evidence:** roediger2006 [1] and rowland2014 [10] (testing beats restudy; recall tests beat recognition tests); kang2007 [166] (with feedback, short-answer quizzing gave the best result; without feedback multiple choice did better); smith2014 [116] ("little or no advantages" of short-answer over multiple-choice); mcdermott2014 [155] (both formats helped in grades 7-12 and format need not match the exam); little2012 [133] (well-built multiple choice triggers real retrieval); vangog2015 [527] (the testing effect shrinks as material gets complex).
    - roediger2006 [1]: "Testing is a powerful means of improving learning, not just assessing it." (College students, prose passages)
    - rowland2014 [10]: "initial recall tests yielding larger testing benefits than recognition tests" (Meta-analysis of testing vs restudy)
    - kang2007 [166]: "having had an intervening MC test led to better performance than an intervening SA test"; "With feedback provided, having had an intervening SA test led to the best performance on the final test" (Undergraduates, journal papers)
    - smith2014 [116]: "little or no advantages of answering short-answer or hybrid format questions over multiple-choice questions" (Undergraduates)
    - mcdermott2014 [155]: "the format of the quiz (multiple-choice or short-answer) did not need to match the format of the criterial test" (Seventh-grade science, high-school history)

    - little2012 [133]: "properly constructed multiple-choice tests can indeed trigger productive retrieval processes" (Two experiments, undergraduates)
    - vangog2015 [527]: "the testing effect decreases as the complexity of learning materials increases" (Review)
- **Verdict: strong** that retrieval of either format beats restudy; **mixed** on type-in over multiple choice (small, feedback-dependent). Code matches: both formats exist, every answer gets feedback.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M21. `preferTypeIn`

- **Where:** `services/learner.js:286,300` (`retrieval > 0.55`, or the scorecard's `typed` flag), `services/sessions.js:81-110` (up to 10 regenerations seeking a type-in item), `services/adaptation.js:546,573`, `services/ai.js:506`.
- **What it does:** for learners whose style leans to recall, or who rush, the generator prefers a typed-answer item when the skill has one; otherwise it returns the multiple-choice item.
- **Evidence:** rowland2014 [10] (recall over recognition); kang2007 [166] (best with feedback); smith2014 [116] (little advantage); mcdermott2014 [155] (equal in classrooms); pashler2008 [31] and rogowsky2015 [317] (preference is not a reason to match).
    - rowland2014 [10]: "initial recall tests yielding larger testing benefits than recognition tests" (Meta-analysis of testing vs restudy)
    - kang2007 [166]: "having had an intervening MC test led to better performance than an intervening SA test"; "With feedback provided, having had an intervening SA test led to the best performance on the final test" (Undergraduates, journal papers)
    - smith2014 [116]: "little or no advantages of answering short-answer or hybrid format questions over multiple-choice questions" (Undergraduates)
    - mcdermott2014 [155]: "the format of the quiz (multiple-choice or short-answer) did not need to match the format of the criterial test" (Seventh-grade science, high-school history)
    - pashler2008 [31]: "the meshing hypothesis, according to which instruction is best provided in a format that matches the preferences of the learner" (Review)
    - rogowsky2015 [317]: "Results failed to show a statistically significant relationship between learning style preference (auditory, visual word) and learning aptitude" (College-educated adults)
- **Verdict: mixed**. The format effect is modest and conditional; choosing it from a preference quiz has no support (M9). The rush-driven trigger is a separate, defensible use. Gap: the typed answer must be graded fairly or learners are punished for typos; see grading sections.
- **Status: Built on 25 October 2026** (switch `SCORECARD_DROP_RUSH_TYPED`, from 03-P1). The rush-driven typed trigger is gone; the style-quiz trigger remains (05-P6, not built). **Open:** the rest of 05-P6, a type-in default independent of the quiz (not built).

05-M22. Neutral "Checking your answer..." state

- **Where:** `public/js/player.js:402-436` (hold at 407, poll at 422-435).
- **What it does:** while the grader is unsure, shows a status (never "right" or "wrong"), lets the learner continue, and replaces it when the check ends (polls every 3 s up to 4 min).
- **Evidence:** nah2004 [528] (feedback prolongs tolerable waiting time); myers1985 [512] (progress and status indicators preferred); wcag-4.1.3 [521] (status via role="status"); ceaparu2004 [529] (waits and error messages are top frustrations).
    - nah2004 [528]: "the presence of feedback prolongs Web users' tolerable waiting time" (Web users, experiment)
    - myers1985 [512]: "people prefer to have progress indicators" (Formal experiment (adults))
    - wcag-4.1.3 [521]: "status messages can be programmatically determined through role or properties such that they can be presented to the user by assistive technologies without receiving focus." (W3C Level AA)
    - ceaparu2004 [529]: "The most-cited causes of frustrating experiences were error messages, dropped network connections, long download times, and hard-to-find features." (111 computer users)
- **Verdict: moderate**. Code matches (role status, announcement).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M23. The coach: asks, never tells, before the answer

- **Where:** `public/js/player.js:278-398` (panel, streamed turns), `services/ai.js:1202-1225` (system prompt: never state the answer, one question per turn, stop when the learner reaches it), `routes/learn.js:787-830` (route, `coach_used` flag, 120 per hour), prime logic `learn.js:755-776`.
- **Evidence:** vanlehn2011 [36] (step-based tutoring d = 0.76, near human tutoring); kulik2016 [37] (median +0.66 SD over 50 evaluations); kestin2025 [530] (a pedagogically designed AI tutor beat in-class active learning); bastani2025 [38] (unguarded GPT lowered later unaided grades by 17%; a guardrailed tutor largely removed the harm); kapur2008 [531] (struggle before support can be productive).
    - vanlehn2011 [36]: "the effect size of intelligent tutoring systems was 0.76, so they are nearly as effective as human tutoring" (Review of tutoring experiments)
    - kulik2016 [37]: "The median effect of intelligent tutoring in the 50 evaluations was to raise test scores 0.66 standard deviations over conventional levels" (50 controlled evaluations)
    - kestin2025 [530]: "students learn significantly more in less time when using the AI tutor, compared with the in-class active learning" (College students, RCT)
    - bastani2025 [38]: "students actually perform worse than those who never had access (17% reduction in grades for GPT Base)"; "These negative learning effects are largely mitigated by the safeguards in GPT Tutor." (Nearly 1,000 high-school maths students)

    - kapur2008 [531]: "engaging students in solving complex, ill-structured problems without the provision of support structures can be a productive exercise in failure" (Eleventh-grade science students)
- **Verdict: strong** for step-based tutoring in general; **moderate** for an LLM coach that withholds the answer (two 2025 trials, different settings). Code matches the guardrail finding. Honest limits: the prompt is a rule, not a guarantee (the route comments say so); no trial of this coach exists.
- **Status:** Existing (in the code before 25 October 2026; not changed that day). **Open:** 05-P7, a graded hint ladder with a try-first step (not built).

05-M24. Help is withheld on scored sittings

- **Where:** `routes/learn.js:749-790` (`heldItem` returns 409), `session.js:1019,1332` and `help.js:171` (`coach: false` on recall and exams).
- **Evidence:** aleven2003 [173] and baker2008 [509] (on-demand help is misused and gamed); roediger2006 [1] (retrieval without aids is what produces the test effect); kalyuga2003 [6] (help that suits novices can hurt more experienced learners).
    - aleven2003 [173]: "although effective help-seeking behavior in ILEs is related to better learning outcomes, learners are not using help facilities effectively" (Selective review)
    - baker2008 [509]: "present evidence on which student behaviors, motivations, and emotions are associated with the choice to game the system" (Students in interactive learning environments)
    - roediger2006 [1]: "Testing is a powerful means of improving learning, not just assessing it." (College students, prose passages)
    - kalyuga2003 [6]: "Instructional techniques that are highly effective with inexperienced learners can lose their effectiveness and even have negative consequences when used with more experienced learners" (Review)
- **Verdict: design choice** with supporting principle. Code matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M25. "Ask about this" after the answer

- **Where:** `player.js:498-530`, `routes/learn.js:948-985`.
- **What it does:** a text box ("What part is unclear?") that returns a direct explanation once the answer is already on screen.
- **Evidence:** chi1994 [532] (self-explanation prompts improve understanding; the box asks the learner to articulate the gap first); butler2013 [17] (explanations aid transfer).
    - chi1994 [532]: "Generating explanations to oneself (self-explaining) facilitates that integration process." (14 eighth graders, one passage)
    - butler2013 [17]: "explanation feedback produced superior performance on the new inference questions" (Prose passages, test two days later)
    - metcalfe2017 [219]: "errorful learning followed by corrective feedback is beneficial to learning"; "Errors committed with high confidence are corrected

more readily than low-confidence errors." (Review of laboratory and classroom studies)

- **Verdict: moderate**. Code matches. Gap: when no model is connected the reply is only the stored rationale (`learn.js:966`), which says so honestly.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M26. "Report a problem."

- **Where:** `player.js:531-540` (uses `window.prompt` at 532), `routes/learn.js:986`.
- **What it does:** removes the item from service while it is checked.
- **Evidence:** lee2004 [533] (trust influences reliance on automation); ceaparu2004 [529] (errors are a leading frustration).
  - lee2004 [533]: "Automation is often problematic because people fail to rely upon it appropriately." (Review)
  - glikson2020 [534]: "The success of integrating AI into organizations critically depends on workers' trust in AI technology." (Review of 20 years of research)
  - ceaparu2004 [529]: "The most-cited causes of frustrating experiences were error messages, dropped network connections, long download times, and hard-to-find features." (111 computer users)
- **Verdict: design choice**. Gap: a native `prompt()` dialog is poor for assistive tech and for small screens; an inline field would match the rest of the player (proposal list, minor).
- **Status: Built on 25 October 2026** (switch `UX_REPORT_DIALOG`). A native, labelled dialog replaces `window.prompt`.

05-M27. "Too fast to count" note and re-served question

- **Where:** `services/effort-monitor.js:5-31` (1.2 s on scored sittings, max 2 re-serves; slow-down note elsewhere for learners with more than 5% rapid answers), off for the owner account.
- **Evidence:** wise2005 [74] (unmotivated examinees answer before reading); wisedemars2005 [523] (low effort lowers measured performance); baker2008 [509] (gaming behaviours).
  - wise2005 [74]: "unmotivated examinees will answer too quickly (i.e., before they have time to read and fully consider the item)" (Computer-based test examinees)
  - wisedemars2005 [523]: "low student motivation is associated with a substantial decrease in test performance" (Low-stakes assessment review)
  - baker2008 [509]: "present evidence on which student behaviors, motivations, and emotions are associated with the choice to game the system" (Students in interactive learning environments)
- **Verdict: moderate**. Code matches the cited principle. Wording is non-punitive ("It is never a penalty").
- **Status: Built on 25 October 2026** (switch `RAPID_ITEM_RELATIVE`, from 03-P8). Long items use an item-relative rapid line once they have 30 or more timed answers.

05-M28. Optional "Give Input" form that never gates progress

- **Where:** `public/js/feedback.js:7-12,42-60`.
- **Evidence:** patall2008 [46] and ryan2000 [47] (choice and autonomy support motivation); cordova1996 [319] (personalization and choice strategies); kluger1996 [342] (task-focused feedback).
    - patall2008 [46]: "providing choice enhanced intrinsic motivation, effort, task performance, and perceived competence" (Meta-analysis of choice studies)
    - ryan2000 [47]: "three innate psychological needs--competence, autonomy, and relatedness--which when satisfied yield enhanced self-motivation and mental health" (Theory review)
    - cordova1996 [319]: "contextualization, personalization, and provision of choices" (School-age children learning with a computer task)
    - kluger1996 [342]: "FIs improved performance on average (d =.41)" (Meta-analysis, 23,663 observations)
- **Verdict: design choice**. Code matches. Tap options and free text lower the cost of answering.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-D. Accessibility, mobile and visual design

05-M29. Touch targets

- **Where:** `public/css/style.css:429` (`.btn` 44 px), `:464` (`.btn-sm` 36 px), `:562-565` and `:973` (`.chip-btn`, confidence chips 36 px), `:2037-2040` (44 px on coarse pointers only inside the study strip), `:1717-1720` (kids variant: buttons 52 px, choices 64 px), tab bar 60 px (`:376`).
- **Evidence:** wcag-2.5.8 [48] (24 by 24 CSS px, AA); wcag-2.5.5 [535] (44 by 44, AAA); parhi2006 [536] (9.2 mm, about 35 CSS px at the 96 dpi reference, sufficient for one-handed thumb discrete taps on 2006 devices).
    - wcag-2.5.8 [48]: "The size of the target for pointer inputs is at least 24 by 24 CSS pixels" (W3C Level AA)
    - wcag-2.5.5 [535]: "The size of the target for pointer inputs is at least 44 by 44 CSS pixels" (W3C Level AAA)
    - parhi2006 [536]: "target size of 9.2 mm for discrete tasks and targets of 9.6 mm for serial tasks should be sufficiently large for one-handed thumb use" (Adults, small touchscreens (2006 devices))
- **Verdict: standard**. Code passes AA everywhere and the 44 px AAA target on the main buttons and primary actions. Gap: the 36 px small buttons and chips (Ask, Report, confidence chips, feedback chips) stay 36 px on phones; the CSS comment at `:2035` names WCAG 2.5.5, which is AAA, not required. Children under 7 are less precise (vatavu2015 [537], title-level only). See P4.

- **Status: Built on 25 October 2026** (switch `UX_TOUCH44`, from 05-P4). Small buttons and chips are 44 px on coarse pointers and phones; the study strip also got 44 px targets (UI fix, no switch).

05-M30. Colour and contrast tokens

- **Where:** `style.css:30-50` (light), `:159-199` (dark), `:68-72`.
- Measured here from the hex values: ink 17.9:1, ink-2 8.1:1, ink-3 5.35:1 on white and 4.9:1 on the well; control outline 3.4:1; dark ink-3 6.0:1 and dark control outline 3.96:1 on the dark page. The right/wrong colours (5.2 and 5.8) are quoted from the CSS comment (oklch values not recomputed).
- **Evidence:** wcag-1.4.3 [538] (4.5:1 text); wcag-1.4.11 [539] (3:1 for UI components); hall2004 [540] (higher contrast ratios gave greater readability, 136 participants).
  - wcag-1.4.3 [538]: "contrast ratio of at least 4.5:1" (W3C Level AA)
  - wcag-1.4.11 [539]: "contrast ratio of at least 3:1 against adjacent color(s)" (W3C Level AA)
  - hall2004 [540]: "Colours with greater contrast ratio generally lead to greater readability" (136 participants, web pages)
- **Verdict: standard**. Code matches on the measured tokens. Dark and light themes both follow the system preference (`app.js:262-292`) with a manual toggle.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M31. Right and wrong are never colour-only

- **Where:** `player.js:19-28,457-480`: tick or cross drawn on the option, a visually hidden "(your answer, correct)" note, and the words "Right." or "Not quite."
- **Evidence:** wcag-1.4.1 [541] (colour not the only means), wcag-1.4.11 [539], wcag-4.1.3 [521].
  - wcag-1.4.1 [541]: "Color is not used as the only visual means of conveying information, indicating an action, prompting a response, or distinguishing a visual element." (W3C Level A)
  - wcag-1.4.11 [539]: "contrast ratio of at least 3:1 against adjacent color(s)" (W3C Level AA)
  - wcag-4.1.3 [521]: "status messages can be programmatically determined through role or properties such that they can be presented to the user by assistive technologies without receiving focus." (W3C Level AA)
- **Verdict: standard**. Code matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M32. Reduced motion

- **Where:** `style.css:1736-1746` (all animations and transitions near zero, marks drawn instantly, spinner kept), `player.js:486-488` (instant scroll when reduced motion is set), non-moving `rise` and `fade` keyframes at `:1733-1734`.

- **Evidence:** wcag-2.3.3 [542] (motion from interaction can be disabled); harp1998 [127] (interesting but irrelevant additions reduced main-idea recall); berney2016 [543] (title-level only: animation meta-analysis); noetel2022 [138] (design principles with positive effects include coherence).
    - wcag-2.3.3 [542]: "Motion animation triggered by interaction can be disabled, unless the animation is essential to the functionality or the information being conveyed." (W3C Level AAA)
    - harp1998 [127]: "recalled significantly fewer main ideas and generated significantly fewer problem-solving transfer solutions" (Four experiments, students reading with interesting but irrelevant additions; indirect background evidence, not a motion study)
    - berney2016 [543]: title-level only.
    - noetel2022 [138]: "We found 11 design principles that demonstrated significant, positive, meta-analytic effects on learning" (Meta-meta-analysis, multimedia design)
- **Verdict: standard** plus **design choice** (no study of reduced-motion CSS on learning). Code matches. The spinner deliberately keeps a slow spin so loading stays visible.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M33. Announcements, toasts and focus

- **Where:** `app.js:52-80` (toast host `role="status"`, error toasts `role="alert"`; a single polite live region), `player.js:493-495` (focus to Next, or the verdict heading), `study.js:162` (focus to the first End button).
- **Evidence:** wcag-4.1.3 [521], wcag-2.4.7 [522], wcag-2.2.1 [544] (time limits).
    - wcag-4.1.3 [521]: "status messages can be programmatically determined through role or properties such that they can be presented to the user by assistive technologies without receiving focus." (W3C Level AA)
    - wcag-2.4.7 [522]: "Any keyboard operable user interface has a mode of operation where the keyboard focus indicator is visible." (W3C Level AA)
    - wcag-2.2.1 [544]: "For each time limit that is set by the content, at least one of the following is true:" (W3C Level A)
- **Verdict: standard**. Code matches 4.1.3 and 2.4.7. Possible gap: toasts disappear after 3.2 s (info) or 5.2 s (error) (`app.js:62`); an error that vanishes can be missed by sighted learners, and WCAG 2.2.1 asks that time limits set by content be adjustable. Whether auto-dismissal counts under 2.2.1 is arguable, so this is flagged, not failed.
- **Status: Built on 25 October 2026** (switch `UX_TOAST_STAY`). Error toasts stay until dismissed (`role="alert"`, a 44 px Dismiss button); info toasts still fade.

05-M34. Skip link and focus ring

- **Where:** `style.css:250-253,300-306`.
- **Evidence:** wcag-2.4.1 [545], wcag-2.4.7 [522], wcag-1.4.11 [539].

    - wcag-2.4.1 [545]: "mechanism is available to bypass blocks of content that are repeated on multiple web pages." (W3C Level A)
    - wcag-2.4.7 [522]: "Any keyboard operable user interface has a mode of operation where the keyboard focus indicator is visible." (W3C Level AA)
    - wcag-1.4.11 [539]: "contrast ratio of at least 3:1 against adjacent color(s)" (W3C Level AA)
- **Verdict: standard**. Code matches (2 px ink ring, offset 2 px).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M35. Mobile layout: bottom tab bar, fixed Next bar, safe areas

- **Where:** `style.css:374-420` (tab bar from 760 px, `env(safe-area-inset-bottom)`, scroll padding), `:951-966` (fixed Next bar in the player), `:480-483` (answer row stacks at 480 px), `player.js:486-488` (verdict scrolls into view).
- **Evidence:** sung2016 [546] (mobile-device learning, moderate effect size 0.523, 110 studies); parhi2006 [536] (thumb reach and target size); wcag-1.4.10 [547] (reflow without two-dimensional scrolling at 320 CSS px).
    - sung2016 [546]: "moderate mean effect size of 0.523 for the application of mobile devices to education" (Meta-analysis, 110 studies, 1993-2013)
    - parhi2006 [536]: "target size of 9.2 mm for discrete tasks and targets of 9.6 mm for serial tasks should be sufficiently large for one-handed thumb use" (Adults, small touchscreens (2006 devices))
    - wcag-1.4.10 [547]: "Content can be presented without loss of information or functionality, and without requiring scrolling in two dimensions" (W3C Level AA)
- **Verdict: moderate** for mobile learning generally; the specific layout is **design choice**. Not verified here: the 320 px reflow behaviour was not tested in a browser in this pass.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M36. Kids audience variant

- **Where:** `app.js:263-300`, `style.css:1709-1721` (larger type, Andika font, 52 px buttons, 64 px choices, larger marks), applied from the consent route (`routes/consent.js:340-341`).
- **Evidence:** parhi2006 [536]; wcag-2.5.5 [535]; hughes2000 [548] (children aged 5 to 7 read larger type faster; children aged 8 to 11 were neither helped nor slowed); vatavu2015 [537] (title-level only: children aged 3 to 6 and touch). The type-size evidence is from children; the button and choice sizes still rest on adult touch data and the AAA minimum.
    - parhi2006 [536]: "target size of 9.2 mm for discrete tasks and targets of 9.6 mm for serial tasks should be sufficiently large for one-handed thumb use" (Adults, small touchscreens (2006 devices))
    - wcag-2.5.5 [535]: "The size of the target for pointer inputs is at least 44 by 44 CSS pixels" (W3C Level AAA)

    - hughes2000 [548]: "The reading speed of children aged 5 to 7 years decreased as the text size decreased" (120 children aged 5 to 11, passages at four type sizes)
    - vatavu2015 [537]: title-level only.
- **Verdict: design choice**. Larger type has direct evidence for children aged 5 to 7; the 52 and 64 px targets are not tested on children.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-E. Calendar, reminders, errors and empty states

05-M37. Retention calendar

- **Where:** `public/js/calendar.js` (month grid as a real table, Monday first, one full sentence per day for screen readers, legend, "Coming up" list).
- **Evidence:** cepeda2006 [22] (spacing intervals matter); lally2009 [45] (consistent context builds automaticity); milkman2011 [549] (planning a specific date and time raised follow-through).
    - cepeda2006 [22]: "This review found 839 assessments of distributed practice in 317 experiments located in 184 articles." (Quantitative synthesis)
    - lally2009 [45]: "Missing one opportunity to perform the behaviour did not materially affect the habit formation process."; "ranged from 18 to 254 days" (96 adult volunteers, 12 weeks, daily behaviour)
    - milkman2011 [549]: "had a 4.2 percentage point higher vaccination rate" (Employees of one large firm, flu vaccination)
- **Verdict: moderate** for visible scheduled practice; the table semantics follow standard patterns. Gap: it is read-only, it asks the learner to plan nothing. See P3.
- **Status: Built on 25 October 2026** (switch `UX_NEXT_STUDY_PLANNER`, from 05-P3). A dashboard card saves the next study day and time. It sends nothing to anyone.

05-M38. Reminders: none exist

- Search found no push, no learner email or SMS (`utils/mailer` is used only by `routes/consent.js`, `routes/ops.js`, `services/error-log.js`).
- **Evidence that reminders help:** milkman2011 [549] (date-and-time implementation-intention prompt, +4.2 points); castleman2015 [550] (personalized text reminders raised enrollment for students with less access to guidance); gollwitzer2006 [551] (title-level only, existing meta-analysis). Evidence of harm: kushlev2016 [552] (alerts raised inattention and hyperactivity symptoms, 221 adults); ico-aadc-13 [553] (no nudges toward weaker privacy).
    - milkman2011 [549]: "had a 4.2 percentage point higher vaccination rate" (Employees of one large firm, flu vaccination)
    - castleman2015 [550]: "The interventions substantially increased college enrollment among students who had less academic-year access to quality college counseling or information." (College-intending low-income US high-school graduates, text-message reminders)

    - gollwitzer2006 [551]: title-level only.
    - kushlev2016 [552]: "Participants reported higher levels of inattention and hyperactivity when alerts were on than when alerts were off." (221 adults, one week per condition)
    - ico-aadc-13 [553]: "Do not use nudge techniques to lead or encourage children to provide unnecessary personal data or turn off privacy protections." (UK statutory code)
- **Verdict: gap**; evidence **moderate** for benefit, with a real cost. Under-13s cannot be contacted directly (COPPA); a reminder would go to the parent account. See P3.
- **Status: Built on 25 October 2026** (switch `UX_NEXT_STUDY_PLANNER`, from 05-P3). Partly closed: the in-app planner only. Emailed or texted reminders are not built (legal review). **Open:** the rest of 05-P3, emailed or texted reminders and the parent digest (owner decision and legal review).

05-M39. Error states

- Toasts (`app.js:52`), inline notices with the server message (`calendar.js:143`, `player.js` ask and report), the honest degrade when the model is down (`learn.js:966`, coach fallback `learn.js:818-830`), session errors (`study.js:177`).
- **Evidence:** ceaparu2004 [529] (error messages are among the most cited frustrations); wcag-3.3.1 [554] (identify the error in text); wcag-3.3.3 [555] (suggest a fix).
    - ceaparu2004 [529]: "The most-cited causes of frustrating experiences were error messages, dropped network connections, long download times, and hard-to-find features." (111 computer users)
    - wcag-3.3.1 [554]: "If an input error is automatically detected, the item that is in error is identified and the error is described to the user in text." (W3C Level A)
    - wcag-3.3.3 [555]: "If an input error is automatically detected and suggestions for correction are known, then the suggestions are provided to the user" (W3C Level AA)
- **Verdict: standard / moderate**. Code matches 3.3.1 where messages are shown; suggestions for a fix are inconsistent (some messages say only "That did not work.", `study.js:177`).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M40. Empty states

- **Where:** `dashboard.js:41` (empty-state block with a route to add a subject), `enforcer.js:270` ("Nothing to lock in..." / "Too short for an enforcer..."), `study.js:170-172` (the toast saying why no enforcer follows), `style.css:1468-1471`.
- **Evidence:** ceaparu2004 [529]; nielsen1994 [556] (title-level only); lazar2006 [603]; taipalus2025 [604].
    - ceaparu2004 [529]: "The most-cited causes of frustrating experiences were error messages, dropped network connections, long download times, and hard-to-find features." (111 computer users)

    - nielsen1994 [556]: title-level only.
    - lazar2006 [603]: "shows high levels of frustration and loss of 1/3–1/2 of time spent" (107 student and 50 workplace computer users). Scope: Shows that computer frustration is common and costly. It did not test empty states.
    - taipalus2025 [604]: "articulating what is wrong in the query with clear natural language, and showing hints on how to fix the error" (165 novices; qualitative analysis of 2,052 answers about SQL error messages). Scope: Shows that novices value a plain statement of what is wrong plus a hint on how to fix it, the pattern an empty state follows. It did not test empty states.
- **Verdict: design choice**. Code matches: each empty state names the reason and the next action.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-F. Owner skips

05-M41. Owner-only skips

- **Where:** `config/capabilities.js:21` (`canSkip` is the owner account only), `services/owner-skips.js:2-12,40` (skip placement, skip a grade, stand on any level; gates not held), `effort-monitor.js:38-43` (owner exempt), `GET can_skip` only decides which buttons are drawn; routes re-check.
- **Evidence:** karich2014 [557] (learner control in educational technology had "almost zero" overall effect, g = 0.05); patall2008 [46] (choice helps motivation, within limits); kalyuga2003 [6] (support that helps novices can hurt experts, a reason placement and gates exist); katz2007 [558] (title-level only).
    - karich2014 [557]: "The overall effect of including learner control within educational technology was almost zero (g = 0.05)" (Meta-analysis, 18 articles, 29 effects)
    - patall2008 [46]: "providing choice enhanced intrinsic motivation, effort, task performance, and perceived competence" (Meta-analysis of choice studies)
    - kalyuga2003 [6]: "Instructional techniques that are highly effective with inexperienced learners can lose their effectiveness and even have negative consequences when used with more experienced learners" (Review)
    - katz2007 [558]: title-level only.
- **Verdict: design choice (product only)**. The evidence says unrestricted learner control is not reliably beneficial, which supports keeping it out of the learner build. Skips are for the owner to inspect the product. Not a learning mechanism, so no learner effect is claimed.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-G. Privacy, consent and minors

05-M42. Neutral age screen

- **Where:** `public/js/age-screen.js:1-46`, `routes/consent.js:82-83`, first answer held for the browser session (`consent.js` header lines 12-14).
- **What it does:** month, year of birth and country, no defaults, no hint of which answer passes; the first answer in a session counts.
- **Evidence:** coppa-312.5a1 [49] (verifiable parental consent before collection from children); gdpr-art8 [559] (16 by default, lower by member state, not below 13); ico-aadc-8 [560] / UK code (age-appropriate application); reyes2018 [561] (a majority of 5,855 children's apps were potentially non-compliant, mainly through third-party SDKs). The FTC FAQ D.7 neutrality guidance is read through the repo's `KIDS-CAMERA-MIC-PRIVACY.md` (the FAQ URL itself is title-level here).
    - coppa-312.5a1 [49]: "An operator is required to obtain verifiable parental consent before any collection , use, or disclosure of personal information from children" (US federal rule, children under 13)
    - gdpr-art8 [559]: "the processing of the personal data of a child shall be lawful where the child is at least 16 years old"; "Member States may provide by law for a lower age for those purposes provided that such lower age is not below 13 years." (EU regulation)
    - ico-aadc-8 [560]: "Settings must be “high privacy” by default" (UK statutory code)
    - reyes2018 [561]: "a majority are potentially in violation of COPPA" (5,855 popular free children's Android apps)
- **Verdict: standard**. Code matches the design in the repo doc. Not legal advice; the repo doc says a lawyer must confirm.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M43. Country age-of-consent table

- **Where:** `consent.js:82-83`: US 13, GB 13, DK 13, SE 13, ES 14, IT 14, FR 15, DE 16, NL 16, default 16.
- **Evidence:** gdpr-art8 [559] (the 13-16 range is member-state law); coppa-312.5a1 [49]; ico-iss-consent [562] (UK: 13 and over may consent); es-lopdgdd-art7 [563] (Spain: "mayor de catorce años", the 14 in the table).
    - gdpr-art8 [559]: "the processing of the personal data of a child shall be lawful where the child is at least 16 years old"; "Member States may provide by law for a lower age for those purposes provided that such lower age is not below 13 years." (EU regulation)
    - coppa-312.5a1 [49]: "An operator is required to obtain verifiable parental consent before any collection , use, or disclosure of personal information from children" (US federal rule, children under 13)
    - ico-iss-consent [562]: "only children aged 13 years and over may lawfully provide their own consent for the processing of their personal data" (UK regulator guidance, UK GDPR Article 8)
    - es-lopdgdd-art7 [563]: "El tratamiento de los datos personales de un menor de edad únicamente podrá fundarse en su consentimiento cuando sea mayor de catorce años" (Spanish national law (Ley Orgánica 3/2018, art. 7))

- **Verdict: standard**. GB 13 and ES 14 now match official texts; US 13 is the COPPA under-13 rule, and the 16 default is the GDPR Article 8 age. Gap: DK 13, SE 13, IT 14, FR 15, DE 16 and NL 16 were not re-checked against national law; treat them as a lawyer item.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M44. Email-plus parental consent with delayed confirmation

- **Where:** `routes/consent.js:20-26,776`.
- **Evidence:** coppa-312.5-email [49] (a confirmatory email to the parent after consent is an additional step in this method); coppa-312.5a1 [49]; reyes2018 [561]; zhao2019 [564] and kumar2017 [565] (young children have gaps in privacy knowledge and rely on parents, which is why the parent is the consenter).
    - coppa-312.5-email [49]: "Sending a confirmatory email to the parent following receipt of consent" (US federal rule (email plus))
    - coppa-312.5a1 [49]: "An operator is required to obtain verifiable parental consent before any collection , use, or disclosure of personal information from children" (US federal rule, children under 13)
    - reyes2018 [561]: "a majority are potentially in violation of COPPA" (5,855 popular free children's Android apps)
    - zhao2019 [564]: "children could identify and articulate certain privacy risks well, such as information oversharing or revealing real identities online" (29 UK children aged 6-10, 12 focus groups)
    - kumar2017 [565]: "children ages 5-7 had gaps in their knowledge" (18 US families, children 5-11)
- **Verdict: standard**. Email plus is permitted only when data is not disclosed externally; the code ties the outside-AI opt-ins to a signed form instead (consent.js header). The tester phase currently grants under-age testers full access (known caveat in the owner's notes; privacy fix needed before investors).
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M45. Camera and microphone off for unconsented child accounts

- **Where:** `middleware/security.js:78` (default `camera=(self), microphone=(self)`), `routes/consent.js:560-567` (`Permissions-Policy: camera=(), microphone=()` unless `camera_mic` is allowed), `public/js/say-aloud.js:116` fallback text.
- **Evidence:** ico-aadc-8 [560] ("high privacy" by default); coppa-312.2 [366] (an image or voice file is personal information); conijn2022 [566] (camera proctoring raised test anxiety, 1,760 students); andrejevic2019 [567] (school-based facial recognition raises social and governance concerns); bipa [568] (title-level only, Illinois biometric statute).
    - ico-aadc-8 [560]: "Settings must be “high privacy” by default" (UK statutory code)
    - coppa-312.2 [366]: "A photograph, video, or audio file where such file contains a child 's image or voice" (US federal rule, definition)

    - conijn2022 [566]: "students reported higher levels of test anxiety" (1,760 Dutch university students across four survey waves studying online exam proctoring)
    - andrejevic2019 [567]: "school-based facial recognition presents a number of other social challenges and concerns that merit specific attention" (Conceptual analysis, schools)
    - bipa [568]: title-level only.
- **Verdict: standard** plus **design choice**. Code matches. Every camera or mic block has a tap fallback.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M46. On-device only camera and mic

- **Where:** `public/engine/basketball/camera-blocks.js:96-97,116` (`audio: false`, `stream.getTracks().forEach(tr => tr.stop())`), `public/engine/piano/piano-io.js:92,152-157`, `public/js/say-aloud.js:13`.
- **Evidence:** gdpr-art5-1c [309] (data minimisation); coppa-312.2 [366] and the repo doc's reading of FTC FAQ F.5 (on-device processing is not collection, an FAQ interpretation not a court holding); hoepman2014 [569] (title-level only: privacy design strategies); acquisti2015 [570] (notice and consent alone are weak because people are uncertain about consequences, so minimising data beats relying on consent).
    - gdpr-art5-1c [309]: "adequate, relevant and limited to what is necessary in relation to the purposes for which they are processed" (EU regulation)
    - coppa-312.2 [366]: "A photograph, video, or audio file where such file contains a child 's image or voice" (US federal rule, definition)
    - hoepman2014 [569]: title-level only.
    - acquisti2015 [570]: "people's uncertainty about the consequences of privacy-related behaviors and their own preferences over those consequences" (review of privacy behaviour; supports uncertainty about privacy consequences, not that minimising data beats consent)
- **Verdict: standard / design choice**. Code matches the contract: video and audio are not uploaded; the camera captures video only; tracks are stopped. Limit: nothing in this pass proves no frame or clip is sent anywhere (third-party SDK risk is among the failures reyes2018 [561] found in children's apps); a network audit is the missing check.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M47. Retention timers

- **Where:** `routes/consent.js:66-68` (pending 14 days, revoked 30 days, parent deletion immediate; a consent stub keeps dates, method, notice version and an email hash).
- **Evidence:** coppa-312.10 [571] (keep only as long as reasonably necessary); gdpr-art5-1c [309]; slade2013 [310] (consent and use of learner data raise ethical issues).
    - coppa-312.10 [571]: "retain personal information collected online from a child for only as long as is reasonably necessary to fulfill the specific purpose(s) for which the information was collected" (US federal rule)

    - gdpr-art5-1c [309]: "adequate, relevant and limited to what is necessary in relation to the purposes for which they are processed" (EU regulation)
    - slade2013 [310]: "face a number of ethical challenges, including location and interpretation of data" (Higher-education learning analytics)
- **Verdict: standard**. Code matches. The day counts are design choices.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M48. Parent review, download, revoke and delete

- **Where:** `consent.js:827-937` (`/parent/children/:id/data`, `settings`, `opt-ins`, `revoke`, `delete`), `public/js/parent.js`.
- **Evidence:** coppa-312.6 [572] (parent's right to review); gdpr-art7-1 [573] (controller must be able to demonstrate consent; withdrawal handled); pardo2014 [574] (trust, accountability and transparency in learner data).
    - coppa-312.6 [572]: "Right of parent to review personal information provided by a child." (US federal rule)
    - gdpr-art7-1 [573]: "the controller shall be able to demonstrate that the data subject has consented to processing of his or her personal data" (EU regulation)
    - pardo2014 [574]: "The privacy and ethical issues that emerge in this context are tightly interconnected with other aspects such as trust, accountability and transparency." (Learning analytics principles)
- **Verdict: standard**. Code matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M49. Pinned notice version

- **Where:** `consent.js:60-64`, `public/js/parent-consent.js:44-47` (an answer to a notice whose text changed is refused).
- **Evidence:** gdpr-art7-1 [573] (demonstrable consent to the words shown); coppa-312.5a1 [49] (consent to the practices in the notice); drachsler2016 [311] (existing entry: trusted learning-analytics checklist, title-level).
    - gdpr-art7-1 [573]: "the controller shall be able to demonstrate that the data subject has consented to processing of his or her personal data" (EU regulation)
    - coppa-312.5a1 [49]: "An operator is required to obtain verifiable parental consent before any collection , use, or disclosure of personal information from children" (US federal rule, children under 13)
    - drachsler2016 [311]: "Privacy and analytics: it's a DELICATE issue a checklist for trusted learning analytics"
- **Verdict: standard**. Code matches.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

05-M50. Trust cues: AI disclosure and coach limits

- **Where:** `session.js:1223,1303` ("Analyzed by the AI tutor (provider)"), `player.js:282` ("The coach asks questions, it will not give you the answer"), `player.js:407-415` (honest "Checking").
- **Evidence:** lee2004 [533] (trust must match reliability); glikson2020 [534] (transparency and reliability shape trust in AI); hoff2015 [575] (trust drives willingness to rely).
    - lee2004 [533]: "Automation is often problematic because people fail to rely upon it appropriately." (Review)
    - glikson2020 [534]: "The success of integrating AI into organizations critically depends on workers' trust in AI technology." (Review of 20 years of research)
    - hoff2015 [575]: "willingness of human operators to rely on automation" (Systematic review (2002-2013))
- **Verdict: moderate**. Code matches: states what the AI does and does not do.
- **Status:** Existing (in the code before 25 October 2026; not changed that day).

Code-versus-evidence gaps found

Status after 25 October 2026: gaps 1 (streak), 5 (touch targets), 6 (toasts) and 7 (Report dialog) were closed on 25 October 2026. So were part of gap 2 (the style copy) and part of gap 4 (the break line and an in-app planner; reminders are still open). Each mechanism's status line has the details.

1. Streak uses UTC dates, not the learner's local day, and has no grace day (M1).
2. Dashboard copy implies learning-style tailoring of formats (M9); `preferTypeIn` is set by a preference quiz (M21).
3. Practice feedback timing follows preference, not evidence (M18).
4. No break guidance and no reminders (M13, M38).
5. Small buttons and chips are 36 px on phones; the 44 px coarse-pointer rule covers only the study strip (M29).
6. Error toasts auto-dismiss in 5.2 s (M33); some errors give no recovery hint (M39).
7. The Report button uses `window.prompt` (M26).
8. Tester-phase access for under-age testers sits outside the consent design (M44); the claim "nothing is uploaded" needs a network audit (M46).

05 proposals (each backed by 3 or more sources; status on each)

Also built, outside the numbered proposals: error toasts that stay until dismissed (flag `UX_TOAST_STAY`), and the Report dialog in place of `window.prompt` (flag `UX_REPORT_DIALOG`).

05-P1. Optional break nudge in the study strip. BUILT (flag BREAK_NUDGE, minutes BREAK_NUDGE_MIN; off for the owner; learners can turn it off on the profile page). After about 25 active minutes without a pause (one setting, `BREAK_NUDGE_MIN`, default 25; the proposal's longer interval for older stages was not built), show a dismissible line: "A short break helps. Pause now?" with Pause as one tap; never block, never touch the streak or XP.

- Sources: albulescu2022 [524] (vigor and fatigue effects d = .36 and .35, performance effect non-significant); biwer2023 [520] (systematic breaks gave mood and efficiency

benefits); sievertsen2016 [199] (20-30 minute break recovered 1.7% of an SD); blasche2018 [525] (a break raised vigor).

- Expected effect: small gain in wellbeing; do not promise a score gain (albulescu2022 [524]).
- Risk: nudges that undermine children's privacy choices (ico-aadc-13 [553], ico-aadc-5 [50]); alerts cost attention (kushlev2016 [552]). Keep it inline, silent, once per stretch, off by a profile setting.
- Files: `public/js/study.js` (`draw`, state), `public/css/style.css` (section 30), `services/study-sessions.js` (`view` could add `stretch_ms` since last pause), a profile setting in `routes/profile.js` and `public/js/profile.js`.

05-P2. Streak on the learner's local day, with one grace day per week. BUILT (flag STREAK_LOCAL_GRACE; services/streak.js; the owner keeps the old rule). Not built: counting a study session or review that is not a finished sitting (needs services/study-sessions.js). Use `zoneFor` from `services/study-sessions.js` in the streak code; allow one missed day without reset; count a qualifying study session or review, not only a finished adaptive session.

- Sources: lally2009 [45] (one miss did not materially affect habit formation); kahneman1979 [498] (loss weighs more than gain); ico-aadc-5 [50] (must be easy to disengage without feeling pressured); kivetz2006 [499].
- Expected effect: fewer abandonments after one missed day; streak numbers that match the learner's calendar.
- Risk: less loss-aversion pull, so possibly lower daily-return rate; this is the trade the ICO text favours for children.
- Files: `services/sessions.js:1604-1613`, `services/review-session.js:256-263`, `services/study-sessions.js` (export a day helper), `public/js/app.js:329-343`, `public/js/dashboard.js:345-349`.

05-P3. "When will you study next?" planner, with reminders only through the parent account for children. PARTLY BUILT (flag UX_NEXT_STUDY_PLANNER): the in-app planner on the dashboard, which sends nothing at any age. DEFERRED: emailed or texted reminders and the parent digest (mailer, consent opt-ins, COPPA review). At End session, offer a date-and-time choice written to the calendar; for adult accounts offer an opt-in email reminder; for under-13s send nothing to the child, and offer the parent an opt-in digest.

- Sources: milkman2011 [549] (+4.2 points with a date-and-time prompt); castleman2015 [550] (text reminders raised enrollment for less-supported students); lally2009 [45] (consistent context); gollwitzer2006 [551] (title-level, meta-analysis); coppa-312.5a1 [49] (consent for any collection or contact).
- Expected effect: moderate for return rate; unknown for learning.
- Risk: attention cost of alerts (kushlev2016 [552]); COPPA and ICO nudge standards; email deliverability.
- Files: `public/js/enforcer.js` (end screen), `public/js/calendar.js`, `routes/retention.js`, `services/enforcers.js`, `utils/mailer.js`, `routes/consent.js` (parent opt-ins), `public/js/parent.js`.

05-P4. 44 px targets on every coarse pointer. BUILT (flag UX_TOUCH44). Extend the rule at `style.css:2037` to `.btn-sm`, `.chip-btn`, `.confidence-row .chip-btn`, `.verdict-actions .btn-sm`, `.back-link`, `.btn-text`.

- Sources: wcag-2.5.8 [48] (24 CSS-pixel minimum), wcag-2.5.5 [535] (44 CSS-pixel enhanced target), parhi2006 [536] (about 9.2 mm is enough for one thumb; 44 px is margin), vatavu2015 [537] (title-level, children).
- Expected effect: fewer mis-taps on phones; no learning effect claimed.
- Risk: layout reflow in the player action row; check at 320 px (wcag-1.4.10 [547]).
- Files: `public/css/style.css` (lines 464, 562, 973, 2037), kids block at 1709.

05-P5. Let learners (or parents) hide XP and streak. BUILT (flag UX_HIDE_STATS). A profile switch "Show my streak and XP".

- Sources: patall2008 [46] (choice raises motivation and effort); ryan2000 [47] (autonomy as a basic need); deci1999 [44] (expected rewards undermined intrinsic motivation); sailer2017 [506] (element-specific effects); ico-aadc-5 [50].
- Expected effect: protects intrinsic motivation for learners harmed by tracking; neutral for others.
- Risk: lower engagement in learners who like the numbers.
- Files: `public/js/app.js:329-344`, `public/js/profile.js`, `routes/profile.js`, user record field, `public/js/dashboard.js:345`, outcome cards (`session.js:1214-1215`, `review.js:254-255`).

05-P6. Make the Retention-check and preference copy honest and ground the format default in retrieval, not style. COPY BUILT (flag UX_STYLE_COPY: dashboard, placement page, quiz intro). DEFERRED: making type-in the default independent of the quiz (services/learner.js, services/sessions.js belong to another implementer). Reword `dashboard.js:354` and `style-quiz.js:87`; make type-in the default where the skill has a reliable grader and feedback follows, independent of the quiz.

- Sources: rowland2014 [10] (recall over recognition); kang2007 [166] (short answer performed best with feedback in its second experiment); mcdermott2014 [155] (grades 7-12); pashler2008 [31], rogowsky2015 [317], rogowsky2020 [517], willingham2015 [518], cuevas2015 [519] (preference matching unsupported). Counter-evidence: smith2014 [116] found little or no advantage over multiple choice in three experiments, but an advantage in a fourth when retrieval success was improved.
- Expected effect: small; mostly removes an unsupported claim.
- Risk: typing friction on phones, grading false negatives, accessibility for learners who cannot type easily.
- Files: `services/learner.js:286`, `services/sessions.js:81-110`, `public/js/dashboard.js:354`, `public/js/style-quiz.js:87`, `public/js/profile.js:153`.

05-P7. A graded hint ladder with a try-first step before the coach opens. DEFERRED: needs routes/learn.js, services/ai.js and services/mastery.js (the XP-with-help line), which implementer D does not own. One free attempt or 20 seconds, then a nudge, then the coach, with a visible "why XP is lower with help" line.

- Sources: kapur2008 [531] (struggle before support); aleven2003 [173] (help is under- and mis-used); aleven2016help [34] (feedback on help seeking changed behaviour, not outcomes); bastani2025 [38] (guardrails matter); kalyuga2003 [6] (fade help for more expert learners).
- Expected effect: moderate for transfer; honest caveat from aleven2016help [34] that outcomes may not move.
- Risk: mismatched support for learners at different expertise levels (kalyuga2003 [6] cuts both ways).
- Files: `public/js/player.js:278-398`, `routes/learn.js:787-830`, `services/ai.js:1202`, `services/mastery.js:679`.

What was not done

- No code changed while this section was written; the implementation phase came after (see each status line). No browser test of the 320 px reflow, the dark-theme ratios in use, or screen reader output; contrast was computed from hex tokens only.
- Country consent ages, BIPA coverage of pose and pitch data, and the COPPA 2025 amendments were taken from `KIDS-CAMERA-MIC-PRIVACY.md` and not re-derived; they are for a lawyer.
- The FTC COPPA FAQ, the Federal Register notice, BIPA and UN General comment 25 are cited by official URL without a quotable passage.

Open items

Every proposal that was not built is listed below with the reason, then the known findings that are still open. Each proposal's full text, sources and files are in its section.

Proposals not built

Proposal	What	Reason
01-P2	Fade worked examples, and add self-explanation prompts	**Owner decision / build:** it changes the lesson format (the lesson schema and prompt in `services/ai.js`, and the lesson render in `public*/js/session.js`).
01-P3 (part)	Extend a near-mark checkpoint by a few questions	**Build:** the quiz client must accept extra questions mid-quiz (`routes/learn.js`, `public*/js/session.js`). The pass decision is ready for it.
01-P4	Calibrate item difficulty and placement from live responses	**Needs data:** answer volume per item.
01-P5 (part)	Move part of the end-of-session recall to the next session	**Overlap:** covered by the retention change 02-P3 (`ENFORCER_DEFER_SITTING_RETENTION`).
01-P7	Nudge toward correctives before a retry	**Build:** UI only (`public*/js/session.js`). No server-only version reaches the learner.
01-P8	3-option multiple choice for practice items	**Build:** the gate is `ai.validateItem` in `services/ai.js`. Objective items already allow 3 options (03-P11).
01-P9	Close the loop on challenge shares	**Needs data:** logged real sessions by role.
02-P2	Choose the forgetting-curve shape on real memory tests	**Needs data:** enough real memory tests (MIN_FIT_TESTS).
03-P7 (run)	Run the split-half reliability check on real learners	**Needs data:** 20 or more learners with 30 or more levels. The check itself is built.
03-P9	Test, then match, the directive caps (6 vs 8)	**Needs data:** model runs with per-directive compliance scoring.
03-P11 (part)	Drop the weakest distractor in the build script	**Build:** today's objective-item files are all 4-option.
04-P4	Model and anchor clips for "film it" and piano self-rating	**Needs media:** licensed or consented clips, or drawn animations.
04-P5 (data)	Scene-set reliability gate	**Built, but waiting on data:** BB1-BB7 stay in shadow until 100 or more learners reach 0.8.
04-P6	Validate camera make/miss and pose per device	**Needs data:** paired hand-logged and camera counts.
04-P8	Run carry-over maintenance as an experiment	**Needs data:** real retention sittings, plus the arm in `services/personalize.js`.
04-P10	First-person decision scenes and response-time bands	**Needs media and data.**
04-P11	A structured neuromuscular warm-up template	**Needs sourced content:** the sources show such programmes work, but do not give the exercises.

Proposal	What	Reason
05-P2 (part)	Count a study session or review toward the streak	**Build:** needs `services/study-sessions.js`.
05-P3 (part)	Emailed or texted reminders, and the parent digest	**Owner decision and legal review:** mailer, consent opt-ins and a lawyer's view on COPPA.
05-P6 (part)	Make type-in the default, independent of the style quiz	**Build:** `services/learner.js` and `services/sessions.js` belong to another implementer.
05-P7	A graded hint ladder with a try-first step	**Build:** needs `routes/learn.js`, `services/ai.js` and `services/mastery.js`.
02 (considered)	Confusable maths concepts side by side	**Needs data:** confusability data the bank does not have.

Known findings still open

- **Reading passages for grades 5-8 are above the CCSS text bands.** The bank's passages are ACT-level prose. CCSS sets text-complexity bands for grades 4-5 and 6-8 (RL/RI.5.10 and 8.10; Appendix A: https://www.thecorestandards.org/assets/Appendix_A.pdf). The fix is grade-band passages written for those grades. The skills themselves fit (`GRADE-FIT-STANDARDS.md`, "Not fixed").
- **The grammar sentences use adult vocabulary** ("crowbar", "kiln", "contract", "bookkeeper"). The rule each one tests now fits the grade, but a first-grader will struggle to decode the words.
- **Some grade pools are small.** Examples: grade 2 comma use (series commas only), grade 4 academic writing (5 items), and grades 2-3 sentence structure (complete sentences only). These repeat sooner instead of reaching a later grade's rule. More items at grade level would help.
- **The daily review's `today()` costs about a second or more per call:** 1.3 to 2.1 s after 25 October 2026, against 1.2 to 1.9 s before. The first version of the item-relative rapid line (03-P8) pushed it to 11.7 s. That was fixed, and a cost-bound test now guards it (`scripts/scorecard-evidence-test.js`). The remaining cost is still worth reducing.
- **"Unsure" placement skills are not shown yet** in `routes/placement.js` (01-P6).
- **Legal items for a lawyer:** the country consent ages, whether BIPA covers pose and pitch data, the COPPA 2025 amendments, and the claim that "nothing is uploaded", which needs a network audit (05-M43, 05-M46).
- **Under-age testers have full access during the test launch.** This sits outside the consent design (05-M44). The privacy fix, a local AI server, is needed before investors.
- **Also from the standards check:** logs (Algebra 2 B) uses very large arguments. Grade 5's sheet introduces square roots early. The Spanish PCIC levels and the ACTFL descriptors could not be re-read.
- **Flaky tests that were there before 25 October 2026:** `self-test` "a real sitting repeats less than 1 question in 4" (about one run in ten), and the English basics fallback check (about 3%).
- **Restart:** the app on 6767 must be restarted to load every change from 25 October 2026.

References

Every distinct source, numbered in order of first citation. **Verified** says how it was checked: "DOI + quote" (DOI resolved, metadata matched, and a verbatim quote found in the abstract or full text), "URL + quote" (an official page, with a verbatim quote found on it), "title-level only" (the DOI or URL is real, but no text could be fetched, so it never counts toward a mechanism's three sources) or "not verified". Author lists come from Crossref where there is a DOI.

1. Roediger HL, Karpicke JD (2006). Test-Enhanced Learning: Taking Memory Tests Improves Long-Term Retention. *Psychological Science*. https://doi.org/10.1111/j.1467-9280.2006.01693.x. **Verified:** DOI + quote. Ids: roediger2006. Cited in: 01, 02, 04, 05.
2. Owen RJ (1975). A Bayesian Sequential Procedure for Quantal Response in the Context of Adaptive Mental Testing. *Journal of the American Statistical Association*. https://doi.org/10.1080/01621459.1975.10479871. **Verified:** DOI + quote. Ids: owen1975. Cited in: 01.
3. Weiss DJ, Kingsbury GG (1984). Application Of Computerized Adaptive Testing To Educational Problems. *Journal of Educational Measurement*. https://doi.org/10.1111/j.1745-3984.1984.tb01040.x. **Verified:** DOI + quote. Ids: weiss1984. Cited in: 01.
4. Kulik CLC, Kulik JA, Bangert-Drowns RL (1990). Effectiveness of Mastery Learning Programs: A Meta-Analysis. *Review of Educational Research*. https://doi.org/10.3102/00346543060002265 (also https://doi.org/10.2307/1170612). **Verified:** DOI + quote. Ids: kulik1990. Cited in: 01, 03, 04.
5. Bloom BS (1984). The 2 Sigma Problem: The Search for Methods of Group Instruction as Effective as One-to-One Tutoring. *Educational Researcher*. https://doi.org/10.3102/0013189x013006004 (also https://doi.org/10.2307/1175554). **Verified:** DOI + quote. Ids: bloom1984. Cited in: 01, 03, 04.
6. Kalyuga S, Ayres P, Chandler P, Sweller J (2003). The Expertise Reversal Effect. *Educational Psychologist*. https://doi.org/10.1207/s15326985ep3801_4. **Verified:** DOI + quote. Ids: kalyuga2003. Cited in: 01, 03, 04, 05.
7. Sweller J, Cooper GA (1985). The Use of Worked Examples as a Substitute for Problem Solving in Learning Algebra. *Cognition and Instruction*. https://doi.org/10.1207/s1532690xci0201_3. **Verified:** DOI + quote. Ids: sweller1985. Cited in: 01.
8. Renkl A (2014). Toward an Instructionally Oriented Theory of Example-Based Learning. *Cognitive Science*. https://doi.org/10.1111/cogs.12086. **Verified:** DOI + quote. Ids: renkl2014. Cited in: 01, 03.
9. Atkinson RK, Derry SJ, Renkl A, Wortham D (2000). Learning from Examples: Instructional Principles from the Worked Examples Research. *Review of Educational Research*. https://doi.org/10.2307/1170661. **Verified:** DOI + quote. Ids: atkinson2000. Cited in: 01.
10. Rowland CA (2014). The effect of testing versus restudy on retention: A meta-analytic review of the testing effect. *Psychological Bulletin*. https://doi.org/10.1037/a0037559. **Verified:** DOI + quote. Ids: rowland2014. Cited in: 01, 02, 03, 04, 05.

11. Adesope OO, Trevisan DA, Sundararajan N (2017). Rethinking the Use of Tests: A Meta-Analysis of Practice Testing. *Review of Educational Research*. https://doi.org/10.3102/0034654316689306. **Verified:** DOI + quote. Ids: adesope2017. Cited in: 01, 02, 03.

12. Rohrer D, Dedrick RF, Stershic S (2015). Interleaved practice improves mathematics learning. *Journal of Educational Psychology*. https://doi.org/10.1037/edu0000001. **Verified:** DOI + quote. Ids: rohrer2015. Cited in: 01, 02, 03, 04.

13. Rohrer D, Dedrick RF, Hartwig MK, Cheung CN (2020). A randomized controlled trial of interleaved mathematics practice. *Journal of Educational Psychology*. https://doi.org/10.1037/edu0000367. **Verified:** DOI + quote. Ids: rohrer2020. Cited in: 01, 02.

14. Brunmair M, Richter T (2019). Similarity matters: A meta-analysis of interleaved learning and its moderators. *Psychological Bulletin*. https://doi.org/10.1037/bul0000209. **Verified:** DOI + quote. Ids: brunmair2019. Cited in: 01, 02, 03.

15. Van der Kleij FM, Feskens RCW, Eggen TJHM (2015). Effects of Feedback in a Computer-Based Learning Environment on Students' Learning Outcomes: A Meta-Analysis. *Review of Educational Research*. https://doi.org/10.3102/0034654314564881. **Verified:** DOI + quote. Ids: vanderkleij2015. Cited in: 01, 02, 05.

16. Shute VJ (2008). Focus on Formative Feedback. *Review of Educational Research*. https://doi.org/10.3102/0034654307313795. **Verified:** DOI + quote. Ids: shute2008. Cited in: 01, 03, 04, 05.

17. Butler AC, Godbole N, Marsh EJ (2013). Explanation feedback is better than correct answer feedback for promoting transfer of learning. *Journal of Educational Psychology*. https://doi.org/10.1037/a0031026. **Verified:** DOI + quote. Ids: butler2013. Cited in: 01, 02, 05.

18. Lewis C, Sheehan K (1990). Using Bayesian Decision Theory to Design a Computerized Mastery Test. *Applied Psychological Measurement*. https://doi.org/10.1177/014662169001400404. **Verified:** DOI + quote. Ids: lewis1990. Cited in: 01.

19. Spray JA, Reckase MD (1996). Comparison of SPRT and Sequential Bayes Procedures for Classifying Examinees Into Two Categories Using a Computerized Test. *Journal of Educational and Behavioral Statistics*. https://doi.org/10.3102/10769986021004405. **Verified:** DOI + quote. Ids: spray1996. Cited in: 01.

20. Karpicke JD, Roediger HL (2008). The Critical Importance of Retrieval for Learning. *Science*. https://doi.org/10.1126/science.1152408. **Verified:** DOI + quote. Ids: karpicke2008. Cited in: 02.

21. Yang C, Luo L, Vadillo MA, Yu R, Shanks DR (2021). Testing (quizzing) boosts classroom learning: A systematic and meta-analytic review. *Psychological Bulletin*. https://doi.org/10.1037/bul0000309. **Verified:** DOI + quote. Ids: yang2021. Cited in: 01, 02, 03.

22. Cepeda NJ, Pashler H, Vul E, Wixted JT, Rohrer D (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*.

https://doi.org/10.1037/0033-2909.132.3.354. **Verified:** DOI + quote. Ids: cepeda2006. Cited in: 01, 02, 05.

23. Cepeda NJ, Vul E, Rohrer D, Wixted JT, Pashler H (2008). Spacing Effects in Learning: A Temporal Ridgeline of Optimal Retention. *Psychological Science*. https://doi.org/10.1111/j.1467-9280.2008.02209.x. **Verified:** DOI + quote. Ids: cepeda2008. Cited in: 01, 02, 04.
24. Rawson KA, Dunlosky J (2011). Optimizing schedules of retrieval practice for durable and efficient learning: How much is enough? *Journal of Experimental Psychology: General*. https://doi.org/10.1037/a0023956. **Verified:** DOI + quote. Ids: rawson2011. Cited in: 01, 02, 04.
25. Settles B, Meeder B (2016). A Trainable Spaced Repetition Model for Language Learning. *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. https://doi.org/10.18653/v1/p16-1174. **Verified:** DOI + quote. Ids: settles2016. Cited in: 01, 02.
26. Tabibian B, Upadhyay U, De A, Zarezade A, Schölkopf B, Gomez-Rodriguez M (2019). Enhancing human learning via spaced repetition optimization. *Proceedings of the National Academy of Sciences*. https://doi.org/10.1073/pnas.1815156116. **Verified:** DOI + quote. Ids: tabibian2019. Cited in: 01, 02.
27. Lindsey RV, Shroyer JD, Pashler H, Mozer MC (2014). Improving Students' Long-Term Knowledge Retention Through Personalized Review. *Psychological Science*. https://doi.org/10.1177/0956797613504302. **Verified:** DOI + quote. Ids: lindsey2014. Cited in: 01, 02.
28. Eglington LG, Pavlik Jr PI (2020). Optimizing practice scheduling requires quantitative tracking of individual item performance. *npj Science of Learning*. https://doi.org/10.1038/s41539-020-00074-4. **Verified:** DOI + quote. Ids: eglington2020. Cited in: 01, 02, 03.
29. Efron B, Morris C (1975). Data Analysis Using Stein's Estimator and its Generalizations. *Journal of the American Statistical Association*. https://doi.org/10.1080/01621459.1975.10479864. **Verified:** DOI + quote. Ids: efron1975. Cited in: 03.
30. Gelman A, Hill J, Yajima M (2012). Why We (Usually) Don't Have to Worry About Multiple Comparisons. *Journal of Research on Educational Effectiveness*. https://doi.org/10.1080/19345747.2011.618213. **Verified:** DOI + quote. Ids: gelman2012. Cited in: 03.
31. Pashler H, McDaniel M, Rohrer D, Bjork R (2008). Learning Styles: Concepts and Evidence. *Psychological Science in the Public Interest*. https://doi.org/10.1111/j.1539-6053.2009.01038.x. **Verified:** DOI + quote. Ids: pashler2008. Cited in: 02, 03, 05.
32. Jansen BR, Louwerse J, Straatemeier M, Van der Ven SH, Klinkenberg S, Van der Maas HL (2013). The influence of experiencing success in math on math anxiety, perceived math competence, and math performance. *Learning and Individual Differences*. https://doi.org/10.1016/j.lindif.2012.12.014. **Verified:** DOI + quote. Ids: jansen2013. Cited in: 03.

33. Koedinger KR, Aleven V (2007). Exploring the Assistance Dilemma in Experiments with Cognitive Tutors. *Educational Psychology Review*. https://doi.org/10.1007/s10648-007-9049-0. **Verified:** DOI + quote. Ids: koedinger2007. Cited in: 03.

34. Aleven V, Roll I, McLaren BM, Koedinger KR (2016). Help Helps, But Only So Much: Research on Help Seeking with Intelligent Tutoring Systems. *International Journal of Artificial Intelligence in Education*. https://doi.org/10.1007/s40593-015-0089-1. **Verified:** DOI + quote. Ids: aleven2016help. Cited in: 01, 03, 05.

35. Wise SL, DeMars CE (2006). An Application of Item Response Time: The Effort-Moderated IRT Model. *Journal of Educational Measurement*. https://doi.org/10.1111/j.1745-3984.2006.00002.x. **Verified:** DOI + quote. Ids: wise2006. Cited in: 01, 02, 03.

36. VanLEHN K (2011). The Relative Effectiveness of Human Tutoring, Intelligent Tutoring Systems, and Other Tutoring Systems. *Educational Psychologist*. https://doi.org/10.1080/00461520.2011.611369. **Verified:** DOI + quote. Ids: vanlehn2011. Cited in: 03, 05.

37. Kulik JA, Fletcher JD (2016). Effectiveness of Intelligent Tutoring Systems: A Meta-Analytic Review. *Review of Educational Research*. https://doi.org/10.3102/0034654315581420. **Verified:** DOI + quote. Ids: kulik2016. Cited in: 03, 05.

38. Bastani H, Bastani O, Sungu A, Ge H, Kabakcı Ö, Mariman R (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*. https://doi.org/10.1073/pnas.2422633122. **Verified:** DOI + quote. Ids: bastani2025. Cited in: 01, 03, 05.

39. Soderstrom NC, Bjork RA (2015). Learning Versus Performance. *Perspectives on Psychological Science*. https://doi.org/10.1177/1745691615569000. **Verified:** DOI + quote. Ids: soderstrom2015. Cited in: 04.

40. Ericsson KA, Krampe RT, Tesch-Römer C (1993). The role of deliberate practice in the acquisition of expert performance. *Psychological Review*. https://doi.org/10.1037/0033-295x.100.3.363. **Verified:** DOI + quote. Ids: ericsson1993. Cited in: 04.

41. Guadagnoli MA, Lee TD (2004). Challenge Point: A Framework for Conceptualizing the Effects of Various Practice Conditions in Motor Learning. *Journal of Motor Behavior*. https://doi.org/10.3200/jmbr.36.2.212-224. **Verified:** DOI + quote. Ids: guadagnoli2004. Cited in: 04.

42. Shea JB, Morgan RL (1979). Contextual interference effects on the acquisition, retention, and transfer of a motor skill. *Journal of Experimental Psychology: Human Learning and Memory*. https://doi.org/10.1037/0278-7393.5.2.179. **Verified:** DOI + quote. Ids: shea1979. Cited in: 04.

43. Sailer M, Homner L (2020). The Gamification of Learning: a Meta-analysis. *Educational Psychology Review*. https://doi.org/10.1007/s10648-019-09498-w. **Verified:** DOI + quote. Ids: sailer2020. Cited in: 04, 05.

44. Deci EL, Koestner R, Ryan RM (1999). A meta-analytic review of experiments examining the effects of extrinsic rewards on intrinsic motivation. *Psychological Bulletin*. https://doi.org/10.1037/0033-2909.125.6.627. **Verified:** DOI + quote. Ids: deci1999. Cited in: 05.

45. Lally P, van Jaarsveld CHM, Potts HWW, Wardle J (2009). How are habits formed: Modelling habit formation in the real world. *European Journal of Social Psychology*. https://doi.org/10.1002/ejsp.674. **Verified:** DOI + quote. Ids: lally2009. Cited in: 05.

46. Patall EA, Cooper H, Robinson JC (2008). The effects of choice on intrinsic motivation and related outcomes: A meta-analysis of research findings. *Psychological Bulletin*. https://doi.org/10.1037/0033-2909.134.2.270. **Verified:** DOI + quote. Ids: patall2008. Cited in: 02, 05.

47. Ryan RM, Deci EL (2000). Self-determination theory and the facilitation of intrinsic motivation, social development, and well-being. *American Psychologist*. https://doi.org/10.1037/0003-066x.55.1.68. **Verified:** DOI + quote. Ids: ryan2000. Cited in: 02, 05.

48. W3C (2023). WCAG 2.2 SC 2.5.8 Target Size (Minimum), Level AA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#target-size-minimum. **Verified:** URL + quote. Ids: wcag-2.5.8. Cited in: 05.

49. FTC (2025). 16 CFR 312.5 (COPPA Rule): parental consent, including the email-plus method. Official standard or regulation. https://www.ecfr.gov/current/title-16/part-312/section-312.5. **Verified:** URL + quote. Ids: coppa-312.5a1, coppa-312.5-email. Cited in: 05.

50. ICO (2020). ICO Age appropriate design code, standard 5: Detrimental use of data. Official standard or regulation. https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/childrens-information/childrens-code-guidance-and-resources/age-appropriate-design-a-code-of-practice-for-online-services/5-detrimental-use-of-data/. **Verified:** URL + quote. Ids: ico-aadc-5, ico-aadc-5b. Cited in: 05.

51. University of MINNESOTA, Babcock B, Weiss D (2012). Termination Criteria in Computerized Adaptive Tests: Do Variable - Length CATs Provide Efficient and Effective Measurement? *Journal of Computerized Adaptive Testing*. https://doi.org/10.7333/1212-0101001. **Verified:** DOI + quote. Ids: babcock2012. Cited in: 01.

52. Ackerman PL, Kanfer R (2009). Test length and cognitive fatigue: An empirical examination of effects on performance and test-taker reactions. *Journal of Experimental Psychology: Applied*. https://doi.org/10.1037/a0015719. **Verified:** DOI + quote. Ids: ackerman2009. Cited in: 01, 05.

53. Debeer D, Buchholz J, Hartig J, Janssen R (2014). Student, School, and Country Differences in Sustained Test-Taking Effort in the 2009 PISA Reading Assessment. *Journal of Educational and Behavioral Statistics*. https://doi.org/10.3102/1076998614558485. **Verified:** DOI + quote. Ids: debeer2014. Cited in: 01.

54. Borgonovi F, Biecek P (2016). An international comparison of students' ability to endure fatigue and maintain motivation during a low-stakes test. *Learning and Individual Differences*. https://doi.org/10.1016/j.lindif.2016.06.001. **Verified:** DOI + quote. Ids: borgonovi2016. Cited in: 01, 05.

55. Wauters K, Desmet P, Van den Noortgate W (2010). Adaptive item-based learning environments based on the item response theory: possibilities and challenges. *Journal of Computer Assisted Learning*. https://doi.org/10.1111/j.1365-2729.2010.00368.x. **Verified:** DOI + quote. Ids: wauters2010. Cited in: 01.

56. Hartig & Buchholz (2012). A multilevel item response model for item position effects and individual persistence. *Psychological test and assessment modeling*. https://openalex.org/W2117924135. **Verified:** URL + quote. Ids: hartig2012. Cited in: 01.

57. Debeer D, Janssen R (2013). Modeling Item-Position Effects Within an IRT Framework. *Journal of Educational Measurement*. https://doi.org/10.1111/jedm.12009. **Verified:** DOI + quote. Ids: debeer2013. Cited in: 01.

58. Wainer H, Kiely GL (1987). Item Clusters and Computerized Adaptive Testing: A Case for Testlets. *Journal of Educational Measurement*. https://doi.org/10.1111/j.1745-3984.1987.tb00274.x. **Verified:** DOI + quote. Ids: wainer1987. Cited in: 01.

59. Wise SL (2017). Rapid-Guessing Behavior: Its Identification, Interpretation, and Implications. *Educational Measurement: Issues and Practice*. https://doi.org/10.1111/emip.12165. **Verified:** DOI + quote. Ids: wise2017. Cited in: 01, 02, 03.

60. Guskey TR (2007). Closing Achievement Gaps: Revisiting Benjamin S. Bloom's "Learning for Mastery". *Journal of Advanced Academics*. https://doi.org/10.4219/jaa-2007-704. **Verified:** DOI + quote. Ids: guskey2007. Cited in: 01, 04.

61. Tetzlaff L, Simonsmeier B, Peters T, Brod G (2025). A cornerstone of adaptivity – A meta-analysis of the expertise reversal effect. *Learning and Instruction*. https://doi.org/10.1016/j.learninstruc.2025.102142. **Verified:** DOI + quote. Ids: tetzlaff2025. Cited in: 01, 03.

62. Pan SC, Carpenter SK (2023). Prequestioning and Pretesting Effects: a Review of Empirical Research, Theoretical Perspectives, and Implications for Educational Practice. *Educational Psychology Review*. https://doi.org/10.1007/s10648-023-09814-5. **Verified:** DOI + quote. Ids: pan2023pretest. Cited in: 01.

63. Morris CN (1983). Parametric Empirical Bayes Inference: Theory and Applications. *Journal of the American Statistical Association*. https://doi.org/10.1080/01621459.1983.10477920. **Verified:** DOI + quote. Ids: morris1983. Cited in: 01, 03.

64. Casella G (1985). An Introduction to Empirical Bayes Data Analysis. *The American Statistician*. https://doi.org/10.1080/00031305.1985.10479400. **Verified:** DOI + quote. Ids: casella1985. Cited in: 01.

65. Bock RD, Mislevy RJ (1982). Adaptive EAP Estimation of Ability in a Microcomputer Environment. *Applied Psychological Measurement*. https://doi.org/10.1177/014662168200600405. **Verified:** DOI + quote. Ids: bock1982. Cited in: 01.

66. Parsons S, Kruijt AW, Fox E (2019). Psychological Science Needs a Standard Practice of Reporting the Reliability of Cognitive-Behavioral Measurements. *Advances in Methods and Practices in Psychological Science*. https://doi.org/10.1177/2515245919879695. **Verified:** DOI + quote. Ids: parsons2019. Cited in: 01, 03.

67. Hedge C, Powell G, Sumner P (2018). The reliability paradox: Why robust cognitive tasks do not produce reliable individual differences. *Behavior Research Methods*. https://doi.org/10.3758/s13428-017-0935-1. **Verified:** DOI + quote. Ids: hedge2018. Cited in: 01, 03.

68. Brown J, Burton R (1978). Diagnostic models for procedural bugs in basic mathematical skills. *Cognitive Science*. https://doi.org/10.1016/s0364-0213(78)80004-4. **Verified:** DOI + quote. Ids: brown1978. Cited in: 01.

69. Kalyuga S, Sweller J (2004). Measuring Knowledge to Optimize Cognitive Load Factors During Instruction. *Journal of Educational Psychology*. https://doi.org/10.1037/0022-0663.96.3.558. **Verified:** DOI + quote. Ids: kalyuga2004. Cited in: 01.

70. Barton MA, Lord FM (1981). An Upper Asymptote For The Three-parameter Logistic Item-response Model*. *ETS Research Report Series*. https://doi.org/10.1002/j.2333-8504.1981.tb01255.x. **Verified:** DOI + quote. Ids: barton1981. Cited in: 01.

71. Baker RSJD, Corbett AT, Aleven V (2008). More Accurate Student Modeling through Contextual Estimation of Slip and Guess Probabilities in Bayesian Knowledge Tracing. *Lecture Notes in Computer Science*. https://doi.org/10.1007/978-3-540-69132-7_44. **Verified:** DOI + quote. Ids: baker2008. Cited in: 01.

72. van der Linden WJ (1998). Bayesian Item Selection Criteria for Adaptive Testing. *Psychometrika*. https://doi.org/10.1007/bf02294775. **Verified:** DOI + quote. Ids: vanderlinden1998. Cited in: 01.

73. Eggen TJHM (1999). Item Selection in Adaptive Testing with the Sequential Probability Ratio Test. *Applied Psychological Measurement*. https://doi.org/10.1177/01466219922031365. **Verified:** DOI + quote. Ids: eggen1999. Cited in: 01.

74. Wise SL, Kong X (2005). Response Time Effort: A New Measure of Examinee Motivation in Computer-Based Tests. *Applied Measurement in Education*. https://doi.org/10.1207/s15324818ame1802_2. **Verified:** DOI + quote. Ids: wise2005. Cited in: 01, 02, 03, 05.

75. Käser T, Alexandron G (2024). Simulated Learners in Educational Technology: A Systematic Literature Review and a Turing-like Test. *International Journal of Artificial Intelligence in Education*. https://doi.org/10.1007/s40593-023-00337-2. **Verified:** DOI + quote. Ids: kaser2024. Cited in: 01.

76. Meyers JL, Miller GE, Way WD (2008). Item Position and Item Difficulty Change in an IRT-Based Common Item Equating Design. *Applied Measurement in Education*. https://doi.org/10.1080/08957340802558342. **Verified:** DOI + quote. Ids: meyers2008. Cited in: 01.

77. Pelánek R (2016). Applications of the Elo rating system in adaptive educational systems. *Computers & Education*. https://doi.org/10.1016/j.compedu.2016.03.017. **Verified:** title-level only. Ids: pelanek2016. Cited in: 01, 04.

78. Klinkenberg S, Straatemeier M, van der Maas H (2011). Computer adaptive practice of Maths ability using a new item response model for on the fly ability and difficulty estimation. *Computers & Education*. https://doi.org/10.1016/j.compedu.2011.02.003. **Verified:** title-level only. Ids: klinkenberg2011. Cited in: 01, 04.

79. Doroudi S, Aleven V, Brunskill E (2017). Robust Evaluation Matrix: Towards a More Principled Offline Exploration of Instructional Policies. *Proceedings of the Fourth (2017) ACM Conference on Learning @ Scale*. https://doi.org/10.1145/3051457.3051463. **Verified:** DOI + quote. Ids: doroudi2017. Cited in: 01.

80. Kalyuga S, Sweller J (2005). Rapid dynamic assessment of expertise to improve the efficiency of adaptive e-learning. *Educational Technology Research and Development*. https://doi.org/10.1007/bf02504800. **Verified:** DOI + quote. Ids: kalyuga2005. Cited in: 01.

81. van de Sande et al. (2013). Properties of the Bayesian Knowledge Tracing Model. *Journal of Educational Data Mining*. https://doi.org/10.5281/zenodo.3554629. **Verified:** DOI + quote. Ids: vandesande2013. Cited in: 01.

82. Corbett AT, Anderson JR (1995). Knowledge tracing: Modeling the acquisition of procedural knowledge. *User Modelling and User-Adapted Interaction*. https://doi.org/10.1007/bf01099821. **Verified:** title-level only. Ids: corbett1995. Cited in: 01, 03.

83. Doroudi S, Brunskill E (2019). Fairer but Not Fair Enough On the Equitability of Knowledge Tracing. *Proceedings of the 9th International Conference on Learning Analytics & Knowledge*. https://doi.org/10.1145/3303772.3303838. **Verified:** DOI + quote. Ids: doroudi2019fair. Cited in: 01.

84. Kalyuga S (2007). Expertise Reversal Effect and Its Implications for Learner-Tailored Instruction. *Educational Psychology Review*. https://doi.org/10.1007/s10648-007-9054-3. **Verified:** DOI + quote. Ids: kalyuga2007. Cited in: 01, 03.

85. Steenbergen-Hu S, Makel MC, Olszewski-Kubilius P (2016). What One Hundred Years of Research Says About the Effects of Ability Grouping and Acceleration on K–12 Students' Academic Achievement. *Review of Educational Research*. https://doi.org/10.3102/0034654316675417. **Verified:** DOI + quote. Ids: steenbergenhu2016. Cited in: 01.

86. Kulik JA, Kulik CLC (1984). Effects of Accelerated Instruction on Students. *Review of Educational Research*. https://doi.org/10.3102/00346543054003409. **Verified:** DOI + quote. Ids: kulik1984accel. Cited in: 01.

87. Major L, Francis GA, Tsapali M (2021). The effectiveness of technology-supported personalised learning in low- and middle-income countries: A meta-analysis. *British Journal of Educational Technology*. https://doi.org/10.1111/bjet.13116. **Verified:** DOI + quote. Ids: major2021. Cited in: 01, 03.

88. VanLehn K (1990). Mind Bugs: The Origins of Procedural Misconceptions. *MIT Press*. https://openalex.org/W1564397683 (catalogue record carrying the publisher's description). **Verified:** URL + quote. Ids: vanlehn1990. Cited in: 01.

89. Koedinger KR, Carvalho PF, Liu R, McLaughlin EA (2023). An astonishing regularity in student learning rate. *Proceedings of the National Academy of Sciences*. https://doi.org/10.1073/pnas.2221311120. **Verified:** DOI + quote. Ids: koedinger2023. Cited in: 01.

90. Bangert-Drowns RL, Kulik JA, Kulik CLC (1991). Effects of Frequent Classroom Testing. *The Journal of Educational Research*. https://doi.org/10.1080/00220671.1991.10702818. **Verified:** DOI + quote. Ids: bangertdrowns1991freq. Cited in: 01.

91. Khanna MM, Brack ASB, Finken LL (2013). Short- and Long-Term Effects of Cumulative Finals on Student Learning. *Teaching of Psychology*.

https://doi.org/10.1177/0098628313487458. **Verified:** DOI + quote. Ids: khanna2013. Cited in: 01, 02.

92. Rohrer D, Taylor K (2007). The shuffling of mathematics problems improves learning. *Instructional Science*. https://doi.org/10.1007/s11251-007-9015-8. **Verified:** DOI + quote. Ids: rohrer2007. Cited in: 01.
93. Taylor K, Rohrer D (2010). The effects of interleaved practice. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.1598. **Verified:** DOI + quote. Ids: taylor2010. Cited in: 01.
94. Wilson RC, Shenhav A, Straccia M, Cohen JD (2019). The Eighty Five Percent Rule for optimal learning. *Nature Communications*. https://doi.org/10.1038/s41467-019-12552-4. **Verified:** DOI + quote. Ids: wilson2019. Cited in: 01, 02, 03.
95. Metcalfe J, Kornell N (2005). A Region of Proximal Learning model of study time allocation. *Journal of Memory and Language*. https://doi.org/10.1016/j.jml.2004.12.001. **Verified:** title-level only. Ids: metcalfe2005. Cited in: 01.
96. Mayfield KH, Chase PN (2002). The Effects Of Cumulative Practice On Mathematics Problem Solving. *Journal of Applied Behavior Analysis*. https://doi.org/10.1901/jaba.2002.35-105. **Verified:** DOI + quote. Ids: mayfield2002. Cited in: 01, 02.
97. Lawrence NK (2013). Cumulative Exams in the Introductory Psychology Course. *Teaching of Psychology*. https://doi.org/10.1177/0098628312465858. **Verified:** DOI + quote. Ids: lawrence2013. Cited in: 01, 02.
98. Szpunar KK, McDermott KB, Roediger HL (2007). Expectation of a final cumulative test enhances long-term retention. *Memory & Cognition*. https://doi.org/10.3758/bf03193473. **Verified:** DOI + quote. Ids: szpunar2007. Cited in: 01, 02.
99. Agarwal PK, Nunes LD, Blunt JR (2021). Retrieval Practice Consistently Benefits Student Learning: a Systematic Review of Applied Research in Schools and Classrooms. *Educational Psychology Review*. https://doi.org/10.1007/s10648-021-09595-9. **Verified:** DOI + quote. Ids: agarwal2021. Cited in: 01, 02, 04.
100. Keller FS (1968). “Good-bye, Teacher …”1. *Journal of Applied Behavior Analysis*. https://doi.org/10.1901/jaba.1968.1-79. **Verified:** title-level only. Ids: keller1968. Cited in: 01.
101. Kulik JA, Kulik CLC, Cohen PA (1979). A meta-analysis of outcome studies of Keller's personalized system of instruction. *American Psychologist*. https://doi.org/10.1037/0003-066x.34.4.307. **Verified:** title-level only. Ids: kulik1979. Cited in: 01.
102. Graham S, Hebert M, Harris KR (2015). Formative Assessment and Writing. *The Elementary School Journal*. https://doi.org/10.1086/681947. **Verified:** DOI + quote. Ids: graham2015. Cited in: 01.
103. Panadero E, Jonsson A (2013). The use of scoring rubrics for formative assessment purposes revisited: A review. *Educational Research Review*. https://doi.org/10.1016/j.edurev.2013.01.002. **Verified:** DOI + quote. Ids: panadero2013. Cited in: 01.
104. Jonsson A, Svingby G (2007). The use of scoring rubrics: Reliability, validity and educational consequences. *Educational Research Review*.

https://doi.org/10.1016/j.edurev.2007.05.002. **Verified:** title-level only. Ids: jonsson2007. Cited in: 01.

105. Kane MT (2013). Validating the Interpretations and Uses of Test Scores. *Journal of Educational Measurement*. https://doi.org/10.1111/jedm.12000. **Verified:** DOI + quote. Ids: kane2013. Cited in: 01, 03.

106. Renkl A, Atkinson RK (2003). Structuring the Transition From Example Study to Problem Solving in Cognitive Skill Acquisition: A Cognitive Load Perspective. *Educational Psychologist*. https://doi.org/10.1207/s15326985ep3801_3. **Verified:** DOI + quote. Ids: renkl2003. Cited in: 01, 04.

107. Mayer RE, Bove W, Bryman A, Mars R, Tapangco L (1996). When less is more: Meaningful learning from visual and verbal summaries of science textbook lessons. *Journal of Educational Psychology*. https://doi.org/10.1037/0022-0663.88.1.64. **Verified:** DOI + quote. Ids: mayer1996. Cited in: 01, 03.

108. Renkl A, Atkinson RK, Maier UH, Staley R (2002). From Example Study to Problem Solving: Smooth Transitions Help Learning. *The Journal of Experimental Education*. https://doi.org/10.1080/00220970209599510. **Verified:** DOI + quote. Ids: renkl2002. Cited in: 01.

109. Atkinson RK, Renkl A, Merrill MM (2003). Transitioning From Studying Examples to Solving Problems: Effects of Self-Explanation Prompts and Fading Worked-Out Steps. *Journal of Educational Psychology*. https://doi.org/10.1037/0022-0663.95.4.774. **Verified:** DOI + quote. Ids: atkinson2003. Cited in: 01.

110. Salden RJCM, Aleven V, Schwonke R, Renkl A (2010). The expertise reversal effect and worked examples in tutored problem solving. *Instructional Science*. https://doi.org/10.1007/s11251-009-9107-8. **Verified:** DOI + quote. Ids: salden2010. Cited in: 01, 03.

111. van Merrienboer JJG, Kirschner PA, Kester L (2003). Taking the Load Off a Learner's Mind: Instructional Design for Complex Learning. *Educational Psychologist*. https://doi.org/10.1207/s15326985ep3801_2. **Verified:** DOI + quote. Ids: vanmerrienboer2003. Cited in: 01.

112. Pollock E, Chandler P, Sweller J (2002). Assimilating complex information. *Learning and Instruction*. https://doi.org/10.1016/s0959-4752(01)00016-0. **Verified:** title-level only. Ids: pollock2002. Cited in: 01.

113. Pan SC, Rickard TC (2018). Transfer of test-enhanced learning: Meta-analytic review and synthesis. *Psychological Bulletin*. https://doi.org/10.1037/bul0000151. **Verified:** DOI + quote. Ids: pan2018. Cited in: 01, 02.

114. Butler AC, Black-Maier AC, Raley ND, Marsh EJ (2017). Retrieving and applying knowledge to different examples promotes transfer of learning. *Journal of Experimental Psychology: Applied*. https://doi.org/10.1037/xap0000142. **Verified:** DOI + quote. Ids: butler2017. Cited in: 01, 02.

115. Hausknecht JP, Halpert JA, Di Paolo NT, Moriarty Gerrard MO (2007). Retesting in selection: A meta-analysis of coaching and practice effects for tests of cognitive ability. *Journal of Applied Psychology*. https://doi.org/10.1037/0021-9010.92.2.373. **Verified:** DOI + quote. Ids: hausknecht2007. Cited in: 01.

116. Smith MA, Karpicke JD (2014). Retrieval practice with short-answer, multiple-choice, and hybrid tests. *Memory*. https://doi.org/10.1080/09658211.2013.831454. **Verified:** DOI + quote. Ids: smith2014. Cited in: 01, 02, 03, 05.

117. Ausubel DP (1960). The use of advance organizers in the learning and retention of meaningful verbal material. *Journal of Educational Psychology*. https://doi.org/10.1037/h0046669. **Verified:** DOI + quote. Ids: ausubel1960. Cited in: 01.

118. Luiten J, Ames W, Ackerson G (1980). A Meta-analysis of the Effects of Advance Organizers on Learning and Retention. *American Educational Research Journal*. https://doi.org/10.3102/00028312017002211. **Verified:** DOI + quote. Ids: luiten1980. Cited in: 01.

119. Porter AC (2002). Measuring the Content of Instruction: Uses in Research and Practice. *Educational Researcher*. https://doi.org/10.3102/0013189x031007003. **Verified:** DOI + quote. Ids: porter2002. Cited in: 01.

120. Hattan C, Alexander PA, Lupo SM (2023). Leveraging What Students Know to Make Sense of Texts: What the Research Says About Prior Knowledge Activation. *Review of Educational Research*. https://doi.org/10.3102/00346543221148478. **Verified:** DOI + quote. Ids: hattan2023. Cited in: 01.

121. Kapur M (2016). Examining Productive Failure, Productive Success, Unproductive Failure, and Unproductive Success in Learning. *Educational Psychologist*. https://doi.org/10.1080/00461520.2016.1155457. **Verified:** DOI + quote. Ids: kapur2016. Cited in: 01.

122. Pashler H, Cepeda NJ, Wixted JT, Rohrer D (2005). When Does Feedback Facilitate Learning of Words? *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/0278-7393.31.1.3. **Verified:** DOI + quote. Ids: pashler2005. Cited in: 01, 02, 05.

123. Hattie J, Timperley H (2007). The Power of Feedback. *Review of Educational Research*. https://doi.org/10.3102/003465430298487. **Verified:** DOI + quote. Ids: hattie2007. Cited in: 01, 04, 05.

124. Liu NF, Lin K, Hewitt J, Paranjape A, Bevilacqua M, Petroni F, et al. (2024). Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics*. https://doi.org/10.1162/tacl_a_00638. **Verified:** DOI + quote. Ids: liu2024lost. Cited in: 01, 03.

125. Levy M, Jacoby A, Goldberg Y (2024). Same Task, More Tokens: the Impact of Input Length on the Reasoning Performance of Large Language Models. *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. https://doi.org/10.18653/v1/2024.acl-long.818. **Verified:** DOI + quote. Ids: levy2024. Cited in: 01, 03.

126. Shi, Chen et al. (2023). Large Language Models Can Be Easily Distracted by Irrelevant Context. *arXiv (Cornell University)*. https://doi.org/10.48550/arxiv.2302.00093. **Verified:** DOI + quote. Ids: shi2023. Cited in: 01, 03.

127. Harp SF, Mayer RE (1998). How seductive details do their damage: A theory of cognitive interest in science learning. *Journal of Educational Psychology*.

https://doi.org/10.1037/0022-0663.90.3.414. **Verified:** DOI + quote. Ids: harp1998. Cited in: 01, 03, 05.

128. Sundararajan N, Adesope O (2020). Keep it Coherent: A Meta-Analysis of the Seductive Details Effect. *Educational Psychology Review*. https://doi.org/10.1007/s10648-020-09522-4. **Verified:** DOI + quote. Ids: sundararajan2020. Cited in: 01, 03.
129. Schroeder NL, Kucera AC (2022). Refutation Text Facilitates Learning: a Meta-Analysis of Between-Subjects Experiments. *Educational Psychology Review*. https://doi.org/10.1007/s10648-021-09656-z. **Verified:** DOI + quote. Ids: schroeder2022. Cited in: 01, 03.
130. McLaren BM, Adams DM, Mayer RE (2015). Delayed Learning Effects with Erroneous Examples: a Study of Learning Decimals with a Web-Based Tutor. *International Journal of Artificial Intelligence in Education*. https://doi.org/10.1007/s40593-015-0064-x. **Verified:** DOI + quote. Ids: mclaren2015. Cited in: 01.
131. Tippett CD (2010). Refutation Text In Science Education: A Review Of Two Decades Of Research. *International Journal of Science and Mathematics Education*. https://doi.org/10.1007/s10763-010-9203-x. **Verified:** DOI + quote. Ids: tippett2010. Cited in: 01, 03.
132. Haladyna TM, Downing SM, Rodriguez MC (2002). A Review of Multiple-Choice Item-Writing Guidelines for Classroom Assessment. *Applied Measurement in Education*. https://doi.org/10.1207/s15324818ame1503_5. **Verified:** DOI + quote. Ids: haladyna2002. Cited in: 01, 03, 04.
133. Little JL, Bjork EL, Bjork RA, Angello G (2012). Multiple-Choice Tests Exonerated, at Least of Some Charges: Fostering Test-Induced Learning and Avoiding Test-Induced Forgetting. *Psychological Science*. https://doi.org/10.1177/0956797612443370. **Verified:** DOI + quote. Ids: little2012. Cited in: 01, 02, 03, 05.
134. Kurdi G, Leo J, Parsia B, Sattler U, Al-Emari S (2020). A Systematic Review of Automatic Question Generation for Educational Purposes. *International Journal of Artificial Intelligence in Education*. https://doi.org/10.1007/s40593-019-00186-y. **Verified:** DOI + quote. Ids: kurdi2020. Cited in: 01.
135. Wang et al. (2022). Self-Consistency Improves Chain of Thought Reasoning in Language Models. *arXiv (ICLR 2023)*. https://doi.org/10.48550/arxiv.2203.11171. **Verified:** DOI + quote. Ids: wang2022sc. Cited in: 01.
136. Rodriguez MC (2005). Three Options Are Optimal for Multiple-Choice Items: A Meta-Analysis of 80 Years of Research. *Educational Measurement: Issues and Practice*. https://doi.org/10.1111/j.1745-3992.2005.00006.x. **Verified:** DOI + quote. Ids: rodriguez2005. Cited in: 01, 03.
137. Rey GD, Beege M, Nebel S, Wirzberger M, Schmitt TH, Schneider S (2019). A Meta-analysis of the Segmenting Effect. *Educational Psychology Review*. https://doi.org/10.1007/s10648-018-9456-4. **Verified:** DOI + quote. Ids: rey2019. Cited in: 01, 03.
138. Noetel M, Griffith S, Delaney O, Harris NR, Sanders T, Parker P, et al. (2022). Multimedia Design for Learning: An Overview of Reviews With Meta-Meta-Analysis. *Review of*

Educational Research. https://doi.org/10.3102/00346543211052329. **Verified:** DOI + quote. Ids: noetel2022. Cited in: 01, 05.

139. Anderson RC (1972). How to Construct Achievement Tests to Assess Comprehension. *Review of Educational Research*. https://doi.org/10.3102/00346543042002145. **Verified:** DOI + quote. Ids: anderson1972. Cited in: 01.

140. Agarwal PK (2019). Retrieval practice & Bloom's taxonomy: Do students need fact knowledge before higher order learning? *Journal of Educational Psychology*. https://doi.org/10.1037/edu0000282. **Verified:** DOI + quote. Ids: agarwal2019. Cited in: 01, 02.

141. Pan SC, Cooke J, Little JL, McDaniel MA, Foster ER, Connor LT, et al. (2019). Online and Clicker Quizzing on Jargon Terms Enhances Definition-Focused but Not Conceptually Focused Biology Exam Performance. *CBE—Life Sciences Education*. https://doi.org/10.1187/cbe.18-12-0248. **Verified:** DOI + quote. Ids: pan2019clicker. Cited in: 01, 02.

142. Ji Z, Lee N, Frieske R, Yu T, Su D, Xu Y, et al. (2023). Survey of Hallucination in Natural Language Generation. *ACM Computing Surveys*. https://doi.org/10.1145/3571730. **Verified:** DOI + quote. Ids: ji2023. Cited in: 01.

143. Huang L, Yu W, Ma W, Zhong W, Feng Z, Wang H, et al. (2025). A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions. *ACM Transactions on Information Systems*. https://doi.org/10.1145/3703155. **Verified:** DOI + quote. Ids: huang2025. Cited in: 01.

144. Dhuliawala S, Komeili M, Xu J, Raileanu R, Li X, Celikyilmaz A, et al. (2024). Chain-of-Verification Reduces Hallucination in Large Language Models. *Findings of the Association for Computational Linguistics ACL 2024*. https://doi.org/10.18653/v1/2024.findings-acl.212. **Verified:** DOI + quote. Ids: dhuliawala2024. Cited in: 01.

145. Min S, Krishna K, Lyu X, Lewis M, Yih WT, Koh P, et al. (2023). FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation. *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. https://doi.org/10.18653/v1/2023.emnlp-main.741. **Verified:** DOI + quote. Ids: min2023. Cited in: 01.

146. Lewis et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *arXiv (NeurIPS 2020)*. https://doi.org/10.48550/arxiv.2005.11401. **Verified:** DOI + quote. Ids: lewis2020. Cited in: 01.

147. Kohavi R, Longbotham R, Sommerfield D, Henne RM (2009). Controlled experiments on the web: survey and practical guide. *Data Mining and Knowledge Discovery*. https://doi.org/10.1007/s10618-008-0114-1. **Verified:** DOI + quote. Ids: kohavi2009. Cited in: 01, 04.

148. Cheung ACK, Slavin RE (2016). How Methodological Features Affect Effect Sizes in Education. *Educational Researcher*. https://doi.org/10.3102/0013189x16656615. **Verified:** DOI + quote. Ids: cheung2016. Cited in: 01.

149. Lortie-Forgues H, Inglis M (2019). Rigorous Large-Scale Educational RCTs Are Often Uninformative: Should We Be Concerned? *Educational Researcher*.

https://doi.org/10.3102/0013189x19832850. **Verified:** DOI + quote. Ids: lortieforgues2019. Cited in: 01.

150. American Educational Research Association; American Psychological Association; National Council on Measurement in Education (2014). Standards for Educational and Psychological Testing (2014 edition). Official standard or regulation. https://www.testingstandards.net/open-access-files.html. **Verified:** title-level only. Ids: aera2014. Cited in: 01, 04.

151. Messick S (1995). Validity of psychological assessment: Validation of inferences from persons' responses and performances as scientific inquiry into score meaning. *American Psychologist*. https://doi.org/10.1037/0003-066x.50.9.741. **Verified:** DOI + quote. Ids: messick1995. Cited in: 01, 03.

152. Walkington CA (2013). Using adaptive learning technologies to personalize instruction to student interests: The impact of relevant contexts on performance and learning outcomes. *Journal of Educational Psychology*. https://doi.org/10.1037/a0031882. **Verified:** DOI + quote. Ids: walkington2013. Cited in: 01, 03.

153. Chan JCK, McDermott KB, Roediger HL (2006). Retrieval-induced facilitation: Initially nontested material can benefit from prior testing of related material. *Journal of Experimental Psychology: General*. https://doi.org/10.1037/0096-3445.135.4.553. **Verified:** DOI + quote. Ids: chan2006. Cited in: 01, 02.

154. McDaniel MA, Anderson JL, Derbish MH, Morrisette N (2007). Testing the testing effect in the classroom. *European Journal of Cognitive Psychology*. https://doi.org/10.1080/09541440701326154. **Verified:** DOI + quote. Ids: mcdaniel2007. Cited in: 01.

155. McDermott KB, Agarwal PK, D'Antonio L, Roediger HL, McDaniel MA (2014). Both multiple-choice and short-answer quizzes enhance later exam performance in middle and high school classes. *Journal of Experimental Psychology: Applied*. https://doi.org/10.1037/xap0000004. **Verified:** DOI + quote. Ids: mcdermott2014. Cited in: 01, 05.

156. Way WD (1998). Protecting the Integrity of Computerized Testing Item Pools. *Educational Measurement: Issues and Practice*. https://doi.org/10.1111/j.1745-3992.1998.tb00632.x. **Verified:** DOI + quote. Ids: way1998. Cited in: 01.

157. Lyle KB, Bego CR, Hopkins RF, Hieb JL, Ralston PAS (2020). How the Amount and Spacing of Retrieval Practice Affect the Short- and Long-Term Retention of Mathematics Knowledge. *Educational Psychology Review*. https://doi.org/10.1007/s10648-019-09489-x. **Verified:** DOI + quote. Ids: lyle2020. Cited in: 01, 02.

158. PastÃ¶tter B, BÃ¤uml KHT (2014). Retrieval practice enhances new learning: the forward effect of testing. *Frontiers in Psychology*. https://doi.org/10.3389/fpsyg.2014.00286. **Verified:** DOI + quote. Ids: pastotter2014. Cited in: 01, 02.

159. Latimier A, Peyre H, Ramus F (2021). A Meta-Analytic Review of the Benefit of Spacing out Retrieval Practice Episodes on Retention. *Educational Psychology Review*. https://doi.org/10.1007/s10648-020-09572-8. **Verified:** DOI + quote. Ids: latimier2021. Cited in: 01, 02.

160. Roediger HL, Butler AC (2011). The critical role of retrieval practice in long-term retention. *Trends in Cognitive Sciences*. https://doi.org/10.1016/j.tics.2010.09.003. **Verified:** DOI + quote. Ids: roediger2011. Cited in: 01.

161. Belland BR, Walker AE, Kim NJ, Lefler M (2017). Synthesizing Results From Empirical Research on Computer-Based Scaffolding in STEM Education: A Meta-Analysis. *Review of Educational Research*. https://doi.org/10.3102/0034654316670999. **Verified:** DOI + quote. Ids: belland2017. Cited in: 01.

162. van de Pol J, Volman M, Beishuizen J (2010). Scaffolding in Teacher–Student Interaction: A Decade of Research. *Educational Psychology Review*. https://doi.org/10.1007/s10648-010-9127-6. **Verified:** DOI + quote. Ids: vandepol2010. Cited in: 01.

163. Firth J, Rivers I, Boyle J (2021). A systematic review of interleaving as a concept learning strategy. *Review of Education*. https://doi.org/10.1002/rev3.3266. **Verified:** DOI + quote. Ids: firth2021. Cited in: 01, 02, 03.

164. Tauber SK, Dunlosky J, Rawson KA, Wahlheim CN, Jacoby LL (2013). Self-regulated learning of a natural category: Do people interleave or block exemplars during study? *Psychonomic Bulletin & Review*. https://doi.org/10.3758/s13423-012-0319-6. **Verified:** DOI + quote. Ids: tauber2013. Cited in: 01.

165. Paas FGWC, Van Merriënboer JJG (1994). Variability of worked examples and transfer of geometrical problem-solving skills: A cognitive-load approach. *Journal of Educational Psychology*. https://doi.org/10.1037/0022-0663.86.1.122. **Verified:** DOI + quote. Ids: paas1994. Cited in: 01.

166. Kang SHK, McDermott KB, Roediger HL (2007). Test format and corrective feedback modify the effect of testing on long-term retention. *European Journal of Cognitive Psychology*. https://doi.org/10.1080/09541440601056620. **Verified:** DOI + quote. Ids: kang2007. Cited in: 01, 02, 03, 05.

167. Pyc MA, Rawson KA (2007). Examining the efficiency of schedules of distributed retrieval practice. *Memory & Cognition*. https://doi.org/10.3758/bf03192925. **Verified:** DOI + quote. Ids: pyc2007. Cited in: 01, 02.

168. Vaughn KE, Rawson KA (2011). Diagnosing Criterion-Level Effects on Memory: What Aspects of Memory Are Enhanced by Repeated Retrieval? *Psychological Science*. https://doi.org/10.1177/0956797611417724. **Verified:** DOI + quote. Ids: vaughn2011. Cited in: 01.

169. Metcalfe J, Finn B (2011). People's hypercorrection of high-confidence errors: Did they know it all along? *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0021962. **Verified:** DOI + quote. Ids: metcalfe2011. Cited in: 01.

170. Carpenter SK, Pashler H, Cepeda NJ (2009). Using tests to enhance 8th grade students' retention of U.S. history facts. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.1507. **Verified:** DOI + quote. Ids: carpenter2009hist. Cited in: 01.

171. Mettler E, Massey CM, Kellman PJ (2016). A comparison of adaptive and fixed schedules of practice. *Journal of Experimental Psychology: General*. https://doi.org/10.1037/xge0000170. **Verified:** DOI + quote. Ids: mettler2016. Cited in: 01, 02.

172. Pelánek R, Řihák J (2017). Experimental Analysis of Mastery Learning Criteria. *Proceedings of the 25th Conference on User Modeling, Adaptation and Personalization*. https://doi.org/10.1145/3079628.3079667. **Verified:** DOI + quote. Ids: pelanek2017mastery. Cited in: 01.

173. Aleven V, Stahl E, Schworm S, Fischer F, Wallace R (2003). Help Seeking and Help Design in Interactive Learning Environments. *Review of Educational Research*. https://doi.org/10.3102/00346543073003277. **Verified:** DOI + quote. Ids: aleven2003. Cited in: 01, 03, 05.

174. Wood D, Bruner JS, Ross G (1976). The Role Of Tutoring In Problem Solving *. *Journal of Child Psychology and Psychiatry*. https://doi.org/10.1111/j.1469-7610.1976.tb00381.x. **Verified:** title-level only. Ids: wood1976. Cited in: 01.

175. Kang SHK, Lindsey RV, Mozer MC, Pashler H (2014). Retrieval practice over the long term: Should spacing be expanding or equal-interval? *Psychonomic Bulletin & Review*. https://doi.org/10.3758/s13423-014-0636-z. **Verified:** DOI + quote. Ids: kang2014. Cited in: 01, 02.

176. National Governors Association Center for Best Practices; Council of Chief State School Officers (2010). Common Core State Standards for Mathematics: Grade 8, Expressions and Equations. Official standard or regulation. https://www.thecorestandards.org/Math/Content/8/EE/. **Verified:** URL + quote. Ids: ccss_8_ee. Cited in: 01.

177. National Governors Association Center for Best Practices; Council of Chief State School Officers (2010). Common Core State Standards for Mathematics: Grade 3, Number and Operations in Base Ten. Official standard or regulation. https://www.thecorestandards.org/Math/Content/3/NBT/. **Verified:** URL + quote. Ids: ccss_3_nbt. Cited in: 01.

178. National Governors Association Center for Best Practices; Council of Chief State School Officers (2010). Common Core State Standards for Mathematics: Kindergarten, Operations and Algebraic Thinking. Official standard or regulation. https://www.thecorestandards.org/Math/Content/K/OA/. **Verified:** URL + quote. Ids: ccss_k_oa. Cited in: 01.

179. National Governors Association Center for Best Practices; Council of Chief State School Officers (2010). Common Core State Standards for Mathematics: Kindergarten, Counting and Cardinality. Official standard or regulation. https://www.thecorestandards.org/Math/Content/K/CC/. **Verified:** URL + quote. Ids: ccss_k_cc. Cited in: 01.

180. Janes JL, Dunlosky J, Rawson KA, Jasnow A (2020). Successive relearning improves performance on a high-stakes exam in a difficult biopsychology course. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.3699. **Verified:** DOI + quote. Ids: janes2020. Cited in: 01.

181. Higham PA, Zengel B, Bartlett LK, Hadwin JA (2022). The benefits of successive relearning on multiple learning outcomes. *Journal of Educational Psychology*. https://doi.org/10.1037/edu0000693. **Verified:** DOI + quote. Ids: higham2022. Cited in: 01, 02.

182. Rawson KA, Vaughn KE, Walsh M, Dunlosky J (2018). Investigating and explaining the effects of successive relearning on long-term retention. *Journal of Experimental Psychology: Applied*. https://doi.org/10.1037/xap0000146. **Verified:** DOI + quote. Ids: rawson2018. Cited in: 01.

183. Guskey TR, Pigott TD (1988). Research on Group-Based Mastery Learning Programs: A Meta-Analysis. *The Journal of Educational Research*. https://doi.org/10.1080/00220671.1988.10885824. **Verified:** DOI + quote. Ids: guskey1988. Cited in: 01.

184. Kulik JA, Kulik CLC (1988). Timing of Feedback and Verbal Learning. *Review of Educational Research*. https://doi.org/10.2307/1170349. **Verified:** DOI + quote. Ids: kulik1988. Cited in: 01, 02, 05.

185. Butler AC, Karpicke JD, Roediger HL (2007). The effect of type and timing of feedback on learning from multiple-choice tests. *Journal of Experimental Psychology: Applied*. https://doi.org/10.1037/1076-898x.13.4.273. **Verified:** DOI + quote. Ids: butler2007. Cited in: 01, 02, 05.

186. Metcalfe J, Kornell N, Finn B (2009). Delayed versus immediate feedback in children's and adults' vocabulary learning. *Memory & Cognition*. https://doi.org/10.3758/mc.37.8.1077. **Verified:** DOI + quote. Ids: metcalfe2009. Cited in: 01, 02.

187. Bisra K, Liu Q, Nesbit JC, Salimi F, Winne PH (2018). Inducing Self-Explanation: a Meta-Analysis. *Educational Psychology Review*. https://doi.org/10.1007/s10648-018-9434-x. **Verified:** title-level only. Ids: bisra2018. Cited in: 01.

188. Rawson KA, Dunlosky J, Sciartelli SM (2013). The Power of Successive Relearning: Improving Performance on Course Exams and Long-Term Retention. *Educational Psychology Review*. https://doi.org/10.1007/s10648-013-9240-4. **Verified:** title-level only. Ids: rawson2013power, rawson2013. Cited in: 01, 04.

189. Manakul P, Liusie A, Gales M (2023). SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. https://doi.org/10.18653/v1/2023.emnlp-main.557. **Verified:** DOI + quote. Ids: manakul2023. Cited in: 01.

190. Lyle KB, Crawford NA (2011). Retrieving Essential Material at the End of Lectures Improves Performance on Statistics Exams. *Teaching of Psychology*. https://doi.org/10.1177/0098628311401587. **Verified:** DOI + quote. Ids: lyle2011. Cited in: 02.

191. Karpicke JD, Bauernschmidt A (2011). Spaced retrieval: Absolute spacing enhances learning regardless of relative spacing. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0023436. **Verified:** DOI + quote. Ids: karpicke2011spaced. Cited in: 02.

192. Karpicke JD, Roediger HL (2007). Expanding retrieval practice promotes short-term retention, but equally spaced retrieval enhances long-term retention. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/0278-7393.33.4.704. **Verified:** DOI + quote. Ids: karpicke2007. Cited in: 02.

193. Cepeda NJ, Coburn N, Rohrer D, Wixted JT, Mozer MC, Pashler H (2009). Optimizing Distributed Practice: Theoretical Analysis and Practical Implications. *Experimental Psychology*. https://doi.org/10.1027/1618-3169.56.4.236. **Verified:** DOI + quote. Ids: cepeda2009. Cited in: 02.

194. Jachimowicz JM, Duncan S, Weber EU, Johnson EJ (2019). When and why defaults influence decisions: a meta-analysis of default effects. *Behavioural Public Policy*. https://doi.org/10.1017/bpp.2018.43. **Verified:** DOI + quote. Ids: jachimowicz2019. Cited in: 02.

195. Johnson EJ, Goldstein D (2003). Do Defaults Save Lives? *Science*. https://doi.org/10.1126/science.1091721. **Verified:** DOI + quote. Ids: johnson2003. Cited in: 02.

196. Trumbo MC, Leiting KA, McDaniel MA, Hodge GK (2016). Effects of reinforcement on test-enhanced learning in a large, diverse introductory college psychology course. *Journal of Experimental Psychology: Applied*. https://doi.org/10.1037/xap0000082. **Verified:** DOI + quote. Ids: trumbo2016. Cited in: 02.

197. Rohrer D, Pashler H (2007). Increasing Retention Without Increasing Study Time. *Current Directions in Psychological Science*. https://doi.org/10.1111/j.1467-8721.2007.00500.x. **Verified:** DOI + quote. Ids: rohrer2007pashler. Cited in: 02.

198. Rohrer D, Taylor K (2006). The effects of overlearning and distributed practise on the retention of mathematics knowledge. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.1266. **Verified:** DOI + quote. Ids: rohrer2006. Cited in: 02, 03.

199. Sievertsen HH, Gino F, Piovesan M (2016). Cognitive fatigue influences students' performance on standardized tests. *Proceedings of the National Academy of Sciences*. https://doi.org/10.1073/pnas.1516947113. **Verified:** DOI + quote. Ids: sievertsen2016. Cited in: 02, 05.

200. Kovanovic V, Gašević D, Dawson S, Joksimovic S, Baker R (2016). Does Time-on-task Estimation Matter? Implications on Validity of Learning Analytics Findings. *Journal of Learning Analytics*. https://doi.org/10.18608/jla.2015.23.6. **Verified:** DOI + quote. Ids: kovanovic2016. Cited in: 02.

201. Nguyen Q (2020). Rethinking time-on-task estimation with outlier detection accounting for individual, time, and task differences. *Proceedings of the Tenth International Conference on Learning Analytics & Knowledge*. https://doi.org/10.1145/3375462.3375538. **Verified:** DOI + quote. Ids: nguyen2020tot. Cited in: 02.

202. Rotelli D, Monreale A (2022). Time-on-Task Estimation by data-driven Outlier Detection based on Learning Activities. *LAK22: 12th International Learning Analytics and Knowledge Conference*. https://doi.org/10.1145/3506860.3506913. **Verified:** DOI + quote. Ids: rotelli2022. Cited in: 02.

203. Goldhammer F, Naumann J, Stelter A, Tóth K, Rölke H, Klieme E (2014). The time on task effect in reading and problem solving is moderated by task difficulty and skill: Insights from a computer-based large-scale assessment. *Journal of Educational Psychology*. https://doi.org/10.1037/a0034716. **Verified:** DOI + quote. Ids: goldhammer2014. Cited in: 02.

204. Kim ASN, Wong-Kee-You AMB, Wiseheart M, Rosenbaum RS (2019). The spacing effect stands up to big data. *Behavior Research Methods*. https://doi.org/10.3758/s13428-018-1184-7. **Verified:** DOI + quote. Ids: kim2019bigdata. Cited in: 02.

205. Vaughn KE, Dunlosky J, Rawson KA (2016). Effects of successive relearning on recall: Does relearning override the effects of initial learning criterion? *Memory & Cognition*. https://doi.org/10.3758/s13421-016-0606-y. **Verified:** DOI + quote. Ids: vaughn2016. Cited in: 02.

206. Rawson KA, Dunlosky J (2013). Relearning attenuates the benefits and costs of spacing. *Journal of Experimental Psychology: General*. https://doi.org/10.1037/a0030498. **Verified:** DOI + quote. Ids: rawson2013relearn. Cited in: 02.

207. Harkin B, Webb TL, Chang BPI, Prestwich A, Conner M, Kellar I, et al. (2016). Does monitoring goal progress promote goal attainment? A meta-analysis of the experimental evidence. *Psychological Bulletin*. https://doi.org/10.1037/bul0000025. **Verified:** DOI + quote. Ids: harkin2016. Cited in: 02.

208. Wissman KT, Rawson KA, Pyc MA (2012). How and when do students use flashcards? *Memory*. https://doi.org/10.1080/09658211.2012.687052. **Verified:** DOI + quote. Ids: wissman2012. Cited in: 02.

209. Kornell N (2009). Optimising learning using flashcards: Spacing is more effective than cramming. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.1537. **Verified:** DOI + quote. Ids: kornell2009cram. Cited in: 02.

210. Bjork RA, Dunlosky J, Kornell N (2013). Self-Regulated Learning: Beliefs, Techniques, and Illusions. *Annual Review of Psychology*. https://doi.org/10.1146/annurev-psych-113011-143823. **Verified:** DOI + quote. Ids: bjork2013. Cited in: 02.

211. Kornell N, Bjork RA (2008). Optimising self-regulated study: The benefits—and costs—of dropping flashcards. *Memory*. https://doi.org/10.1080/09658210701763899. **Verified:** DOI + quote. Ids: kornell2008spacing. Cited in: 02.

212. Kornell N, Son LK (2009). Learners' choices and beliefs about self-testing. *Memory*. https://doi.org/10.1080/09658210902832915. **Verified:** DOI + quote. Ids: kornell2009son. Cited in: 02.

213. Pavlik PI, Anderson JR (2008). Using a model to compute the optimal schedule of practice. *Journal of Experimental Psychology: Applied*. https://doi.org/10.1037/1076-898x.14.2.101. **Verified:** DOI + quote. Ids: pavlik2008. Cited in: 02.

214. Ye J, Su J, Cao Y (2022). A Stochastic Shortest Path Algorithm for Optimizing Spaced Repetition Scheduling. *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. https://doi.org/10.1145/3534678.3539081. **Verified:** DOI + quote. Ids: ye2022. Cited in: 02.

215. Pyc MA, Rawson KA (2009). Testing the retrieval effort hypothesis: Does greater difficulty correctly recalling information lead to higher levels of memory? *Journal of Memory and Language*. https://doi.org/10.1016/j.jml.2009.01.004. **Verified:** title-level only. Ids: pyc2009. Cited in: 02.

216. Fazio LK, Marsh EJ (2019). Retrieval-Based Learning in Children. *Current Directions in Psychological Science*. https://doi.org/10.1177/0963721418806673. **Verified:** DOI + quote. Ids: fazio2019. Cited in: 02.

217. Kornell N, Hays MJ, Bjork RA (2009). Unsuccessful retrieval attempts enhance subsequent learning. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0015729. **Verified:** DOI + quote. Ids: kornell2009. Cited in: 02, 04.

218. Butterfield B, Metcalfe J (2001). Errors committed with high confidence are hypercorrected. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/0278-7393.27.6.1491. **Verified:** DOI + quote. Ids: butterfield2001. Cited in: 02.

219. Metcalfe J (2017). Learning from Errors. *Annual Review of Psychology*. https://doi.org/10.1146/annurev-psych-010416-044022. **Verified:** DOI + quote. Ids: metcalfe2017. Cited in: 02, 05.

220. Butler AC, Karpicke JD, Roediger HL (2008). Correcting a metacognitive error: Feedback increases retention of low-confidence correct responses. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/0278-7393.34.4.918. **Verified:** DOI + quote. Ids: butler2008. Cited in: 02, 05.

221. Eglington LG, Pavlik PI (2023). How to Optimize Student Learning Using Student Models That Adapt Rapidly to Individual Differences. *International Journal of Artificial Intelligence in Education*. https://doi.org/10.1007/s40593-022-00296-0. **Verified:** DOI + quote. Ids: eglington2023. Cited in: 02.

222. Zerr CL, Berg JJ, Nelson SM, Fishell AK, Savalia NK, McDermott KB (2018). Learning Efficiency: Identifying Individual Differences in Learning Rate and Retention in Healthy Adults. *Psychological Science*. https://doi.org/10.1177/0956797618772540. **Verified:** DOI + quote. Ids: zerr2018. Cited in: 02.

223. Sense F, Behrens F, Meijer RR, van Rijn H (2016). An Individual's Rate of Forgetting Is Stable Over Time but Differs Across Materials. *Topics in Cognitive Science*. https://doi.org/10.1111/tops.12183. **Verified:** DOI + quote. Ids: sense2016. Cited in: 02.

224. Vaughn KE, Rawson KA, Pyc MA (2013). Repeated retrieval practice and item difficulty: Does criterion learning eliminate item difficulty effects? *Psychonomic Bulletin & Review*. https://doi.org/10.3758/s13423-013-0434-z. **Verified:** DOI + quote. Ids: vaughn2013. Cited in: 02.

225. Rawson KA, Dunlosky J, Janes JL (2020). All Good Things Must Come to an End: a Potential Boundary Condition on the Potency of Successive Relearning. *Educational Psychology Review*. https://doi.org/10.1007/s10648-020-09528-y. **Verified:** DOI + quote. Ids: rawson2020end. Cited in: 02.

226. Finn B (2010). Ending on a high note: Adding a better end to effortful study. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0020605. **Verified:** DOI + quote. Ids: finn2010. Cited in: 02.

227. Finn B, Miele DB (2016). Hitting a high note on math tests: Remembered success influences test preferences. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/xlm0000150. **Verified:** DOI + quote. Ids: finn2016. Cited in: 02.

228. Kahneman D, Fredrickson BL, Schreiber CA, Redelmeier DA (1993). When More Pain Is Preferred to Less: Adding a Better End. *Psychological Science*.

https://doi.org/10.1111/j.1467-9280.1993.tb00589.x. **Verified:** DOI + quote. Ids: kahneman1993. Cited in: 02.

229. Fredrickson BL, Kahneman D (1993). Duration neglect in retrospective evaluations of affective episodes. *Journal of Personality and Social Psychology*. https://doi.org/10.1037/0022-3514.65.1.45. **Verified:** DOI + quote. Ids: fredrickson1993. Cited in: 02.

230. Yeo DJ, Fazio LK (2019). The optimal learning strategy depends on learning goals and processes: Retrieval practice versus worked examples. *Journal of Educational Psychology*. https://doi.org/10.1037/edu0000268. **Verified:** DOI + quote. Ids: yeo2019. Cited in: 02.

231. Dunlosky J, Rawson KA (2012). Overconfidence produces underachievement: Inaccurate self evaluations undermine students' learning and retention. *Learning and Instruction*. https://doi.org/10.1016/j.learninstruc.2011.08.003. **Verified:** DOI + quote. Ids: dunlosky2012. Cited in: 02, 05.

232. Khanna MM (2015). Ungraded Pop Quizzes: Test-Enhanced Learning Without All the Anxiety. *Teaching of Psychology*. https://doi.org/10.1177/0098628315573144. **Verified:** DOI + quote. Ids: khanna2015. Cited in: 02.

233. Agarwal PK, D'Antonio L, Roediger HL, McDermott KB, McDaniel MA (2014). Classroom-based programs of retrieval practice reduce middle school and high school students' test anxiety. *Journal of Applied Research in Memory and Cognition*. https://doi.org/10.1016/j.jarmac.2014.07.002. **Verified:** DOI + quote. Ids: agarwal2014anx. Cited in: 02.

234. Smith AM, Floerke VA, Thomas AK (2016). Retrieval practice protects memory against acute stress. *Science*. https://doi.org/10.1126/science.aah5067. **Verified:** DOI + quote. Ids: smith2016stress. Cited in: 02.

235. Woźniak P, Gorzelańczyk E (1994). Optimization of repetition spacing in the practice of learning. *Acta Neurobiologiae Experimentalis*. https://doi.org/10.55782/ane-1994-1003. **Verified:** DOI + quote. Ids: wozniak1994. Cited in: 02.

236. Bell MC, Kawadri N, Simone PM, Wiseheart M (2014). Long-term memory, sleep, and the spacing effect. *Memory*. https://doi.org/10.1080/09658211.2013.778294. **Verified:** DOI + quote. Ids: bell2014. Cited in: 02.

237. Arnold KM, McDermott KB (2013). Test-potentiated learning: Distinguishing between direct and indirect effects of tests. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0029199. **Verified:** DOI + quote. Ids: arnold2013. Cited in: 02.

238. Hays MJ, Kornell N, Bjork RA (2013). When and why a failed test potentiates the effectiveness of subsequent study. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0028468. **Verified:** DOI + quote. Ids: hays2013. Cited in: 02.

239. Mazza S, Gerbier E, Gustin MP, Kasikci Z, Koenig O, Toppino TC, et al. (2016). Relearn Faster and Retain Longer: Along With Practice, Sleep Makes Perfect. *Psychological Science*. https://doi.org/10.1177/0956797616659930. **Verified:** DOI + quote. Ids: mazza2016. Cited in: 02, 04.

240. Chan JCK, Meissner CA, Davis SD (2018). Retrieval potentiates new learning: A theoretical and meta-analytic review. *Psychological Bulletin*. https://doi.org/10.1037/bul0000166. **Verified:** DOI + quote. Ids: chan2018. Cited in: 02.

241. Yang C, Potts R, Shanks DR (2018). Enhancing learning and retrieval of new information: a review of the forward testing effect. *npj Science of Learning*. https://doi.org/10.1038/s41539-018-0024-y. **Verified:** DOI + quote. Ids: yang2018fwd. Cited in: 02.

242. Szpunar KK, McDermott KB, Roediger HL (2008). Testing during study insulates against the buildup of proactive interference. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0013082. **Verified:** DOI + quote. Ids: szpunar2008. Cited in: 02.

243. Emeny WG, Hartwig MK, Rohrer D (2021). Spaced mathematics practice improves test scores and reduces overconfidence. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.3814. **Verified:** DOI + quote. Ids: emeny2021. Cited in: 02.

244. Kapler IV, Weston T, Wiseheart M (2015). Spacing in a simulated undergraduate classroom: Long-term benefits for factual and higher-level learning. *Learning and Instruction*. https://doi.org/10.1016/j.learninstruc.2014.11.001. **Verified:** title-level only. Ids: kapler2015. Cited in: 02.

245. Szpunar KK, Khan NY, Schacter DL (2013). Interpolated memory tests reduce mind wandering and improve learning of online lectures. *Proceedings of the National Academy of Sciences*. https://doi.org/10.1073/pnas.1221764110. **Verified:** DOI + quote. Ids: szpunar2013. Cited in: 02, 05.

246. Mullet HG, Butler AC, Verdin B, von Borries R, Marsh EJ (2014). Delaying feedback promotes transfer of knowledge despite student preferences to receive feedback immediately. *Journal of Applied Research in Memory and Cognition*. https://doi.org/10.1016/j.jarmac.2014.05.001. **Verified:** DOI + quote. Ids: mullet2014. Cited in: 02, 05.

247. Corral D, Carpenter SK, Clingan-Siverly S (2021). The effects of immediate versus delayed feedback on complex concept learning. *Quarterly Journal of Experimental Psychology*. https://doi.org/10.1177/1747021820977739. **Verified:** DOI + quote. Ids: corral2021. Cited in: 02.

248. Butler AC, Roediger HL (2008). Feedback enhances the positive effects and reduces the negative effects of multiple-choice testing. *Memory & Cognition*. https://doi.org/10.3758/mc.36.3.604. **Verified:** DOI + quote. Ids: butler2008mc. Cited in: 02.

249. Smith TA, Kimball DR (2010). Learning from feedback: Spacing and the delay–retention effect. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/a0017407. **Verified:** DOI + quote. Ids: smith2010kimball. Cited in: 02.

250. Berres S, Erdfelder E (2021). The sleep benefit in episodic memory: An integrative review and a meta-analysis. *Psychological Bulletin*. https://doi.org/10.1037/bul0000350. **Verified:** DOI + quote. Ids: berres2021. Cited in: 02.

251. Gais S, Lucas B, Born J (2006). Sleep after learning aids memory recall. *Learning & Memory*. https://doi.org/10.1101/lm.132106. **Verified:** DOI + quote. Ids: gais2006. Cited in: 02.

252. Bäuml KHT, Holterman C, Abel M (2014). Sleep can reduce the testing effect: It enhances recall of restudied items but can leave recall of retrieved items unaffected. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/xlm0000025. **Verified:** DOI + quote. Ids: bauml2014. Cited in: 02.

253. Wilhelm I, Diekelmann S, Born J (2008). Sleep in children improves memory performance on declarative but not procedural tasks. *Learning & Memory*. https://doi.org/10.1101/lm.803708. **Verified:** DOI + quote. Ids: wilhelm2008. Cited in: 02, 04.

254. Pavlik PI, Anderson JR (2005). Practice and Forgetting Effects on Vocabulary Memory: An Activation-Based Model of the Spacing Effect. *Cognitive Science*. https://doi.org/10.1207/s15516709cog0000_14. **Verified:** DOI + quote. Ids: pavlik2005. Cited in: 02.

255. Wixted JT, Ebbesen EB (1991). On the Form of Forgetting. *Psychological Science*. https://doi.org/10.1111/j.1467-9280.1991.tb00175.x. **Verified:** DOI + quote. Ids: wixted1991. Cited in: 02.

256. Rubin DC, Wenzel AE (1996). One hundred years of forgetting: A quantitative description of retention. *Psychological Review*. https://doi.org/10.1037/0033-295x.103.4.734. **Verified:** DOI + quote. Ids: rubin1996. Cited in: 02.

257. Murre JMJ, Dros J (2015). Replication and Analysis of Ebbinghaus' Forgetting Curve. *Plos One*. https://doi.org/10.1371/journal.pone.0120644. **Verified:** DOI + quote. Ids: murre2015. Cited in: 02.

258. Averell L, Heathcote A (2011). The form of the forgetting curve and the fate of memories. *Journal of Mathematical Psychology*. https://doi.org/10.1016/j.jmp.2010.08.009. **Verified:** title-level only. Ids: averell2011. Cited in: 02.

259. Donovan JJ, Radosevich DJ (1999). A meta-analytic review of the distribution of practice effect: Now you see it, now you don't. *Journal of Applied Psychology*. https://doi.org/10.1037/0021-9010.84.5.795. **Verified:** DOI + quote. Ids: donovan1999. Cited in: 02, 04.

260. Khajah MM, Lindsey RV, Mozer MC (2014). Maximizing Students' Retention via Spaced Review: Practical Guidance From Computational Models of Memory. *Topics in Cognitive Science*. https://doi.org/10.1111/tops.12077. **Verified:** DOI + quote. Ids: khajah2014. Cited in: 02.

261. Rohrer D, Taylor K, Pashler H, Wixted JT, Cepeda NJ (2005). The effect of overlearning on long-term retention. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.1083. **Verified:** DOI + quote. Ids: rohrer2005over. Cited in: 02, 03, 04.

262. Anderson MC, Bjork RA, Bjork EL (1994). Remembering can cause forgetting: Retrieval dynamics in long-term memory. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/0278-7393.20.5.1063. **Verified:** DOI + quote. Ids: anderson1994. Cited in: 02.

263. Murayama K, Miyatsu T, Buchli D, Storm BC (2014). Forgetting as a consequence of retrieval: A meta-analytic review of retrieval-induced forgetting. *Psychological Bulletin*. https://doi.org/10.1037/a0037505. **Verified:** DOI + quote. Ids: murayama2014. Cited in: 02.

264. Greving S, Richter T (2018). Examining the Testing Effect in University Teaching: Retrievability and Question Format Matter. *Frontiers in Psychology*. https://doi.org/10.3389/fpsyg.2018.02412. **Verified:** DOI + quote. Ids: greving2018. Cited in: 03.

265. Wise SL, Bhola DS, Yang S (2006). Taking the Time to Improve the Validity of Low-Stakes Tests: The Effort-Monitoring CBT. *Educational Measurement: Issues and Practice*. https://doi.org/10.1111/j.1745-3992.2006.00054.x. **Verified:** DOI + quote. Ids: wise2006cbt. Cited in: 03.

266. Wise SL, Kuhfeld MR, Soland J (2019). The Effects of Effort Monitoring With Proctor Notification on Test-Taking Engagement, Test Performance, and Validity. *Applied Measurement in Education*. https://doi.org/10.1080/08957347.2019.1577248. **Verified:** DOI + quote. Ids: wise2019. Cited in: 03.

267. He Q, Zeng J, Huang W, Chen L, Xiao J, He Q, et al. (2024). Can Large Language Models Understand Real-World Complex Instructions? *Proceedings of the AAAI Conference on Artificial Intelligence*. https://doi.org/10.1609/aaai.v38i16.29777. **Verified:** DOI + quote. Ids: he2024cello. Cited in: 03.

268. Lomas D, Patel K, Forlizzi JL, Koedinger KR (2013). Optimizing challenge in an educational game using large-scale design experiments. *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. https://doi.org/10.1145/2470654.2470668. **Verified:** DOI + quote. Ids: lomas2013. Cited in: 03.

269. Kirschner PA, van Merriënboer JJ (2013). Do Learners Really Know Best? Urban Legends in Education. *Educational Psychologist*. https://doi.org/10.1080/00461520.2013.804395. **Verified:** DOI + quote. Ids: kirschner2013. Cited in: 03.

270. Koo TK, Li MY (2016). A Guideline of Selecting and Reporting Intraclass Correlation Coefficients for Reliability Research. *Journal of Chiropractic Medicine*. https://doi.org/10.1016/j.jcm.2016.02.012. **Verified:** DOI + quote. Ids: koo2016. Cited in: 03.

271. Cicchetti DV (1994). Guidelines, criteria, and rules of thumb for evaluating normed and standardized assessment instruments in psychology. *Psychological Assessment*. https://doi.org/10.1037/1040-3590.6.4.284. **Verified:** DOI + quote. Ids: cicchetti1994. Cited in: 03.

272. Lance CE, Butts MM, Michels LC (2006). The Sources of Four Commonly Reported Cutoff Criteria. *Organizational Research Methods*. https://doi.org/10.1177/1094428105284919. **Verified:** DOI + quote. Ids: lance2006. Cited in: 03.

273. Brown W (1910). Some Experimental Results In The Correlation Of Mental Abilities 1. *British Journal of Psychology, 1904-1920*. https://doi.org/10.1111/j.2044-8295.1910.tb00207.x. **Verified:** title-level only. Ids: brown1910. Cited in: 03.

274. Spearman C (1910). Correlation Calculated From Faulty Data. *British Journal of Psychology, 1904-1920*. https://doi.org/10.1111/j.2044-8295.1910.tb00206.x. **Verified:** title-level only. Ids: spearman1910. Cited in: 03.

275. Newton MA (2004). Detecting differential gene expression with a semiparametric hierarchical mixture method. *Biostatistics*. https://doi.org/10.1093/biostatistics/5.2.155. **Verified:** DOI + quote. Ids: newton2004. Cited in: 03.

276. Efron B (2008). Microarrays, Empirical Bayes and the Two-Groups Model. *Statistical Science*. https://doi.org/10.1214/07-sts236. **Verified:** DOI + quote. Ids: efron2008. Cited in: 03.

277. Kass RE, Raftery AE (1995). Bayes Factors. *Journal of the American Statistical Association*. https://doi.org/10.1080/01621459.1995.10476572. **Verified:** DOI + quote. Ids: kass1995. Cited in: 03.

278. Berry DA (2006). Bayesian clinical trials. *Nature Reviews Drug Discovery*. https://doi.org/10.1038/nrd1927. **Verified:** DOI + quote. Ids: berry2006. Cited in: 03.

279. Haldane BJBS (1956). The Estimation And Significance Of The Logarithm Of A Ratio Of Frequencies. *Annals of Human Genetics*. https://doi.org/10.1111/j.1469-1809.1955.tb01285.x. **Verified:** title-level only. Ids: haldane1956. Cited in: 03.

280. Anscombe FJ (1956). On estimating binomial response relations. *Biometrika*. https://doi.org/10.1093/biomet/43.3-4.461. **Verified:** title-level only. Ids: anscombe1956. Cited in: 03.

281. Gart JJ, Zweifel JR (1967). On the bias of various estimators of the logit and its variance with application to quantal bioassay. *Biometrika*. https://doi.org/10.1093/biomet/54.1-2.181. **Verified:** DOI + quote. Ids: gart1967. Cited in: 03.

282. Warton DI, Hui FKC (2011). The arcsine is asinine: the analysis of proportions in ecology. *Ecology*. https://doi.org/10.1890/10-0340.1. **Verified:** DOI + quote. Ids: warton2011. Cited in: 03.

283. Baker RS, Corbett AT, Koedinger KR, Wagner AZ (2004). Off-task behavior in the cognitive tutor classroom: when students "game the system". *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. https://doi.org/10.1145/985692.985741. **Verified:** DOI + quote. Ids: baker2004. Cited in: 03.

284. Usher EL, Pajares F (2008). Sources of Self-Efficacy in School: Critical Review of the Literature and Future Directions. *Review of Educational Research*. https://doi.org/10.3102/0034654308321456. **Verified:** DOI + quote. Ids: usher2008. Cited in: 03.

285. Byars-Winston A, Diestelmann J, Savoy JN, Hoyt WT (2017). Unique effects and moderators of effects of sources on self-efficacy: A model-based meta-analysis. *Journal of Counseling Psychology*. https://doi.org/10.1037/cou0000219. **Verified:** DOI + quote. Ids: byarswinston2017. Cited in: 03.

286. Bandura A, Schunk DH (1981). Cultivating competence, self-efficacy, and intrinsic interest through proximal self-motivation. *Journal of Personality and Social Psychology*. https://doi.org/10.1037/0022-3514.41.3.586. **Verified:** DOI + quote. Ids: bandura1981. Cited in: 03.

287. Schunk DH (1983). Ability versus effort attributional feedback: Differential effects on self-efficacy and achievement. *Journal of Educational Psychology*. https://doi.org/10.1037/0022-0663.75.6.848. **Verified:** DOI + quote. Ids: schunk1983. Cited in: 03.

288. Talsma K, Schüz B, Schwarzer R, Norris K (2018). I believe, therefore I achieve (and vice versa): A meta-analytic cross-lagged panel analysis of self-efficacy and academic performance. *Learning and Individual Differences*. https://doi.org/10.1016/j.lindif.2017.11.015. **Verified:** DOI + quote. Ids: talsma2018. Cited in: 03.

289. Kendeou P, Butterfuss R, Kim J, Van Boekel M (2018). Knowledge revision through the lenses of the three-pronged approach. *Memory & Cognition*. https://doi.org/10.3758/s13421-018-0848-y. **Verified:** DOI + quote. Ids: kendeou2018. Cited in: 03.

290. Shtulman A, Valcarcel J (2012). Scientific knowledge suppresses but does not supplant earlier intuitions. *Cognition*. https://doi.org/10.1016/j.cognition.2012.04.005. **Verified:** DOI + quote. Ids: shtulman2012. Cited in: 03.

291. Payne SJ, Squibb HR (1990). Algebra Mal-Rules and Cognitive Accounts of Error. *Cognitive Science*. https://doi.org/10.1207/s15516709cog1403_4. **Verified:** DOI + quote. Ids: payne1990. Cited in: 03.

292. Smith III JP, diSessa AA, Roschelle J (1994). Misconceptions Reconceived: A Constructivist Analysis of Knowledge in Transition. *Journal of the Learning Sciences*. https://doi.org/10.1207/s15327809jls0302_1. **Verified:** DOI + quote. Ids: smith1994. Cited in: 03.

293. D'Mello S, Graesser A (2012). Dynamics of affective states during complex learning. *Learning and Instruction*. https://doi.org/10.1016/j.learninstruc.2011.10.001. **Verified:** DOI + quote. Ids: dmello2012. Cited in: 03.

294. Camacho-Morles J, Slemp GR, Pekrun R, Loderer K, Hou H, Oades LG (2021). Activity Achievement Emotions and Academic Performance: A Meta-analysis. *Educational Psychology Review*. https://doi.org/10.1007/s10648-020-09585-3. **Verified:** DOI + quote. Ids: camachomorles2021. Cited in: 03.

295. Botelho AF, Varatharaj A, Patikorn T, Doherty D, Adjei SA, Beck JE (2019). Developing Early Detectors of Student Attrition and Wheel Spinning Using Deep Learning. *IEEE Transactions on Learning Technologies*. https://doi.org/10.1109/tlt.2019.2912162. **Verified:** DOI + quote. Ids: botelho2019. Cited in: 03.

296. Ventura M, Shute V, Zhao W (2013). The relationship between video game use and a performance-based measure of persistence. *Computers & Education*. https://doi.org/10.1016/j.compedu.2012.07.003. **Verified:** DOI + quote. Ids: ventura2013ce. Cited in: 03.

297. Mace FC, Hock ML, Lalli JS, West BJ, Belfiore P, Pinter E, et al. (1988). Behavioral Momentum In The Treatment Of Noncompliance. *Journal of Applied Behavior Analysis*. https://doi.org/10.1901/jaba.1988.21-123. **Verified:** DOI + quote. Ids: mace1988. Cited in: 03.

298. Lee DL (2005). Increasing Compliance: A Quantitative Synthesis of Applied Research on High-Probability Request Sequences. *Exceptionality*. https://doi.org/10.1207/s15327035ex1303_1. **Verified:** DOI + quote. Ids: lee2005. Cited in: 03.

299. Aamodt MG, McShane T (1992). A Meta-Analytic Investigation of the Effect of Various Test Item Characteristics on Test Scores and Test Completion Times. *Public Personnel Management*. https://doi.org/10.1177/009102609202100203. **Verified:** DOI + quote. Ids: aamodt1992. Cited in: 03.

300. Benjamini Y, Hochberg Y (1995). Controlling the False Discovery Rate: A Practical and Powerful Approach to Multiple Testing. *Journal of the Royal Statistical Society Series B: Statistical Methodology*. https://doi.org/10.1111/j.2517-6161.1995.tb02031.x. **Verified:** DOI + quote. Ids: benjamini1995. Cited in: 03.

301. Rouder JN, Haaf JM (2019). A psychometrics of individual differences in experimental tasks. *Psychonomic Bulletin & Review*. https://doi.org/10.3758/s13423-018-1558-y. **Verified:** DOI + quote. Ids: rouder2019. Cited in: 03.

302. Enkavi AZ, Eisenberg IW, Bissett PG, Mazza GL, MacKinnon DP, Marsch LA, et al. (2019). Large-scale analysis of test–retest reliabilities of self-regulation measures. *Proceedings of the National Academy of Sciences*. https://doi.org/10.1073/pnas.1818430116. **Verified:** DOI + quote. Ids: enkavi2019. Cited in: 03.

303. Schmidt FL, Hunter JE (1996). Measurement error in psychological research: Lessons from 26 research scenarios. *Psychological Methods*. https://doi.org/10.1037/1082-989x.1.2.199. **Verified:** DOI + quote. Ids: schmidt1996. Cited in: 03.

304. Raudenbush SW, Bryk AS (1985). Empirical Bayes Meta-Analysis. *Journal of Educational Statistics*. https://doi.org/10.2307/1164836. **Verified:** DOI + quote. Ids: raudenbush1985. Cited in: 03.

305. DerSimonian R, Laird N (1986). Meta-analysis in clinical trials. *Controlled Clinical Trials*. https://doi.org/10.1016/0197-2456(86)90046-2. **Verified:** DOI + quote. Ids: dersimonian1986. Cited in: 03.

306. Jaroslawicz, Whiting et al. (2025). How Many Instructions Can LLMs Follow at Once? *arXiv*. https://arxiv.org/abs/2507.11538. **Verified:** URL + quote. Ids: jaroslawicz2025. Cited in: 03.

307. Jiang Y, Wang Y, Zeng X, Zhong W, Li L, Mi F, et al. (2024). FollowBench: A Multi-level Fine-grained Constraints Following Benchmark for Large Language Models. *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. https://doi.org/10.18653/v1/2024.acl-long.257. **Verified:** DOI + quote. Ids: jiang2024follow. Cited in: 03.

308. Qin Y, Song K, Hu Y, Yao W, Cho S, Wang X, et al. (2024). InFoBench: Evaluating Instruction Following Ability in Large Language Models. *Findings of the Association for Computational Linguistics ACL 2024*. https://doi.org/10.18653/v1/2024.findings-acl.772. **Verified:** DOI + quote. Ids: qin2024. Cited in: 03.

309. EU (2016). GDPR Article 5(1)(c): data minimisation. Official standard or regulation. https://gdpr-info.eu/art-5-gdpr/. **Verified:** URL + quote. Ids: gdpr-art5-1c. Cited in: 03, 05.

310. Slade S, Prinsloo P (2013). Learning Analytics: Ethical Issues and Dilemmas. *American Behavioral Scientist*. https://doi.org/10.1177/0002764213479366. **Verified:** DOI + quote. Ids: slade2013. Cited in: 03, 05.

311. Drachsler H, Greller W (2016). Privacy and analytics: it's a DELICATE issue a checklist for trusted learning analytics. *Proceedings of the Sixth International Conference on Learning Analytics & Knowledge - LAK '16*. https://doi.org/10.1145/2883851.2883893. **Verified:** DOI + quote. Ids: drachsler2016. Cited in: 03, 05.

312. Dwork C, Roth A (2013). The Algorithmic Foundations of Differential Privacy. *now Publishers Inc*. https://doi.org/10.1561/9781601988195. **Verified:** DOI + quote. Ids: dwork2013. Cited in: 03.

313. Saltzer J, Schroeder M (1975). The protection of information in computer systems. *Proceedings of the IEEE*. https://doi.org/10.1109/proc.1975.9939. **Verified:** DOI + quote. Ids: saltzer1975. Cited in: 03.

314. Brown A, Maydeu-Olivares A (2011). Item Response Modeling of Forced-Choice Questionnaires. *Educational and Psychological Measurement*. https://doi.org/10.1177/0013164410375112. **Verified:** DOI + quote. Ids: brown2011fc. Cited in: 03.

315. Cao M, Drasgow F (2019). Does forcing reduce faking? A meta-analytic review of forced-choice personality measures in high-stakes situations. *Journal of Applied Psychology*. https://doi.org/10.1037/apl0000414. **Verified:** DOI + quote. Ids: cao2019. Cited in: 03.

316. Nisbett RE, Wilson TD (1977). Telling more than we can know: Verbal reports on mental processes. *Psychological Review*. https://doi.org/10.1037/0033-295x.84.3.231. **Verified:** DOI + quote. Ids: nisbett1977. Cited in: 03.

317. Rogowsky BA, Calhoun BM, Tallal P (2015). Matching learning style to instructional method: Effects on comprehension. *Journal of Educational Psychology*. https://doi.org/10.1037/a0037478. **Verified:** DOI + quote. Ids: rogowsky2015. Cited in: 03, 05.

318. Bernacki ML, Walkington C (2018). The role of situational interest in personalized learning. *Journal of Educational Psychology*. https://doi.org/10.1037/edu0000250. **Verified:** DOI + quote. Ids: bernacki2018. Cited in: 03.

319. Cordova DI, Lepper MR (1996). Intrinsic motivation and the process of learning: Beneficial effects of contextualization, personalization, and choice. *Journal of Educational Psychology*. https://doi.org/10.1037/0022-0663.88.4.715. **Verified:** DOI + quote. Ids: cordova1996. Cited in: 03, 05.

320. Lomas JD, Koedinger K, Patel N, Shodhan S, Poonwala N, Forlizzi JL (2017). Is Difficulty Overrated?: The Effects of Choice, Novelty and Suspense on Intrinsic Motivation in Educational Games. *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*. https://doi.org/10.1145/3025453.3025638. **Verified:** DOI + quote. Ids: lomas2017. Cited in: 03.

321. Pekrun R (2006). The Control-Value Theory of Achievement Emotions: Assumptions, Corollaries, and Implications for Educational Research and Practice. *Educational Psychology Review*. https://doi.org/10.1007/s10648-006-9029-9. **Verified:** DOI + quote. Ids: pekrun2006. Cited in: 03.

322. Loderer K, Pekrun R, Lester JC (2020). Beyond cold technology: A systematic review and meta-analysis on emotions in technology-based learning environments. *Learning and Instruction*. https://doi.org/10.1016/j.learninstruc.2018.08.002. **Verified:** title-level only. Ids: loderer2020. Cited in: 03.

323. Ocumpaugh J, Baker R, Gowda S, Heffernan N, Heffernan C (2014). Population validity for educational data mining models: A case study in affect detection. *British Journal of Educational Technology*. https://doi.org/10.1111/bjet.12156. **Verified:** DOI + quote. Ids: ocumpaugh2014. Cited in: 03.

324. Kizilcec RF, Piech C, Schneider E (2013). Deconstructing disengagement: analyzing learner subpopulations in massive open online courses. *Proceedings of the Third International Conference on Learning Analytics and Knowledge*. https://doi.org/10.1145/2460296.2460330. **Verified:** DOI + quote. Ids: kizilcec2013. Cited in: 03.

325. Beck JE, Gong Y (2013). Wheel-Spinning: Students Who Fail to Master a Skill. *Lecture Notes in Computer Science*. https://doi.org/10.1007/978-3-642-39112-5_44. **Verified:** DOI + quote. Ids: beck2013. Cited in: 03.

326. Kai, Almeda et al. (2018). Decision Tree Modeling of Wheel- Spinning and Productive Persistence in Skill Builders. *Journal of Educational Data Mining*. https://doi.org/10.5281/zenodo.3344810. **Verified:** DOI + quote. Ids: kai2018. Cited in: 03.

327. Miller JB, Sanjurjo A (2018). Surprised by the Hot Hand Fallacy? A Truth in the Law of Small Numbers. *Econometrica*. https://doi.org/10.3982/ecta14943. **Verified:** DOI + quote. Ids: miller2018. Cited in: 03.

328. Wise SL, Pastor DA, Kong XJ (2009). Correlates of Rapid-Guessing Behavior in Low-Stakes Testing: Implications for Test Development and Measurement Practice. *Applied Measurement in Education*. https://doi.org/10.1080/08957340902754650. **Verified:** DOI + quote. Ids: wise2009. Cited in: 03.

329. Ryan AM, Pintrich PR, Midgley C (2001). Avoiding Seeking Help in the Classroom: Who and Why? *Educational Psychology Review*. https://doi.org/10.1023/a:1009013420053. **Verified:** DOI + quote. Ids: ryan2001. Cited in: 03.

330. Razzaq L, Heffernan NT (2010). Hints: Is It Better to Give or Wait to Be Asked? *Lecture Notes in Computer Science*. https://doi.org/10.1007/978-3-642-13388-6_39. **Verified:** title-level only. Ids: razzaq2010. Cited in: 03.

331. McLaren BM, van Gog T, Ganoe C, Karabinos M, Yaron D (2016). The efficiency of worked examples compared to erroneous examples, tutored problem solving, and problem solving in computer-based learning environments. *Computers in Human Behavior*. https://doi.org/10.1016/j.chb.2015.08.038. **Verified:** DOI + quote. Ids: mclaren2016. Cited in: 03.

332. van Dinther M, Dochy F, Segers M (2011). Factors affecting students' self-efficacy in higher education. *Educational Research Review*. https://doi.org/10.1016/j.edurev.2010.10.003. **Verified:** DOI + quote. Ids: vandinther2011. Cited in: 03.

333. Ecker UKH, Lewandowsky S, Cook J, Schmid P, Fazio LK, Brashier N, et al. (2022). The psychological drivers of misinformation belief and its resistance to correction. *Nature*

Reviews Psychology. https://doi.org/10.1038/s44159-021-00006-y. **Verified:** title-level only. Ids: ecker2022. Cited in: 03.

334. Bertsch S, Pesta BJ, Wiscott R, McDaniel MA (2007). The generation effect: A meta-analytic review. *Memory & Cognition*. https://doi.org/10.3758/bf03193441. **Verified:** DOI + quote. Ids: bertsch2007. Cited in: 03.
335. Callender AA, McDaniel MA (2009). The limited benefits of rereading educational texts. *Contemporary Educational Psychology*. https://doi.org/10.1016/j.cedpsych.2008.07.001. **Verified:** title-level only. Ids: callender2009. Cited in: 03.
336. Cohen SA (1987). Instructional Alignment: Searching for a Magic Bullet. *Educational Researcher*. https://doi.org/10.3102/0013189x016008016. **Verified:** DOI + quote. Ids: cohen1987. Cited in: 03.
337. Martone A, Sireci SG (2009). Evaluating Alignment Between Curriculum, Assessment, and Instruction. *Review of Educational Research*. https://doi.org/10.3102/0034654309341375. **Verified:** DOI + quote. Ids: martone2009. Cited in: 03.
338. Kong XJ, Wise SL, Bhola DS (2007). Setting the Response Time Threshold Parameter to Differentiate Solution Behavior From Rapid-Guessing Behavior. *Educational and Psychological Measurement*. https://doi.org/10.1177/0013164406294779. **Verified:** DOI + quote. Ids: kong2007. Cited in: 03.
339. Soland J, Kuhfeld M, Rios J (2021). Comparing different response time threshold setting methods to detect low effort on a large-scale assessment. *Large-scale Assessments in Education*. https://doi.org/10.1186/s40536-021-00100-w. **Verified:** DOI + quote. Ids: soland2021. Cited in: 03.
340. Lee YH, Jia Y (2014). Using response time to investigate students' test-taking behaviors in a NAEP computer-based study. *Large-scale Assessments in Education*. https://doi.org/10.1186/s40536-014-0008-1. **Verified:** DOI + quote. Ids: lee2014rt. Cited in: 03.
341. Rios JA, Guo H, Mao L, Liu OL (2017). Evaluating the Impact of Careless Responding on Aggregated-Scores: To Filter Unmotivated Examinees or Not? *International Journal of Testing*. https://doi.org/10.1080/15305058.2016.1231193. **Verified:** DOI + quote. Ids: rios2017. Cited in: 03.
342. Kluger AN, DeNisi A (1996). The effects of feedback interventions on performance: A historical review, a meta-analysis, and a preliminary feedback intervention theory. *Psychological Bulletin*. https://doi.org/10.1037/0033-2909.119.2.254. **Verified:** DOI + quote. Ids: kluger1996. Cited in: 03, 05.
343. Chi MT, Siler SA, Jeong H, Yamauchi T, Hausmann RG (2001). Learning from human tutoring. *Cognitive Science*. https://doi.org/10.1207/s15516709cog2504_1. **Verified:** DOI + quote. Ids: chi2001. Cited in: 03.
344. Graesser AC, Lu S, Jackson GT, Mitchell HH, Ventura M, Olney A, et al. (2004). AutoTutor: A tutor with dialogue in natural language. *Behavior Research Methods, Instruments, & Computers*. https://doi.org/10.3758/bf03195563. **Verified:** DOI + quote. Ids: graesser2004. Cited in: 03.
345. Nye BD, Graesser AC, Hu X (2014). AutoTutor and Family: A Review of 17 Years of Natural Language Tutoring. *International Journal of Artificial Intelligence in Education*.

https://doi.org/10.1007/s40593-014-0029-5. **Verified:** DOI + quote. Ids: nye2014. Cited in: 03.

346. Feng M, Heffernan N, Koedinger K (2009). Addressing the assessment challenge with an online system that tutors as it assesses. *User Modeling and User-Adapted Interaction*. https://doi.org/10.1007/s11257-009-9063-7. **Verified:** DOI + quote. Ids: feng2009. Cited in: 03.

347. Beck JE, Chang KM, Mostow J, Corbett A (2008). Does Help Help? Introducing the Bayesian Evaluation and Assessment Methodology. *Lecture Notes in Computer Science*. https://doi.org/10.1007/978-3-540-69132-7_42. **Verified:** DOI + quote. Ids: beck2008help. Cited in: 03.

348. Haladyna TM, Downing SM (2004). Construct-Irrelevant Variance in High-Stakes Testing. *Educational Measurement: Issues and Practice*. https://doi.org/10.1111/j.1745-3992.2004.tb00149.x. **Verified:** DOI + quote. Ids: haladyna2004. Cited in: 03.

349. Downing SM (2005). The Effects of Violating Standard Item Writing Principles on Tests and Students: The Consequences of Using Flawed Test Items on Achievement Examinations in Medical Education. *Advances in Health Sciences Education*. https://doi.org/10.1007/s10459-004-4019-5. **Verified:** DOI + quote. Ids: downing2005. Cited in: 03, 04.

350. Burrows S, Gurevych I, Stein B (2015). The Eras and Trends of Automatic Short Answer Grading. *International Journal of Artificial Intelligence in Education*. https://doi.org/10.1007/s40593-014-0026-8. **Verified:** DOI + quote. Ids: burrows2015. Cited in: 03.

351. National Governors Association Center for Best Practices; Council of Chief State School Officers (2010). CCSS Math, Grade 4, Number and Operations—Fractions (4.NF.A.1 and 4.NF.C.6). Official standard or regulation. https://www.thecorestandards.org/Math/Content/4/NF/. **Verified:** URL + quote. Ids: ccss-4nf-c6, ccss-4nf-a1. Cited in: 03.

352. National Governors Association Center for Best Practices; Council of Chief State School Officers (2010). CCSS Math, Grade 3, Number and Operations—Fractions (3.NF.A.3). Official standard or regulation. https://www.thecorestandards.org/Math/Content/3/NF/. **Verified:** URL + quote. Ids: ccss-3nf-a3. Cited in: 03.

353. Otheguy R, Zentella AC, Livert D (2007). Language and Dialect Contact in Spanish in New York: Toward the Formation of a Speech Community. *Language*. https://doi.org/10.1353/lan.2008.0019. **Verified:** DOI + quote. Ids: otheguy2007. Cited in: 03.

354. Real Academia Española (2005). Real Academia Española, Diccionario panhispánico de dudas: "tilde" (diacritic accent). Official standard or regulation. https://www.rae.es/dpd/tilde. **Verified:** URL + quote. Ids: rae-tilde. Cited in: 03.

355. Greshake K, Abdelnabi S, Mishra S, Endres C, Holz T, Fritz M (2023). Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection. *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*. https://doi.org/10.1145/3605764.3623985. **Verified:** DOI + quote. Ids: greshake2023. Cited in: 03.

356. Shi J, Yuan Z, Liu Y, Huang Y, Zhou P, Sun L, et al. (2024). Optimization-based Prompt Injection Attack to LLM-as-a-Judge. *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security*. https://doi.org/10.1145/3658644.3690291. **Verified:** DOI + quote. Ids: shi2024judge. Cited in: 03.

357. OWASP Foundation (2025). OWASP Top 10 for LLM Applications 2025: LLM01 Prompt Injection. Official standard or regulation. https://genai.owasp.org/llmrisk/llm01-prompt-injection/. **Verified:** URL + quote. Ids: owasp-llm01. Cited in: 03.

358. Zheng L, Chiang WL, Sheng Y, Zhuang S, Wu Z, Zhuang Y, et al. (2023). Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *Advances in Neural Information Processing Systems 36*. https://doi.org/10.52202/075280-2020. **Verified:** DOI + quote. Ids: zheng2023. Cited in: 03.

359. Tarrant M, Ware J (2008). Impact of item-writing flaws in multiple-choice questions on student achievement in high-stakes nursing assessments. *Medical Education*. https://doi.org/10.1111/j.1365-2923.2007.02957.x. **Verified:** DOI + quote. Ids: tarrant2008. Cited in: 03, 04.

360. Gierl MJ, Bulut O, Guo Q, Zhang X (2017). Developing, Analyzing, and Using Distractors for Multiple-Choice Tests in Education: A Comprehensive Review. *Review of Educational Research*. https://doi.org/10.3102/0034654317726529. **Verified:** DOI + quote. Ids: gierl2017. Cited in: 03.

361. Pardos ZA, Bhandari S (2024). ChatGPT-generated help produces learning gains equivalent to human tutor-authored help on mathematics skills. *Plos One*. https://doi.org/10.1371/journal.pone.0304013. **Verified:** DOI + quote. Ids: pardos2024. Cited in: 03.

362. National Governors Association Center for Best Practices; Council of Chief State School Officers (2010). Ccss Ela W.k.2. Official standard or regulation. https://www.thecorestandards.org/ELA-Literacy/W/K/. **Verified:** URL + quote. Ids: ccss-wk2. Cited in: 03.

363. Puranik CS, Lonigan CJ (2009). From scribbles to scrabble: preschool children's developing knowledge of written language. *Reading and Writing*. https://doi.org/10.1007/s11145-009-9220-8. **Verified:** DOI + quote. Ids: puranik2009. Cited in: 03.

364. Reed DK, Vaughn S (2011). Retell as an Indicator of Reading Comprehension. *Scientific Studies of Reading*. https://doi.org/10.1080/10888438.2010.538780. **Verified:** DOI + quote. Ids: reed2011. Cited in: 03.

365. Rittle-Johnson B (2006). Promoting Transfer: Effects of Self-Explanation and Direct Instruction. *Child Development*. https://doi.org/10.1111/j.1467-8624.2006.00852.x. **Verified:** DOI + quote. Ids: rittlejohnson2006. Cited in: 03.

366. FTC (2025). 16 CFR 312.2 (COPPA Rule): definition of personal information. Official standard or regulation. https://www.ecfr.gov/current/title-16/part-312/section-312.2. **Verified:** URL + quote. Ids: coppa-312.2. Cited in: 03, 05.

367. Livingstone, Stoilova & Nandagiri (2019). Children's data and privacy online: growing up in a digital age: an evidence review. *London School of Economics and Political Science*

Research Online (London School of Economics and Political Science). https://openalex.org/W2965426950. **Verified:** URL + quote. Ids: livingstone2019. Cited in: 03.

368. WHATWG (n.d.). WHATWG HTML Living Standard: the iframe sandbox attribute. Official standard or regulation. https://html.spec.whatwg.org/multipage/iframe-embed-object.html. **Verified:** URL + quote. Ids: whatwg-iframe-sandbox. Cited in: 03.

369. W3C (n.d.). W3C Content Security Policy Level 3. Official standard or regulation. https://www.w3.org/TR/CSP3/. **Verified:** URL + quote. Ids: w3c-csp3. Cited in: 03.

370. Weichselbaum L, Spagnuolo M, Lekies S, Janc A (2016). CSP Is Dead, Long Live CSP! On the Insecurity of Whitelists and the Future of Content Security Policy. *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*. https://doi.org/10.1145/2976749.2978363. **Verified:** DOI + quote. Ids: weichselbaum2016. Cited in: 03.

371. Mueller CM, Dweck CS (1998). Praise for intelligence can undermine children's motivation and performance. *Journal of Personality and Social Psychology*. https://doi.org/10.1037/0022-3514.75.1.33. **Verified:** DOI + quote. Ids: mueller1998. Cited in: 03, 05.

372. Kantak SS, Winstein CJ (2012). Learning–performance distinction and memory processes for motor skills: A focused review and perspective. *Behavioural Brain Research*. https://doi.org/10.1016/j.bbr.2011.11.028. **Verified:** DOI + quote. Ids: kantak2012. Cited in: 04.

373. Schmidt RA, Bjork RA (1992). New Conceptualizations of Practice: Common Principles in Three Paradigms Suggest New Concepts for Training. *Psychological Science*. https://doi.org/10.1111/j.1467-9280.1992.tb00029.x. **Verified:** DOI + quote. Ids: schmidt1992. Cited in: 04.

374. Lee TD, Genovese ED (1988). Distribution of Practice in Motor Skill Acquisition: Learning and Performance Effects Reconsidered. *Research Quarterly for Exercise and Sport*. https://doi.org/10.1080/02701367.1988.10609373. **Verified:** DOI + quote. Ids: lee1988. Cited in: 04.

375. Adams JA (1961). The second facet of forgetting: A review of warm-up decrement. *Psychological Bulletin*. https://doi.org/10.1037/h0044798. **Verified:** title-level only. Ids: adams1961. Cited in: 04.

376. Walker MP, Brakefield T, Morgan A, Hobson J, Stickgold R (2002). Practice with Sleep Makes Perfect. *Neuron*. https://doi.org/10.1016/s0896-6273(02)00746-8. **Verified:** DOI + quote. Ids: walker2002. Cited in: 04.

377. Simmons AL, Duke RA (2006). Effects of Sleep on Performance of a Keyboard Melody. *Journal of Research in Music Education*. https://doi.org/10.1177/002242940605400308. **Verified:** DOI + quote. Ids: simmons2006. Cited in: 04.

378. Abushanab B, Bishara AJ (2013). Memory and metacognition for piano melodies: Illusory advantages of fixed- over random-order practice. *Memory & Cognition*. https://doi.org/10.3758/s13421-013-0311-z. **Verified:** DOI + quote. Ids: abushanab2013. Cited in: 04.

379. Rickard TC, Cai DJ, Rieth CA, Jones J, Ard MC (2008). Sleep does not enhance motor sequence learning. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/0278-7393.34.4.834. **Verified:** DOI + quote. Ids: rickard2008. Cited in: 04.

380. Brashers-Krug T, Shadmehr R, Bizzi E (1996). Consolidation in human motor memory. *Nature*. https://doi.org/10.1038/382252a0. **Verified:** DOI + quote. Ids: brashers1996. Cited in: 04.

381. Gelman A, Carlin JB, Stern HS, Dunson DB, Vehtari A, Rubin DB (2013). Bayesian Data Analysis. *Chapman and Hall/CRC*. https://doi.org/10.1201/b16018. **Verified:** DOI + quote. Ids: gelman2013. Cited in: 04.

382. Brown LD, Cai TT, DasGupta A (2001). Interval Estimation for a Binomial Proportion. *Statistical Science*. https://doi.org/10.1214/ss/1009213286. **Verified:** DOI + quote. Ids: brown2001. Cited in: 04.

383. Hopkins WG (2000). Measures of Reliability in Sports Medicine and Science. *Sports Medicine*. https://doi.org/10.2165/00007256-200030010-00001. **Verified:** DOI + quote. Ids: hopkins2000. Cited in: 04.

384. Wald A (1945). Sequential Tests of Statistical Hypotheses. *The Annals of Mathematical Statistics*. https://doi.org/10.1214/aoms/1177731118. **Verified:** title-level only. Ids: wald1945. Cited in: 04.

385. de Cheveigné A, Kawahara H (2002). YIN, a fundamental frequency estimator for speech and music. *The Journal of the Acoustical Society of America*. https://doi.org/10.1121/1.1458024. **Verified:** DOI + quote. Ids: decheveigne2002. Cited in: 04.

386. Emiya V, Badeau R, David B (2010). Multipitch Estimation of Piano Sounds Using a New Probabilistic Spectral Smoothness Principle. *IEEE Transactions on Audio, Speech, and Language Processing*. https://doi.org/10.1109/tasl.2009.2038819. **Verified:** DOI + quote. Ids: emiya2010. Cited in: 04.

387. Kong Q, Li B, Song X, Wan Y, Wang Y (2021). High-Resolution Piano Transcription With Pedals by Regressing Onset and Offset Times. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*. https://doi.org/10.1109/taslp.2021.3121991. **Verified:** DOI + quote. Ids: kong2021. Cited in: 04.

388. Washabaugh EP, Shanmugam TA, Ranganathan R, Krishnan C (2022). Comparing the accuracy of open-source pose estimation methods for measuring gait kinematics. *Gait & Posture*. https://doi.org/10.1016/j.gaitpost.2022.08.008. **Verified:** DOI + quote. Ids: washabaugh2022. Cited in: 04.

389. Stenum J, Rossi C, Roemmich RT (2021). Two-dimensional video-based analysis of human gait using pose estimation. *PLOS Computational Biology*. https://doi.org/10.1371/journal.pcbi.1008935. **Verified:** DOI + quote. Ids: stenum2021. Cited in: 04.

390. Benetos E, Dixon S, Giannoulis D, Kirchhoff H, Klapuri A (2013). Automatic music transcription: challenges and future directions. *Journal of Intelligent Information Systems*. https://doi.org/10.1007/s10844-013-0258-3. **Verified:** title-level only. Ids: benetos2013. Cited in: 04.

391. Kruger J, Dunning D (1999). Unskilled and unaware of it: How difficulties in recognizing one's own incompetence lead to inflated self-assessments. *Journal of Personality and Social Psychology*. https://doi.org/10.1037/0022-3514.77.6.1121. **Verified:** DOI + quote. Ids: kruger1999. Cited in: 04.

392. Bergee MJ (1997). Relationships among Faculty, Peer, and Self-Evaluations of Applied Performances. *Journal of Research in Music Education*. https://doi.org/10.2307/3345425. **Verified:** DOI + quote. Ids: bergee1997. Cited in: 04.

393. Hewitt MP (2001). The Effects of Modeling, Self-Evaluation, and Self-Listening on Junior High Instrumentalists' Music Performance and Practice Attitude. *Journal of Research in Music Education*. https://doi.org/10.2307/3345614. **Verified:** DOI + quote. Ids: hewitt2001. Cited in: 04.

394. Andrade HL (2019). A Critical Review of Research on Student Self-Assessment. *Frontiers in Education*. https://doi.org/10.3389/feduc.2019.00087. **Verified:** DOI + quote. Ids: andrade2019. Cited in: 04.

395. Arthur Jr. W, Bennett Jr. W, Stanush PL, McNelly TL (1998). Factors That Influence Skill Decay and Retention: A Quantitative Review and Analysis. *Human Performance*. https://doi.org/10.1207/s15327043hup1101_3. **Verified:** DOI + quote. Ids: arthur1998. Cited in: 04.

396. Kim JW, Ritter FE, Koubek RJ (2013). An integrated theory for improved skill acquisition and retention in the three stages of learning. *Theoretical Issues in Ergonomics Science*. https://doi.org/10.1080/1464536x.2011.573008. **Verified:** DOI + quote. Ids: kim2013. Cited in: 04.

397. Downing SM, Haladyna TM (1997). Test Item Development: Validity Evidence From Quality Assurance Procedures. *Applied Measurement in Education*. https://doi.org/10.1207/s15324818ame1001_4. **Verified:** DOI + quote. Ids: downing1997. Cited in: 04.

398. Rösch D, Schultz F, Höner O (2021). Decision-Making Skills in Youth Basketball Players: Diagnostic and External Validation of a Video-Based Assessment. *International Journal of Environmental Research and Public Health*. https://doi.org/10.3390/ijerph18052331. **Verified:** DOI + quote. Ids: rosch2021. Cited in: 04.

399. Cortina JM (1993). What is coefficient alpha? An examination of theory and applications. *Journal of Applied Psychology*. https://doi.org/10.1037/0021-9010.78.1.98. **Verified:** DOI + quote. Ids: cortina1993. Cited in: 04.

400. Duke RA, Simmons AL, Cash CD (2009). It's Not How Much; It's How. *Journal of Research in Music Education*. https://doi.org/10.1177/0022429408328851. **Verified:** DOI + quote. Ids: duke2009. Cited in: 04.

401. Platz F, Kopiez R, Lehmann AC, Wolf A (2014). The influence of deliberate practice on musical achievement: a meta-analysis. *Frontiers in Psychology*. https://doi.org/10.3389/fpsyg.2014.00646. **Verified:** DOI + quote. Ids: platz2014. Cited in: 04.

402. Driskell JE, Willis RP, Copper C (1992). Effect of overlearning on retention. *Journal of Applied Psychology*. https://doi.org/10.1037/0021-9010.77.5.615. **Verified:** DOI + quote. Ids: driskell1992. Cited in: 04.

403. Shibata K, Sasaki Y, Bang JW, Walsh EG, Machizawa MG, Tamaki M, et al. (2017). Overlearning hyperstabilizes a skill by rapidly making neurochemical processing inhibitory-dominant. *Nature Neuroscience*. https://doi.org/10.1038/nn.4490. **Verified:** DOI + quote. Ids: shibata2017. Cited in: 04.

404. Wiseheart M, D'Souza AA, Chae J (2017). Lack of spacing effects during piano learning. *Plos One*. https://doi.org/10.1371/journal.pone.0182986. **Verified:** DOI + quote. Ids: wiseheart2017. Cited in: 04.

405. Charness N, Tuffiash M, Krampe R, Reingold E, Vasyukova E (2005). The role of deliberate practice in chess expertise. *Applied Cognitive Psychology*. https://doi.org/10.1002/acp.1106. **Verified:** DOI + quote. Ids: charness2005. Cited in: 04.

406. Macnamara BN, Hambrick DZ, Oswald FL (2014). Deliberate Practice and Performance in Music, Games, Sports, Education, and Professions: A Meta-Analysis. *Psychological Science*. https://doi.org/10.1177/0956797614535810. **Verified:** DOI + quote. Ids: macnamara2014ps. Cited in: 04.

407. Macnamara BN, Moreau D, Hambrick DZ (2016). The Relationship Between Deliberate Practice and Performance in Sports. *Perspectives on Psychological Science*. https://doi.org/10.1177/1745691616635591. **Verified:** DOI + quote. Ids: macnamara2016. Cited in: 04.

408. Bergeron MF, Mountjoy M, Armstrong N, Chia M, Côté J, Emery CA, et al. (2015). International Olympic Committee consensus statement on youth athletic development. *British Journal of Sports Medicine*. https://doi.org/10.1136/bjsports-2015-094962. **Verified:** DOI + quote. Ids: bergeron2015. Cited in: 04.

409. DiFiori JP, Güllich A, Brenner JS, Côté J, Hainline B, Ryan E, et al. (2018). The NBA and Youth Basketball: Recommendations for Promoting a Healthy and Positive Experience. *Sports Medicine*. https://doi.org/10.1007/s40279-018-0950-0. **Verified:** DOI + quote. Ids: difiori2018. Cited in: 04.

410. International Test Commission (ITC) (2015). The International Test Commission Guidelines on the Security of Tests, Examinations, and Other Assessments. *International Journal of Testing*. https://doi.org/10.1080/15305058.2015.1111221. **Verified:** DOI + quote. Ids: itc2015security. Cited in: 04.

411. Gobet F, Campitelli G (2007). The role of domain-specific practice, handedness, and starting age in chess. *Developmental Psychology*. https://doi.org/10.1037/0012-1649.43.1.159. **Verified:** DOI + quote. Ids: gobet2007. Cited in: 04.

412. Bilalić M, McLeod P, Gobet F (2008). Why good thoughts block better ones: The mechanism of the pernicious Einstellung (set) effect. *Cognition*. https://doi.org/10.1016/j.cognition.2008.05.005. **Verified:** DOI + quote. Ids: bilalic2008. Cited in: 04.

413. Van Der Maas HLJ, Wagenmakers EJ (2005). A Psychometric Analysis of Chess Expertise. *The American Journal of Psychology*. https://doi.org/10.2307/30039042. **Verified:** DOI + quote. Ids: vandermaas2005. Cited in: 04.

414. Burns BD (2004). The Effects of Speed on Skilled Chess Performance. *Psychological Science*. https://doi.org/10.1111/j.0956-7976.2004.00699.x. **Verified:** DOI + quote. Ids: burns2004. Cited in: 04.

415. Glickman ME (1999). Parameter Estimation in Large Dynamic Paired Comparison Experiments. *Journal of the Royal Statistical Society Series C: Applied Statistics*. https://doi.org/10.1111/1467-9876.00159. **Verified:** DOI + quote. Ids: glickman1999. Cited in: 04.

416. Glickman ME (2001). Dynamic paired comparison models with stochastic variances. *Journal of Applied Statistics*. https://doi.org/10.1080/02664760120059219. **Verified:** DOI + quote. Ids: glickman2001. Cited in: 04.

417. Eggen TJHM, Verschoor AJ (2006). Optimal Testing With Easy or Difficult Items in Computerized Adaptive Testing. *Applied Psychological Measurement*. https://doi.org/10.1177/0146621606288890. **Verified:** DOI + quote. Ids: eggen2006. Cited in: 04.

418. Ash DW, Holding DH (1990). Backward versus Forward Chaining in the Acquisition of a Keyboard Skill. *Human Factors: The Journal of the Human Factors and Ergonomics Society*. https://doi.org/10.1177/001872089003200202. **Verified:** DOI + quote. Ids: ash1990. Cited in: 04.

419. Pinder RA, Davids K, Renshaw I, Araújo D (2011). Representative Learning Design and Functionality of Research and Practice in Sport. *Journal of Sport and Exercise Psychology*. https://doi.org/10.1123/jsep.33.1.146. **Verified:** DOI + quote. Ids: pinder2011. Cited in: 04.

420. Thompson K (1986). Retrograde Analysis of Certain Endgames. *ICGA Journal*. https://doi.org/10.3233/icg-1986-9302. **Verified:** title-level only. Ids: thompson1986. Cited in: 04.

421. Bourzutschky M, Tamplin J, Haworth G (2005). Chess endgames: 6-man data and strategy. *Theoretical Computer Science*. https://doi.org/10.1016/j.tcs.2005.09.043. **Verified:** title-level only. Ids: bourzutschky2005. Cited in: 04.

422. Regan K, Haworth G (2011). Intrinsic Chess Ratings. *Proceedings of the AAAI Conference on Artificial Intelligence*. https://doi.org/10.1609/aaai.v25i1.7951. **Verified:** DOI + quote. Ids: regan2011. Cited in: 04.

423. McIlroy-Young R, Sen S, Kleinberg J, Anderson A (2020). Aligning Superhuman AI with Human Behavior. *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*. https://doi.org/10.1145/3394486.3403219. **Verified:** DOI + quote. Ids: mcilroy2020. Cited in: 04.

424. Guid M, Bratko I (2006). Computer Analysis Of World Chess Champions1. *ICGA Journal*. https://doi.org/10.3233/icg-2006-29203. **Verified:** title-level only. Ids: guid2006. Cited in: 04.

425. Gobet F, Simon HA (1996). Templates in Chess Memory: A Mechanism for Recalling Several Boards. *Cognitive Psychology*. https://doi.org/10.1006/cogp.1996.0011. **Verified:** DOI + quote. Ids: gobet1996. Cited in: 04.

426. Chase WG, Simon HA (1973). Perception in chess. *Cognitive Psychology*. https://doi.org/10.1016/0010-0285(73)90004-2. **Verified:** title-level only. Ids: chase1973. Cited in: 04.

427. de Bruin AB, Rikers RM, Schmidt HG (2007). The effect of self-explanation and prediction on the development of principled understanding of chess in novices. *Contemporary*

Educational Psychology. https://doi.org/10.1016/j.cedpsych.2006.01.001. **Verified:** title-level only. Ids: debruin2007. Cited in: 04.

428. Repp BH (2005). Sensorimotor synchronization: A review of the tapping literature. *Psychonomic Bulletin & Review*. https://doi.org/10.3758/bf03206433. **Verified:** DOI + quote. Ids: repp2005. Cited in: 04.
429. Repp BH, Su YH (2013). Sensorimotor synchronization: A review of recent research (2006–2012). *Psychonomic Bulletin & Review*. https://doi.org/10.3758/s13423-012-0371-2. **Verified:** DOI + quote. Ids: repp2013. Cited in: 04.
430. Friberg A, Sundberg J (1995). Time discrimination in a monotonic, isochronous sequence. *The Journal of the Acoustical Society of America*. https://doi.org/10.1121/1.413218. **Verified:** DOI + quote. Ids: friberg1995. Cited in: 04.
431. Dannenberg RB, Raphael C (2006). Music score alignment and computer accompaniment. *Communications of the ACM*. https://doi.org/10.1145/1145287.1145311. **Verified:** DOI + quote. Ids: dannenberg2006. Cited in: 04.
432. Jabusch H, Vauth H, Altenmüller E (2004). Quantification of focal dystonia in pianists using scale analysis. *Movement Disorders*. https://doi.org/10.1002/mds.10671. **Verified:** DOI + quote. Ids: jabusch2004. Cited in: 04.
433. Van Vugt FT, Jabusch HC, Altenmüller E (2013). Individuality That is Unheard of: Systematic Temporal Deviations in Scale Playing Leave an Inaudible Pianistic Fingerprint. *Frontiers in Psychology*. https://doi.org/10.3389/fpsyg.2013.00134. **Verified:** DOI + quote. Ids: vanvugt2013. Cited in: 04.
434. Jabusch HC, Alpers H, Kopiez R, Vauth H, Altenmüller E (2009). The influence of practice on the development of motor skills in pianists: A longitudinal study in a selected motor task. *Human Movement Science*. https://doi.org/10.1016/j.humov.2008.08.001. **Verified:** DOI + quote. Ids: jabusch2009. Cited in: 04.
435. Brady F (2004). Contextual Interference: A Meta-Analytic Study. *Perceptual and Motor Skills*. https://doi.org/10.2466/pms.99.1.116-126. **Verified:** DOI + quote. Ids: brady2004. Cited in: 04.
436. Wulf G, Shea CH (2002). Principles derived from the study of simple skills do not generalize to complex skill learning. *Psychonomic Bulletin & Review*. https://doi.org/10.3758/bf03196276. **Verified:** DOI + quote. Ids: wulf2002. Cited in: 04.
437. Magill RA, Hall KG (1990). A review of the contextual interference effect in motor skill acquisition. *Human Movement Science*. https://doi.org/10.1016/0167-9457(90)90005-x. **Verified:** title-level only. Ids: magill1990. Cited in: 04.
438. Mishra J (2014). Improving sightreading accuracy: A meta-analysis. *Psychology of Music*. https://doi.org/10.1177/0305735612463770. **Verified:** DOI + quote. Ids: mishra2014a. Cited in: 04.
439. Mishra J (2014). Factors Related to Sight-Reading Accuracy. *Journal of Research in Music Education*. https://doi.org/10.1177/0022429413508585. **Verified:** DOI + quote. Ids: mishra2014b. Cited in: 04.
440. Kopiez R, In Lee J (2006). Towards a dynamic model of skills involved in sight reading music. *Music Education Research*. https://doi.org/10.1080/14613800600570785. **Verified:** DOI + quote. Ids: kopiez2006. Cited in: 04.

441. McPherson GE, Bailey M, Sinclair KE (1997). Path Analysis of a Theoretical Model to Describe the Relationship among Five Types of Musical Performance. *Journal of Research in Music Education*. https://doi.org/10.2307/3345469. **Verified:** DOI + quote. Ids: mcpherson1997. Cited in: 04.

442. Lehmann AC, Ericsson KA (1996). Performance without preparation: Structure and acquisition of expert sight-reading and accompanying performance. *Psychomusicology: A Journal of Research in Music Cognition*. https://doi.org/10.1037/h0094082. **Verified:** title-level only. Ids: lehmann1996. Cited in: 04.

443. Chaffin R, Imreh G (2002). Practicing Perfection: Piano Performance as Expert Memory. *Psychological Science*. https://doi.org/10.1111/1467-9280.00462. **Verified:** DOI + quote. Ids: chaffin2002. Cited in: 04.

444. Williamon A, Valentine E (2002). The Role of Retrieval Structures in Memorizing Music. *Cognitive Psychology*. https://doi.org/10.1006/cogp.2001.0759. **Verified:** DOI + quote. Ids: williamon2002. Cited in: 04.

445. Chaffin R, Lisboa T, Logan T, Begosh KT (2010). Preparing for memorized cello performance: the role of performance cues. *Psychology of Music*. https://doi.org/10.1177/0305735608100377. **Verified:** DOI + quote. Ids: chaffin2010. Cited in: 04.

446. Goebl W (2001). Melody lead in piano performance: Expressive device or artifact? *The Journal of the Acoustical Society of America*. https://doi.org/10.1121/1.1376133. **Verified:** DOI + quote. Ids: goebl2001. Cited in: 04.

447. Lerch A, Arthur C, Pati A, Gururani S (2020). An Interdisciplinary Review of Music Performance Analysis. *Transactions of the International Society for Music Information Retrieval*. https://doi.org/10.5334/tismir.53. **Verified:** DOI + quote. Ids: lerch2020. Cited in: 04.

448. Repp BH (1996). Pedal Timing and Tempo in Expressive Piano Performance: A Preliminary Investigation. *Psychology of Music*. https://doi.org/10.1177/0305735696242011. **Verified:** DOI + quote. Ids: repp1996. Cited in: 04.

449. Maxwell J, Masters R, Kerr E, Weedon E (2001). The implicit benefit of learning without errors. *The Quarterly Journal of Experimental Psychology Section A*. https://doi.org/10.1080/713756014. **Verified:** DOI + quote. Ids: maxwell2001. Cited in: 04.

450. Fitts PM (1954). The information capacity of the human motor system in controlling the amplitude of movement. *Journal of Experimental Psychology*. https://doi.org/10.1037/h0055392. **Verified:** title-level only. Ids: fitts1954. Cited in: 04.

451. Potdevin F, Vors O, Huchez A, Lamour M, Davids K, Schnitzler C (2018). How can video feedback be used in physical education to support novice learning in gymnastics? Effects on motor learning, self-assessment and motivation. *Physical Education and Sport Pedagogy*. https://doi.org/10.1080/17408989.2018.1485138. **Verified:** DOI + quote. Ids: potdevin2018. Cited in: 04.

452. McPherson GE (2005). From child to musician: skill development during the beginning stages of learning an instrument. *Psychology of Music*.

https://doi.org/10.1177/0305735605048012. **Verified:** DOI + quote. Ids: mcpherson2005. Cited in: 04.

453. Woody RH, Lehmann AC (2010). Student Musicians' Ear-Playing Ability as a Function of Vernacular Music Experiences. *Journal of Research in Music Education*. https://doi.org/10.1177/0022429410370785. **Verified:** DOI + quote. Ids: woody2010. Cited in: 04.

454. Drake C, Jones MR, Baruch C (2000). The development of rhythmic attending in auditory sequences: attunement, referent period, focal attending. *Cognition*. https://doi.org/10.1016/s0010-0277(00)00106-2. **Verified:** DOI + quote. Ids: drake2000. Cited in: 04.

455. Provasi J, Bobin-Bègue A (2003). Spontaneous motor tempo and rhythmical synchronisation in 2½- and 4-year-old children. *International Journal of Behavioral Development*. https://doi.org/10.1080/01650250244000290. **Verified:** DOI + quote. Ids: provasi2003. Cited in: 04.

456. Hallam S, Rinta T, Varvarigou M, Creech A, Papageorgi I, Gomes T, et al. (2012). The development of practising strategies in young people. *Psychology of Music*. https://doi.org/10.1177/0305735612443868. **Verified:** DOI + quote. Ids: hallam2012. Cited in: 04.

457. Ste-Marie DM, Law B, Rymal AM, Jenny O, Hall C, McCullagh P (2012). Observation interventions for motor skill learning and performance: an applied model for the use of observation. *International Review of Sport and Exercise Psychology*. https://doi.org/10.1080/1750984x.2012.665076. **Verified:** DOI + quote. Ids: stemarie2012. Cited in: 04.

458. Wulf G, Shea C, Lewthwaite R (2010). Motor skill learning and performance: a review of influential factors. *Medical Education*. https://doi.org/10.1111/j.1365-2923.2009.03421.x. **Verified:** DOI + quote. Ids: wulf2010. Cited in: 04.

459. Agresti A, Coull BA (1998). Approximate is Better than "Exact" for Interval Estimation of Binomial Proportions. *The American Statistician*. https://doi.org/10.1080/00031305.1998.10480550. **Verified:** DOI + quote. Ids: agresti1998. Cited in: 04.

460. FIBA (2026). Official Basketball Rules 2026 v1.1. *FIBA (official standard)*. https://assets.fiba.basketball/image/upload/documents-corporate-fiba-official-rules-2026-v1-1.pdf. **Verified:** URL + quote. Ids: fiba2026rules. Cited in: 04.

461. Kubatko J, Oliver D, Pelton K, Rosenbaum DT (2007). A Starting Point for Analyzing Basketball Statistics. *Journal of Quantitative Analysis in Sports*. https://doi.org/10.2202/1559-0410.1070. **Verified:** DOI + quote. Ids: kubatko2007. Cited in: 04.

462. Gilovich T, Vallone R, Tversky A (1985). The hot hand in basketball: On the misperception of random sequences. *Cognitive Psychology*. https://doi.org/10.1016/0010-0285(85)90010-6. **Verified:** title-level only. Ids: gilovich1985. Cited in: 04.

463. Cobley S, Baker J, Wattie N, McKenna J (2009). Annual Age-Grouping and Athlete Development. *Sports Medicine*. https://doi.org/10.2165/00007256-200939030-00005. **Verified:** DOI + quote. Ids: cobley2009. Cited in: 04.

464. Lloyd RS, Oliver JL (2012). The Youth Physical Development Model. *Strength & Conditioning Journal*. https://doi.org/10.1519/ssc.0b013e31825760ea. **Verified:** DOI + quote. Ids: lloyd2012. Cited in: 04.

465. Crane J, Temple V (2015). A systematic review of dropout from organized sport among children and youth. *European Physical Education Review*. https://doi.org/10.1177/1356336x14555294. **Verified:** DOI + quote. Ids: crane2015. Cited in: 04.

466. Mann DT, Williams AM, Ward P, Janelle CM (2007). Perceptual-Cognitive Expertise in Sport: A Meta-Analysis. *Journal of Sport and Exercise Psychology*. https://doi.org/10.1123/jsep.29.4.457. **Verified:** DOI + quote. Ids: mann2007. Cited in: 04.

467. Zhu R, Zheng M, Liu S, Guo J, Cao C (2024). Effects of Perceptual-Cognitive Training on Anticipation and Decision-Making Skills in Team Sports: A Systematic Review and Meta-Analysis. *Behavioral Sciences*. https://doi.org/10.3390/bs14100919. **Verified:** DOI + quote. Ids: zhu2024. Cited in: 04.

468. Broadbent DP, Causer J, Williams AM, Ford PR (2015). Perceptual-cognitive skill training and its transfer to expert performance in the field: Future research directions. *European Journal of Sport Science*. https://doi.org/10.1080/17461391.2014.957727. **Verified:** DOI + quote. Ids: broadbent2015. Cited in: 04.

469. Hadlow SM, Panchuk D, Mann DL, Portus MR, Abernethy B (2018). Modified perceptual training in sport: A new classification framework. *Journal of Science and Medicine in Sport*. https://doi.org/10.1016/j.jsams.2018.01.011. **Verified:** DOI + quote. Ids: hadlow2018. Cited in: 04.

470. Abad Robles MT, Collado-Mateo D, Fernández-Espínola C, Castillo Viera E, Giménez Fuentes-Guerra FJ (2020). Effects of Teaching Games on Decision Making and Skill Execution: A Systematic Review and Meta-Analysis. *International Journal of Environmental Research and Public Health*. https://doi.org/10.3390/ijerph17020505. **Verified:** DOI + quote. Ids: abadrobles2020. Cited in: 04.

471. Harvey S, Jarrett K (2014). A review of the game-centred approaches to teaching and coaching literature since 2006. *Physical Education and Sport Pedagogy*. https://doi.org/10.1080/17408989.2012.754005. **Verified:** DOI + quote. Ids: harvey2014. Cited in: 04.

472. Bangert-Drowns RL, Kulik CLC, Kulik JA, Morgan M (1991). The Instructional Effect of Feedback in Test-Like Events. *Review of Educational Research*. https://doi.org/10.3102/00346543061002213. **Verified:** DOI + quote. Ids: bangert1991. Cited in: 04.

473. Cao Z, Hidalgo G, Simon T, Wei SE, Sheikh Y (2021). OpenPose: Realtime Multi-Person 2D Pose Estimation Using Part Affinity Fields. *IEEE Transactions on Pattern Analysis and Machine Intelligence*. https://doi.org/10.1109/tpami.2019.2929257. **Verified:** DOI + quote. Ids: cao2021. Cited in: 04.

474. Hartley R, Zisserman A (2004). Multiple View Geometry in Computer Vision. *Cambridge University Press*. https://doi.org/10.1017/cbo9780511811685. **Verified:** DOI + quote. Ids: hartley2004. Cited in: 04.

475. Mödinger M, Woll A, Wagner I (2022). Video-based visual feedback to enhance motor learning in physical education—a systematic review. *German Journal of Exercise and Sport Research*. https://doi.org/10.1007/s12662-021-00782-y. **Verified:** DOI + quote. Ids: modinger2022. Cited in: 04.

476. Chua LK, Jimenez-Diaz J, Lewthwaite R, Kim T, Wulf G (2021). Superiority of external attentional focus for motor performance and learning: Systematic reviews and meta-analyses. *Psychological Bulletin*. https://doi.org/10.1037/bul0000335. **Verified:** DOI + quote. Ids: chua2021. Cited in: 04.

477. Wulf G (2013). Attentional focus and motor learning: a review of 15 years. *International Review of Sport and Exercise Psychology*. https://doi.org/10.1080/1750984x.2012.723728. **Verified:** DOI + quote. Ids: wulf2013. Cited in: 04.

478. Wulf G, Lewthwaite R (2016). Optimizing performance through intrinsic motivation and attention for learning: The OPTIMAL theory of motor learning. *Psychonomic Bulletin & Review*. https://doi.org/10.3758/s13423-015-0999-9. **Verified:** DOI + quote. Ids: wulf2016. Cited in: 04.

479. Tran CM, Silverberg LM (2008). Optimal release conditions for the free throw in men's basketball. *Journal of Sports Sciences*. https://doi.org/10.1080/02640410802004948. **Verified:** DOI + quote. Ids: tran2008. Cited in: 04.

480. Soligard T, Myklebust G, Steffen K, Holme I, Silvers H, Bizzini M, et al. (2008). Comprehensive warm-up programme to prevent injuries in young female footballers: cluster randomised controlled trial. *Bmj*. https://doi.org/10.1136/bmj.a2469. **Verified:** DOI + quote. Ids: soligard2008. Cited in: 04.

481. Longo UG, Loppini M, Berton A, Marinozzi A, Maffulli N, Denaro V (2012). The FIFA 11+ Program Is Effective in Preventing Injuries in Elite Male Basketball Players. *The American Journal of Sports Medicine*. https://doi.org/10.1177/0363546512438761. **Verified:** DOI + quote. Ids: longo2012. Cited in: 04.

482. Emery CA, Roy TO, Whittaker JL, Nettel-Aguirre A, van Mechelen W (2015). Neuromuscular training injury prevention strategies in youth sport: a systematic review and meta-analysis. *British Journal of Sports Medicine*. https://doi.org/10.1136/bjsports-2015-094639. **Verified:** DOI + quote. Ids: emery2015. Cited in: 04.

483. DiFiori JP, Benjamin HJ, Brenner J, Gregory A, Jayanthi N, Landry GL, et al. (2014). Overuse Injuries and Burnout in Youth Sports. *Clinical Journal of Sport Medicine*. https://doi.org/10.1097/jsm.0000000000000060. **Verified:** DOI + quote. Ids: difiori2014. Cited in: 04.

484. Renshaw I, Chow JY, Davids K, Hammond J (2010). A constraints-led perspective to understanding skill acquisition and game play: a basis for integration of motor learning theory and physical education praxis? *Physical Education and Sport Pedagogy*. https://doi.org/10.1080/17408980902791586. **Verified:** DOI + quote. Ids: renshaw2010. Cited in: 04.

485. Haladyna TM, Downing SM (1989). Validity of a Taxonomy of Multiple-Choice Item-Writing Rules. *Applied Measurement in Education*. https://doi.org/10.1207/s15324818ame0201_4. **Verified:** DOI + quote. Ids: haladyna1989. Cited in: 04.

486. Attali Y, Bar-Hillel M (2003). Guess Where: The Position of Correct Answers in Multiple-Choice Test Items as a Psychometric Variable. *Journal of Educational Measurement*. https://doi.org/10.1111/j.1745-3984.2003.tb01099.x. **Verified:** DOI + quote. Ids: attali2003. Cited in: 04.

487. Millman J, Bishop CH, Ebel R (1965). An Analysis of Test-Wiseness. *Educational and Psychological Measurement*. https://doi.org/10.1177/001316446502500304. **Verified:** title-level only. Ids: millman1965. Cited in: 04.

488. Wisniewski B, Zierer K, Hattie J (2020). The Power of Feedback Revisited: A Meta-Analysis of Educational Feedback Research. *Frontiers in Psychology*. https://doi.org/10.3389/fpsyg.2019.03087. **Verified:** DOI + quote. Ids: wisniewski2020. Cited in: 04, 05.

489. Leary LF, Dorans NJ (1985). Implications for Altering the Context in Which Test Items Appear: A Historical Perspective on an Immediate Concern. *Review of Educational Research*. https://doi.org/10.3102/00346543055003387. **Verified:** DOI + quote. Ids: leary1985. Cited in: 04.

490. Hamari J, Shernoff DJ, Rowe E, Coller B, Asbell-Clarke J, Edwards T (2016). Challenging games help students learn: An empirical study on engagement, flow and immersion in game-based learning. *Computers in Human Behavior*. https://doi.org/10.1016/j.chb.2015.07.045. **Verified:** title-level only. Ids: hamari2016. Cited in: 04.

491. Engeser S, Rheinberg F (2008). Flow, performance and moderators of challenge-skill balance. *Motivation and Emotion*. https://doi.org/10.1007/s11031-008-9102-4. **Verified:** title-level only. Ids: engeser2008. Cited in: 04.

492. Biggs J (1996). Enhancing teaching through constructive alignment. *Higher Education*. https://doi.org/10.1007/bf00138871. **Verified:** title-level only. Ids: biggs1996. Cited in: 04.

493. Wouters P, van Nimwegen C, van Oostendorp H, van der Spek ED (2013). A meta-analysis of the cognitive and motivational effects of serious games. *Journal of Educational Psychology*. https://doi.org/10.1037/a0031311. **Verified:** DOI + quote. Ids: wouters2013. Cited in: 04.

494. Clark DB, Tanner-Smith EE, Killingsworth SS (2016). Digital Games, Design, and Learning. *Review of Educational Research*. https://doi.org/10.3102/0034654315582065. **Verified:** DOI + quote. Ids: clark2016. Cited in: 04.

495. Sherwood DE (1988). Effect of Bandwidth Knowledge of Results on Movement Consistency. *Perceptual and Motor Skills*. https://doi.org/10.2466/pms.1988.66.2.535. **Verified:** DOI + quote. Ids: sherwood1988. Cited in: 04.

496. Salmoni AW, Schmidt RA, Walter CB (1984). Knowledge of results and motor learning: A review and critical reappraisal. *Psychological Bulletin*. https://doi.org/10.1037/0033-2909.95.3.355. **Verified:** title-level only. Ids: salmoni1984. Cited in: 04.

497. Winstein CJ, Schmidt RA (1990). Reduced frequency of knowledge of results enhances motor skill learning. *Journal of Experimental Psychology: Learning, Memory, and Cognition*. https://doi.org/10.1037/0278-7393.16.4.677. **Verified:** title-level only. Ids: winstein1990. Cited in: 04.

498. Kahneman D, Tversky A (1979). Prospect Theory: An Analysis of Decision under Risk. *Econometrica*. https://doi.org/10.2307/1914185. **Verified:** DOI + quote. Ids: kahneman1979. Cited in: 05.

499. Kivetz R, Urminsky O, Zheng Y (2006). The Goal-Gradient Hypothesis Resurrected: Purchase Acceleration, Illusionary Goal Progress, and Customer Retention. *Journal of Marketing Research*. https://doi.org/10.1509/jmkr.43.1.39. **Verified:** DOI + quote. Ids: kivetz2006. Cited in: 05.

500. Shortt M, Tilak S, Kuznetcova I, Martens B, Akinkuolie B (2021). Gamification in mobile-assisted language learning: a systematic review of Duolingo literature from public release of 2012 to early 2020. *Computer Assisted Language Learning*. https://doi.org/10.1080/09588221.2021.1933540. **Verified:** DOI + quote. Ids: shortt2021. Cited in: 05.

501. Henderlong J, Lepper MR (2002). The effects of praise on children's intrinsic motivation: A review and synthesis. *Psychological Bulletin*. https://doi.org/10.1037/0033-2909.128.5.774. **Verified:** DOI + quote. Ids: henderlong2002. Cited in: 05.

502. Zeng J, Sun D, Looi C, Fan ACW (2024). Exploring the impact of gamification on students' academic performance: A comprehensive meta-analysis of studies from the year 2008 to 2023. *British Journal of Educational Technology*. https://doi.org/10.1111/bjet.13471. **Verified:** DOI + quote. Ids: zeng2024. Cited in: 05.

503. Cameron J, Pierce WD (1994). Reinforcement, Reward, and Intrinsic Motivation: A Meta-Analysis. *Review of Educational Research*. https://doi.org/10.3102/00346543064003363. **Verified:** DOI + quote. Ids: cameron1994. Cited in: 05.

504. Lepper MR, Greene D, Nisbett RE (1973). Undermining children's intrinsic interest with extrinsic reward: A test of the "overjustification" hypothesis. *Journal of Personality and Social Psychology*. https://doi.org/10.1037/h0035519. **Verified:** DOI + quote. Ids: lepper1973. Cited in: 05.

505. Eisenberger R, Cameron J (1996). Detrimental effects of reward: Reality or myth? *American Psychologist*. https://doi.org/10.1037/0003-066x.51.11.1153. **Verified:** DOI + quote. Ids: eisenberger1996. Cited in: 05.

506. Sailer M, Hense JU, Mayr SK, Mandl H (2017). How gamification motivates: An experimental study of the effects of specific game design elements on psychological need satisfaction. *Computers in Human Behavior*. https://doi.org/10.1016/j.chb.2016.12.033. **Verified:** DOI + quote. Ids: sailer2017. Cited in: 05.

507. Hamari J, Koivisto J, Sarsa H (2014). Does Gamification Work? -- A Literature Review of Empirical Studies on Gamification. *2014 47th Hawaii International Conference on System Sciences*. https://doi.org/10.1109/hicss.2014.377. **Verified:** DOI + quote. Ids: hamari2014. Cited in: 05.

508. Przybylski AK, Rigby CS, Ryan RM (2010). A Motivational Model of Video Game Engagement. *Review of General Psychology*. https://doi.org/10.1037/a0019440. **Verified:** DOI + quote. Ids: przybylski2010. Cited in: 05.

509. Baker RS, Walonoski JA, Heffernan NT, Roll I, Corbett AT, Koedinger KR (2008). Why Students Engage in "Gaming the System" Behavior in Interactive Learning Environments.

Journal of Interactive Learning Research. https://openalex.org/W2115486496. **Verified:** URL + quote. Ids: baker2008. Cited in: 05.

510. Nunes JC, Drèze X (2006). The Endowed Progress Effect: How Artificial Advancement Increases Effort. *Journal of Consumer Research*. https://doi.org/10.1086/500480. **Verified:** DOI + quote. Ids: nunes2006. Cited in: 05.

511. Cheema A, Bagchi R (2011). The Effect of Goal Visualization on Goal Pursuit: Implications for Consumers and Managers. *Journal of Marketing*. https://doi.org/10.1509/jm.75.2.109. **Verified:** DOI + quote. Ids: cheema2011. Cited in: 05.

512. Myers BA (1985). The importance of percent-done progress indicators for computer-human interfaces. *Proceedings of the SIGCHI conference on Human factors in computing systems - CHI '85*. https://doi.org/10.1145/317456.317459. **Verified:** DOI + quote. Ids: myers1985. Cited in: 05.

513. Locke EA, Latham GP (2002). Building a practically useful theory of goal setting and task motivation: A 35-year odyssey. *American Psychologist*. https://doi.org/10.1037/0003-066x.57.9.705. **Verified:** DOI + quote. Ids: locke2002. Cited in: 05.

514. Bodily R, Verbert K (2017). Review of Research on Student-Facing Learning Analytics Dashboards and Educational Recommender Systems. *IEEE Transactions on Learning Technologies*. https://doi.org/10.1109/tlt.2017.2740172. **Verified:** DOI + quote. Ids: bodily2017. Cited in: 05.

515. Matcha W, Uzir NA, Gasevic D, Pardo A (2020). A Systematic Review of Empirical Studies on Learning Analytics Dashboards: A Self-Regulated Learning Perspective. *IEEE Transactions on Learning Technologies*. https://doi.org/10.1109/tlt.2019.2916802. **Verified:** DOI + quote. Ids: matcha2020. Cited in: 05.

516. Sedrakyan G, Malmberg J, Verbert K, Järvelä S, Kirschner PA (2020). Linking learning behavior analytics and learning science concepts: Designing a learning analytics dashboard for feedback to support learning regulation. *Computers in Human Behavior*. https://doi.org/10.1016/j.chb.2018.05.004. **Verified:** title-level only. Ids: sedrakyan2020. Cited in: 05.

517. Rogowsky BA, Calhoun BM, Tallal P (2020). Providing Instruction Based on Students' Learning Style Preferences Does Not Improve Learning. *Frontiers in Psychology*. https://doi.org/10.3389/fpsyg.2020.00164. **Verified:** DOI + quote. Ids: rogowsky2020. Cited in: 05.

518. Willingham DT, Hughes EM, Dobolyi DG (2015). The Scientific Status of Learning Styles Theories. *Teaching of Psychology*. https://doi.org/10.1177/0098628315589505. **Verified:** DOI + quote. Ids: willingham2015. Cited in: 05.

519. Cuevas J (2015). Is learning styles-based instruction effective? A comprehensive analysis of recent research on learning styles. *Theory and Research in Education*. https://doi.org/10.1177/1477878515606621. **Verified:** DOI + quote. Ids: cuevas2015. Cited in: 05.

520. Biwer F, Wiradhany W, oude Egbrink MGA, de Bruin ABH (2023). Understanding effort regulation: Comparing 'Pomodoro' breaks and self-regulated breaks. *British Journal of Educational Psychology*. https://doi.org/10.1111/bjep.12593. **Verified:** DOI + quote. Ids: biwer2023. Cited in: 05.

521. W3C (2023). WCAG 2.2 SC 4.1.3 Status Messages, Level AA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#status-messages. **Verified:** URL + quote. Ids: wcag-4.1.3. Cited in: 05.

522. W3C (2023). WCAG 2.2 SC 2.4.7 Focus Visible, Level AA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#focus-visible. **Verified:** URL + quote. Ids: wcag-2.4.7. Cited in: 05.

523. Wise SL, DeMars CE (2005). Low Examinee Effort in Low-Stakes Assessment: Problems and Potential Solutions. *Educational Assessment*. https://doi.org/10.1207/s15326977ea1001_1. **Verified:** DOI + quote. Ids: wisedemars2005. Cited in: 05.

524. Albulescu P, Macsinga I, Rusu A, Sulea C, Bodnaru A, Tulbure BT (2022). "Give me a break!" A systematic review and meta-analysis on the efficacy of micro-breaks for increasing well-being and performance. *Plos One*. https://doi.org/10.1371/journal.pone.0272460. **Verified:** DOI + quote. Ids: albulescu2022. Cited in: 05.

525. Blasche G, Szabo B, Wagner-Menghin M, Ekmekcioglu C, Gollner E (2018). Comparison of rest-break interventions during a mentally demanding task. *Stress and Health*. https://doi.org/10.1002/smi.2830. **Verified:** DOI + quote. Ids: blasche2018. Cited in: 05.

526. Stankov L, Lee J, Luo W, Hogan DJ (2012). Confidence: A better predictor of academic achievement than self-efficacy, self-concept and anxiety? *Learning and Individual Differences*. https://doi.org/10.1016/j.lindif.2012.05.013. **Verified:** DOI + quote. Ids: stankov2012. Cited in: 05.

527. van Gog T, Sweller J (2015). Not New, but Nearly Forgotten: the Testing Effect Decreases or even Disappears as the Complexity of Learning Materials Increases. *Educational Psychology Review*. https://doi.org/10.1007/s10648-015-9310-x. **Verified:** DOI + quote. Ids: vangog2015. Cited in: 05.

528. Nah FFH (2004). A study on tolerable waiting time: how long are Web users willing to wait? *Behaviour & Information Technology*. https://doi.org/10.1080/01449290410001669914. **Verified:** DOI + quote. Ids: nah2004. Cited in: 05.

529. Ceaparu I, Lazar J, Bessiere K, Robinson J, Shneiderman B (2004). Determining Causes and Severity of End-User Frustration. *International Journal of Human-Computer Interaction*. https://doi.org/10.1207/s15327590ijhc1703_3. **Verified:** DOI + quote. Ids: ceaparu2004. Cited in: 05.

530. Kestin G, Miller K, Klales A, Milbourne T, Ponti G (2025). AI tutoring outperforms in-class active learning: an RCT introducing a novel research-based design in an authentic educational setting. *Scientific Reports*. https://doi.org/10.1038/s41598-025-97652-6. **Verified:** DOI + quote. Ids: kestin2025. Cited in: 05.

531. Kapur M (2008). Productive Failure. *Cognition and Instruction*. https://doi.org/10.1080/07370000802212669. **Verified:** DOI + quote. Ids: kapur2008. Cited in: 05.

532. Chi MT, De Leeuw N, Chiu M, Lavancher C (1994). Eliciting Self-Explanations Improves Understanding. *Cognitive Science*. https://doi.org/10.1207/s15516709cog1803_3. **Verified:** DOI + quote. Ids: chi1994. Cited in: 05.

533. Lee JD, See KA (2004). Trust in Automation: Designing for Appropriate Reliance. *Human Factors: The Journal of the Human Factors and Ergonomics Society*. https://doi.org/10.1518/hfes.46.1.50_30392. **Verified:** DOI + quote. Ids: lee2004. Cited in: 05.

534. Glikson E, Woolley AW (2020). Human Trust in Artificial Intelligence: Review of Empirical Research. *Academy of Management Annals*. https://doi.org/10.5465/annals.2018.0057. **Verified:** DOI + quote. Ids: glikson2020. Cited in: 05.

535. W3C (2023). WCAG 2.2 SC 2.5.5 Target Size (Enhanced), Level AAA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#target-size-enhanced. **Verified:** URL + quote. Ids: wcag-2.5.5. Cited in: 05.

536. Parhi P, Karlson AK, Bederson BB (2006). Target size study for one-handed thumb use on small touchscreen devices. *Proceedings of the 8th conference on Human-computer interaction with mobile devices and services*. https://doi.org/10.1145/1152215.1152260. **Verified:** DOI + quote. Ids: parhi2006. Cited in: 05.

537. Vatavu RD, Cramariuc G, Schipor DM (2015). Touch interaction for children aged 3 to 6 years: Experimental findings and relationship to motor skills. *International Journal of Human-Computer Studies*. https://doi.org/10.1016/j.ijhcs.2014.10.007. **Verified:** title-level only. Ids: vatavu2015. Cited in: 05.

538. W3C (2023). WCAG 2.2 SC 1.4.3 Contrast (Minimum), Level AA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#contrast-minimum. **Verified:** URL + quote. Ids: wcag-1.4.3. Cited in: 05.

539. W3C (2023). WCAG 2.2 SC 1.4.11 Non-text Contrast, Level AA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#non-text-contrast. **Verified:** URL + quote. Ids: wcag-1.4.11. Cited in: 05.

540. Hall RH, Hanna P (2004). The impact of web page text-background colour combinations on readability, retention, aesthetics and behavioural intention. *Behaviour & Information Technology*. https://doi.org/10.1080/01449290410001669932. **Verified:** DOI + quote. Ids: hall2004. Cited in: 05.

541. W3C (2023). WCAG 2.2 SC 1.4.1 Use of Color, Level A. Official standard or regulation. https://www.w3.org/TR/WCAG22/#use-of-color. **Verified:** URL + quote. Ids: wcag-1.4.1. Cited in: 05.

542. W3C (2023). WCAG 2.2 SC 2.3.3 Animation from Interactions, Level AAA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#animation-from-interactions. **Verified:** URL + quote. Ids: wcag-2.3.3. Cited in: 05.

543. Berney S, Bétrancourt M (2016). Does animation enhance learning? A meta-analysis. *Computers & Education*. https://doi.org/10.1016/j.compedu.2016.06.005. **Verified:** title-level only. Ids: berney2016. Cited in: 05.

544. W3C (2023). WCAG 2.2 SC 2.2.1 Timing Adjustable, Level A. Official standard or regulation. https://www.w3.org/TR/WCAG22/#timing-adjustable. **Verified:** URL + quote. Ids: wcag-2.2.1. Cited in: 05.

545. W3C (2023). WCAG 2.2 SC 2.4.1 Bypass Blocks, Level A. Official standard or regulation. https://www.w3.org/TR/WCAG22/#bypass-blocks. **Verified:** URL + quote. Ids: wcag-2.4.1. Cited in: 05.

546. Sung YT, Chang KE, Liu TC (2016). The effects of integrating mobile devices with teaching and learning on students' learning performance: A meta-analysis and research synthesis. *Computers & Education*. https://doi.org/10.1016/j.compedu.2015.11.008. **Verified:** DOI + quote. Ids: sung2016. Cited in: 05.

547. W3C (2023). WCAG 2.2 SC 1.4.10 Reflow, Level AA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#reflow. **Verified:** URL + quote. Ids: wcag-1.4.10. Cited in: 05.

548. Hughes L, Wilkins A (2000). Typography in children's reading schemes may be suboptimal: Evidence from measures of reading rate. *Journal of Research in Reading*. https://doi.org/10.1111/1467-9817.00126. **Verified:** DOI + quote. Ids: hughes2000. Cited in: 05.

549. Milkman KL, Beshears J, Choi JJ, Laibson D, Madrian BC (2011). Using implementation intentions prompts to enhance influenza vaccination rates. *Proceedings of the National Academy of Sciences*. https://doi.org/10.1073/pnas.1103170108. **Verified:** DOI + quote. Ids: milkman2011. Cited in: 05.

550. Castleman BL, Page LC (2015). Summer nudging: Can personalized text messages and peer mentor outreach increase college going among low-income high school graduates? *Journal of Economic Behavior & Organization*. https://doi.org/10.1016/j.jebo.2014.12.008. **Verified:** DOI + quote. Ids: castleman2015. Cited in: 05.

551. Gollwitzer PM, Sheeran P (2006). Implementation Intentions and Goal Achievement: A Meta-analysis of Effects and Processes. *Advances in Experimental Social Psychology*. https://doi.org/10.1016/s0065-2601(06)38002-1. **Verified:** title-level only. Ids: gollwitzer2006. Cited in: 05.

552. Kushlev K, Proulx J, Dunn EW (2016). "Silence Your Phones". *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. https://doi.org/10.1145/2858036.2858359. **Verified:** DOI + quote. Ids: kushlev2016. Cited in: 05.

553. ICO (2020). ICO Age appropriate design code, standard 13 Nudge techniques. Official standard or regulation. https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/childrens-information/childrens-code-guidance-and-resources/age-appropriate-design-a-code-of-practice-for-online-services/13-nudge-techniques/. **Verified:** URL + quote. Ids: ico-aadc-13. Cited in: 05.

554. W3C (2023). WCAG 2.2 SC 3.3.1 Error Identification, Level A. Official standard or regulation. https://www.w3.org/TR/WCAG22/#error-identification. **Verified:** URL + quote. Ids: wcag-3.3.1. Cited in: 05.

555. W3C (2023). WCAG 2.2 SC 3.3.3 Error Suggestion, Level AA. Official standard or regulation. https://www.w3.org/TR/WCAG22/#error-suggestion. **Verified:** URL + quote. Ids: wcag-3.3.3. Cited in: 05.

556. Nielsen J (1994). Enhancing the explanatory power of usability heuristics. *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. https://doi.org/10.1145/191666.191729. **Verified:** title-level only. Ids: nielsen1994. Cited in: 05.

557. Karich AC, Burns MK, Maki KE (2014). Updated Meta-Analysis of Learner Control Within Educational Technology. *Review of Educational Research*. https://doi.org/10.3102/0034654314526064. **Verified:** DOI + quote. Ids: karich2014. Cited in: 05.

558. Katz I, Assor A (2007). When Choice Motivates and When It Does Not. *Educational Psychology Review*. https://doi.org/10.1007/s10648-006-9027-y. **Verified:** title-level only. Ids: katz2007. Cited in: 05.

559. EU (2016). GDPR Article 8: conditions applicable to child's consent. Official standard or regulation. https://gdpr-info.eu/art-8-gdpr/. **Verified:** URL + quote. Ids: gdpr-art8, gdpr-art8-3. Cited in: 05.

560. ICO (2020). ICO Age appropriate design code, summary of the standards. Official standard or regulation. https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/childrens-information/childrens-code-guidance-and-resources/age-appropriate-design-a-code-of-practice-for-online-services/. **Verified:** URL + quote. Ids: ico-aadc-8. Cited in: 05.

561. Reyes I, Wijesekera P, Reardon J, On AEB, Razaghpanah A, Vallina-Rodriguez N, et al. (2018). “Won’t Somebody Think of the Children?” Examining COPPA Compliance at Scale. *Proceedings on Privacy Enhancing Technologies*. https://doi.org/10.1515/popets-2018-0021. **Verified:** DOI + quote. Ids: reyes2018. Cited in: 05.

562. ICO (n.d.). ICO guidance: What are the rules about an ISS and consent? (UK GDPR Article 8). Official standard or regulation. https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/childrens-information/children-and-the-uk-gdpr/what-are-the-rules-about-an-iss-and-consent/. **Verified:** URL + quote. Ids: ico-iss-consent. Cited in: 05.

563. Spain (BOE) (2018). Ley Orgánica 3/2018 (LOPDGDD), artículo 7: Consentimiento de los menores de edad. Official standard or regulation. https://www.boe.es/buscar/act.php?id=BOE-A-2018-16673. **Verified:** URL + quote. Ids: es-lopdgdd-art7. Cited in: 05.

564. Zhao J, Wang G, Dally C, Slovak P, Edbrooke-Childs J, Van Kleek M, et al. (2019). ‘I make up a silly name’. *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*. https://doi.org/10.1145/3290605.3300336. **Verified:** DOI + quote. Ids: zhao2019. Cited in: 05.

565. Kumar P, Naik SM, Devkar UR, Chetty M, Clegg TL, Vitak J (2017). 'No Telling Passcodes Out Because They're Private'. *Proceedings of the ACM on Human-Computer Interaction*. https://doi.org/10.1145/3134699. **Verified:** DOI + quote. Ids: kumar2017. Cited in: 05.

566. Conijn R, Kleingeld A, Matzat U, Snijders C (2022). The fear of big brother: The potential negative side-effects of proctored exams. *Journal of Computer Assisted Learning*. https://doi.org/10.1111/jcal.12651. **Verified:** DOI + quote. Ids: conijn2022. Cited in: 05.

567. Andrejevic M, Selwyn N (2019). Facial recognition technology in schools: critical questions and concerns. *Learning, Media and Technology*. https://doi.org/10.1080/17439884.2020.1686014. **Verified:** DOI + quote. Ids: andrejevic2019. Cited in: 05.

568. Illinois General Assembly (2008). Illinois Biometric Information Privacy Act, 740 ILCS 14. Official standard or regulation. http://www.ilga.gov/legislation/ilcs/ilcs3.asp?ActID=3004&ChapterID=57. **Verified:** title-level only. Ids: bipa. Cited in: 05.

569. Hoepman JH (2014). Privacy Design Strategies. *IFIP Advances in Information and Communication Technology*. https://doi.org/10.1007/978-3-642-55415-5_38. **Verified:** title-level only. Ids: hoepman2014. Cited in: 05.

570. Acquisti A, Brandimarte L, Loewenstein G (2015). Privacy and human behavior in the age of information. *Science*. https://doi.org/10.1126/science.aaa1465. **Verified:** DOI + quote. Ids: acquisti2015. Cited in: 05.

571. FTC (2025). 16 CFR 312.10 (COPPA Rule): data retention and deletion. Official standard or regulation. https://www.ecfr.gov/current/title-16/part-312/section-312.10. **Verified:** URL + quote. Ids: coppa-312.10. Cited in: 05.

572. FTC (2025). 16 CFR 312.6 (COPPA Rule): right of parent to review. Official standard or regulation. https://www.ecfr.gov/current/title-16/part-312/section-312.6. **Verified:** URL + quote. Ids: coppa-312.6. Cited in: 05.

573. EU (2016). GDPR Article 7(1): conditions for consent. Official standard or regulation. https://gdpr-info.eu/art-7-gdpr/. **Verified:** URL + quote. Ids: gdpr-art7-1. Cited in: 05.

574. Pardo A, Siemens G (2014). Ethical and privacy principles for learning analytics. *British Journal of Educational Technology*. https://doi.org/10.1111/bjet.12152. **Verified:** DOI + quote. Ids: pardo2014. Cited in: 05.

575. Hoff KA, Bashir M (2015). Trust in Automation. *Human Factors: The Journal of the Human Factors and Ergonomics Society*. https://doi.org/10.1177/0018720814547570. **Verified:** DOI + quote. Ids: hoff2015. Cited in: 05.

576. Anderson A, Huttenlocher D, Kleinberg J, Leskovec J (2013). Steering user behavior with badges. *Proceedings of the 22nd international conference on World Wide Web*. https://doi.org/10.1145/2488388.2488398. **Verified:** DOI + quote. Ids: anderson2013. Cited in: 05. Listed in the section's sources file; not cited in the text.

577. Keith N, Frese M (2008). Effectiveness of error management training: A meta-analysis. *Journal of Applied Psychology*. https://doi.org/10.1037/0021-9010.93.1.59. **Verified:** DOI + quote. Ids: keith2008. Cited in: 05. Listed in the section's sources file; not cited in the text.

578. Mekler ED, Brühlmann F, Tuch AN, Opwis K (2017). Towards understanding the effects of individual gamification elements on intrinsic motivation and performance. *Computers in Human Behavior*. https://doi.org/10.1016/j.chb.2015.08.048. **Verified:** title-level only. Ids: mekler2017. Cited in: 05. Listed in the section's sources file; not cited in the text.

579. Hanus MD, Fox J (2015). Assessing the effects of gamification in the classroom: A longitudinal study on intrinsic motivation, social comparison, satisfaction, effort, and academic performance. *Computers & Education*. https://doi.org/10.1016/j.compedu.2014.08.019. **Verified:** title-level only. Ids: hanus2015. Cited in: 05. Listed in the section's sources file; not cited in the text.

580. Bai S, Hew KF, Huang B (2020). Does gamification improve student learning outcome? Evidence from a meta-analysis and synthesis of qualitative data in educational contexts. *Educational Research Review*. https://doi.org/10.1016/j.edurev.2020.100322. **Verified:** title-level only. Ids: bai2020. Cited in: 05. Listed in the section's sources file; not cited in the text.

581. Toda AM, Valle PHD, Isotani S (2018). The Dark Side of Gamification: An Overview of Negative Effects of Gamification in Education. *Communications in Computer and Information Science*. https://doi.org/10.1007/978-3-319-97934-2_9. **Verified:** title-level only. Ids: toda2018. Cited in: 05. Listed in the section's sources file; not cited in the text.

582. FTC (2020). Complying with COPPA: Frequently Asked Questions. Official standard or regulation. https://www.ftc.gov/business-guidance/resources/complying-coppa-frequently-asked-questions. **Verified:** title-level only. Ids: ftc-coppa-faq. Cited in: 05. Listed in the section's sources file; not cited in the text.

583. FTC (2025). Children's Online Privacy Protection Rule (final rule amending 16 CFR 312). Official standard or regulation. https://www.federalregister.gov/documents/2025/04/22/2025-05904/childrens-online-privacy-protection-rule. **Verified:** title-level only. Ids: fr-coppa-2025. Cited in: 05. Listed in the section's sources file; not cited in the text.

584. UN Committee on the Rights of the Child (2021). General comment No. 25 on children's rights in relation to the digital environment. Official standard or regulation. https://docs.un.org/en/CRC/C/GC/25. **Verified:** title-level only. Ids: un-crc-gc25. Cited in: 05. Listed in the section's sources file; not cited in the text.

585. Lievens F, Reeve CL, Heggestad ED (2007). An examination of psychometric bias due to retesting on cognitive ability tests in selection settings. *Journal of Applied Psychology*. https://doi.org/10.1037/0021-9010.92.6.1672. **Verified:** DOI + quote. Ids: lievens2007. Cited in: 01.

586. Calamia M, Markon K, Tranel D (2012). Scoring Higher the Second Time Around: Meta-Analyses of Practice Effects in Neuropsychological Assessment. *The Clinical Neuropsychologist*. https://doi.org/10.1080/13854046.2012.680913. **Verified:** DOI + quote. Ids: calamia2012. Cited in: 01.

587. Zell E, Krizan Z (2014). Do People Have Insight Into Their Abilities? A Metasynthesis. *Perspectives on Psychological Science*. https://doi.org/10.1177/1745691613518075. **Verified:** DOI + quote. Ids: zell2014. Cited in: 03.

588. Bernes GA, Villodas M, Coles CD, Kable JA, May PA, Kalberg WO, Sowell ER, Jones KL, Riley EP, Mattson SN (2021). Validity and Reliability of Executive Function Measures in Children With Heavy Prenatal Alcohol Exposure: Correspondence Between Multiple Raters and Laboratory Measures. *Alcoholism: Clinical and Experimental Research*. https://doi.org/10.1111/acer.14547. **Verified:** DOI + quote. Ids: bernes2021. Cited in: 03.

589. Gobet F, Simon HA (1996). Recall of rapidly presented random chess positions is a function of skill. *Psychonomic Bulletin & Review*. https://doi.org/10.3758/bf03212414. **Verified:** DOI + quote. Ids: gobet1996b. Cited in: 04.

590. Astudillo-Sandoval I, Sanchez-Lopez J, Berchicci M, Moreno-Aguirre AJ, Tellez-Alanis B, Maldonado-Paz G (2026). Cognitive foundations of chess performance in novice players. *Journal of Sports Sciences*. https://doi.org/10.1080/02640414.2026.2632505. **Verified:** DOI + quote. Ids: astudillo2026. Cited in: 04.

591. Chabris CF, Hearst ES (2003). Visualization, pattern recognition, and forward search: effects of playing speed and sight of the position on grandmaster chess errors. *Cognitive Science*. https://doi.org/10.1207/s15516709cog2704_3. **Verified:** DOI + quote. Ids: chabris2003. Cited in: 04.

592. Liang B, Fazekas G, Sandler M (2018). Measurement, Recognition, and Visualization of Piano Pedaling Gestures and Techniques. *Journal of the Audio Engineering Society*. https://doi.org/10.17743/jaes.2018.0035. **Verified:** DOI + quote. Ids: liang2018. Cited in: 04.

593. Poolton JM, Masters RSW, Maxwell JP (2005). The relationship between initial errorless learning conditions and subsequent performance. *Human Movement Science*. https://doi.org/10.1016/j.humov.2005.06.006. **Verified:** DOI + quote. Ids: poolton2005. Cited in: 04.

594. Furuya S, Nakamura A, Nagata N (2013). Transfer of piano practice in fast performance of skilled finger movements. *BMC Neuroscience*. https://doi.org/10.1186/1471-2202-14-133. **Verified:** DOI + quote. Ids: furuya2013. Cited in: 04.

595. Thall PF, Simon R (1994). Practical Bayesian Guidelines for Phase IIB Clinical Trials. *Biometrics*. https://doi.org/10.2307/2533377. **Verified:** DOI + quote. Ids: thall1994. Cited in: 04.

596. Tuyl F, Gerlach R, Mengersen K (2008). A Comparison of Bayes–Laplace, Jeffreys, and Other Priors. *The American Statistician*. https://doi.org/10.1198/000313008x267839. **Verified:** DOI + quote. Ids: tuyl2008. Cited in: 04.

597. Shortridge A, Goldsberry K, Adams M (2014). Creating space to shoot: quantifying spatial relative field goal efficiency in basketball. *Journal of Quantitative Analysis in Sports*. https://doi.org/10.1515/jqas-2013-0094. **Verified:** DOI + quote. Ids: shortridge2014. Cited in: 04.

598. Wong-Toi E, Yang HC, Shen W, Hu G (2023). A Joint Analysis for Field Goal Attempts and Percentages of Professional Basketball Players: Bayesian Nonparametric Resource. *Journal of Data Science*. https://doi.org/10.6339/22-jds1062. **Verified:** DOI + quote. Ids: wongtoi2023. Cited in: 04.

599. Hambleton RK, Traub RE (1974). The Effects of Item Order on Test Performance and Stress. *The Journal of Experimental Education*. https://doi.org/10.1080/00220973.1974.10806302. **Verified:** DOI + quote. Ids: hambleton1974. Cited in: 04.

600. Wang TH, Kao CH, Chen HC (2021). Factors Associated with the Equivalence of the Scores of Computer-Based Test and Paper-and-Pencil Test: Presentation Type, Item

Difficulty and Administration Order. *Sustainability*. https://doi.org/10.3390/su13179548. **Verified:** DOI + quote. Ids: wang2021. Cited in: 04.

601. Skipper Y, Douglas K (2012). Is no praise good praise? Effects of positive feedback on children's and university students' responses to subsequent failures. *British Journal of Educational Psychology*. https://doi.org/10.1111/j.2044-8279.2011.02028.x. **Verified:** DOI + quote. Ids: skipper2012. Cited in: 05.
602. Morris BJ, Zentall SR (2014). High fives motivate: the effects of gestural and ambiguous verbal praise on motivation. *Frontiers in Psychology*. https://doi.org/10.3389/fpsyg.2014.00928. **Verified:** DOI + quote. Ids: morris2014. Cited in: 05.
603. Lazar J, Jones A, Hackley M, Shneiderman B (2006). Severity and impact of computer user frustration: A comparison of student and workplace users. *Interacting with Computers*. https://doi.org/10.1016/j.intcom.2005.06.001. **Verified:** DOI + quote. Ids: lazar2006. Cited in: 05.
604. Taipalus T, Grahn H, Ritonummi S, Siitonen V, Vartiainen T, Zhidkikh D (2025). Novice Perceptions on Effective Elements of PostgreSQL Error Messages. *ACM Transactions on Computing Education*. https://doi.org/10.1145/3732790. **Verified:** DOI + quote. Ids: taipalus2025. Cited in: 05.
605. Sinha T, Kapur M (2021). When Problem Solving Followed by Instruction Works: Evidence for Productive Failure. *Review of Educational Research*. https://doi.org/10.3102/00346543211019105. **Verified:** DOI + quote. Ids: sinha2021. Cited in: 03.
606. Kapur M (2014). Productive Failure in Learning Math. *Cognitive Science*. https://doi.org/10.1111/cogs.12107. **Verified:** DOI + quote. Ids: kapur2014. Cited in: 03.
607. Roll I, Aleven V, McLaren BM, Koedinger KR (2011). Improving students' help-seeking skills using metacognitive feedback in an intelligent tutoring system. *Learning and Instruction*. https://doi.org/10.1016/j.learninstruc.2010.07.004. **Verified:** DOI + quote. Ids: roll2011. Cited in: 05.